

Lots of authors

Static Single Assignment Book

Tuesday 8th July, 2014
18:11

Springer

Contents

Part I Vanilla SSA

1	Introduction — (<i>J. Singer</i>)	5
1.1	Definition of SSA	
1.2	Informal Semantics of SSA	
1.3	Comparison with Classical Data-flow Analysis	
1.4	SSA in Context	
1.5	About the Rest of this Book	
1.5.1	Benefits of SSA	
1.5.2	Fallacies about SSA	
2	Properties and flavours — (<i>P. Brisk</i>)	15
2.1	Preliminaries	
2.2	Def-Use and Use-Def Chains	
2.3	Minimality	
2.4	Strict SSA Form and Dominance Property	
2.5	Pruned SSA Form	
2.6	Conventional and Transformed SSA Form	
2.7	A Stronger Definition of Interference	
2.8	Further readings	
3	Standard Construction and Destruction Algorithms — (<i>J. Singer, F. Rastello</i>)	27
3.1	Construction	
3.1.1	Join Sets and Dominance Frontiers	
3.1.2	ϕ -function Insertion	
3.1.3	Variable Renaming	
3.1.4	Summary	
3.2	Destruction	
3.3	SSA Property Transformations	
3.3.1	Pessimistic ϕ -function Insertion	
3.4	Additional reading	

4 Advanced Construction Algorithms for SSA — (<i>D. Das, U. Ramakrishna, V. Sreedhar</i>)	45
4.1 Introduction	
4.2 Basic Algorithm	
4.3 Placing ϕ -functions using DJ-graphs	
4.3.1 Key Observation	
4.3.2 Main Algorithm	
4.4 Placing ϕ -functions using Merge Sets	
4.4.1 Merge Set and Merge Relation	
4.4.2 Iterative Merge Set Computation	
4.5 Computing Iterated Dominance Frontier Using Loop Nesting Forests	
4.5.1 Loop nesting forest	
4.5.2 Main Algorithm	
4.6 Concluding Remarks	
4.7 Further Reading	
5 SSA Reconstruction — (<i>S. Hack</i>)	61
5.1 Introduction	
5.1.1 Live-Range Splitting	
5.1.2 Jump Threading	
5.2 General Considerations	
5.3 Reconstruction based on the Dominance Frontier	
5.4 Search-based Reconstruction	
5.5 Conclusions	
6 Functional representations of SSA — (<i>L. Beringer</i>)	71
6.1 Low-level functional program representations	
6.1.1 Variable assignment versus binding	
6.1.2 Control flow: continuations	
6.1.3 Control flow: direct style	
6.1.4 Let-normal form	
6.2 Functional construction and destruction of SSA	
6.2.1 Initial construction using liveness analysis	
6.2.2 λ -dropping	
6.2.3 Nesting, dominance, loop-closure	
6.2.4 Destruction of SSA	
6.3 Refined block sinking and loop nesting forests	
6.4 Pointers to the literature	
6.5 Concluding remarks	

Part II Analysis

7	Introduction — (<i>A. Ertl</i>)	101
7.1	benefits of SSA	
7.2	properties: suffixed representation can use dense vector, vs DF needs (x,PP) needs hashtable	
7.3	phi node construction precomputes DF merge	
7.4	dominance properties imply that simple forward analysis works	
7.5	Flow insensitive somewhat improved by SSA	
7.6	building defuse chains: for each SSA variable can get pointer to next var	
7.7	improves SSA DF-propagation	
8	Propagating Information using SSA— (<i>F. Brandner, D. Novillo</i>)	103
8.1	Overview	
8.2	Preliminaries	
8.2.1	Solving Data-flow Problems	
8.3	Data-flow Propagation under SSA Form	
8.3.1	Program Representation	
8.3.2	Sparse Data-flow Propagation	
8.3.3	Discussion	
8.3.4	Limitations	
8.4	Example – Copy Propagation	
8.5	Further Reading	
9	Liveness — (<i>B. Boissinot, F. Rastello</i>)	117
9.1	TODO	
9.2	SSA tells nothing about liveness	
9.3	Computing liveness sets	
9.4	Liveness check	
10	Alias analysis — (???)	119
10.1	TODO	
11	Loop tree and induction variables — (<i>S. Pop, A. Cohen</i>)	121
11.1	Part of the CFG and Loop Tree can be exposed from the SSA	
11.1.1	An SSA representation without the CFG	
11.1.2	Natural loop structures on the SSA	
11.1.3	Improving the SSA pretty printer for loops	
11.2	Analysis of Induction Variables	
11.2.1	Stride detection	
11.2.2	Translation to chains of recurrences	
11.2.3	Instantiation of symbols and region parameters	
11.2.4	Number of iterations and computation of the end of loop value	
11.3	Further reading	

12 Redundancy Elimination— (<i>F. Chow</i>)	131
12.1 Introduction	
12.2 Why PRE and SSA are related	
12.3 How SSAPRE Works	
12.3.1 The Safety Criterion	
12.3.2 The Computational Optimality Criterion	
12.3.3 The Lifetime Optimality Criterion	
12.4 Speculative PRE	
12.5 Register Promotion via PRE	
12.5.1 Register Promotion as Placement Optimization	
12.5.2 Load Placement Optimization	
12.5.3 Store Placement Optimization	
12.6 Value-based Redundancy Elimination	
12.6.1 Value Numbering	
12.6.2 Redundancy Elimination under Value Numbering	
12.7 Additional Reading	

Part III Extensions

13 Introduction — (<i>V. Sarkar, F. Rastello</i>)	157
14 Static Single Information Form — (<i>F. Pereira, F. Rastello</i>)	163
14.1 Introduction	
14.2 Static Single Information	
14.2.1 Sparse Analysis	
14.2.2 Partitioned Lattice per Variable (PLV) data-flow Problems	
14.2.3 The Static Single Information Property	
14.2.4 Special instructions used to split live ranges	
14.2.5 Propagating Information Forwardly and Backwardly	
14.2.6 Examples of sparse data-flow analyses	
14.3 Construction and Destruction of the Intermediate Program Representation	
14.3.1 Splitting strategy	
14.3.2 Splitting live ranges	
14.3.3 Variable Renaming	
14.3.4 Dead and Undefined Code Elimination	
14.3.5 Implementation Details	
14.4 Further Reading	
15 Hashed SSA form: HSSA — (<i>M. Mantione, F. Chow</i>)	183
15.1 Introduction	
15.2 SSA and aliasing: μ and χ -functions	
15.3 Introducing “zero versions” to limit the explosion of the number of variables	
15.4 SSA and indirect memory operations: virtual variables	
15.5 GVN and indirect memory operations: HSSA	

15.6	Building HSSA
15.7	Using HSSA
15.8	Further readings
16	Array SSA Form — (<i>Vivek Sarkar, Kathleen Knobe, Stephen Fink</i>) 197
16.1	Introduction
16.2	Array SSA Form
16.3	Sparse Constant Propagation of Array Elements
16.3.1	Array Lattice for Sparse Constant Propagation
16.3.2	Beyond Constant Indices
16.4	Extension to Objects: Redundant Load Elimination
16.4.1	Analysis Framework
16.4.2	Definitely-Same and Definitely-Different Analyses for Heap Array Indices
16.4.3	Scalar Replacement Algorithm
16.5	Further Reading
17	Psi-SSA Form — (<i>F. de Ferrière</i>) 213
17.1	Overview
17.2	Definition and Construction
17.3	SSA algorithms
17.4	Psi-SSA algorithms
17.5	Psi-SSA destruction
17.5.1	Psi-normalize
17.5.2	Psi-web
17.6	Additional reading
18	Graphs and gating functions — (<i>J. Stanier</i>) 227
18.1	Introduction
18.2	Data-flow Graphs
18.3	The SSA Graph
18.3.1	Finding induction variables with the SSA Graph
18.4	Program Dependence Graph
18.4.1	Detecting parallelism with the PDG
18.5	Gating functions and GSA
18.5.1	Backwards symbolic analysis with GSA
18.6	Value State Dependence Graph
18.6.1	Definition of the VSDG
18.6.2	Nodes
18.6.3	γ -Nodes
18.6.4	θ -Nodes
18.6.5	State Nodes
18.6.6	Dead node elimination with the VSDG
18.7	Further readings

Part IV Machine code generation and optimization

19 SSA form and code generation — (<i>B. Dupont de Dinechin</i>)	249
19.1 SSA form engineering issues	
19.1.1 About instructions, operands, operations, and operators	
19.1.2 Representation of instruction semantics	
19.1.3 Operand naming constraints	
19.1.4 Non-kill target operands	
19.1.5 Program representation invariants	
19.2 Code generation phases and the SSA form	
19.2.1 Instruction selection	
19.2.2 If-conversion	
19.2.3 Inner loop optimizations	
19.2.4 Code cleanups and instruction re-selection	
19.2.5 Pre-pass instruction scheduling	
19.2.6 SSA form destruction	
19.2.7 Register allocation	
20 Instruction Code Selection — (<i>D. Ebner, A. Krall, B. Scholz</i>)	263
20.1 Introduction	
20.2 Instruction Code Selection for Tree Patterns on SSA-Graphs	
20.2.1 Partitioned Boolean Quadratic Problem	
20.2.2 Instruction Code Selection with PBQP	
20.3 Extensions and Generalizations	
20.3.1 Instruction Code Selection for DAG Patterns	
20.4 Summary and Further Reading	
21 If-Conversion — (<i>C. Bruel</i>)	277
21.1 Introduction	
21.1.1 Architectural requirements	
21.2 Basic Transformations	
21.2.1 SSA operations on Basic Blocks	
21.2.2 Handling of predicated execution model	
21.3 Global Analysis and Transformations	
21.3.1 SSA Incremental if-conversion algorithm	
21.3.2 Tail Duplication	
21.3.3 Profitability	
21.4 Conclusion	
21.5 Additional reading	
22 SSA destruction for machine code — (<i>F. Rastello</i>)	293
22.1 Introduction	
22.2 Correctness	
22.3 Code Quality	
22.4 Speed and Memory Footprint	
22.5 Further Readings	

23 Register Allocation— (<i>F. Bouchez, S. Hack</i>)	311
23.1 Introduction	
23.1.1 Non-SSA Register Allocators	
23.1.2 SSA form to the rescue of register allocation	
23.1.3 Maxlive and colorability of graphs	
23.2 Spilling	
23.2.1 Spilling under the SSA Form	
23.2.2 Finding Spilling Candidates	
23.2.3 Spilling under SSA	
23.3 Coloring and coalescing	
23.3.1 Greedy coloring scheme	
23.3.2 SSA graphs are colorable with a greedy scheme	
23.3.3 A light tree-scan coloring algorithm	
23.3.4 Coalescing under SSA form	
23.4 Practical and advanced discussions	
23.4.1 Handling registers constraints	
23.4.2 Out-of-SSA and edge splitting	
23.4.3 Parallel copy motion	
23.5 Further Reading & References	
24 Hardware Compilation using SSA— (<i>Pedro C. Diniz, Philip Brisk</i>) ...	331
24.1 Brief history and overview	
24.2 Why use SSA for hardware compilation?	
24.3 Mapping a control-flow graph to hardware	
24.3.1 Basic block mapping	
24.3.2 Basic control-flow graph mapping	
24.3.3 Control-flow graph mapping using SSA	
24.3.4 ϕ -function and multiplexer optimizations	
24.3.5 Implications of using SSA-form in floor-planing	
24.4 Existing SSA-based hardware compilation efforts	
25 Building SSA form in a compiler for PHP	
— (<i>Paul Biggar and David Gregg, Lero@TCD, Trinity College Dublin</i>)	347
25.1 Introduction	
25.2 SSA form in “traditional” languages	
25.3 PHP and aliasing	
25.4 Our whole-program analysis	
25.4.1 The analysis algorithm	
25.4.2 Analysis results	
25.4.3 Building def-use sets	
25.4.4 HSSA	
25.5 Other challenging PHP features	
25.5.1 Run-time symbol tables	
25.5.2 Execution of arbitrary code	
25.5.3 Object handlers	

25.6 Discussion, implications and conclusion	
Index	360
References	363

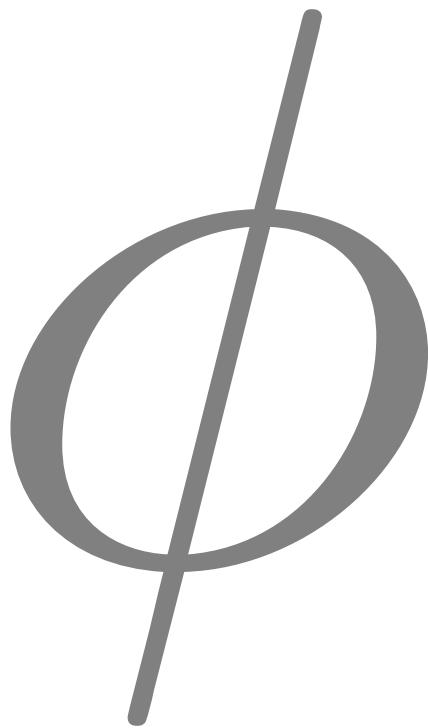
Foreword

Author: Zadeck

Preface

TODO: Roadmap - for students, compiler engineers, etc

TODO: Pre-requisites for reading this book (Tiger book, dragon book, muchnick, etc)?



Part I

Vanilla SSA

Progress: 91%

Progress: 91%

TODO: (tutorial style -
not too heavy)

CHAPTER 1

Introduction

J. Singer

Progress: 95%



Final reading in progress

In computer programming, as in real life, names are useful handles for concrete entities. The key message of this book is that having *unique names* for *distinct entities* reduces uncertainty and imprecision.

For example, consider overhearing a conversation about ‘Homer.’ Without any more contextual clues, you cannot disambiguate between Homer Simpson and Homer the classical Greek poet; or indeed, any other people called Homer that you may know. As soon as the conversation mentions Springfield (rather than Smyrna), you are fairly sure that the Simpsons television series (rather than Greek poetry) is the subject. On the other hand, if everyone had a *unique* name, then there would be no possibility of confusing 20th century American cartoon characters with ancient Greek literary figures.

This book is about the *static single assignment form* (SSA), which is a naming convention for storage locations (variables) in low-level representations of computer programs. The term *static* indicates that SSA relates to properties and analysis of program text (code). The term *single* refers to the uniqueness property of variable names that SSA imposes. As illustrated above, this enables a greater degree of precision. The term *assignment* means variable definitions. For instance, in the code

$$x = y + 1;$$

the variable x is being assigned the value of expression $(y + 1)$. This is a definition, or assignment statement, for x . A compiler engineer would interpret the above assignment statement to mean that the lvalue of x (i.e., the memory location labelled as x) should be modified to store the value $(y + 1)$.

referential transparency

1.1 Definition of SSA

The simplest, least constrained, definition of SSA can be given using the following informal prose:

“A program is defined to be in SSA form if each variable is a target of exactly one assignment statement in the program text.”

However there are various, more specialized, varieties of SSA, which impose further constraints on programs. Such constraints may relate to graph-theoretic properties of variable definitions and uses, or the encapsulation of specific control-flow or data-flow information. Each distinct SSA variety has specific characteristics. Basic varieties of SSA are discussed in Chapter 2. Part III of this book presents more complex extensions.

One important property that holds for all varieties of SSA, including the simplest definition above, is *referential transparency*^{*}: i.e., since there is only a single definition for each variable in the program text, a variable’s value is *independent of its position* in the program. We may refine our knowledge about a particular variable based on branching conditions, e.g. we know the value of x in the conditionally executed block following an `if` statement beginning:

```
if (x == 0)
```

however the *underlying value* of x does not change at this `if` statement. Programs written in pure functional languages are referentially transparent. Referentially transparent programs are more amenable to formal methods and mathematical reasoning, since the meaning of an expression depends only on the meaning of its subexpressions and not on the order of evaluation or side-effects of other expressions. For a referentially opaque program, consider the following code fragment.

```
x = 1;
y = x + 1;
x = 2;
z = x + 1;
```

A naive (and incorrect) analysis may assume that the values of y and z are equal, since they have identical definitions of $(x + 1)$. However the value of variable x depends on whether the current code position is before or after the second definition of x , i.e., variable values depend on their *context*. When a compiler transforms this program fragment to SSA code, it becomes referentially transparent. The translation process involves renaming to eliminate multiple assignment statements for the same variable. Now it is apparent that y and z are equal if and only if x_1 and x_2 are equal.

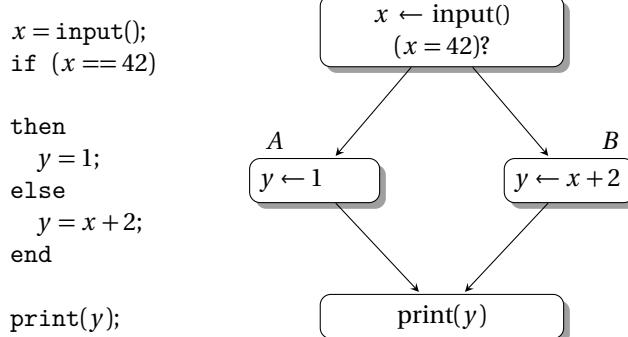
$x_1 = 1;$	<i>variable name</i>
$y = x_1 + 1;$	ϕ -function
$x_2 = 2;$	ϕ -function, as multiplexer
$z = x_2 + 1;$	control-flow graph

1.2 Informal Semantics of SSA

In the previous section, we saw how straightline sequences of code can be transformed to SSA by simple renaming of variable definitions. The *target* of the definition is the variable being defined, on the left-hand side of the assignment statement. In SSA, each definition target must be a unique variable name.^{*} Conversely variable names can be used multiple times on the right-hand side of any assignment statements, as *source* variables for definitions. Throughout this book, renaming is generally performed by adding integer subscripts to original variable names. In general this is an unimportant implementation feature, although it can prove useful for compiler debugging purposes.

The ϕ -function^{*} is the most important SSA concept to grasp. It is a special statement, known as a *pseudo-assignment* function. Some call it a “notational fiction.”¹ The purpose of a ϕ -function^{*} is to merge values from different incoming paths, at control-flow merge points.

Consider the following code example and its corresponding control-flow graph (CFG) representation:



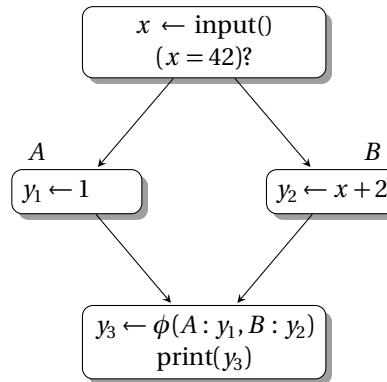
There is a distinct definition of y in each branch of the `if` statement. So multiple definitions of y reach the `print` statement at the control-flow merge point. When a compiler transforms this program to SSA, the multiple definitions of y are renamed as y_1 and y_2 . However the `print` statement could use either variable, dependent on the outcome of the `if` conditional test. A ϕ -function intro-

¹ Kenneth Zadeck reports that ϕ -functions were originally known as *phoney*-functions, during the development of SSA at IBM Research. Although this was an in-house joke, it did serve as the basis for the eventual name.

ϕ -function, as multiplexer
 ϕ -function, parallel semantic of
 ϕ_{if} -function
 γ -function

duces a new variable y_3 , which takes the value of either y_1 or y_2 . Thus the SSA version of the program is:

```
x = input();
if (x == 42)
then
    y1 = 1;
else
    y2 = x + 2;
end
y3 = φ(y1, y2);
print(y3);
```



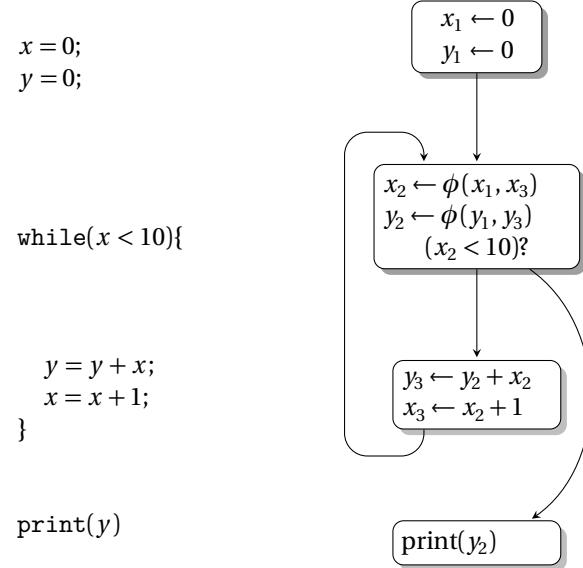
In terms of their position, ϕ -functions^{*} are generally placed at control-flow merge points, i.e., at the heads of basic blocks that have multiple predecessors in control-flow graphs. A ϕ -function at block b has n parameters if there are n incoming control-flow paths to b . The behavior of the ϕ -function is to select dynamically the value of the parameter associated with the actually executed control-flow path into b . This parameter value is assigned to the fresh variable name, on the left-hand-side of the ϕ -function. Such pseudo-functions are required to maintain the SSA property of unique variable definitions, in the presence of branching control flow. Hence, in the above example, y_3 is set to y_1 if control flows from basic block A , and set to y_2 if it flows from basic block B . Notice that the CFG representation here adopts a more expressive syntax for ϕ -functions than the standard one, as it associates predecessor basic block labels B_i with corresponding SSA variable names a_i , i.e., $a_0 = \phi(B_1 : a_1, \dots, B_n : a_n)$. Throughout this book, basic block labels will be omitted from ϕ -function operands when the omission does not cause ambiguity.

It is important to note that, if there are multiple ϕ -functions^{*} at the head of a basic block, then these are executed in parallel, i.e., simultaneously *not* sequentially. This distinction becomes important if the target of a ϕ -function is the same as the source of another ϕ -function, perhaps after optimizations such as copy propagation (see Chapter 8). When ϕ -functions are eliminated in the SSA destruction phase, they are sequentialized using conventional copy operations, as described in Chapters 3 and 22. This subtlety is particularly important in the context of register allocated code (see Chapter 23, e.g., Figure 23.11).

Strictly speaking, ϕ -functions are not directly executable in software, since the dynamic control-flow path leading to the ϕ -function is not explicitly encoded as an input to ϕ -function. This is tolerable, since ϕ -functions are generally only used during static analysis of the program. They are removed before any program interpretation or execution takes place. However, there are various executable extensions of ϕ -functions, such as ϕ_{if}^* or γ^* functions (see Chapter 18), which take an extra predicate parameter to replace the implicit control

dependence that dictates the argument the ϕ -function should select by a condition dependence. Such extensions are useful for program interpretation (see Chapter 18), if-conversion (see Chapter 21), or hardware synthesis (see Chapter 24).

We present one further example in this section, to illustrate how a loop control-flow structure appears in SSA. Here is the non-SSA version of the program and its corresponding control-flow graph SSA version:



The SSA code has two ϕ -functions in the compound loop header statement that merge incoming definitions from before the loop for the first iteration, and from the loop body for subsequent iterations.

It is important to outline that SSA should not be confused with (*dynamic*) single assignment (DSA or simply SA) form used in automatic parallelization. *Static* single assignment does not prevent multiple assignments to a variable during program execution. For instance, in the SSA code fragment above, variables y_3 and x_3 in the loop body are redefined dynamically with fresh values at each loop iteration.

Full details of the SSA construction algorithm are given in Chapter 3. For now, it is sufficient to see that:

1. A ϕ -function has been inserted at the appropriate control-flow merge point where multiple reaching definitions of the same variable converged in the original program.
2. Integer subscripts have been used to rename variables x and y from the original program.

*data-flow analysis
data-flow analysis,
non-zero value
top, \top
data-flow analysis, dense*

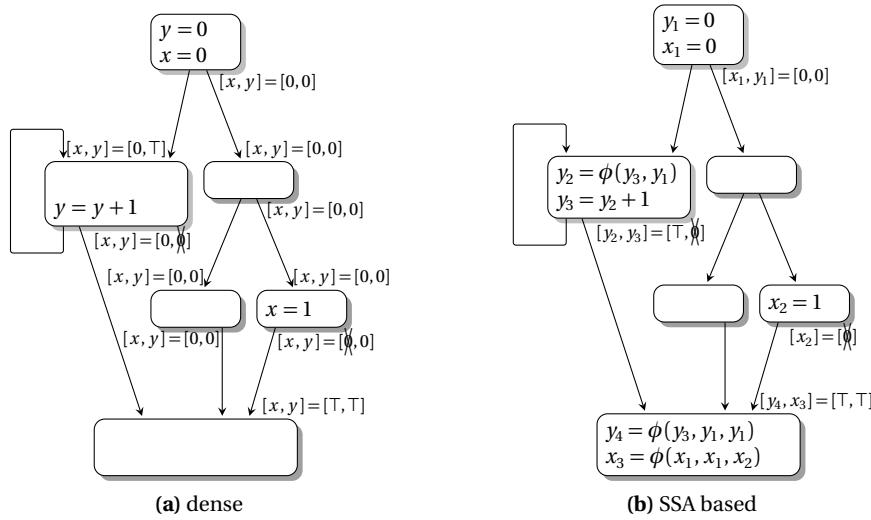


Fig. 1.1 Example control-flow graph for non-zero value analysis, only showing relevant definition statements for variables x and y .

1.3 Comparison with Classical Data-flow Analysis

As we will discover further in Chapters 14 and 8, one of the major advantages of SSA form concerns data-flow analysis.^{*} Data-flow analysis collects information about programs at compile time in order to make optimizing code transformations. During actual program execution, information flows between variables. Static analysis captures this behaviour by propagating *abstract* information, or data-flow facts, using an operational representation of the program such as the control-flow graph (CFG). This is the approach used in classical data-flow analysis.

Often, data-flow information can be propagated more efficiently using a *functional*, or *sparse*, representation of the program such as SSA. When a program is translated into SSA form, variables are renamed at definition points. For certain data-flow problems (e.g. constant propagation) this is exactly the set of program points where data-flow facts may change. Thus it is possible to associate data-flow facts directly with variable names, rather than maintaining a vector of data-flow facts indexed over all variables, at each program point.

Figure 1.1 illustrates this point through an example of non-zero value analysis.^{*} For each variable in a program, the aim is to determine statically whether that variable can contain a zero integer value (i.e., null) at runtime. Here 0 represents the fact that the variable is null, \emptyset the facts that it is non-null, and \top the fact that it is maybe-null. With classical dense data-flow analysis on the CFG in Figure 1.1(a), we would compute information about variables x and y for each of the

entry and exit points of the six basic-blocks in the CFG, using suitable data-flow equations. Using sparse SSA-based^{*} data-flow analysis on Figure 1.1(b), we compute information about each variable based on a simple analysis of its definition statement. This gives us six data-flow facts, one for each SSA version of variables x and y .

*data-flow analysis,
sparse
def-use chains*

For other data-flow problems, properties may change at points that are not variable definitions. These problems can be accommodated in a sparse analysis framework by inserting additional pseudo-definition functions at appropriate points to induce additional variable renaming. See Chapter 14 for one such example. However, this example illustrates some key advantages of the SSA-based analysis.

1. Data-flow information *propagates directly* from definition statements to uses, via the def-use links implicit in the SSA naming scheme. In contrast, the classical data-flow framework propagates information throughout the program, including points where the information does not change, or is not relevant.
2. The results of the SSA data-flow analysis are *more succinct*. In the example, there are fewer data-flow facts associated with the sparse (SSA) analysis than with the dense (classical) analysis.

Part ?? of this textbook gives a comprehensive treatment of some SSA-based data-flow analysis.

1.4 SSA in Context

Historical Context. Throughout the 1980s, as optimizing compiler technology became more mature, various intermediate representations (IRs) were proposed to encapsulate data dependence in a way that enabled fast and accurate data-flow analysis. The motivation behind the design of such IRs was the exposure of direct links between variable definitions and uses, known as *def-use chains*, enabling efficient propagation of data-flow information. Example IRs include the program dependence graph [118] and program dependence web [212]. Chapter 18 gives further details on dependence graph style IRs.

Static single assignment form was one such IR, which was developed at IBM Research, and announced publicly in several research papers in the late 1980s [241, 11, 88]. SSA rapidly acquired popularity due to its intuitive nature and straightforward construction algorithm. The SSA property gives a standardized shape for variable def-use chains, which simplifies data-flow analysis techniques.

Current Usage. The majority of current commercial and open-source compilers, including GCC, Open64, Sun’s HotSpot JVM, IBM’s Java Jikes RVM, Chromium V8

*just-in-time compiler, JIT
back-end, compiler
register allocation
functional language*

JavaScript engine, Mono, and LLVM, use SSA as a key intermediate representation for program analysis. As optimizations in SSA are fast and powerful, SSA is increasingly used in just-in-time (JIT) compilers^{*} that operate on a high-level target-independent program representation such as Java byte-code, CLI byte-code (.NET MSIL), or LLVM bitcode.

Initially developed to facilitate the development of high-level program transformations, SSA form has gained much interest due to its favorable properties that often enable the simplification of algorithms and reduction of computational complexity. Today, SSA form is even adopted for the final code generation phase (see Part ??), i.e., the back-end.^{*} Several industrial and academic compilers, static or just-in-time, use SSA in their back-ends, e.g., LLVM, Java HotSpot, LAO, libFirm, Mono. Many industrial compilers that use SSA form perform SSA elimination before register allocation, including GCC, HotSpot, Jikes, and LLVM. Recent research on register allocation (see Chapter 23) even allows the retention of SSA form until the very end of the code generation process.

SSA for High-Level Languages. So far, we have presented SSA as a useful feature for compiler-based analysis of low-level programs. It is interesting to note that some high-level languages enforce the SSA property. The SISAL language is defined in such a way that programs automatically have referential transparency, since multiple assignments are not permitted to variables. Other languages allow the SSA property to be applied on a per-variable basis, using special annotations like `final` in Java, or `const` and `readonly` in C#.

The main motivation for allowing the programmer to enforce SSA in an explicit manner in high-level programs is that *immutability simplifies concurrent programming*. Read-only data can be shared freely between multiple threads, without any data dependence problems. This is becoming an increasingly important issue, with the trend of multi- and many-core processors.

High-level functional languages^{*} claim referential transparency as one of the cornerstones of their programming paradigm. Thus functional programming supports the SSA property implicitly. Chapter ?? explains the dualities between SSA and functional programming.

1.5 About the Rest of this Book

In this chapter, we have introduced the notion of SSA. The rest of this book presents various aspects of SSA, from the pragmatic perspective of compiler engineers and code analysts. The ultimate goals of this book are:

1. To demonstrate clearly the *benefits* of SSA-based analysis.
2. To dispell the *fallacies* that prevent people from using SSA.

This section gives pointers to later parts of the book that deal with specific topics.

1.5.1 Benefits of SSA

*SSA!myths
myths*

SSA imposes a strict discipline on variable naming in programs, so that each variable has a unique definition. Fresh variable names are introduced at assignment statements, and control-flow merge points. This serves to simplify the structure of variable *def-use* relationships (see Section 2.2) and *live ranges* (see Section 2.4), which underpin data-flow analysis. Part II of this book focus on data-flow analysis using SSA. There are three major advantages to SSA:

Compile time benefit. Certain compiler optimizations can be more efficient when operating on SSA programs, since referential transparency means that data-flow information can be associated directly with variables, rather than with variables at each program point. We have illustrated this simply with the non-zero value analysis in Section 1.3.

Compiler development benefit. Program analyses and transformations can be easier to express in SSA. This means that compiler engineers can be more productive, in writing new compiler passes, and debugging existing passes. For example, the *dead code elimination* pass in GCC 4.x, which relies on an underlying SSA-based intermediate representation, takes only 40% as many lines of code as the equivalent pass in GCC 3.x, which does not use SSA. The SSA version of the pass is simpler, since it relies on the general-purpose, factored-out, data-flow propagation engine.

Program runtime benefit. Conceptually, any analysis and optimization that can be done under SSA form can also be done identically out of SSA form. Because of the compiler development mentioned above, several compiler optimizations show to be more effective when operating on programs in SSA form. These include the class of *control-flow insensitive analyses*, e.g. [139].

1.5.2 Fallacies about SSA

Some people believe that SSA is too cumbersome to be an effective program representation. This book aims to convince the reader that such a concern is unnecessary, given the application of suitable techniques. The table below presents some common myths about SSA, and references in this first part of the book containing material to dispell these myths.*

Myth	Reference
SSA greatly increases the number of variables	Chapter 2 reviews the main varieties of SSA, some of which introduce far fewer variables than the original SSA formulation.
SSA property is difficult to maintain	Chapters 3 and 5 discuss simple techniques for the repair of SSA invariants that have been broken by optimization rewrites.
SSA destruction generates many copy operations	Chapters 3 and 22 present efficient and effective SSA destruction algorithms.

CHAPTER 2

Properties and flavours

P. Brisk

Progress: 95%



Final reading in progress

2.1 Preliminaries

Recall from the previous chapter that a procedure is in SSA form if every variable is defined once, and every use of a variable refers to exactly one definition. Many variations, or flavors, of SSA form that satisfy these criteria can be defined, each offering its own considerations. For example, different flavors vary in terms of the number of ϕ -functions, which affects the size of the intermediate representation; some are more difficult to construct, maintain, and deconstruct compared to others. This chapter explores these SSA flavors and provides insight onto the contexts that favor some over others.

2.2 Def-Use and Use-Def Chains

Under SSA form, each variable is defined once. Def-use chains are data structures that provide for the single definition of a variable the set of all its uses. A use-def chain which under SSA consists of a single name, in turn, uniquely specifies the definition that reaches the use. As we will illustrate further in the book (see Chapter 8) def-use chains are useful for forward data-flow analysis as they provide direct connections that shorten the propagation distance between nodes that generate and use data-flow information.

forward control-flow graph

Because of its single definition per variable property, SSA form simplifies def-use and use-def chains in several ways. First, SSA form simplifies def-use chains as it combines the information as early as possible. This is illustrated by Figure 2.1 where the def-use chains in the non-SSA program requires as many merge as x is used while the corresponding SSA form allows early and more efficient combination.

Second, as it is easy to associate to each variable its single defining operation, use-def chains can be explicitly represented and maintained almost for free. As this constitutes the skeleton of the so called SSA graph (see Chapter 18), when considering a program under SSA form, use-def chains are implicitly considered as a given. The explicit representation of use-def chains simplifies backward propagation, which favors algorithms such as dead-code elimination.

For forward propagation, def-use chains being just the reverse of use-def chains, computing them is also easy, and maintaining them can be done without much efforts either. However, even without def-use chains, some lightweight forward propagation algorithms such as copy-folding are possible: using a single pass that processes operations along a topological order traversal of a forward CFG¹, most definitions are processed prior to their uses. When processing an operation, the use-def chain allows to carry the already computed value of an argument. Conservative merge is performed when, at some loop headers, a ϕ -function encounters an unprocessed argument. Such a lightweight propagation engine shows to be already quite efficient.

2.3 Minimality

SSA construction is a two-phase process: placement of ϕ -functions, followed by renaming. The goal of the first phase is to generate a code that fulfills the further defined single reaching-definition property. Minimality is an additional property of the code with ϕ -functions inserted, but prior to renaming; Chapter 3 describes the classical SSA construction algorithm in detail, while this section focuses primarily on describing the minimality property.

A definition D of variable v reaches a point p in the CFG if there exists a path from D to p that does not pass through another definition of v . We say that a code has the *single reaching-definition property* iff no program point can be reached by two definitions of the same variable. Under the assumption that the single reaching-definition property is fulfilled, the *minimality property* states the minimality of the number of inserted ϕ -functions.

This property can be characterized using the notion of join sets that we introduce next. Let n_1 and n_2 be distinct basic blocks in a CFG. A basic block n_3 , which may or may not be distinct from n_1 or n_2 , is a *join node* of n_1 and n_2 if

¹ A forward control-flow graph is an acyclic reduction of the CFG obtained by removing back-edges.

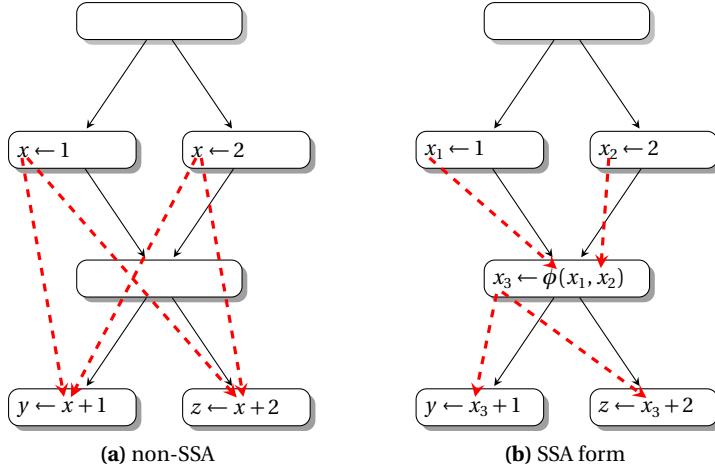


Fig. 2.1 Def-use chains (dashed) for non-SSA form and its corresponding SSA form program

there exist at least two non-empty paths, i.e., paths containing at least one CFG edge, from n_1 to n_3 and from n_2 to n_3 , respectively, such that n_3 is the only basic block that occurs on both of the paths. In other words, the two paths converge at n_3 and no other CFG node. Given a set S of basic blocks, n_3 is a join node of S if it is the join node of at least two basic blocks in S . The set of join nodes of set S is denoted by the set $\mathcal{J}(S)$.

Intuitively, a join set corresponds to the placement of ϕ -functions. In other words, if n_1 and n_2 are basic blocks that both contain the definition of a variable v , then we ought to instantiate ϕ -functions for v at every basic block in $\mathcal{J}(n_1, n_2)$. Generalizing this statement, if D_v is the set of basic blocks containing definitions of v , then ϕ -functions should be instantiated in every basic block in $\mathcal{J}(D_v)$. As inserted ϕ -functions are themselves definition points, some new ϕ -functions should be inserted at $\mathcal{J}(D_v \cup \mathcal{J}(D_v))$. Actually it turns out that $\mathcal{J}(S \cup \mathcal{J}(S)) = \mathcal{J}(S)$, so the join set of the set of definition points of a variable in the original program characterizes exactly the minimum set of program points where ϕ -functions should be inserted.

We are not aware of any optimizations that require a strict enforcement of minimality property. However, placing ϕ -functions only at the join sets can be done easily using a simple topological traversal of the CFG as described in Section 4.5. Classical techniques place ϕ -functions of a variable v at $\mathcal{J}(D_v, r)$, with r the entry node of the CFG. There are good reasons for that as we will explain further. Finally, as explained in Section 3.3 for reducible flow graphs, some copy-propagation engines can easily turn a non-minimal SSA code into a minimal one.

2.4 Strict SSA Form and Dominance Property

A procedure is defined to be *strict* if every variable is defined before it is used along every path from the entry to exit point; otherwise, it is *non-strict*. Some languages, such as Java, impose strictness as part of the language definition; others, such as C/C++, impose no such restrictions. The code in Figure 2.2(a) is non-strict as there exists a path from the entry to the use of a that does not go through the definition. If this path is taken through the CFG during the execution, then a will be used without ever being assigned a value. Although this may be permissible in some cases, it is usually indicative of a programmer error or poor software design.

Under SSA, because there is only a single (static) definition per variable, strictness is equivalent to the *dominance property*: each use of a variable is dominated by its definition. In a CFG, basic block n_1 *dominates* basic block n_2 if every path in the CFG from the entry point to n_2 includes n_1 . By convention, every basic block in a CFG dominates itself. Basic block n_1 *strictly dominates* n_2 if n_1 dominates n_2 and $n_1 \neq n_2$. We use the symbols $n_1 \text{ dom } n_2$ and $n_1 \text{ sdom } n_2$ to denote dominance and strict dominance respectively.

Adding a (undefined) pseudo-definition of each variable to the procedure's entry point ensures strictness. The single reaching-definition property discussed previously mandates that each program point be reachable by exactly one definition (or pseudo-definition) of each variable. If a program point U is a use of variable v , then the reaching definition D of v will dominate U ; otherwise, there would be a path from the CFG entry node to U that does not include D . If such a path existed, then the program would not be in SSA form, and a ϕ -function would need to be inserted somewhere in $\mathcal{J}(r, D)$ as in our example of Figure 2.2(b) where \perp represents the undefined pseudo-definition. The so called *minimal SSA form* is a variant of SSA form that satisfies both the minimality and dominance properties. As shall be seen in Chapter 3, minimal SSA form is obtained by placing the ϕ -functions of variable v at $\mathcal{J}(D_v, r)$ using the formalism of dominance frontier. If the original procedure is non-strict, conversion to minimal SSA will create a strict SSA-based representation. Here, strictness refers solely to the SSA representation; if the input program is non-strict, conversion to and from strict SSA form cannot address errors due to uninitialized variables. To finish with, the use of an implicit pseudo-definition in the CFG entry node to enforce strictness does not change the semantics of the program by any means.

SSA with dominance property is useful for many reasons that directly originate from the structural properties of the variable live-ranges. The immediate dominator or idom of a node n is the unique node that strictly dominates n but does not strictly dominate any other node that strictly dominates n . All nodes but the entry node have immediate dominators. A dominator tree is a tree where each node's children are those nodes it immediately dominates. Because the immediate dominator is unique, it is a tree with the entry node as root. For each variable, its live-range^{*}, i.e., the set of program points where it is live, is a sub-tree

*liveness
interference graph*

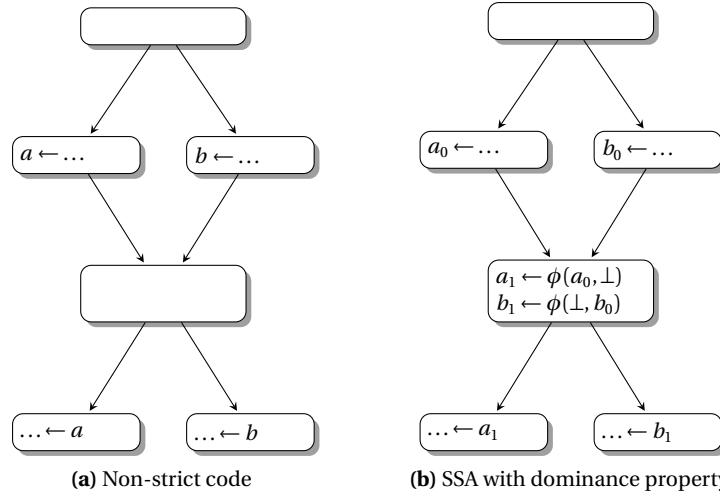


Fig. 2.2 A non-strict code and its corresponding strict SSA form. The presence of \perp indicates a use of an undefined value.

of the dominator tree. Among other consequences of this property, we can cite the ability to design a fast and efficient method to query whether a variable is live at point q or an iteration free algorithm to computes liveness sets (see Chapter 9); this property allows also efficient algorithms to test whether two variables interfere (see Chapter 22). Note that in the general case, a variable is considered to be live^{*} at a program point if there exists a definition of that variable that can reach this point (reaching definition analysis), *and* if there exists a definition free path to a use (upward-exposed use analysis). For strict programs, any program point from which you can reach a use without going through a definition is necessarily reachable from a definition.

Another elegant consequence is that the intersection graph of live-ranges belongs to a special class of graphs called chordal graphs. Chordal graphs are significant because several problems that are NP-complete on general graphs have efficient linear-time solutions on chordal graphs, including graph coloring. Graph coloring plays an important role in register allocation, as the register assignment problem can be expressed as a coloring problem of the interference graph, the graph that links two variables with an interference edge if they cannot be assigned the same register (color). If we suppose that two variables interfere if their live-ranges intersect (see Section 2.7 for further discussions about this hypothesis), the underlying chordal property highly simplifies the assignment problem otherwise considered NP-complete. In particular, a traversal of the dominator tree called a tree scan can color all of the variables in the pro-

gram, without requiring the explicit construction of an interference graph. The trees scan algorithm can be used for register allocation, which is discussed in greater detail in Chapter 23.

As we have already mentioned, most ϕ -function placement algorithms are based on the notion of dominance frontier (see chapters 3 and 4) consequently do provide the dominance property. As we will see in Chapter 3, this property can be broken by copy-propagation: in our example of Figure 2.2(b), the argument a_1 of the copy represented by $a_2 = \phi(a_1, \perp)$ can be propagated and every occurrence of a_2 can be safely replaced by a_1 ; the now identity ϕ -function can then be removed obtaining the initial SSA but non strict code. Making a non-strict SSA code, strict is somehow as “difficult” as SSA construction (actually we need a pruned version as described below). Still the “strictification” usually concerns only a few variables and a restricted region of the CFG: the incremental update described in Chapter 5 will do the work with less efforts.

2.5 Pruned SSA Form

One drawback of Minimal SSA form is that it may place ϕ -functions for a variable at a point in the control-flow graph where the variable was not actually live prior to SSA. Many program analyses and optimizations, including register allocation, are only concerned with the region of a program where a given variable is live. The primary advantage of eliminating those dead ϕ -functions over minimal SSA form is that it has far fewer ϕ -functions in most cases. It is possible to construct such a form while still maintaining the minimality and dominance properties otherwise. The new constraint is that every *use point* for a given variable must be reached by exactly one definition, as opposed to all program points. Pruned SSA form satisfies these properties.

Under minimal SSA, ϕ -functions for variable v are placed at the entry points of basic blocks belonging to the set $\mathcal{J}(S, r)$. Under pruned SSA, we suppress the instantiation of a ϕ -function at the beginning of a basic block if v is not live at the entry point of that block. One possible way to do this is to perform liveness analysis prior to SSA construction, and then use the liveness information to suppress the placement of ϕ -functions as described above; another approach is to construct minimal SSA and then remove the dead ϕ -functions using dead code elimination. Details can be found in Chapter 3.

Figure 2.3(a) shows an example of minimal non pruned SSA. The corresponding pruned SSA form would remove the dead ϕ -function that defines Y_3 .

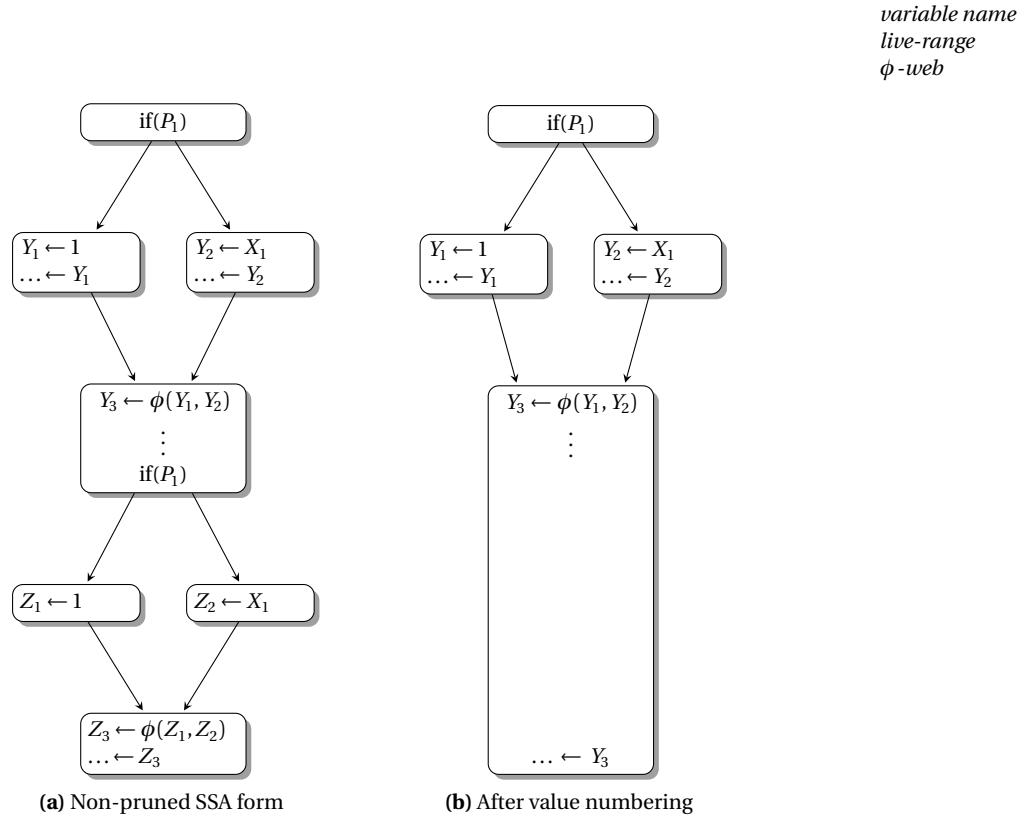


Fig. 2.3 Non pruned SSA form allows value numbering to determine that Y_3 and Z_3 have the same value.

2.6 Conventional and Transformed SSA Form

In many non-SSA and graph coloring based register allocation schemes, register assignment is done at the granularity of webs. In this context, a web is the maximum unions of def-use chains that have either a use or a def in common. As an example, the code of Figure 2.4(a) leads to two separate webs for variable a . The conversion to minimal SSA form replaces each web of a variable v in the pre-SSA program with some variable names v_i^* . In pruned SSA, these variable names partition the live-range^{*} of the web: at every point in the procedure where the web is live, *exactly* one variable v_i is also live; and none of the v_i are live at any point where the web is not.

Based on this observation, we can partition the variables in a program that has been converted to SSA form into ϕ -equivalence classes that we will refer as ϕ -webs^{*}. We say that x and y are ϕ -related to one another if they are refer-

enced by the same ϕ -function, i.e., if x and y are either parameters or defined by the ϕ -function. The transitive closure of this relation defines an equivalence relation that partitions the variables defined locally in the procedure into equivalence classes, the ϕ -webs. Intuitively, the ϕ -equivalence class of a resource represents a set of resources “connected” via ϕ -functions. For any freshly constructed SSA code, the ϕ -webs exactly correspond to the register web of the original non-SSA code.

Conventional SSA form (C-SSA) is defined as SSA form for which each ϕ -web is interference free. Many program optimizations such as copy-propagation may transform a procedure from conventional to a non-conventional (T-SSA for Transformed-SSA) form, in which some variables belonging to the same ϕ -web interfere with one another. Figure 2.4(c) shows the corresponding transformed SSA form of our previous example: here variable a_1 interferes with variables a_2 , a_3 , and a_4 .

Bringing back the conventional property of a T-SSA code is as “difficult” as destructing SSA (see Chapter 3). The translation out of conventional SSA form is straightforward: each ϕ -web can be replaced with a single variable; all definitions and uses are renamed to use the new variable, and all ϕ -functions involving this equivalence class are removed. SSA destruction starting from non-conventional SSA form can be performed through a conversion to conventional SSA form as an intermediate step. This conversion is achieved by inserting copy operations that dissociate interfering variables from the connecting ϕ -functions. As those copy instructions will have to be inserted at some points to get rid of ϕ -functions, for machine level transformations such as register allocation or scheduling, T-SSA provides an inaccurate view of the resource usage. Another motivation for sticking to C-SSA is that the names used in the original program might help capturing some properties otherwise difficult to discover. Lexical partial redundancy elimination (PRE) as described in Chapter 12 illustrates this point.

Apart from those specific examples most current compilers choose not to maintain the conventional property. Still, we should outline that, as later described in Chapter 22, checking if a given ϕ -web is (and if necessary turning it back to) interference free can be done in linear time (instead of the naive quadratic time algorithm) in the size of the ϕ -web.

2.7 A Stronger Definition of Interference

Throughout this chapter, two variables have been said to interfere if their live ranges intersect. Intuitively, two variables with overlapping lifetimes will require two distinct storage locations; otherwise, a write to one variable will overwrite the value of the other. In particular, this definition has applied to the discussion

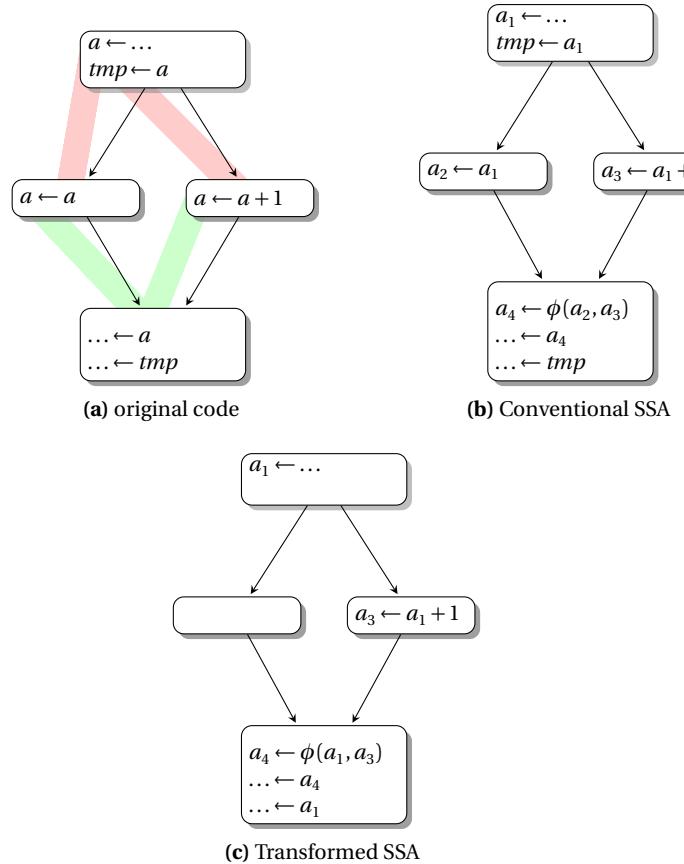


Fig. 2.4 (a) Non-SSA register webs of variable a and (b) corresponding SSA ϕ -webs $\{a_1\}$ and $\{a_2, a_3, a_4\}$; (b) conventional and (c) corresponding transformed SSA form after copy propagating a_1

of interference graphs and the definition of conventional SSA form, as described above.

Although it suffices for correctness, this is a fairly restrictive definition of interference, based on static considerations. The ultimate notion of interference, that is obviously undecidable because of a reduction to the Halting problem, should decide for two distinct variables whether there exists an execution for which they simultaneously hold two different values. Several “static” extensions to our simple definition are still possible, in which, under very specific conditions, variables whose live ranges overlap one another may not interfere. We present two examples.

live-range

Firstly, consider the double-diamond graph of Figure 2.2(a) again, which although non-strict, is correct as soon as the two if-conditions are the same. Even if a and b are unique variables with overlapping live ranges, the paths along which a and b are respectively used and defined are mutually exclusive with one another. In this case, the program will either pass through the definition of a and the use of a , or the definition of b and the use of b , since all statements involved are controlled by the same condition, albeit at different conditional statements in the program. Since only one of the two paths will ever execute, it suffices to allocate a single storage location that can be used for a or b . Thus, a and b do not actually interfere with one another. A simple way to refine the interference test is to check if one of the variable is live at the definition point of the other. This relaxed but correct notion of interference would not make a and b of Figure 2.2(a) interfere while variables a_1 and b_1 of Figure 2.2(b) would still interfere. This example illustrates the fact that live-range splitting required here to make the code fulfill the dominance property may lead to less accurate analysis results. As far as the interference is concerned, for a SSA code with dominance property, the two notions are strictly equivalent: two live-ranges intersect iff one contains the definition of the other.

Secondly, consider two variables u and v , whose live ranges overlap. If we can prove that u and v will always hold the same value at every place where both are live, then they do not actually interfere with one another. Since they always have the same value, a single storage location can be allocated for both variables, because there is only one unique value between them. Of course, this new criteria is in general undecidable. Still, a technique such as global value numbering that is straightforward to implement under SSA (See Chapter 12.6.1) can make a fairly good job, especially in the presence of a code with many variable-to-variable copies, such as one obtained after a naive SSA destruction pass (See Chapter 3). In that case (See Chapter 22), the difference between the refined notion of interference and non-value based one is significant.

This refined notion of interference has significant implications if applied to SSA form. In particular, the interference graph of a procedure is no longer chordal, as any edge between two variables whose lifetimes overlap could be eliminated by this property.

2.8 Further readings

The advantages of def-use and use-def chains provided for almost free under SSA are well illustrated in Chapters 8 and 14.

The notion of minimal SSA and a corresponding efficient algorithm to compute it were introduced by Cytron et al. [89]. For this purpose they extensively develop the notion of dominance frontier of a node n , $\mathcal{DF}(n) = \mathcal{J}(n, r)$. The fact that $\mathcal{J}^+(S) = \mathcal{J}(S)$ has been actually discovered later, with a simple proof

by Wolfe [291]. More details about the theory on (iterated) dominance frontier can be found in Chapters 3 and 4. The post-dominance frontier, which is its symmetric notion, also known as the control dependence graph finds many applications. Further discussions on control dependence graph can be found in Chapter 18.

Most SSA papers implicitly consider the SSA form to fulfill the dominance property. The first technique that really exploits the structural properties of the strictness is the fast SSA destruction algorithm developed by Budimlić et al. [59] and revisited in Chapter 22.

The notion of pruned SSA has been introduced by Choi, Cytron, and Ferrante [71]. The example of Figure 2.3 to illustrate the difference between pruned and non pruned SSA have been borrowed from Cytron et al. [89]. The notions of conventional and transformed SSA were introduced by Sreedhar et al. in their seminal paper [255] for destructing SSA form. The description of the existing techniques to turn a general SSA into either a minimal, a pruned, a conventional, or a strict SSA is provided in Chapter 3.

The ultimate notion of interference was first discussed by Chaitin in his seminal paper [67] that presents the graph coloring approach for register allocation. His interference test is similar to the refined test presented in this chapter. In the context of SSA destruction, Chapter 22 addresses the issue of taking advantage of the dominance property with this refined notion of interference.

CHAPTER 3

Standard Construction and Destruction Algorithms

J. Singer
F. Rastello

Progress: 95%

Final reading in progress

This chapter^{*} describes the standard algorithms for construction and destruction of SSA form. SSA *construction* refers to the process of translating a non-SSA program into one that satisfies the SSA constraints. In general, this transformation occurs as one of the earliest phases in the middle-end of an optimizing compiler, when the program has been converted to three-address intermediate code. SSA *destruction* is sometimes called out-of-SSA translation. This step generally takes place in an optimizing compiler after all SSA optimizations have been performed, and prior to code generation. Note however that there are specialized code generation techniques that can work directly on SSA-based intermediate representations such as instruction selection (see Chapter 20), if-conversion (see Chapter 21), and register allocation (see Chapter 23).

The algorithms presented in this chapter are based on material from the seminal research papers on SSA. These original algorithms are straightforward to implement and have acceptable efficiency. Therefore such algorithms are widely implemented in current compilers. Note that more efficient, albeit more complex, alternative algorithms have been devised. These are described further in Chapters 4 and 22.

Figure 3.1 shows the control-flow graph (CFG) of an example program. The set of nodes is $\{r, A, B, C, D, E\}$, and the variables used are $\{x, y, \text{tmp}\}$. Note that the program shows the complete control-flow structure, denoted by directed edges between the nodes. However, the program only shows statements that define relevant variables, together with the unique `return` statement at the exit point of the CFG. All of the program variables are undefined on entry. On certain control-flow paths, some variables may be used without being defined, e.g. x

*insertion, of ϕ -function
live-range splitting
reachable, by a definition
live-range
single definition property
renaming, of variables
variable name
minimal SSA form*

on the path $r \rightarrow A \rightarrow C$. We discuss this issue later in the chapter. We intend to use this program as a running example throughout the chapter, to demonstrate various aspects of SSA construction.

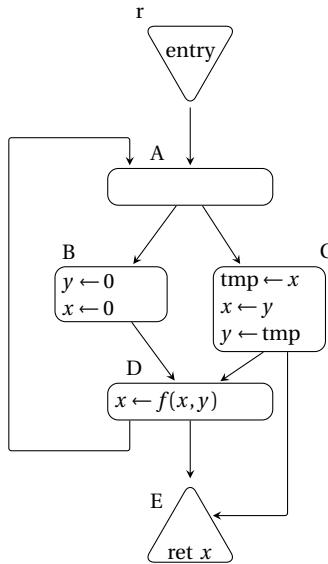


Fig. 3.1 Example control-flow graph, before SSA construction occurs.

3.1 Construction

The original construction algorithm for SSA form consists of two distinct phases.

1. **ϕ -function insertion*** performs *live-range splitting** to ensure that any use of a given variable v is reached¹ by exactly one definition of v . The resulting live-ranges exhibit the property of having a single definition, which occurs at the beginning of each live-range.*
2. **Variable renaming*** assigns a unique variable name* to each live-range. This second phase rewrites variable names in program statements such that the program text contains only one definition of each variable, and every use refers to its corresponding unique reaching definition.

As already outlined in Chapter 2, there are different flavors of SSA with distinct properties. In this chapter, we focus on the *minimal SSA form**.

¹ A program point p is said to be *reachable* by a definition of v if there exists a path in the CFG from that definition to p that does not contain any other definition of v .

3.1.1 Join Sets and Dominance Frontiers

In order to explain how ϕ -function insertions occur, it will be helpful to review the related concepts of *join sets* and *dominance frontiers*.

For a given set of nodes S in a CFG, the join set $\mathcal{J}(S)$ is the set of *join nodes* of S , i.e., nodes in the CFG that can be reached by two (or more) distinct elements of S using disjoint paths. Join sets were introduced in Chapter 2, Section 2.3.

Let us consider some join set examples from the program in Figure 3.1.

1. $\mathcal{J}(\{B, C\}) = \{D\}$, since it is possible to get from B to D and from C to D along different, non-overlapping, paths.
2. Again, $\mathcal{J}(\{r, A, B, C, D, E\}) = \{A, D, E\}$, since the nodes A , D , and E are the only nodes with multiple predecessors in the program.

The *dominance frontier*^{*} of a node n , $DF(n)$, is the border of the CFG region that is dominated by n . More formally,

- node x *strictly dominates*^{*} node y if x dominates y and $x \neq y$;
- the set of nodes $DF(n)$ contains all nodes x such that n dominates a predecessor of x but n does not strictly dominate x .

For instance, in our Figure 3.1, the dominance frontier of the y defined in block B is the first operation of D, while the DF of the y defined in block C would be the first operations of D and E.

Note that DF is defined over individual nodes, but for simplicity of presentation, we overload it to operate over sets of nodes too, i.e., $DF(S) = \bigcup_{s \in S} DF(s)$. The *iterated dominance frontier* $DF^+(S)$ is obtained by iterating the computation of DF until reaching a fixed point, i.e., it is the limit $DF_{i \rightarrow \infty}(S)$ of the sequence:

$$\begin{aligned} DF_1(S) &= DF(S) \\ DF_{i+1}(S) &= DF(S \cup DF_i(S)) \end{aligned}$$

Construction of minimal SSA requires for each variable v the insertion of ϕ -functions at $\mathcal{J}(\text{Defs}(v))$, where $\text{Defs}(v)$ is the set of nodes that contain definitions of v . The original construction algorithm for SSA form uses the iterated dominance frontier $DF^+(\text{Defs}(v))$. This is an over-approximation of join set, since $DF^+(S) = \mathcal{J}(S \cup \{r\})$, i.e., the original algorithm assumes an *implicit* definition of every variable at the entry node r .

<i>join set</i>
<i>dominance frontier</i> , DF
<i>dominates, strictly</i>
<i>iterated dominance frontier</i> , DF^+

3.1.2 ϕ -function Insertion

This concept of dominance frontiers naturally leads to a straightforward approach that places ϕ -functions on a per-variable basis. For a given variable v , we place ϕ -functions at the iterated dominance frontier $DF^+(\text{Defs}(v))$ where

dominance property

$\text{Defs}(v)$ is the set of nodes containing definitions of v . This leads to the construction of SSA form that has the dominance property*, i.e., where each the definition of each renamed variable dominates its entire live-range.

Consider again our running example from Figure 3.1. The set of nodes containing definitions of variable x is $\{B, C, D\}$. The iterated dominance frontier of this set is $\{A, D, E\}$ (it is also the DF, no iteration needed here). Hence we need to insert ϕ -functions for x at the beginning of nodes A , D , and E . Figure 3.2 shows the example CFG program with ϕ -functions for x inserted.

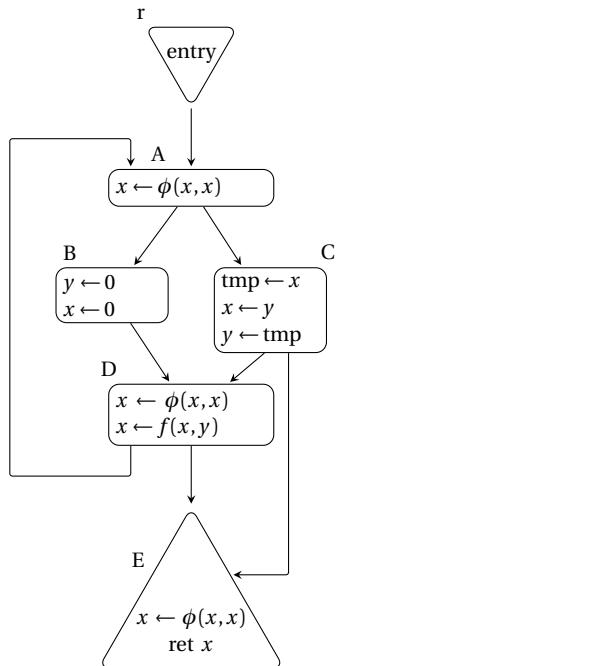


Fig. 3.2 Example control-flow graph, including inserted ϕ -functions for variable x .

As far as the actual algorithm for ϕ -functions insertion is concerned, we will assume that the dominance frontier of each CFG node is pre-computed and that the iterated dominance frontier is computed on the fly, as the algorithm proceeds. The algorithm works by inserting ϕ -functions iteratively using a worklist of definition points, and flags (to avoid multiple insertions). The corresponding pseudo-code for ϕ -function insertion is given in Algorithm 3.1. The worklist of nodes W is used to record definition points that the algorithm has not yet processed, i.e., it has not yet inserted ϕ -functions at their dominance frontiers. Because a ϕ -function is itself a definition, it may require further ϕ -functions to be inserted. This is the cause of node insertions into the worklist W during iterations of the inner loop in Algorithm 3.1. Effectively, we compute the iterated dominance frontier on the fly. The set F is used to avoid repeated insertion of

ϕ -functions on a single block. Dominance frontiers of distinct nodes may intersect, e.g., in the example CFG in Figure 3.1, $\text{DF}(B)$ and $\text{DF}(C)$ both contain D , but once a ϕ -function for a particular variable has been inserted at a node, there is no need to insert another, since a single ϕ -function per variable handles all incoming definitions of that variable to that node.

Algorithm 3.1: Standard algorithm for inserting ϕ -functions for a variable v .

```

1  $F \leftarrow \{\}$ ;                                /* set of basic blocks where  $\phi$  is added */
2 for  $v$ : variable names in original program do
3    $W \leftarrow \{\}$ ;                                /* set of basic blocks */
4   for  $d \in \text{Defs}(v)$  do
5     let  $B$  be the basic block containing  $d$ ;
6      $W \leftarrow W \cup \{B\}$ ;
7   end
8   while  $W \neq \{\}$  do
9     remove a basic block  $X$  from  $W$ ;
10    for  $Y$ : basic block  $\in \text{DF}(X)$  do
11      if  $Y \notin F$  then
12        add  $v \leftarrow \phi(\dots)$  at entry of  $Y$ ;
13         $F \leftarrow F \cup \{Y\}$ ;
14        if  $Y \notin \text{Defs}(v)$  then
15           $W \leftarrow W \cup \{Y\}$ ;
16        end
17      end
18    end
19  end
20 end
```

We give gives a walkthrough example of Algorithm 3.1 in Table 3.1. It shows the stages of execution for a single iteration of the outermost for loop, inserting ϕ -functions for variable x . Each row represents a single iteration of the while loop that iterates over the worklist W . The table shows the values of X , F and W at the start of each while loop iteration. At the beginning, the CFG looks like Figure 3.1. At the end, when all the ϕ -functions for x have been placed, then the CFG looks like Figure 3.2.

while loop #	X	$\text{DF}(X)$	F	W
-	-	-	$\{\}$	$\{B, C, D\}$
1	B	$\{D\}$	$\{D\}$	$\{C, D\}$
2	C	$\{D, E\}$	$\{D, E\}$	$\{D, E\}$
3	D	$\{E, A\}$	$\{D, E, A\}$	$\{E, A\}$
4	E	$\{\}$	$\{D, E, A\}$	$\{A\}$
5	A	$\{A\}$	$\{D, E, A\}$	$\{\}$

Table 3.1 Walkthrough of placement of ϕ -functions for variable x in example CFG.

*DJ-graph@DJ-graph
dominance edge, D-edge
join edge, J-edge
dominance frontier edge,
DF-edge
def-use chains*

Providing the dominator tree is given, the computation of the dominance frontier is quite straightforward. As illustrated by Figure 3.3, this can be understood using the DJ-graph notation. The skeleton of the DJ-graph is the dominator tree of the CFG that makes the D-edges^{*} (dominance edges). This is augmented with J-edges (join edges) that correspond to all edges of the CFG whose source does not strictly dominate its destination. A DF-edge (dominance frontier edge) is an edge whose destination is in the dominance frontier of its source. By definition, there is a DF-edge (a, b) between every CFG nodes a, b such that a dominates a predecessor of b , but does not strictly dominate b . In other words, for each J-edge (a, b) , all ancestors of a (including a) that do not strictly dominate b have b in their dominance frontier. For example, in Figure 3.3, (F, G) is a J-edge, so $\{(F, G), (E, G), (B, G)\}$ are DF-edges. This leads to the pseudo-code given in Algorithm 3.2, where for every edge (a, b) we visit all ancestors of a to add b to their dominance frontier.

Since the iterated dominance frontier is simply the transitive closure of the dominance frontier, we can define the DF^+ -graph as the transitive closure of the DF-graph. In our example, as $\{(C, E), (E, G)\}$ are DF-edges, (C, G) is a DF^+ -edge. Hence, a definition of x in C will lead to inserting ϕ -functions in E and G . We can compute the iterated dominance frontier for each variable independently, as outlined in this chapter, or “cache” it to avoid repeated computation of the iterated dominance frontier of the same node. This leads to more sophisticated algorithms detailed in Chapter 4.

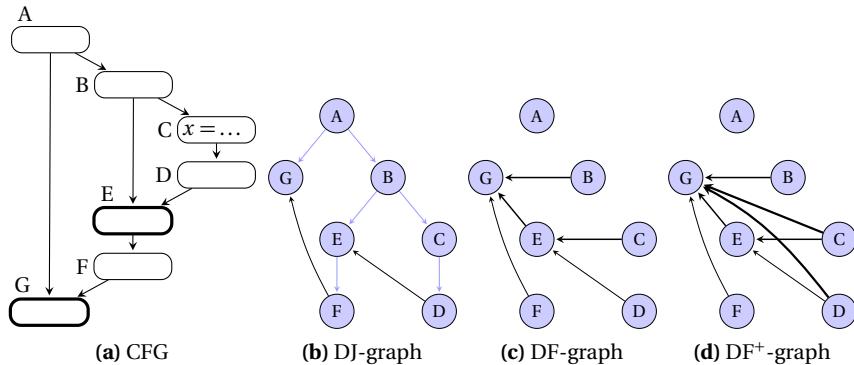


Fig. 3.3 An example CFG and its corresponding DJ-graph (D-edges are top-down), DF-graph and DF^+ -graph.

Once ϕ -functions have been inserted using this algorithm, the program usually still contains several definitions per variable, however, now there is a single definition statement in the CFG that reaches each use. For each variable use in a ϕ -function, it is conventional to treat them as if the use actually occurs on the corresponding incoming edge or at the end of the corresponding predecessor node. If we follow this convention, then def-use chains^{*} are aligned with the

Algorithm 3.2: Algorithm for computing the dominance frontier of each CFG node.

```

1 for  $(a, b) \in$  CFG edges do
2    $x \leftarrow a;$ 
3   while  $x$  does not strictly dominate  $b$  do
4      $\text{DF}(x) \leftarrow \text{DF}(x) \cup b;$ 
5      $x \leftarrow \text{immediate dominator}(x);$ 
6   end
7 end
*
```

*renaming, of variables
live-range splitting
variable!version
variable!name*

CFG dominator tree. In other words, *the single definition that reaches each use dominates that use*.

3.1.3 Variable Renaming

* To obtain the desired property of a static single assignment per variable, it is necessary to perform variable renaming, which is the second phase in the SSA construction process. ϕ -function insertions have the effect of splitting the live-range(s)* of each original variable into pieces. The variable renaming phase associates to each individual live-range a new variable name, also called a *version*.^{*} The pseudo-code for this process is presented in Algorithm 3.3. Because of the dominance property outlined above, it is straightforward to rename variables using a depth-first traversal of the dominator tree. During the traversal, for each variable v , it is necessary to remember the version of its unique reaching definition at some point p in the graph. This corresponds to the closest definition that dominates p . In Algorithm 3.3, we compute and cache the reaching definition for v in the per-variable slot “ $v.\text{reachingDef}$ ” that is updated as the algorithm traverses the dominator tree of the SSA graph. This per-variable slot stores the in-scope, “new” variable name (version) for the equivalent variable at the same point in the un-renamed program.

The variable renaming algorithm translates our running example from Figure 3.1 into the SSA form of Figure 3.4a. The table in Figure 3.4b gives a walk-through example of Algorithm 3.3, only considering variable x . The labels l_i mark instructions in the program that mention x , shown in Figure 3.4a. The table records (1) when $x.\text{reachingDef}$ is updated from x_{old} into x_{new} due to a call of updateReachingDef, and (2) when $x.\text{reachingDef}$ is x_{old} before a definition statement, then x_{new} afterwards.

The original presentation of the renaming algorithm uses a per-variable stack that stores all variable versions that dominate the current program point, rather than the slot-based approach outlined above. In the stack-based algorithm, a

Algorithm 3.3: Renaming algorithm for second phase of SSA construction

```

/* rename variable definitions and uses to have one definition per
   variable name
*/
1 foreach v : Variable do
2   |   v.reachingDef  $\leftarrow \perp$ ;
3 end
4 foreach BB : basic Block in depth-first search preorder traversal of the dominance tree
do
5   |   foreach i : instruction in linear code sequence of BB do
6     |     foreach v : variable used by non- $\phi$ -function i do
7       |       updateReachingDef( $v, i$ );
8       |       replace this use of  $v$  by  $v$ .reachingDef in  $i$ ;
9     |     end
10    |     foreach v : variable defined by i (may be a  $\phi$ -function) do
11      |       updateReachingDef( $v, i$ );
12      |       create fresh variable  $v'$ ;
13      |       replace this definition of  $v$  by  $v'$  in  $i$ ;
14      |        $v'$ .reachingDef  $\leftarrow v$ .reachingDef;
15      |        $v$ .reachingDef  $\leftarrow v'$ ;
16    |     end
17  |   end
18  |   foreach  $\phi$  :  $\phi$ -function in a successor of BB do
19    |     foreach v : variable used by  $\phi$  do
20      |       updateReachingDef( $v, \phi$ );
21      |       replace this use of  $v$  by  $v$ .reachingDef in  $\phi$ ;
22    |   end
23  | end
24 end
*
```

Procedure updateReachingDef(v,i) Utility function for SSA renaming

Data: v : variable from program
Data: i : instruction from program

```

/* search through chain of definitions for  $v$  until we find the closest
   definition that dominates  $i$ , then update  $v$ .reachingDef in-place with
   this definition
*/
1  $r \leftarrow v$ .reachingDef;
2 while not ( $r == \perp$  or definition( $r$ ) dominates  $i$ ) do
3   |    $r \leftarrow r$ .reachingDef;
4 end
5  $v$ .reachingDef  $\leftarrow r$ ;
*
```

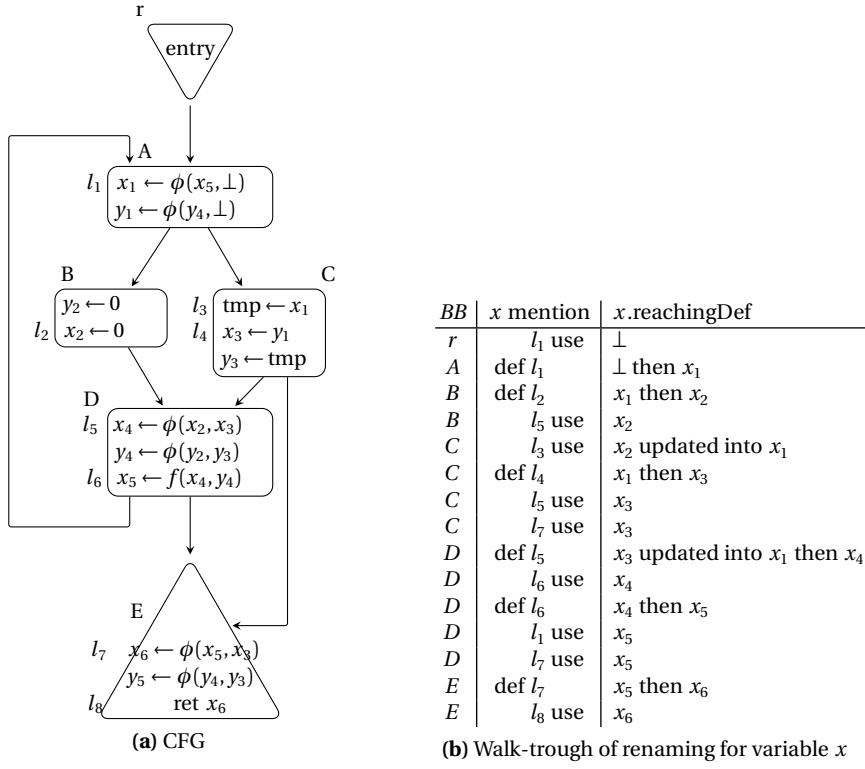


Fig. 3.4 SSA form of the example of Figure 3.1.

stack value is pushed when a variable definition is processed (while we explicitly update the reachingDef field at this point). The top stack value is peeked when a variable use is encountered (we read from the reachingDef field at this point). Multiple stack values may be popped when moving to a different node in the dominator tree (we always check whether we need to update the reachingDef field before we read from it). While the slot-based algorithm requires more memory, it can take advantage of an existing working field for a variable, and be more efficient in practice.

3.1.4 Summary

Now let us review the flavour of SSA form that this simple construction algorithm produces. We refer back to several SSA properties that were introduced in Chapter 2.

minimal SSA form
pruned SSA form
conventional SSA form,
 CSSA
strict SSA form
dominance property, SSA
form with
undefined variable
SSA destruction
conventional SSA form
coalescing, of variables
 ϕ -web
conventional SSA

- It is *minimal* (see Section 2.3).^{*} After the ϕ -function insertion phase, but before variable renaming, the CFG contains the minimal number of inserted ϕ -functions to achieve the property that exactly one definition of each variable v reaches every point in the graph.
- It is *not pruned* (see Section 2.5).^{*} Some of the inserted ϕ -functions may be dead, i.e., there is not always an explicit use of the variable subsequent to the ϕ -function (e.g., y_5 in Figure 3.4a).
- It is *conventional* (see Section 2.6).^{*} The transformation that renames all ϕ -related variables into a unique representative name and then removes all ϕ -functions is a correct SSA-destruction algorithm.
- Finally, it has the *dominance* property (see Section 2.4).^{*} Each variable use is dominated by its unique definition. This is due to the use of iterated dominance frontiers during the ϕ -placement phase, rather than join sets. Whenever the iterated dominance frontier of the set of definition points of a variable differs from its join set, then at least one program point can be reached both by r (the entry of the CFG) and one of the definition points. In other words, as in Figure 3.1, one of the uses of the ϕ -function inserted in block A for x does not have any actual reaching definition that dominates it. This corresponds to the \perp value used to initialize each `reachingDef` slot in Algorithm 3.3. Actual implementation code can use a `NULL` value, create a fake undefined variable at the entry of the CFG, or create undefined pseudo-operations on-the-fly just before the particular use.

3.2 Destruction

*

SSA form is a sparse representation of program information, which enables simple, efficient code analysis and optimization. Once we have completed SSA based optimization passes, and certainly before code generation, it is necessary to eliminate ϕ -functions since these are not executable machine instructions. This elimination phase is known as *SSA destruction*.

When freshly constructed, an untransformed SSA code is conventional and its destruction is straightforward.^{*} One simply has to rename all ϕ -related variables (source and destination operands of the same ϕ -function) into a unique representative variable. Then, each ϕ -function should have syntactically identical names for all its operands, and thus can be removed to coalesce the related live-ranges.^{*}

We refer to a set of ϕ -related variables as a ϕ -web.^{*} We recall from Chapter 2 that conventional SSA is defined as a flavor under which each ϕ -web is free from interferences. Hence, if all variables of a ϕ -web^{*} have non-overlapping live-ranges then the SSA form is conventional. The discovery of ϕ -webs can be performed efficiently using the classical *union-find* algorithm with a disjoint-set

data structure, which keeps track of a set of elements partitioned into a number of disjoint (non-overlapping) subsets. The ϕ -webs discovery algorithm is presented in Algorithm 3.4.

Algorithm 3.4: The ϕ -webs discovery algorithm, based on the union-find pattern

```

1 begin
2   for each variable  $v$  do
3     | phiweb( $v$ )  $\leftarrow \{v\}$ ;
4   end
5   for each instruction of the form  $a_{\text{dest}} = \phi(a_1, \dots, a_n)$  do
6     | for each source operand  $a_i$  in instruction do
7       |   | union(phiweb( $a_{\text{dest}}$ ), phiweb( $a_i$ ))
8     | end
9   end
10 end
*
```

parallel copy
critical edge
edge splitting
parallel copy
transformed SSA, T-SSA
conventional SSA, C-SSA
edge splitting

While freshly constructed SSA code is conventional, this may not be the case after performing some optimizations such as copy propagation. Going back to conventional SSA form requires the insertion of copies. The simplest (although not the most efficient) way to destroy non-conventional SSA form is to split all *critical edges*, and then replace ϕ -functions by copies^{*} at the end of predecessor basic blocks. A critical edge is an edge from a node with several successors to a node with several predecessors.^{*} The process of splitting an edge, say (b_1, b_2) , involves replacing edge (b_1, b_2) by (i) an edge from b_1 to a freshly created basic block and by (ii) another edge from this fresh basic block to b_2 . As ϕ -functions have a parallel semantic, i.e., have to be executed simultaneously not sequentially, the same holds for the corresponding copies inserted at the end of predecessor basic blocks. To this end, a pseudo instruction called a *parallel copy*^{*} is created to represent a set of copies that have to be executed in parallel. The replacement of parallel copies by sequences of simple copies is handled later on. Algorithm 3.5 presents the corresponding pseudo-code that makes non-conventional SSA conventional.^{*} As already mentioned, SSA destruction of such form is straightforward. However, Algorithm 3.5 can be slightly modified to directly destruct SSA by deleting line 17, replacing a'_i by a_0 in the following lines, and adding “remove the ϕ -function” after them.

We stress that the above destruction technique has several drawbacks: first because of specific architectural constraints, region boundaries, or exception handling code, the compiler might not permit the splitting of a given edge; second, the resulting code contains many temporary-to-temporary copy operations. In theory, reducing the frequency of these copies is the role of the coalescing during the register allocation phase. A few memory- and time-consuming coalescing heuristics mentioned in Chapter 23 can handle the removal of these copies effectively. Coalescing can also, with less effort, be performed prior to the

*conservative coalescing
register allocation
aggressive coalescing
parallel copy,
sequentialization of*

Algorithm 3.5: Critical Edge Splitting Algorithm for making non-conventional SSA form conventional.

```

1 begin
2   | foreach  $B$ : basic block of the CFG do
3     |   | let  $(E_1, \dots, E_n)$  be the list of incoming edges of  $B$ ;
4     |   | foreach  $E_i = (B_i, B)$  do
5       |     |   let  $PC_i$  be an empty parallel copy instruction;
6       |     |   if  $B_i$  has several outgoing edges then
7         |       |     create fresh empty basic block  $B'_i$ ;
8         |       |     replace edge  $E_i$  by edges  $B_i \rightarrow B'_i$  and  $B'_i \rightarrow B$ ;
9         |       |     insert  $PC_i$  in  $B'_i$ ;
10        |     |   end
11        |     |   else
12          |       |     append  $PC_i$  at the end of  $B_i$ ;
13        |     |   end
14   |   end
15   | foreach  $\phi$ -function at the entry of  $B$  of the form  $a_0 = \phi(B_1 : a_1, \dots, B_n : a_n)$  do
16     |   | foreach  $a_i$  (argument of the  $\phi$ -function corresponding to  $B_i$ ) do
17       |     |   let  $a'_i$  be a freshly created variable ;
18       |     |   add copy  $a'_i \leftarrow a_i$  to  $PC_i$ ;
19       |     |   replace  $a_i$  by  $a'_i$  in the  $\phi$ -function;
20     |   | end
21   |   end
22 end
23 end

```

register allocation phase. As opposed to a (so-called conservative) coalescing^{*} during register allocation, this *aggressive* coalescing^{*} would not cope with the interference graph colorability. Further, the process of copy insertion itself might take a substantial amount of time and might not be suitable for dynamic compilation. The goal of Chapter 22 is to cope both with non-splittable edges and difficulties related to SSA destruction at machine code level, but also aggressive coalescing in the context of resource constrained compilation.

Once ϕ -functions have been replaced by parallel copies, we need to sequentialize the parallel copies, i.e., replace them by a sequence of simple copies. This phase can be performed immediately after SSA destruction or later on, perhaps even after register allocation (see Chapter 23). It might be useful to postpone the copy sequentialization since it introduces arbitrary interference between variables. As an example, $a_1 \leftarrow a_2 \parallel b_1 \leftarrow b_2$ (where $inst_1 \parallel inst_2$ represents two instructions $inst_1$ and $inst_2$ to be executed simultaneously) can be sequentialized into $a_1 \leftarrow a_2; b_1 \leftarrow b_2$ which would make b_2 interfere with a_1 while the other way round $b_1 \leftarrow b_2; a_1 \leftarrow a_2$ would make a_2 interfere with b_1 instead.

If we still decide to replace parallel copies into a sequence of simple copies immediately after SSA destruction, this can be done as shown in Algorithm 3.6.^{*} To see that this algorithm converges, one can visualize the parallel copy as a graph where nodes represent resources and edges represent transfer of values: the

number of steps is exactly the number of cycles plus the number of non-self edges of this graph. The correctness comes from the invariance of the behavior of *seq*, *pcopy*. An optimized implementation of this algorithm will be presented in Chapter 22.

*conventional, making
SSA
strict, making SSA*

Algorithm 3.6: Replacement of parallel copies with sequences of sequential copy operations.

```

1 begin
2   let pcopy denote the parallel copy to be sequentialized;
3   let seq=() denote the sequence of copies;
4   while  $\neg[\forall(b \leftarrow a) \in pcopy, a = b]$  do
5     if  $\exists(b \leftarrow a) \in pcopy$  s.t.  $\nexists(c \leftarrow b) \in pcopy$  then
6       /* b is not live-in of pcopy */
7       append b  $\leftarrow$  a to seq;
8       remove copy b  $\leftarrow$  a from pcopy;
9     end
10    else
11      /* pcopy is only made-up of cycles; Break one of them */
12      let b  $\leftarrow$  a  $\in$  pcopy s.t. a  $\neq$  b;
13      let a' be a freshly created variable;
14      append a'  $\leftarrow$  a to seq;
15      replace in pcopy b  $\leftarrow$  a into b  $\leftarrow$  a';
16    end
17  end
18 end

```

3.3 SSA Property Transformations

As discussed in Chapter 2, SSA comes in different flavors. This section describes algorithms that transform arbitrary SSA code into the desired flavor. Making SSA *conventional** corresponds exactly to the first phase of SSA destruction (described in Section 3.2) that splits critical edges and introduces parallel copies (sequentialized later in bulk or on-demand) around ϕ -functions. As already discussed, this straightforward algorithm has several drawbacks addressed in Chapter 22.

Making SSA *strict*, i.e., fulfill *the dominance property**, is as “hard” as constructing SSA. Of course, a pre-pass through the graph can detect the offending variables that have definitions that do not dominate their uses. Then there are several possible single-variable ϕ -function insertion algorithms (see Chapter 4) that can be used to patch up the SSA, by restricting attention to the set of non-conforming variables. The renaming phase can also be applied with the same filtering process. As the number of variables requiring repair might be a

*pruned SSA, building
dead ϕ -functions
dead code elimination
def-use chains
use-def chains
semi-pruned SSA*

small proportion of all variables, a costly traversal of the whole program can be avoided by building the def-use chains (for non-conforming variables) during the detection pre-pass. Renaming can then be done on a per-variable basis or better (if pruned SSA is preferred) the reconstruction algorithm presented in Chapter 5 can be used for both ϕ -functions placement and renaming.

The construction algorithm described above does not build *pruned SSA* form.^{*} If available, liveness information can be used to filter out the insertion of ϕ -functions wherever the variable is not live: the resulting SSA form is pruned. Alternatively, pruning SSA form is equivalent to a dead-code elimination pass^{*} after SSA construction. As use-def chains are implicitly provided by SSA form, dead- ϕ -function elimination simply relies on marking actual uses (non- ϕ -function ones) as *useful* and propagating *usefulness* backward through ϕ -functions. Algorithm 3.7 presents the relevant pseudo-code for this operation. Here, *stack* is used to store useful and unprocessed variables defined by ϕ -functions.

Algorithm 3.7: ϕ -function pruning algorithm.

```

1 begin
2   |   stack  $\leftarrow$  ();
3   |   /* -- initial marking phase -- */ 
4   |   foreach I : instruction of the CFG in dominance order do
5   |   |   if I is  $\phi$ -function defining variable a then mark a as useless;
6   |   |   else I is not  $\phi$ -function
7   |   |   |   foreach x : source operand of I do
8   |   |   |   |   if x is defined by some  $\phi$ -function then
9   |   |   |   |   |   mark x as useful; stack.push(x);
10  |   |   |   |   end
11  |   |   end
12  |   end
13  |   /* -- usefulness propagation phase -- */ 
14  |   while stack not empty do
15  |   |   a  $\leftarrow$  stack.pop();
16  |   |   let I be the  $\phi$ -function that defines a;
17  |   |   foreach x : source operand of I do
18  |   |   |   if x is marked as useless then
19  |   |   |   |   mark x as useful; stack.push(x);
20  |   |   |   end
21  |   |   end
22  |   |   /* -- final pruning phase -- */ 
23  |   |   foreach I:  $\phi$ -function do
24  |   |   |   if destination operand of I marked as useless then delete I;
25 end

```

To construct pruned SSA form via dead code elimination, it is generally much faster to first build *semi-pruned SSA form*^{*} rather than minimal SSA form, and

then apply dead code elimination. Semi-pruned SSA form is based on the observation that many variables are *local*, i.e., have a small live-range that is within a single basic block. Consequently, pruned SSA would not instantiate any ϕ -functions for these variables. Such variables can be identified by a linear traversal over each basic block of the CFG. All of these variables can be filtered out: minimal SSA form restricted to the remaining variables gives rise to the so called semi-pruned SSA form.

local variable insertion of ϕ -function, pessimistic minimal, making SSA non-copy-propagation strict, making SSA non-reducible CFG minimal, making SSA

3.3.1 Pessimistic ϕ -function Insertion

Construction of minimal SSA, as outlined in Section 3.1, comes at the price of a sophisticated algorithm involving the computation of iterated dominance frontiers. Alternatively, ϕ -functions may be inserted in a *pessimistic* fashion, as detailed below.^{*} In the pessimistic approach, a ϕ -function is inserted at the start of each CFG node (basic block) for each variable that is live at the start of that node. A less sophisticated, or *crude*, strategy is to insert a ϕ -function for each variable at the start of each CFG node; the resulting SSA will not be pruned. When a pessimistic approach is used, many inserted ϕ -functions are redundant. Code transformations such as code motion, or other CFG modifications, can also introduce redundant ϕ -functions, i.e., make a minimal SSA program become non-minimal.^{*}

The application of *copy propagation*^{*} and rule-based ϕ -function rewriting can remove many of these redundant ϕ -functions. As already mentioned in Chapter 2, copy propagation can break the dominance property by propagating variable a_j through ϕ -functions of the form $a_i = \phi(a_{x_1}, \dots, a_{x_k})$ where all the source operands are syntactically equal to either a_j or \perp . If we want to avoid breaking the dominance property we simply have to avoid applying copy propagations that involve \perp . A more interesting rule is the one that propagates a_j through a ϕ -function of the form $a_i = \phi(a_{x_1}, \dots, a_{x_k})$ where all the source operands are syntactically equal to either a_i or a_j . These ϕ -functions turn out to be “identity” operations, where all the source operands become syntactically identical and equivalent to the destination operand. As such, they can be trivially eliminated from the program. Identity ϕ -functions can be simplified this way from inner to outer loops. To be efficient, variable def-use chains should be pre-computed—otherwise as many iterations as the maximum loop depth might be necessary for copy propagation. When the CFG is reducible,² this simplification produces minimal SSA.^{*} The underlying reason is that the def-use chain of a given ϕ -web is, in this context, isomorphic with a subgraph of the reducible CFG: all nodes except those that can be reached by two different paths (from actual non- ϕ -function definitions) can be simplified by iterative application of $T1$ (removal of

² A CFG is reducible if there are no jumps into the middle of loops from the outside, so the only entry to a loop is through its header. Section 3.4 gives a fuller discussion of reducibility, with pointers to further reading.

self edge) and $T2$ (merge of a node with its unique predecessor) graph reductions. Figure 3.5 illustrates these graph transformations. To be precise, we can simplify an SSA program as follows:

1. Remove any ϕ -function $a_i = \phi(a_{x_1}, \dots, a_{x_k})$ where all a_{x_n} are a_i . This corresponds to $T1$ reduction.
2. Remove any ϕ -function $a_i = \phi(a_{x_1}, \dots, a_{x_k})$ where all $a_{x_n} \in \{a_i, a_j\}$. Replace all occurrences of a_i by a_j . This corresponds to $T2$ possibly followed by $T1$ reduction.

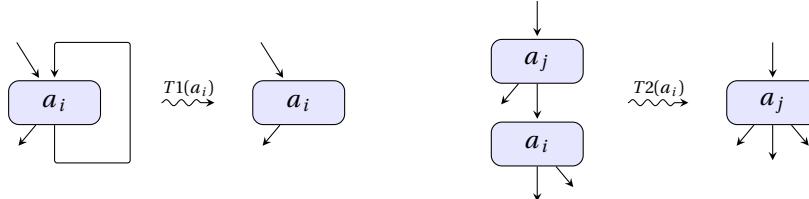


Fig. 3.5 $T1$ and $T2$ rewrite rules for graph reduction, applied to def-use relations between SSA variables.

This approach can be implemented using a worklist, which stores the candidate nodes for simplification. Using the graph made up of def-use chains (see Chapter 18), the worklist can be initialized with successors of non- ϕ -functions. However, for simplicity, we may initialize it with all ϕ -functions. Of course, if loop nesting forest information is available, the worklist can be avoided by traversing the CFG in a single pass from inner to outer loops, and in a topological order within each loop (header excluded). But since we believe the main motivation for this approach to be its simplicity, the pseudo-code shown in Algorithm 3.8 uses a work queue.

This algorithm is guaranteed to terminate in a fixed number of steps. At every iteration of the while loop, it removes a ϕ -function from the work queue W . Whenever it adds new ϕ -functions to W , it removes a ϕ -function from the program. The number of ϕ -functions in the program is bounded so the number of insertions to W is bounded. The queue could be replaced by a worklist, and the insertions/removals done at random. The algorithm would be less efficient, but the end result would be the same.

3.4 Additional reading

The early literature on SSA form [88, 90] introduces the two phases of the construction algorithm we have outlined in this chapter, and discusses algorithmic complexity on common and worst-case inputs. These initial presentations trace

Algorithm 3.8: Removal of redundant ϕ -functions using rewriting rules and work queue.

```

1 begin
2   |    $W \leftarrow ()$ ;
3   |   foreach  $I$ :  $\phi$ -function in program in reverse post-order do
4   |   |    $W.\text{enqueue}(I)$ ; mark  $I$  ;
5   |   end
6   while  $W$  not empty do
7   |   |    $I \leftarrow W.\text{dequeue}()$ ; unmark  $I$ ;
8   |   |   let  $I$  be of the form  $a_i = \phi(a_{x_1}, \dots, a_{x_k})$ ;
9   |   |   if all source operands of the  $\phi$ -function are  $\in \{a_i, a_j\}$  then
10  |   |   |   Remove  $I$  from program;
11  |   |   |   foreach  $I'$ : instruction different from  $I$  that uses  $a_i$  do
12  |   |   |   |   replace  $a_i$  by  $a_j$  in  $I'$ ;
13  |   |   |   |   if  $I'$  is a  $\phi$ -function and  $I'$  is not marked then
14  |   |   |   |   |   mark  $I'$ ;  $W.\text{enqueue}(I')$ ;
15  |   |   |   |   end
16  |   |   end
17  |   end
18 end
19 end

```

the ancestry of SSA form back to early work on data-flow representations by Shapiro and Saint [246].

Briggs et al. [51] discuss pragmatic refinements to the original algorithms for SSA construction and destruction, with the aim of reducing execution time. They introduce the notion of semi-pruned form, show how to improve the efficiency of the stack-based renaming algorithm, and describe how copy propagation must be constrained to preserve correct code during SSA destruction.

There are numerous published descriptions of alternative algorithms for SSA construction, in particular for the ϕ -function insertion phase. The pessimistic approach that first inserts ϕ -functions at all control-flow merge points and then removes unnecessary ones using simple $T1/T2$ rewrite rules was proposed by Aycock and Horspool [21]. Brandis and Mössenböck [47] describe a simple, syntax-directed approach to SSA construction from well structured high-level source code. Throughout this textbook, we consider the more general case of SSA construction from arbitrary CFGs.

A *reducible* CFG is one that will collapse to a single node when it is transformed using repeated application of $T1/T2$ rewrite rules. Aho et al. [6] describe the concept of reducibility and trace its history in early compilers literature.

Sreedhar and Gao [251] pioneer linear-time complexity ϕ -function insertion algorithms based on DJ-graphs. These approaches have been refined by other researchers. Chapter 4 explores these alternative construction algorithms in depth.

Blech et al. [34] formalize the semantics of SSA, in order to verify the correctness of SSA destruction algorithms. Boissinot et al. [41] review the history

of SSA destruction approaches, and highlight misunderstandings that led to incorrect destruction algorithms. Chapter 22 presents more details on alternative approaches to SSA destruction.

There are instructive dualisms between concepts in SSA form and functional programs, including construction, dominance and copy propagation. Chapter ?? explores these issues in more detail.

CHAPTER 4

Advanced Construction Algorithms for SSA *D. Das*
 U. Ramakrishna
 V. Sreedhar

Progress: 80%

Minor polishing in progress

4.1 Introduction

The placement of ϕ -functions is an important step in the construction of Static Single Assignment (SSA) form. In SSA form every variable is assigned only once. At control-flow merge points ϕ -functions are added to ensure that every use of a variable corresponds to exactly one definition. In the rest of this chapter we will present three different approaches for placing ϕ -functions at appropriate merge nodes. Recall that SSA construction falls into two phases of ϕ -placement and variable renaming. Here we present different approaches for the first phase for minimal SSA form. We first present Sreedhar and Gao's algorithm for placing ϕ -functions using the DJ-graph^{*} representation of a CFG. Then we will discuss the notions of merge set and merge relation, and present Das and Ramakrishna's algorithm for placing ϕ -functions based on merge relation and using DJ-graph. Finally we describe another approach for placing ϕ -functions based on the loop nesting forest^{*} proposed by Ramalingam.

dominance frontier

4.2 Basic Algorithm

The original algorithm for placing ϕ -functions is based on computing the dominance frontier (DF) set for the given control-flow graph. The dominance frontier^{*} $DF(x)$ of a node x is the set of all nodes z such that x dominates a predecessor of z , without strictly dominating z . Thus, from the point of view of x , the DF nodes are the earliest appearance of paths joining the flow without passing through x . For example, $DF(8) = \{6, 8\}$ in Figure 4.1. A straightforward algorithm for computing DF for each node takes $O(|V|^2)$, where $|V|$ is the number of nodes in the CFG, since the size of the full DF set in the worst case can be $O(|V|^2)$. Cytron et al.'s algorithm for the placement of ϕ -functions consists of computing the iterated dominance frontier (DF^+) for a set of all definition points (or nodes where variables are defined). Let $Defs(x)$ be the set of nodes where variable x are defined. Given that the dominance frontier for a set of nodes is just the union of the DF set of each node, we can compute $DF^+(Defs(x))$ as a limit of the following recurrence equation (where S is initially $Defs(x)$):

$$\begin{aligned} DF^+_1(S) &= DF(S) \\ DF^+_{i+1}(S) &= DF(S \cup DF^+_i(S)) \end{aligned}$$

A ϕ -function is then placed at each merge node in the $DF^+(Defs(x))$ set. Although Cytron et al.'s algorithm works very well in practice, in the worst case the time complexity of ϕ -placement algorithm for a single variable is quadratic in the number of nodes in the original control-flow graph.

4.3 Placing ϕ -functions using DJ-graphs

Sreedhar and Gao proposed the first linear time algorithm for computing the DF^+ set without the need for explicitly computing the full DF set. Sreedhar and Gao's original algorithm was implemented using the DJ-graph (See Section ?? and Figure 3.3b) representation of a CFG. The DJ-graph for the example CFG is also shown in Figure 4.1(b). Rather than explicitly computing the DF set, Sreedhar and Gao's algorithm uses a DJ-graph to compute the $DF^+(Defs(x))$ on the fly. Even though the time complexity of Sreedhar and Gao's algorithm is linear in the size of the DJ-graph, in practice it sometimes performs worse than the Cytron et al. algorithm. The main reason for this is that the size of the DF set is linear in the number of nodes in the CFG and in practice sometimes smaller than the size of the DJ-graph.

4.3.1 Key Observation

Now let us try to understand how to compute the DF set for a single node using the DJ-graph. Consider the DJ-graph shown in Figure 4.1(b). To compute $DF(8)$ we simply walk down the dominator (D) tree edges from node 8 and identify all join (J) edges $y \xrightarrow{J} z$ such that $z.level \leq 8.level$, where $level$ of a node is the depth of the node from the root of the dominator tree. The $level$ of each node is also shown in Figure 4.1(b). For our example the J edges that satisfy this condition are $9 \xrightarrow{J} 6$ and $10 \xrightarrow{J} 8$. Therefore $DF(8) = \{6, 8\}$. To generalize the example, we can compute the DF of a node x using:

$$DF(x) = \{z \mid \exists y \in \text{dominated}(x) \wedge y \xrightarrow{J} z \in \text{J-edges} \wedge z.level \leq x.level\}$$

where

$$\text{dominated}(x) = \{y \mid x \text{ dom } y\}$$

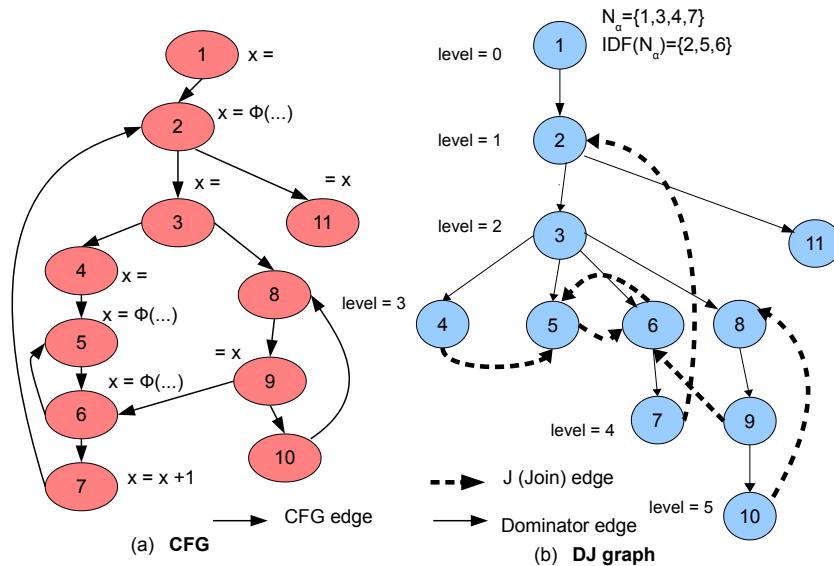


Fig. 4.1 A motivating example (a) CFG (b) DJ-graph.

Now we can extend the above idea to compute the DF^+ for a set of nodes, and hence the placement of ϕ -functions. This algorithm does not precompute DF; given a set of initial nodes $\text{Defs}(x)$ for which we want to compute the relevant set of ϕ -functions, a key observation can be made. Let w be an ancestor node of a node x on the dominator tree. If $DF(x)$ has already been computed before the

computation of $\text{DF}(w)$, the traversal of $\text{dominated}(x)$ can be avoided and $\text{DF}(x)$ directly used for the computation of $\text{DF}(w)$. This is because nodes reachable from $\text{dominated}(x)$ are already in $\text{DF}(x)$. However, the converse may not be true, and therefore the order of the computation of DF is crucial.

To illustrate the key observation consider the example DJ-graph in Figure 4.1(b), and let us compute $\text{DF}^+(\{3, 8\})$. It is clear from the recursive definition of DF^+ that we have to compute $\text{DF}(3)$ and $\text{DF}(8)$ as a first step. Now supposing we start with node 3 and compute $\text{DF}(3)$. The resulting DF set is $\text{DF}(3) = \{2\}$. Now supposing we next compute the DF set for node 8, and the resulting set is $\text{DF}(8) = \{6, 8\}$. Notice here that we have already visited node 8 and its subtree when visiting node 3. We can avoid such duplicate visits by ordering the computation of DF set so that we first compute $\text{DF}(8)$ and then during the computation of $\text{DF}(3)$ we avoid visiting the sub-tree of node 8, and use the result $\text{DF}(8)$ that was previously computed.

Thus, to compute $\text{DF}(x)$, where x is an ancestor of y in the DJ-graph, we do not need to compute it from scratch as we can re-use the information computed as part of $\text{DF}(y)$ as shown. For this, we need to compute the DF of deeper (based on level) nodes (here, y), before computing the DF of a shallower node (here, x).

$$\begin{aligned}\text{DF}(x) = & \{w \mid w \in \text{DF}(y) \wedge w.\text{level} \leq x.\text{level}\} \cup \\ & \{w \mid t \in \text{Subtree}(x) \setminus \text{Subtree}(y) \wedge t \xrightarrow{L} w \wedge w.\text{level} \leq x.\text{level}\}\end{aligned}$$

4.3.2 Main Algorithm

In this section we present Sreedhar and Gao's algorithm for computing DF^+ . Let $x.\text{level}$ be the depth of the node from the root node, with $\text{root.level} = 0$. To ensure that the nodes are processed according to the key observation we use a simple array of sets OrderedBucket , and two functions defined over this array of sets: (1) $\text{InsertNode}(n)$ that inserts the node n in the set $\text{OrderedBucket}[n.\text{level}]$, and (2) $\text{GetDeepestNode}()$ that returns a node from the OrderedBucket with the deepest level number.

In Algorithm 4.1, at first we insert all nodes belonging to $\text{Defs}(x)$ in the OrderedBucket . Then the nodes are processed in a bottom-up fashion over the DJ-graph from deepest node level to least node level by calling $\text{Visit}(x)$. The procedure $\text{Visit}(x)$ essentially walks top-down in the DJ-graph avoiding already visited nodes. During this traversal it also peeks at destination nodes of J edges. Whenever it notices that the level number of the destination node of a J edge is less than or equal to the level number of the currentRoot , the destination node is added to the DF^+ set (Line 5) if it is not present in DF^+ already. Notice that at Line 6 the destination node is also inserted in the OrderedBucket if it was never inserted before. Finally at Line 13 we continue to process the nodes in the sub-

Algorithm 4.1: Sreedhar and Gao's algorithm for computing DF^+ set.

Input: DJ-graph representation of a program
Input: $Defs(x)$: nodes that define x
Output: $DF^+(Defs(x))$

```

1  $DF^+ \leftarrow \emptyset;$ 
2 foreach node  $x \in Defs(x)$  do
3   | InsertNode( $x$ );
4   |  $x.visited \leftarrow \text{False};$ 
5 end
6 while  $z \leftarrow GetDeepestNode()$  do           /* Get node from the deepest level */
7   |  $currentRoot \leftarrow z;$ 
8   |  $z.visited \leftarrow \text{True};$ 
9   | Visit( $z$ );                                /* Find  $DF(z)$  */
10 end

```

Procedure Visit(x)

```

1 foreach node  $y \in Succ(x)$  do
2   | if  $x \rightarrow y$  is a J edge then
3     |   | if  $y.level \leq currentRoot.level$  then
4       |     | if  $y \notin DF^+$  then
5         |       |  $DF^+ \leftarrow DF^+ \cup y;$ 
6         |       | if  $y \notin Defs(x)$  then InsertNode( $y$ );
7       |     | end
8     |   | end
9   | end
10  | else                                /* visit D edges */
11    |   | if  $y.visited = \text{False}$  then
12      |     |  $y.visited \leftarrow \text{True};$ 
13      |     | /*  $\text{if}(y.boundary = \text{False})$  // See the section on Further
14      |     |   Reading for details
15      |     | Visit( $y$ );
16    |   | end
17  | end
18 end

```

tree by visiting over the D edges. When the algorithm terminates the set DF^+ contains the iterated dominance frontier for the initial $Defs(x)$.

In Figure 4.2, some of the phases of the algorithm are depicted for clarity. The *OrderedBucket* is populated with the nodes 1, 3, 4 and 7 corresponding to $Defs(x) = \{1, 3, 4, 7\}$. The nodes are placed in the buckets corresponding to the levels at which they appear. Hence, node 1 which appears at level 0 is in the zero-th bucket, node 3 is in the second bucket and so on. Since the nodes are processed bottom-up, the first node that is visited is node 7. The successor of node 7 is node 2 and since there exists a J edge $7 \xrightarrow{J} 2$, and the DF^+ set is empty, the DF^+ set is updated to hold node 2 according to Line 5 of the **Visit** procedure. In addition, **InsertNode(2)** is invoked and node 2 is placed in the second bucket.

merge graph

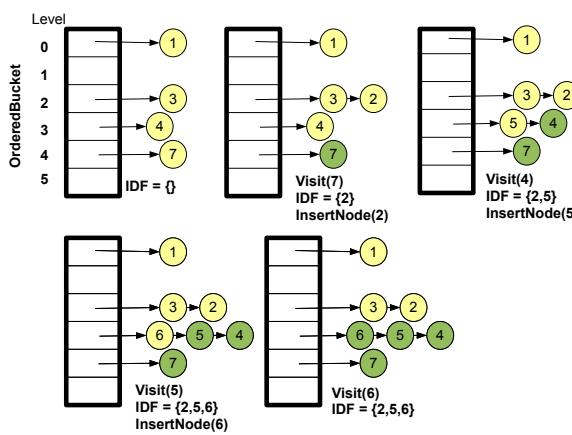


Fig. 4.2 Phases of Sreedhar and Gao's algorithm for $\text{Defs}(x) = \{1, 3, 4, 7\}$.

The next node visited is node 4. The successor of node 4 which is node 5 has an incoming J edge $4 \xrightarrow{J} 5$ which results in the new $\text{DF}^+ = \{2, 5\}$. The final DF^+ set converges to $\{2, 5, 6\}$ when node 5 is visited. Subsequent visits of other nodes do not add anything to the DF^+ set. An interesting case arises when node 3 is visited. Node 3 finally causes nodes 9 and 10 also to be visited. However, when node 10 is visited, its successor node which is node 8 and which also corresponds to the J edge $10 \xrightarrow{J} 8$, does not result in an update of the DF^+ set as the level of node 8 is deeper than that of node 3.

4.4 Placing ϕ -functions using Merge Sets

In the previous section we described ϕ -function placement algorithm in terms of iterated dominance frontiers. There is another way to approach the problem of ϕ -function placement using the concepts of merge relation and merge set. In this section we will first introduce the notion of merge set and merge relation and then show how the merge relation is related to the DF relation. We will then show how to compute the M (merge) graph^{*} using the DJ-graph and then use the M-graph for placing ϕ -functions.

4.4.1 Merge Set and Merge Relation

First, let us define the notion of a *join* set $J(S)$ for a given set of nodes S in a control-flow graph.¹ Consider two nodes u and v and distinct non-empty (denoted by \dashv) paths from $u \dashv w$ and $v \dashv w$, where w is some node in the CFG. If the two paths meet only at w then w is in the join set of the nodes $\{u, v\}$. For instance, consider nodes 1 and 9 in Figure 4.1(a). The paths $1 \rightarrow 2 \rightarrow 3 \rightarrow 8$ and $9 \rightarrow 10 \rightarrow 8$ meet at 8 for the first time and so $\{8\} \in J(\{1, 9\})$.

Now let us define the *merge* relation as a relation $v = M(u)$ that holds between two nodes u and v whenever $v \in J(\{\text{root}, u\})$. We insert a ϕ -function at v for a variable that is assigned only at *root* and u . One can show that $J(S) = \bigcup_{u \in S} M(u)$ where $S \subseteq V$. Also, for any node $u \in V$, $v \in M(u)$ if and only if there is a path $u \dashv v$ that does not contain $\text{idom}(v)$. This relationship between dominance and merge can be conveniently encoded using a DJ-graph and can be used for placing ϕ -functions. First, we will construct the DF (Dominance Frontier) graph using the DJ-graph. For each J edge $u \rightarrow v$ in the DJ-graph insert a new $w \dashv v$ where $w = \text{idom}(u)$ and w does not strictly dominate v . Another way to look at this is to insert a new J edge $w \rightarrow v$ if $w = \text{idom}(u)$ and $w.\text{level} \geq v.\text{level}$.² We repeatedly insert new J edges in a bottom-up fashion over the DJ-graph until no more J edge can be inserted. The resulting DJ-graph is a “fully cached” DJ-graph. We will call the fully cached DJ-graph the DF-graph.

Next, we can compute the M relation from the DF-graph. Using the DF-graph we compute the M-graph as a transitive closure using only the DF edges of the DJ-graph. In Figure 4.3, $M(4) = \{2, 5, 6\}$ as the nodes 2, 5 and 6 are the only nodes reachable from node 4 following the DF edges. Similarly, for node 8, $M(8) = \{2, 5, 6, 8\}$ due to these nodes being reachable from node 8 via the DF edges. Given the M-graph of a CFG we can place ϕ -functions for a set $S \subseteq V$ at $M(S)$.

The algorithms for placing ϕ -functions based on the construction of the DF-graph or M relation can be quadratic in the worst case. The merge relation is due to Pingali and Bilardi. The DF-graph and the M sets for the running example are given in Figure 4.3.

4.4.2 Iterative Merge Set Computation

In this section we describe a method to iteratively compute the *merge* relation using a data-flow formulation. In the previous section we saw that the *merge* relation can be computed using a transitive closure of the DF-graph, which in turn can be computed from the DJ-graph. In the algorithm proposed by Das and Ra-

¹ In English, “join” and “merge” are synonyms, but unfortunately in the literature, due to lack of better terms, these two synonyms are used to mean distinct but related concepts.

² Recall that $w.\text{level}$ is the depth of w from the *root* of the dominator tree.

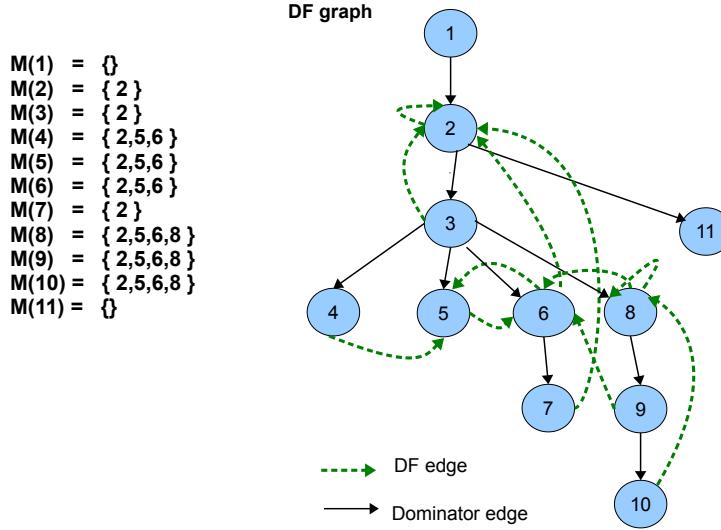


Fig. 4.3 DF-graph and M sets.

makrishna, explicit DF-graph construction or the transitive closure formation are not necessary. Instead, the same result can be achieved by formulating the *merge* relation as a data-flow problem and solving it iteratively. For several applications, this approach has been found to be a fast and effective method to construct $M(x)$ for each node x and the corresponding ϕ -function placement using the $M(x)$ sets.

Consider a J edge $u \xrightarrow{J} v$. Also consider all nodes w such that w dominates u and $w.level \geq v.level$. For every w we can set up the data-flow equation as $M(w) = M(w) \cup M(v) \cup \{v\}$. The set of data-flow equations for each node n in the DJ-graph can be solved iteratively using a top-down pass over the DJ-graph. To check whether multiple-passes are required over the DJ-graph before a fixed-point is reached for the data-flow equations, we devise an “inconsistency condition” stated as follows:

Inconsistency Condition:

*For a J edge, $u \xrightarrow{J} v$, if u does not satisfy $M(u) \supseteq M(v)$,
then the node u is said to be inconsistent.*

The algorithm described in the next section is directly based on the method of building up the $M(x)$ sets of the nodes as each J edge is encountered in an iterative fashion by traversing the DJ-graph top-down. If no node is found to be *inconsistent* after a single top-down pass, all the nodes are supposed to have

reached fixed-point solutions. If some node is found to be inconsistent, multiple passes are required until a fixed-point solution is reached.

4.4.2.1 Top Down Merge Set Computation

Input: A DJ-graph representation of a program. **Output:** The merge sets for the nodes.

```

Procedure TDMSCMain {
1:    $\forall$  node  $x \in$  DJ-graph set  $M(x) = \{\}$ 
2:    $done = false;$ 
3:   while ( $! done$ ) do
4:      $done = TDMSC-I(DJ\text{-graph})$ ;
5:   end while
}
Procedure TDMSC-I(DJ-graph) {
6:   RequireAnotherPass = False;
7:   while ( $n =$  next node in B(readth) F(irst) S(earch) order of DJ-graph) do
8:     for ( $\forall$  incoming edge to  $n$ ) do
9:       Let  $e = s \xrightarrow{L} n$ , be an incoming J edge
10:      if ( $e$  not visited) then
11:        Mark  $e$  as visited
12:         $tmp = s$ ;
13:         $lnode = Null$ ;
14:        while ( $level(tmp) \geq level(n)$ ) do
15:           $M(tmp) = M(tmp) \cup M(n) \cup \{n\}$ ;
16:           $lnode = tmp$ ;
17:           $tmp = parent(tmp)$ ; //dominator tree parent
18:        end while
19:        for ( $\forall$  incoming edge to  $lnode$ ) do
20:          Let  $e' = s' \xrightarrow{L} lnode$ , be an incoming J edge
21:          if ( $e'$  visited) then
22:            if ( $M(s') \not\supseteq M(lnode)$ ) then //Check inconsistency
23:              RequireAnotherPass = true;
24:            end if
25:          end if
26:        end for
27:      end if
28:    end for
29:  end while
30:  return RequireAnotherPass;
}

```

Fig. 4.4 Top Down Merge Set Computation algorithm for computing Merge sets.

The first and direct variant of the approach laid out above is poetically termed TDMSC-I. This variant works by scanning the DJ-graph in a top-down fashion as shown in Line 7 of Figure 4.4. All $M(x)$ sets are set to the empty set before the initial pass of TDMSC-I. The $M(x)$ sets computed in a previous pass are carried over if a subsequent pass is required.

The DJ-graph is visited level by level. During this process, for each node n encountered, if there is an *incoming* J edge $s \xrightarrow{J} n$ as in Line 8, then a separate bottom-up pass starts at Line 14. This bottom-up pass traverses all nodes w such that w dominates s and $w.level \geq s.level$, updating the $M(w)$ values using the aforementioned data-flow equation. Line 19 is used for the inconsistency check. *RequireAnotherPass* is set to true only if a fixed point is not reached and the inconsistency check succeeds for some node.

There are some subtleties in the algorithm that should be noted. Line 19 of the algorithm visits incoming edges to $lnode$ only when $lnode$ is at the same level as n , which is the current level of inspection and the incoming edges to $lnode$'s posterity are at a level greater than that of node n and are unvisited yet.

Here, we will briefly walk through TDMSC-I using the DJ-graph of Figure 4.1(b). Moving top-down over the graph, the first J edge encountered is $7 \xrightarrow{J} 2$. As a result, a bottom-up climbing of the nodes happens, starting at node 7 and ending at node 2 and the merge sets of these nodes are updated so that $M(7) = M(6) = M(3) = M(2) = \{2\}$. The next J edge to be visited can be any of $4 \xrightarrow{J} 5$, $5 \xrightarrow{J} 6$ or $6 \xrightarrow{J} 5$ at $level = 3$. Assume it is $5 \xrightarrow{J} 6$. This results in $M(5) = M(5) \cup M(6) \cup \{6\} = \{2, 6\}$. Now, let $6 \xrightarrow{J} 5$ be visited. Hence, $M(6) = M(6) \cup M(5) \cup \{5\} = \{2, 5, 6\}$. At this point, the *inconsistency check* comes into picture for the edge $6 \xrightarrow{J} 5$, as $5 \xrightarrow{J} 6$ is another J edge that is already visited and is an incoming edge of node 6. Checking for $M(5) \supseteq M(6)$ fails, implying that the $M(5)$ needs to be computed again. This will be done in a succeeding pass as suggested by the *RequireAnotherPass* value of true. In a second iterative pass, the J edges are visited in the same order. Now, when $5 \xrightarrow{J} 6$ is visited, $M(5) = M(5) \cup M(6) \cup \{6\} = \{2, 5, 6\}$ as this time $M(5) = \{2, 6\}$ and $M(6) = \{2, 5, 6\}$. On a subsequent visit of $6 \xrightarrow{J} 5$, $M(6)$ is also set to $\{2, 5, 6\}$. The inconsistency does not appear any more and the algorithm proceeds to handle the edges $4 \xrightarrow{J} 5$, $9 \xrightarrow{J} 6$ and $10 \xrightarrow{J} 8$ which have also been visited in the earlier pass. TDMSC-I is invoked repeatedly by a different function which calls it in a loop till *RequireAnotherPass* is returned as False as shown in the procedure TDMSC-Main.

Once the *Merge* relation is computed for the entire CFG, placing ϕ is a straightforward application of the $M(x)$ values for a given $\text{Defs}(x)$, as shown in Figure 4.5.

Input: A CFG with Merge sets computed and $\text{Defs}(x)$. **Output:** Blocks augmented by ϕ -functions for the given $\text{Defs}(x)$.

```

Procedure  $\phi$ -placement for  $\text{Defs}(x)$  {
1:   for ( $\forall n \in \text{Defs}(x)$ ) do
2:     for ( $\forall n' \in M(n)$ ) do
3:       Add a  $\phi$  for  $n$  in  $n'$  if not placed already
4:     end for
5:   end for
}

```

Fig. 4.5 ϕ -placement using *Merge* sets.

4.5 Computing Iterated Dominance Frontier Using Loop Nesting Forests

This section illustrates the use of *loop nesting forests* to construct the iterated dominance frontier (DF^+) of a set of vertices in a CFG. This method works with reducible as well as irreducible loops and is based on Ramalingam's work.

4.5.1 Loop nesting forest

A loop nesting forest is a data structure that represents the loops in a CFG and the containment relation between them. In the example shown in Figure 4.6(a) the loops with back edges $11 \rightarrow 9$ and $12 \rightarrow 2$ are both reducible loops. The corresponding loop nesting forest is shown in Figure 4.6(b) and consists of two loops whose header nodes are 2 and 9. The loop with header node 2 contains the loop with header node 9.

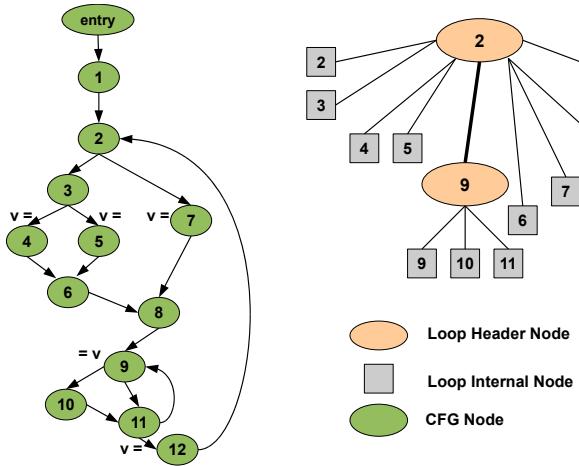


Fig. 4.6 An example (a) CFG and (b) Loop Nesting Forest.

4.5.2 Main Algorithm

The idea is to use an acyclic version of the control-flow graph (i.e., without back edge) and construct the DF^+ for a variable in this context: whenever two def-

initiations reach a join point, it belongs to the DF^+ . Then, we take into account the back edges using the loop nesting forest: if a loop contains a definition, its header also belongs to the DF^+ .

A definition node d “reaches” another node u if there is a path in the graph from d to u which does not contain any re-definition. If at least two definitions reach a node u , then u belongs to $\text{DF}^+(X)$ where X consists of these definition nodes. This suggests the Algorithm 4.2 which works for *acyclic* graphs. For a given $\text{Defs}(x)$, we can compute $\text{DF}^+(\text{Defs}(x))$ as follows:

- Initialize DF^+ to the empty set ;
- Using a topological order, compute the subset of $\text{DF}^+(\text{Defs}(x))$ that can reach a node using forward data-flow analysis ;
- Add a node to DF^+ if it is reachable from multiple nodes;

For Figure 4.6, the acyclic version of the graph G , termed G_{ac} , is formed by dropping the back edges $11 \rightarrow 9$ and $12 \rightarrow 2$. Also, $\text{entry}(G) = \{\text{entry}\}$ where entry is a specially designated node that is the root of the CFG. For the definitions of v in nodes 4, 5, 7 and 12 in Figure 4.6, the subsequent nodes reached by multiple definitions are 6 and 8: node 6 can be reached by any one of two definitions in nodes 4 or 5, and node 8 by either the definition from node 7 or one of 4 or 5. Note that the back edges do not exist in the acyclic graph and hence node 2 is not part of the DF^+ set. We will see later how the DF^+ set for the entire graph is computed by considering the contribution of the back edges.

Let us walk through this algorithm computing DF^+ for variable v , i.e., $X = \{4, 5, 7, 12\}$. The nodes in Figure 4.6 are already numbered in topological order. Nodes 1 to 5 have only one predecessor, none of them being in X , so their *UniqueReachingDef* stays entry , and DF^+ is still empty. For node 6, its two predecessor belong to X , hence *ReachingDefs* = $\{4, 5\}$, and 6 is added to DF^+ . Nothing changes for 7, then for 8 its predecessors 6 and 7 are respectively in DF^+ and X : they are added to *ReachingDefs*, and 8 is then added to DF^+ . Finally, for nodes 8 to 12, their *UniqueReachingDef* will be updated to node 8, but this will not change DF^+ anymore which will end up being $\{6, 8\}$.

DF⁺ on reducible graphs

A reducible graph can be decomposed into an acyclic graph and a set of back edges. The contribution of back edges to the iterated dominance frontier can be identified by using the loop nesting forest. If a vertex u is contained in a loop, then $\text{DF}^+(u)$ will contain the loop header. For any vertex u , let $\text{HLC}(u)$ denote the set of loop headers of the loops containing u . Given a set of vertices X , it turns out that $\text{DF}^+(X) = \text{HLC}(X) \cup \text{DF}_{ac}^+(X \cup \text{HLC}(X))$ where DF_{ac}^+ denote the DF^+ restricted to the acyclic graph G_{ac} .

Reverting back to Figure 4.6, we see that in order to find the DF^+ for the nodes defining v , we need to evaluate $\text{DF}_{ac}^+(\{4, 5, 7, 12\} \cup \text{HLC}(\{4, 5, 7, 12\}))$. As all

Algorithm 4.2: Ramalingam's algorithm for DF^+ in an acyclic graph.

Input: G_{ac} : acyclic CFG
Input: X : subset of nodes
Output: $\text{DF}^+(X)$

```

1   $\text{DF}^+ \leftarrow \emptyset;$ 
2  foreach node  $n$  do  $\text{UniqueReachingDef}(n) \leftarrow \{\text{entry}\};$ 
3  foreach node  $u \neq \text{entry}$  in topological order do
4     $\text{ReachingDefs} \leftarrow \emptyset;$ 
5    foreach  $v$  predecessor of  $u$  do
6      if  $v \in \text{DF}^+ \cup X \cup \{\text{entry}\}$  then
7         $\text{ReachingDefs} \leftarrow \text{ReachingDefs} \cup \{v\};$ 
8      end
9      else
10      $\text{ReachingDefs} \leftarrow \text{ReachingDefs} \cup \text{UniqueReachingDef}(v);$ 
11   end
12 end
13 if  $|\text{ReachingDefs}| = 1$  then
14    $\text{UniqueReachingDef}(u) \leftarrow \text{ReachingDefs};$ 
15 end
16 else
17    $\text{DF}^+ \leftarrow \text{DF}^+ \cup \{u\};$ 
18 end
19 end
20 return  $\text{DF}^+$ 

```

these nodes are contained in a single loop with header 2 $\text{HLC}(\{4, 5, 7, 12\}) = \{2\}$. Computing $\text{DF}^+_{ac}(\{2, 4, 5, 7, 12\})$ gives us $\{6, 8\}$, and finally, $\text{DF}^+(\{4, 5, 7, 12\}) = \text{HLC}(\{4, 5, 7, 12\}) \cup \{6, 8\} = \{2, 6, 8\}$.

DF⁺ on irreducible graphs

We will now briefly explain how graphs containing irreducible loops can be handled. The insight behind the implementation is to transform the irreducible loop in such a way that an acyclic graph is created from the loop without changing the dominance properties of the nodes.

The loop in the graph of Figure 4.7(a) is irreducible as it has two headers, v and w . It can be transformed to the acyclic graph (c) by removing the edges between nodes v and w that create the irreducible loop. We create a dummy node, θ , to which we connect all predecessors of the headers (u and s), and that we connect to all the headers of the loop (v and w), creating graph (d).

TODO: subfigure

Following this transformation, the graph is now acyclic, and computing the DF^+ for the nodes in the original irreducible graph translates to computing DF^+ using the transformed graph to get DF^+_{ac} and using the loop forest of the original graph (b).

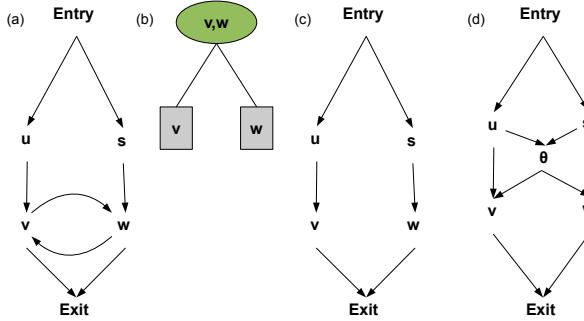


Fig. 4.7 (a) An irreducible graph (b) The Loop Nesting Forest (c) The acyclic subgraph (d) Transformed graph.

The crucial observation that allows this transformation to create an equivalent acyclic graph is the fact that the dominator tree of the transformed graph remains identical to the original graph containing an irreducible cycle. One of the drawbacks of the transformation is the possible explosion in the number of dummy nodes and edges for graphs with many irreducible cycles as a unique dummy node needs to be created for each irreducible loop.

4.6 Concluding Remarks

Although all these algorithms claim to be better than the original algorithm by Cytron et al., they are difficult to compare due to the unavailability of these algorithms in a common compiler framework.

In particular, while constructing the whole DF⁺ set seems very costly in the classical construction algorithm, its cost is actually amortized as it will serve to place ϕ -functions for many variables. It is however interesting not to pay this cost whenever we only have a few variables to consider, for instance when repairing SSA as in the next chapter.

Note also that people have observed in production compilers that, during SSA construction, what seems to be the most expensive part is the renaming of the variables and not the placement of ϕ -functions.

4.7 Further Reading

Cytron's approach for ϕ -placement involves a fully eager approach of constructing the entire DF-graph. On the other hand Sreedhar and Gao's approach is a fully lazy approach as it constructs DF on-the-fly only when a query is encountered. Pingali and Bilardi [?] suggested a middle-ground by combining both approaches. They proposed a new representation called ADT (Augmented Dominator Tree). The ADT representation can be thought as a DJ-graph, where the DF sets are pre-computed for certain nodes called “boundary nodes” using an eager approach. For the rest of the nodes, termed “interior nodes,” the DF needs to be computed on-the-fly as in the Sreedhar-Gao algorithm. The nodes which act as “boundary nodes” are detected in a separate pass. A factor β is used to determine the partitioning of the nodes of a CFG into boundary or interior nodes by dividing the CFG into zones. β is a number that represents space/query-time tradeoff. $\beta \ll 1$ denotes a fully eager approach where storage requirement for DF is maximum but query-time is faster while $\beta \gg 1$ denotes a fully lazy approach where storage requirement is zero but query is slower.

Given the ADT of a control-flow graph, it is straight forward to modify Sreedhar and Gao's algorithm for computing ϕ -functions in linear time. The only modification that is needed is to ensure that we need not visit all the nodes of a sub-tree rooted at a node y when y is a boundary node whose DF set is already known. This change is reflected in Line 12 of Figure ??, where a subtree rooted at y is visited or not visited based on whether it is a boundary node or not.

For iterative merge set computation, TDMSC-II is an improvement to algorithm TDMSC-I. This improvement is fueled by the observation that for an inconsistent node u , the merge sets of all nodes w such that w dominates u and $w.level \geq u.level$, can be locally corrected for some special cases. This heuristic works very well for certain class of problems—especially for CFGs with DF-graphs having cycles consisting of a few edges. This eliminates extra passes as an inconsistent node is made consistent immediately on being detected. For details of this algorithm refer to [?].

TODO: citations

- SSA form [?]
- Sreedhar & gao [?]
- Das & Ramakrishna no explicit DF-graph construction [?]
- Ramalingam loop nesting forest [?]
- computing DF^+ set without full DF set [?, ?]
- [?]

CHAPTER 5

SSA Reconstruction

S. Hack

Progress: 90%

Text and figures formating in progress



5.1 Introduction

Some optimizations break the single-assignment property of the SSA form by inserting additional definitions for a single SSA value. A common example is live-range splitting by inserting copy operations or inserting spill and reload code during register allocation. Other optimizations, such as loop unrolling or jump threading, might duplicate code, thus adding additional variable definitions, *and* modify the control flow of the program. We will first go through two examples before we present algorithms to properly repair SSA.

5.1.1 Live-Range Splitting

The first example is depicted in Figure 5.1. Our spilling pass decided to spill a part of the live range of variable x_0 in the right block in (a), resulting in the code shown in (b), where it inserted a store and a load instruction. This is indicated by assigning to the memory location X . The load is now a second definition of x_0 , violating the SSA form that has to be reconstructed as shown in (c). This example shows that maintaining SSA also involves placing new ϕ -functions.

Many optimizations perform such program modifications, and maintaining SSA is often one of the more complicated and error-prone parts in such opti-

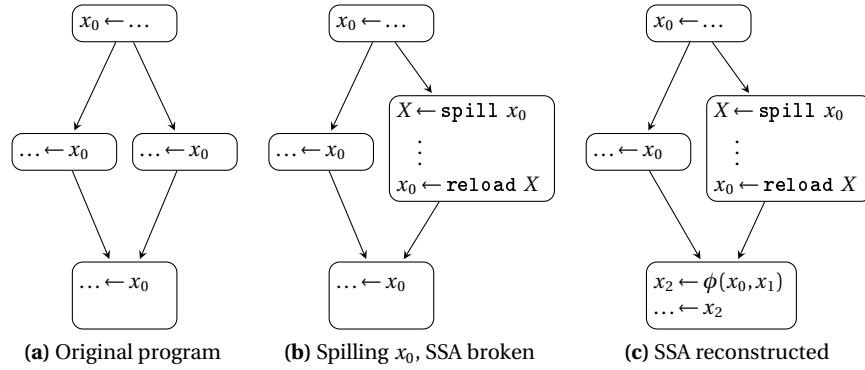


Fig. 5.1 Adding a second definition as a side-effect of spilling.

mizations, owing to the insertion of additional ϕ -functions and the correct redirection of the uses of the variable.

5.1.2 Jump Threading

Jump threading is a transformation performed by many popular compilers such as GCC and LLVM. It partially evaluates branch targets to save branches and provide more context to other optimizations. Consider the situation where a block contains a conditional branch that depends on some variable x . In the example shown in Figure 5.2, the conditional branch tests if $x > 0$.

Assume that the block containing that conditional branch has multiple predecessors and x can be proven constant for one of the predecessors. In our example, this is shown by the assignment $x_1 \leftarrow 0$. Jump threading now partially evaluates the conditional branch by directly making the corresponding successor of the branch block a successor of the *predecessor* of the branch. To this end, the code of the branch block is duplicated and attached to the predecessor. This also duplicates potential definitions of variables in the branch block. Hence, SSA is destroyed for those variables and has to be reconstructed. However, in contrast to the previous spill example, the control flow has changed which poses an additional challenge for an efficient SSA reconstruction algorithm.

5.2 General Considerations

In this chapter, we will discuss two algorithms. The first is an adoption of the classical dominance-frontier based algorithm. The second performs a search

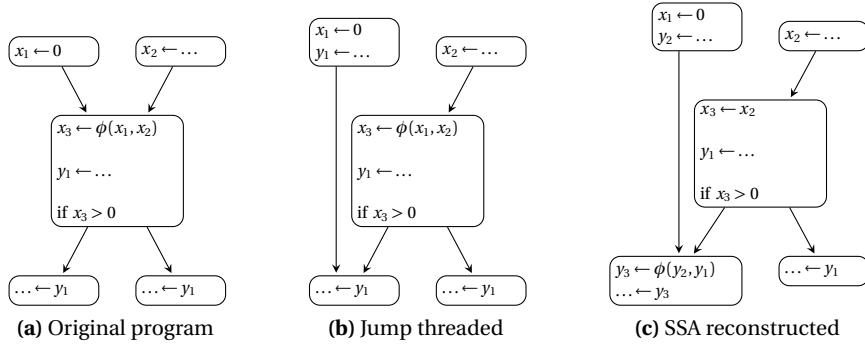


Fig. 5.2 Jump Threading can break SSA, which has to be reconstructed.

from the uses of the variables to the definition and places ϕ -functions on demand at appropriate places. In contrast to the first, the second algorithm might not construct minimal SSA form in general; However, it does not need to update its internal data structures when the CFG is modified.

We consider the following scenario: The program is represented as a control-flow graph (CFG) and is in SSA form with dominance property. For the sake of simplicity, we assume that each instruction in the program only writes to a single variable. An optimization or transformation violates SSA by inserting additional definitions for an existing SSA variable, like in the examples above. The original variable and the additional definitions can be seen as a single non-SSA variable that has multiple definitions and uses.

In the following, v will be such a non-SSA variable. When reconstructing SSA for v , we will first create fresh variables for every definition of v to establish the single-assignment property. What remains is associating every use of v with a suitable definition. In the algorithms, $v.\text{defs}$ denotes the set of all instructions that define v . A use of a variable is a pair consisting of a program point (an instruction) and an integer denoting the index of the operand at this instruction.

Both algorithms presented in this chapter share the same driver routine described in Algorithm 5.1. First, we scan all definitions of v so that for every basic block b we have the list $b.\text{defs}$ that contains all instructions in the block which define one of the variables in $v.\text{defs}$. It is best to sort this according to the schedule of the instructions in the block from back to front, making, the latest definition the first in the list.

Then, all uses of the variable v are scanned. Here we have to differentiate whether the use is in a ϕ -function or not. If so, the use occurs at the end of the predecessor block that corresponds to the position of the variable in the ϕ 's argument list. In that case, we start looking for the reaching definition from the end of that block. Otherwise, we scan the instructions of the block backwards until we reach the first definition that is before the use (Line 16). If there is no such definition, we have to find one that reaches this block from *outside*.

TODO: reprendre paragraphe

We use two functions, `FindDefFromTop` and `FindDefFromBottom` that search the reaching definition respectively from the beginning or the end of a block. `FindDefFromBottom` actually just returns the last definition in the block, or call `FindDefFromTop` if there is none.

The two presented approaches to SSA repairing differ in the implementation of the function `FindDefFromTop`. The differences are described in the next two sections.

Algorithm 5.1: SSA Reconstruction Driver

```

Input:  $v$ , a variable that breaks SSA property
1 foreach  $d \in v.\text{defs}$  do
2   | create fresh variable  $v'$ ;
3   | rewrite def of  $v$  by  $v'$  in  $d$ ;
4   |  $b \leftarrow d.\text{block}$ ;
5   | insert  $d$  in  $b.\text{defs}$ ;
6 end
7 foreach use (inst, index) of  $v$  do
8   | if inst is a  $\phi$ -function then
9   |   |  $b \leftarrow \text{inst.block.pred}(index)$ ;
10  |   |  $d \leftarrow \text{FindDefFromBottom}(v, b)$ ;
11  | end
12  | else
13  |   |  $d \leftarrow \perp$ ;
14  |   |  $b \leftarrow \text{inst.block}$ ;
15  |   | foreach  $l \in b.\text{defs}$  from bottom to top do
16  |   |   | if  $l$  is before inst in  $b$  then
17  |   |   |   |  $d \leftarrow l$ ;
18  |   |   |   | break;
19  |   | end
20  |   | end
21  |   | if  $d = \perp$  then /* no local def. found, searching in the preds */
22  |   |   |  $d \leftarrow \text{FindDefFromTop}(v, b)$ ;
23  |   | end
24  | end
25  |  $v' \leftarrow$  version of  $v$  defined by  $d$ ;
26  | rewrite use of  $v$  by  $v'$  in inst;
27 end

```

Procedure `FindDefFromBottom(v, b)`

```

1 if  $b.\text{defs} \neq \emptyset$  then
2   | return latest instruction in  $b.\text{defs}$ ;
3 end
4 else
5   | return FindDefFromTop( $v, b$ );
6 end

```

5.3 Reconstruction based on the Dominance Frontier

This algorithm follows the same principles as the classical SSA construction algorithm by Cytron et al. as described in Chapter 3. We first compute the iterated dominance frontier (DF^+) of v . This set is a sound approximation of the set where ϕ -functions must be placed—it might contain blocks where a ϕ -function would be dead. Then, we search for each use u the corresponding reaching definition. This search starts at the block of u . If that block b is in the DF^+ of v , a ϕ -function needs to be placed at its entrance. This ϕ -function becomes a new definition of v and has to be inserted in $v.defs$ and in $b.defs$. The operands of the newly created ϕ -function will query their reaching definitions by recursive calls to `FindDefFromBottom` on predecessors of b . Because we inserted the ϕ -function into $b.defs$ *before* searching for the arguments, no infinite recursion can occur (otherwise, it could happen for instance with a loop back edge).

If the block is not in the DF^+ , the search continues in the immediate dominator of the block. This is because in SSA, every use of a variable must be dominated by its definition.¹ Therefore, the reaching definition is the same for all predecessors of the block, and hence for the immediate dominator of this block.

Procedure `FindDefFromTop(v, b)`

```

/* SSA Reconstruction based on Dominance Frontiers */  

1 if  $b \in DF^+(v.defs)$  then  

2    $v' \leftarrow$  fresh variable;  

3    $d \leftarrow$  new  $\phi$ -function in  $b$ :  $v' \leftarrow \phi(\dots)$ ;  

4   append  $d$  to  $b.defs$ ;  

5   foreach  $p \in b.preds$  do  

6      $o \leftarrow$  FindDefFromBottom( $v, p$ );  

7      $v' \leftarrow$  version of  $v$  defined by  $o$ ;  

8     set corresponding operand of  $d$  to  $v'$ ;  

9   end  

10 end  

11 else  

12    $| d \leftarrow$  FindDefFromBottom( $v, b.idom$ ); /* search in immediate dominator */  

13 end  

14 return  $d$ 

```

¹ The definition of an operand of a ϕ -function has to dominate the corresponding predecessor block.

5.4 Search-based Reconstruction

The second algorithm presented here is adapted from an algorithm designed to construct SSA from the abstract syntax tree, but it also works well on control-flow graphs. Its major advantage over the algorithm presented in the last section is that it does neither require dominance information nor dominance frontiers. Thus it is well suited to be used in transformations that change the control-flow graph. Its disadvantage is that potentially more blocks have to be visited during the reconstruction. The principle idea is to start a search from every use to find the corresponding definition, inserting ϕ -functions on the fly while caching the SSA variable alive at the beginning of basic blocks. As in the last section, we only consider the reconstruction for a single variable called v in the following. If multiple variables have to be reconstructed, the algorithm can be applied to each variable separately.

Search on an acyclic CFG.

The algorithm performs a backward depth-first search in the CFG to collect the reaching definitions of v in question at each block, recording the SSA variable that is alive at the beginning of a block in the “beg” field of this block. If the CFG is an acyclic graph (DAG), all predecessors of a block can be visited before the block itself is processed, as we are using a post-order traversal following edges backward. Hence, we know all the definitions that reach a block b : if there is more than one definition, we need to place a ϕ -function in b , otherwise it is not necessary.

Search on a general CFG.

If the CFG has loops, there are blocks for which not all reaching definitions can be computed before we can decide whether a ϕ -function has to be placed or not. In a loop, recursively computing the reaching definitions for a block b will end up at b itself. To avoid infinite recursion, when we enter a block during the traversal, we first create a ϕ -function without arguments, “`pending_phi`.” This creates a new definition for v , v_ϕ , which is the variable alive at the beginning of this block.

When we return to b after traversing the rest of the CFG, we decide whether a ϕ -function has to be placed in b by looking at the reaching definition for every predecessor. These reaching definitions can be either v_ϕ itself (loop in the CFG without a definition of v), or some other definitions of v . If there is only one such other definition, say w , then this definition cannot be in a loop around b or the program would not be strict. In that case, it means `pending_phi` is not necessary

TODO: strictness not necessary

and we can remove it, propagating w downward instead of v_ϕ . Otherwise, we keep pending $_\phi$ and fill its missing operands with the reaching definitions.

In this version of function **FindDefFromTop**, this check is done by the function **Phi-Necessary**.

Procedure FindDefFromTop(b)

```

/* Search-based SSA Reconstruction
Input:  $b$ , a basic block
1 if  $b.\text{top} \neq \perp$  then
2   | return  $b.\text{top}$ 
3 end
4 pending $_\phi$   $\leftarrow$  new  $\phi$ -function in  $b$ ;
5  $v_\phi$   $\leftarrow$  result of pending $_\phi$ ;
6  $b.\text{top} \leftarrow v_\phi$ ;
7 reaching_defs  $\leftarrow []$ ;
8 foreach  $p \in b.\text{preds}$  do
9   | reaching_defs  $\leftarrow$  reaching_defs  $\cup$  FindDefFromBottom( $v, p$ )
10 end
11  $v_{\text{def}} \leftarrow \text{Phi-Necessary}(v_\phi, \text{reaching\_defs})$ ;
12 if  $v_{\text{def}} = v_\phi$  then
13   | set arguments of pending $_\phi$  to reaching_defs;
14 end
15 else
16   | remove pending $_\phi$ ;
17   |  $b.\text{top} \leftarrow v_{\text{def}}$ ;
18 end
19 return  $v_{\text{def}}$ 
```

Removing more unnecessary ϕ -functions.

In programs with loops, it can be the case that the local optimization performed when function **FindDefFromTop** calls **Phi-Necessary** does not remove all unnecessary ϕ -functions. This can happen in loops where ϕ -functions can become unnecessary because other ϕ -functions are optimized away. Consider the example in Figure 5.3. We look for a definition of x from block E . If Algorithm **FindDefFromTop** considers the blocks in an unfavorable order, e.g., E, D, C, B, A, D, C , some unnecessary ϕ -functions can not be removed by **Phi-Necessary**, as shown in Figure 5.3b. While the ϕ -function in block C can be eliminated by the local criterion applied by **Phi-Necessary**, the ϕ -function in block B remains. This is because the depth-first search carried out by **FindDefFromTop** will not visit block B a second time. To remove the remaining ϕ -functions, the local criterion can be iteratively applied to all placed ϕ -functions until a fixpoint is reached. For reducible control flow, this then produces the minimal number of placed

Procedure Phi-Necessary(v_ϕ , reaching_defs)

```

/* Checks if the set reaching_defs makes pending_φ necessary. This is
   the case if it is a subset of { $v_\phi$ , other} */
```

Input: v_ϕ , the variable defined by pending_φ

Input: reaching_defs, list of variables

Output: v_ϕ if pending_φ is necessary, the variable to use instead otherwise

```

1 other  $\leftarrow \perp$ ;
2 foreach  $v' \in$  reaching_defs do
3   if  $v' = v_\phi$  then continue;
4   if other  $= \perp$  then
5     | other  $\leftarrow v'$ 
6   end
7   else if  $d \neq$  other then
8     | return  $v_\phi$ 
9   end
10 end
/* this assertion is violated if reaching_defs contains only pending_φ
   which never can happen */
```

```

11 assert(other  $\neq \perp$ );
12 return other

```

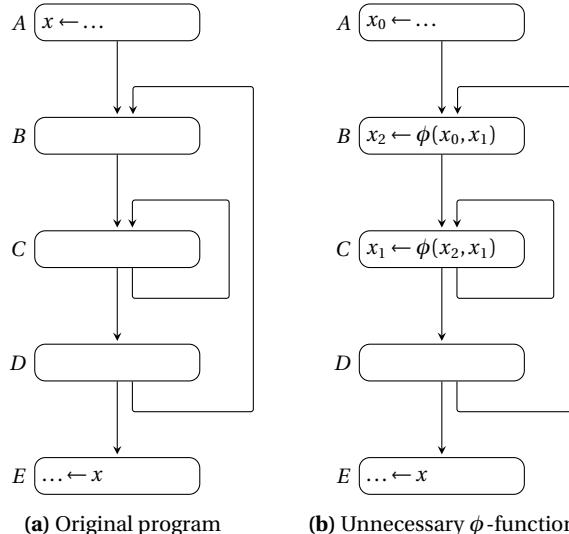


Fig. 5.3 Removal of unnecessary ϕ -functions.

ϕ -functions. The classical $(*)$ -graph in Figure 5.4 illustrates that this does not hold for the irreducible case.

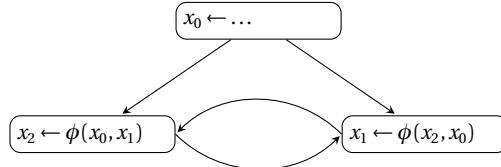


Fig. 5.4 The irreducible $(*)$ -Graph.

5.5 Conclusions

The algorithms presented in this chapter are independent of the transformation that violated SSA and can be used as a black box: For every variable for which SSA was violated, a routine is called that restores SSA. Both algorithms rely on computed def-use chains because they traverse all uses from a SSA variable to find suitable definitions; However, they differ in their prerequisites and their runtime behavior:

First algorithm. Choi et al. [70] based this algorithm on the iterated dominance frontiers like the classical SSA construction algorithm by Cytron et al. [87]. Hence, it is less suited for optimizations that also change the flow of control since that would require recomputing the iterated dominance frontiers. On the other hand, by using the iterated dominance frontiers, the algorithm can find the reaching definitions quickly by scanning the dominance tree upwards. Furthermore, one could also envision applying incremental algorithms to construct the dominance tree [230, 252] and the dominance frontier [253] to account for changes in the control flow. This has not yet been done and no practical experiences have been reported so far.

Second algorithm. It does not depend on additional analysis information such as iterated dominance frontiers or the dominance tree. Thus, it is well suited for transformations that change the CFG because no information needs to be recomputed. On the other hand, it might be slower to find the reaching definitions because they are searched by a depth-first search in the CFG.

Both approaches construct *pruned* SSA, i.e., they do not add dead ϕ -functions. The first approach produces minimal SSA by the very same arguments Cytron et al. [87] give. The second is similar to the construction algorithm that Click describes in his thesis [80] to construct SSA from an abstract syntax tree, and only produce minimal SSA (in the sense of Cytron) in the case of reducible control

flow. This follows from the iterative application of the function **Phi-Necessary** which implements the two simplification rules presented by Aycock and Hor-spool [22], who showed that their iterative application yields minimal SSA on re-duicible graphs. These two local rules can be extended to a non-local one which has to find strongly connected ϕ -components that all refer to the same exterior variable. Such a non-local check also eliminates unnecessary ϕ -functions in the irreducible case.

TODO: citation?

CHAPTER 6

Functional representations of SSA

L. Beringer

Progress: 95%

Final reading in progress

This chapter discusses alternative representations of SSA using the terminology and structuring mechanisms of functional programming languages. The reading of SSA as a discipline of functional programming arises from a correspondence between dominance and syntactic scope that subsequently extends to numerous aspects of control and data-flow structure.

The development of functional representations of SSA is motivated by the following considerations:

1. Relating the core ideas of SSA to concepts from other areas of compiler and programming language research provides conceptual insight into the SSA discipline and thus contributes to a better understanding of the practical appeal of SSA to compiler writers;
2. Reformulating SSA as a functional program makes explicit some of the syntactic conditions and semantic invariants that are implicit in the definition and use of SSA. Indeed, the introduction of SSA itself was motivated by a similar goal: to represent aspects of program structure—namely the def-use relationships—explicitly in syntax, by enforcing a particular naming discipline. In a similar way, functional representations directly enforce invariants such as “all ϕ -functions in a block must be of the same arity,” “the variables assigned to by these ϕ -functions must be distinct,” “ ϕ -functions are only allowed to occur at the beginning of a basic block,” or “each use of a variable should be dominated by its (unique) definition.” Constraints such as these would typically have to be validated or (re-)established after each optimization phase of an SSA-based compiler, but are typically enforced *by construction* if a functional representation is chosen. Consequently, less code is required, improving the robustness, maintainability, and code readability of the compiler;

formal parameters

3. The intuitive meaning of “unimplementable” ϕ -instructions is complemented by a concrete execution model, facilitating the rapid implementation of interpreters. This enables the compiler developers to experimentally validate SSA-based analyses and transformations at their genuine language level, without requiring SSA destruction. Indeed, functional intermediate code can often be directly emitted as a program in a high-level mainstream functional language, giving the compiler writer access to existing interpreters and compilation frameworks. Thus, rapid prototyping is supported and high-level evaluation of design decisions is enabled;
4. Formal frameworks of program analysis that exist for functional languages become applicable. Type systems provide a particularly attractive formalism due to their declarativeness and compositional nature. As type systems for functional languages typically support higher-order functions, they can be expected to generalize more easily to interprocedural analyses than other static analysis formalisms;
5. We obtain a formal basis for comparing variants of SSA—such as the variants discussed elsewhere in this book—for translating between these variants, and for constructing and destructing SSA. Correctness criteria for program analyses and associated transformations can be stated in a uniform manner, and can be proved to be satisfied using well-established reasoning principles.

Rather than discussing all these considerations in detail, the purpose of the present chapter is to informally highlight particular aspects of the correspondence and then point the reader to some more advanced material. Our exposition is example-driven but leads to the identification of concrete correspondence pairs between the imperative/SSA world and the functional world.

Like the remainder of the book, our discussion is restricted to code occurring in a single procedure.

6.1 Low-level functional program representations

Functional languages represent code using declarations of the form

$$\mathbf{function} \ f(x_0, \dots, x_n) = e \tag{6.1}$$

where the syntactic category of expression e conflates the notions of expressions and commands of imperative languages. Typically, e may contain further nested or (mutually) recursive function declarations. A declaration of the form (6.1) binds the formal parameters^{*} x_i and the function name f within e .

6.1.1 Variable assignment versus binding

A language construct provided by almost all functional languages is the *let-binding*:

`let $x = e_1$ in e_2 end.`

*let-binding
scope!static
variable!free
free|see variable
referential transparency*

The effect of this expression is to evaluate e_1 and bind the resulting value to variable x for the duration of the evaluation of e_2 . The code affected by this binding, e_2 , is called the *static scope** of x and is easily syntactically identifiable. In the following, we occasionally indicate scopes by code-enclosing boxes, and list the variables that are in scope using subscripts.

For example, the scope associated with the top-most binding of v to 3 in code

(6.2)

$$\begin{aligned} &\text{let } v = 3 \text{ in} \\ &\quad \boxed{\text{let } y = (\text{let } v = 2 \times v \text{ in } \boxed{4 \times v}_v \text{ end})} \\ &\quad \text{in } \boxed{y \times v}_{v,y} \text{ end} \\ &\quad \end{aligned}$$

spans both inner let-bindings, the scopes of which are themselves not nested inside one other as the inner binding of v occurs in the e_1 position of the let-binding for y .

In contrast to an assignment in an imperative language, a let-binding for variable x hides any previous value bound to x for the duration of evaluating e_2 but does not permanently overwrite it. Bindings are treated in a stack-like fashion, resulting in a tree-shaped nesting structure of boxes in our code excerpts. For example, in the above code, the inner binding of v to value $2 \times 3 = 6$ shadows the outer binding of v to value 3 precisely for the duration of the evaluation of the expression $4 \times v$. Once this evaluation has terminated (resulting in the binding of y to 24), the binding of v to 3 comes back in the game, yielding the overall result of 72.

The concepts of binding and static scope ensure that functional programs enjoy the characteristic feature of SSA, namely the fact that each use of a variable is uniquely associated with a point of definition. Indeed, the point of definition for a use of x is given by the *nearest enclosing binding of x* . Occurrences of variables in an expression that are not enclosed by a binding are called *free**. A well-formed procedure declaration contains all free variables of its body amongst its formal parameters. Thus, the notion of scope makes explicit a crucial invariant of SSA that is often left implicit: each use of a variable should be dominated by its (unique) definition.

In contrast to SSA, functional languages achieve the association of definitions to uses without imposing the global uniqueness of variables, as witnessed by the duplicate binding occurrences for v in the above code. As a consequence of this decoupling, functional languages enjoy a strong notion of *referential transparency**: the choice of x as the variable holding the result of e_1 depends only on the free variables of e_2 . For example, we may rename the inner v in code (6.2) to

*alpha-renaming@ α -renaming
compositional equational reasoning
continuation
CPS|seecontinuation*

z without altering the meaning of the code:

```
let  $v = 3$  in
  let  $y = (\text{let } z = 2 \times v \text{ in } [4 \times z]_{v,z} \text{ end})$ 
    in  $[y \times v]_{v,y} \text{ end}$ 
  end
```

(6.3)

Note that this conversion formally makes the outer v visible for the expression $4 \times z$, as indicated by the index v, z decorating its surrounding box.

In order to avoid altering the meaning of the program, the choice of the newly introduced variable has to be such that confusion with other variables is avoided. Formally, this means that a renaming

$\text{let } x = e_1 \text{ in } e_2 \text{ end}$ to $\text{let } y = e_1 \text{ in } e_2[y \leftrightarrow x] \text{ end}$

can only be carried out if y is not a *free* variable of e_2 . Moreover, in case that e_2 already contains some preexisting bindings to y , the *substitution* of x by y in e_2 (denoted by $e_2[y \leftrightarrow x]$ above) first renames these preexisting bindings in a suitable manner. Also note that the renaming only affects e_2 —any occurrences of x or y in e_1 refer to conceptually *different* but identically named variables, but the static scoping discipline ensures these will never be confused with the variables involved in the renaming. In general, the semantics-preserving renaming of bound variables is called α -renaming.^{*} Typically, program analyses for functional languages are compatible with α -renaming in that they behave equivalently for fragments that differ only in their choice of bound variables, and program transformations α -rename bound variables whenever necessary.

A consequence of referential transparency, and thus a property typically enjoyed by functional languages, is *compositional equational reasoning*:^{*} the meaning of a piece of code e is only dependent of its free variables, and can be calculated from the meaning of its subexpressions. For example, the meaning of a phrase $\text{let } x = e_1 \text{ in } e_2 \text{ end}$ only depends on the free variables of e_1 and on the free variables of e_2 *other than* x . Hence, languages with referential transparency allow one to replace a subexpression by some semantically equivalent phrase without altering the meaning of the surrounding code. Since semantic preservation is a core requirement of program transformations, the suitability of SSA for formulating and implementing such transformations can be explained by the proximity of SSA to functional languages.

6.1.2 Control flow: continuations

The correspondence between let-bindings and points of variable definition in assignments extends to other aspects of program structure, in particular to code in *continuation-passing-style*^{*}(CPS), a program representation routinely used in compilers for functional languages.

*lambda-notation@λ-
notation*

Satisfying a roughly similar purpose as return addresses or function pointers in imperative languages, a continuation specifies how the execution should proceed once the evaluation of the current code fragment has terminated. Syntactically, continuations are expressions that may occur in functional position (i.e., are typically applied to argument expressions), as is the case for the variable k in the following modification from code (6.2):

```
let v = 3 in
  let y = (let v = 2 × v in 4 × v end)
    in k(y × v) end
end
```

(6.4)

In effect, k represents any function that may be applied to the result of expression (6.2).

Surrounding code may specify the concrete continuation by binding k to a suitable expression. It is common practice to write these continuation-defining expressions in λ -notation*, i.e., in the form $\lambda x.e$ where x typically occurs free in e . The effect of the expression is to act as the (unnamed) function that sends x to $e(x)$, i.e., formal parameter x represents the place-holder for the argument to which the continuation is applied. Note that x is α -renameable, as λ acts as a binder. For example, a client of the above code fragment wishing to multiply the result by 2 may insert code (6.4) in the e_2 position of a let-binding for k that contains $\lambda x.2 \times x$ in its e_1 -position, as in the following code:

```
let k =  $\lambda x.2 \times x$ 
in let v = 3 in
  let y = (let z = 2 × v in 4 × z end)
    in k(y × v) end
  end
end
```

(6.5)

When the continuation k is applied to the argument $y \times v$, the (dynamic) value $y \times v$ (i.e., 72) is substituted for x in the expression $2 \times x$, just like in an ordinary function application.

Alternatively, the client may wrap fragment (6.4) in a function definition with formal argument k and construct the continuation in the calling code, where he would be free to choose a different name for the continuation-representing variable:

```
function f(k)=
  let v = 3 in
    let y = (let z = 2 × v in 4 × z end)
      in k(y × v) end
    end
  in let k =  $\lambda x.2 \times x$  in f(k) end
end.
```

(6.6)

Typically, the caller of f is itself parametric in *its* continuation, as in

```
function g(k)=
  let k' =  $\lambda x.k(x + 7)$  in f(k') end.
```

(6.7)

where f is invoked with a newly constructed continuation k' that applies the addition of 7 to its formal argument x (which at runtime will hold the result of f) before passing the resulting value on as an argument to the outer continuation k . In a similar way, the function

```
function  $h(y, k) =$ 
  let  $x = 4$  in
    let  $k' = \lambda z. k(z \times x)$ 
    in if  $y > 0$ 
      then let  $z = y \times 2$  in  $k'(z)$  end
      else let  $z = 3$  in  $k'(z)$  end
    end
  end
```

(6.8)

constructs from k a continuation k' that is invoked (with different arguments) in each branch of the conditional. In effect, the sharing of k' amounts to the definition of a control-flow merge point, as indicated by the CFG corresponding to h in Figure 6.1a. Contrary to the functional representation, the top-level continuation parameter k is not explicitly visible in the CFG—it roughly corresponds to the frame slot that holds the return address in an imperative procedure call.

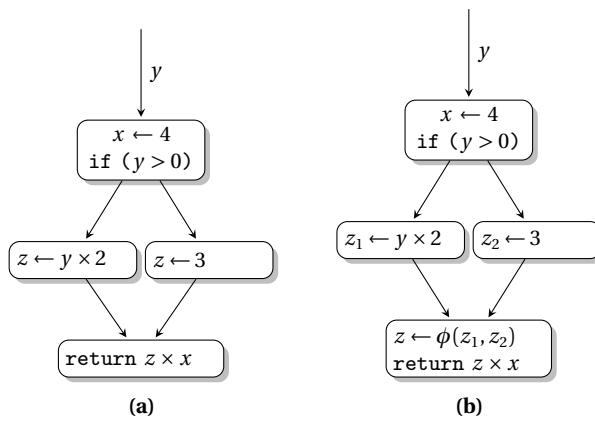


Fig. 6.1 Control-flow graph for code (6.8) (a), and SSA representation (b).

The SSA form of this CFG is shown in Figure 6.1b. If we apply similar renamings of z to z_1 and z_2 in the two branches of (6.8), we obtain the following fragment:

<pre>function h(y, k)= let x = 4 in let k' = $\lambda z. k(z \times x)$ in if y > 0 then let z₁ = y × 2 in k'(z₁) end else let z₂ = 3 in k'(z₂) end end end</pre>	<i>direct style</i> <i>tail-recursive</i>
---	--

(6.9)

We observe that the role of the formal parameter z of continuation k' is exactly that of a ϕ -function: to unify the arguments stemming from various calls sites by binding them to a common name for the duration of the ensuing code fragment—in this case just the return expression. As expected from the above understanding of scope and dominance, the scopes of the bindings for z_1 and z_2 coincide with the dominance regions of the identically named imperative variables: both terminate at the point of function invocation / jump to the control-flow merge point.

The fact that transforming (6.8) into (6.9) only involves the referentially transparent process of α -renaming indicates that program (6.8) already contains the essential structural properties that SSA distills from an imperative program.

Programs in CPS equip *all* functions declarations with continuation arguments. By interspersing ordinary code with continuation-forming expressions as shown above, they model the flow of control exclusively by communicating, constructing, and invoking continuations.

6.1.3 Control flow: direct style

An alternative to the explicit passing of continuation terms via additional function arguments is the *direct style*, in which we represent code as a set of locally named tail-recursive^{*} functions, for which the last operation is a call to a function, eliminating the need to save the return address.

In direct style, no continuation terms are constructed dynamically and then passed as function arguments, and we hence exclusively employ λ -free function definitions in our representation. For example, code (6.8) may be represented as

<pre>function h(y)= let x = 4 in function h'(z)= $z \times x \boxed{_{x,y,z,h,h'}}$ in if y > 0 then let z = y × 2 in $\boxed{h'(z)} \boxed{_{x,y,z,h,h'}}$ end else let z = 3 in $\boxed{h'(z)} \boxed{_{x,y,z,h,h'}}$ end end end</pre>	x, y, h
---	---

(6.10)

where the local function h' plays a similar role as the continuation k' and is jointly called from both branches. In contrast to the CPS representation, however, the body of h' returns its result directly rather than by passing it on as an argument to some continuation. Also note that neither the declaration of h nor that of h' contain additional continuation parameters. Thus, rather than handing its result directly over to some caller-specified receiver (as communicated by the continuation argument k), h simply returns control back to the caller, who is then responsible for any further execution. Roughly speaking, the effect is similar to the imperative compilation discipline of always setting the return address of a procedure call to the instruction pointer immediately following the call instruction.

A stricter format is obtained if the granularity of local functions is required to be that of basic blocks:

```
function  $h(y) =$ 
  let  $x = 4$  in
    function  $h'(z) = z \times x$ 
    in if  $y > 0$ 
      then function  $h_1() = \text{let } z = y \times 2 \text{ in } h'(z) \text{ end}$ 
      in  $h_1()$  end
    else function  $h_2() = \text{let } z = 3 \text{ in } h'(z) \text{ end}$ 
      in  $h_2()$  end
  end
end
```

(6.11)

Now, function invocations correspond precisely to jumps, reflecting more directly the CFG from Figure 6.1.

The choice between CPS and direct style is orthogonal to the granularity level of functions: both CPS and direct style are compatible with the strict notion of basic blocks but also with more relaxed formats such as extended basic blocks. In the extreme case, *all* control-flow points are explicitly named, i.e., one local function or continuation is introduced per instruction. The exploration of the resulting design space is ongoing, and has so far not yet led to a clear consensus in the literature. In our discussion below, we employ the arguably easier-to-read direct style, but the gist of the discussion applies equally well to CPS.

Independent of the granularity level of local functions, the process of moving from the CFG to the SSA form is again captured by suitably α -renaming the bindings of z in h_1 and h_2 :

```
function  $h(y) =$ 
  let  $x = 4$  in
    function  $h'(z) = z \times x$ 
    in if  $y > 0$ 
      then function  $h_1() = \text{let } z_1 = y \times 2 \text{ in } h'(z_1) \text{ end}$ 
      in  $h_1()$  end
    else function  $h_2() = \text{let } z_2 = 3 \text{ in } h'(z_2) \text{ end}$ 
      in  $h_2()$  end
  end
end
```

(6.12)

Again, the role of the formal parameter z of the control-flow merge point function h' is identical to that of a ϕ -function. In accordance with the fact that the basic blocks representing the arms of the conditional do not contain ϕ -functions, the local functions h_1 and h_2 have empty parameter lists—the free occurrence of y in the body of h_1 is bound at the top level by the formal argument of h .

let-normal form

6.1.4 Let-normal form

For both direct style and CPS the correspondence to SSA is most pronounced for code in *let-normal form*: each intermediate result must be explicitly named by a variable, and function arguments must be names or constants. Syntactically, let-normal form isolates basic instructions in a separate category of primitive terms a and then requires let-bindings to be of the form `let $x = a$ in e end`. In particular, neither jumps (conditional or unconditional) nor let-bindings are primitive. Let-normalized form is obtained by repeatedly rewriting code as follows:

$$\begin{array}{ccc} \text{let } x = \text{let } y = e \text{ in } \boxed{e'}_y \text{ end} & \text{into} & \text{let } y = e \\ \text{in } \boxed{e''}_x \text{ } & & \text{in } \boxed{\text{let } x = e' \text{ in } \boxed{e''}_{x,y} \text{ end}}_y \\ \text{end} & & \text{end}, \end{array}$$

subject to the side condition that y is not free in e'' . For example, let-normalizing code (6.3) pulls the let-binding for z to the outside of the binding for y , yielding

$$\begin{array}{c} \text{let } v = 3 \text{ in} \\ \quad \boxed{\text{let } z = 2 \times v \text{ in} \\ \quad \quad \boxed{\text{let } y = 4 \times z \text{ in } \boxed{y \times v}_{v,y,z} \text{ end}}_{v,z} \\ \quad \quad \text{end}}_v \\ \quad \text{end} \end{array} \tag{6.13}$$

Programs in let-normal form thus do not contain let-bindings in the e_1 -position of outer let-expressions. The stack discipline in which let-bindings are managed is simplified as scopes are nested inside each other. While still enjoying referential transparency, let-normal code is in closer correspondence to imperative code than non-normalized code as the chain of nested let-bindings directly reflects the sequence of statements in a basic block, interspersed occasionally by the definition of continuations or local functions. Exploiting this correspondence, we apply a simplified notation in the rest of this chapter where the scope-identifying boxes (including their indices) are omitted and chains of let-normal bindings are abbreviated by single comma-separated (and ordered) let-blocks. Using this convention, code (6.13) becomes

```
let  $v = 3,$ 
 $z = 2 \times v,$ 
 $y = 4 \times z$ 
in  $y \times v$  end
```

(6.14)

Summarizing our discussion up to this point, Table 6.1 collects some correspondences between functional and imperative/SSA concepts.

<i>Functional concept</i>	<i>Imperative/SSA concept</i>
variable binding in let	assignment (point of definition)
α -renaming	variable renaming
unique association of binding occurrences to uses	unique association of defs to uses
formal parameter of continuation/local function	ϕ -function (point of definition)
lexical scope of bound variable	dominance region

Table 6.1 Correspondence pairs between functional form and SSA (part I).

6.2 Functional construction and destruction of SSA

The relationship between SSA and functional languages is extended by the correspondences shown in Table 6.2. We discuss some of these aspects by considering the translation into SSA, using the program in Figure 6.2 as a running example.

<i>Functional concept</i>	<i>Imperative/SSA concept</i>
subterm relationship	control-flow successor relationship
arity of function f_i	number of ϕ -functions at beginning of b_i
distinctness of formal param. of f_i	distinctness of LHS-variables in the ϕ -block of b_i
number of call sites of function f_i	arity of ϕ -functions in block b_i
parameter lifting/dropping	addition/removal of ϕ -function
block floating/sinking	reordering according to dominator tree structure
potential nesting structure	dominator tree

Table 6.2 Correspondence pairs between functional form and SSA: program structure.

6.2.1 Initial construction using liveness analysis

A simple way to represent this program in let-normalized direct style is to introduce one function f_i for each basic block b_i . The body of each f_i arises by intro-

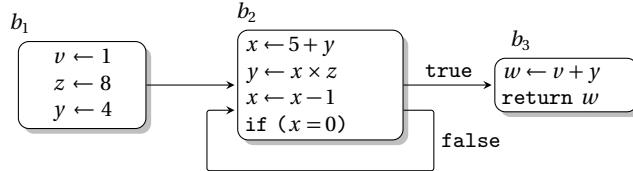


Fig. 6.2 Functional construction of SSA: running example.

ducing one let-binding for each assignment and converting jumps into function calls. In order to determine the formal parameters of these functions we perform a liveness analysis. For each basic block b_i , we choose an arbitrary enumeration of its live-in variables. We then use this enumeration as the list of formal parameters in the declaration of the function f_i , and also as the list of actual arguments in calls to f_i . We collect all function definitions in a block of mutually tail-recursive functions at the top level:

```

function  $f_1()$  = let  $v = 1, z = 8, y = 4$ 
    in  $f_2(v, z, y)$  end
and  $f_2(v, z, y)$  = let  $x = 5 + y, y = x \times z, x = x - 1$ 
    in if  $x = 0$  then  $f_3(y, v)$  else  $f_2(v, z, y)$  end
and  $f_3(y, v)$  = let  $w = y + v$  in  $w$  end
    in  $f_1()$  end
  
```

(6.15)

The resulting program has the following properties:

- All function declarations are *closed*: the free variables of their bodies are contained in their formal parameter lists¹;
- Variable names are not unique, but the unique association of definitions to uses is satisfied;
- Each subterm e_2 of a let-binding **let** $x = e_1$ **in** e_2 **end** corresponds to the control-flow successor of the assignment to x .

If desired, we may α -rename to make names globally unique. As the function declarations in code (6.15) are closed, all variable renamings are independent from each other. The resulting code (6.16) corresponds precisely to the SSA-program shown in Figure 6.3 (see also Table 6.2): each formal parameter of a function f_i is the target of one ϕ -function for the corresponding block b_i . The arguments of these ϕ -functions are the arguments in the corresponding positions in the calls to f_i . As the number of arguments in each call to f_i coincides with the number of formal parameters of f_i , the ϕ -functions in b_i are all of the same arity, namely the number of call sites to f_i . In order to coordinate the relative positioning of the arguments of the ϕ -functions, we choose an arbitrary enumeration of these call sites.

¹ Apart from the function identifiers f_i , which can always be chosen distinct from the variables.

*lambda-dropping@ λ -
dropping
 λ
block!sinking*

```

function  $f_1()$  = let  $v_1 = 1$ ,  $z_1 = 8$ ,  $y_1 = 4$ 
    in  $f_2(v_1, z_1, y_1)$  end
and  $f_2(v_2, z_2, y_2)$  = let  $x_1 = 5 + y_2$ ,  $y_3 = x_1 \times z_2$ ,  $x_2 = x_1 - 1$ 
    in if  $x_2 = 0$  then  $f_3(y_3, v_2)$  else  $f_2(v_2, z_2, y_3)$  end
and  $f_3(y_4, v_3)$  = let  $w_1 = y_4 + v_3$  in  $w_1$  end
in  $f_1()$  end

```

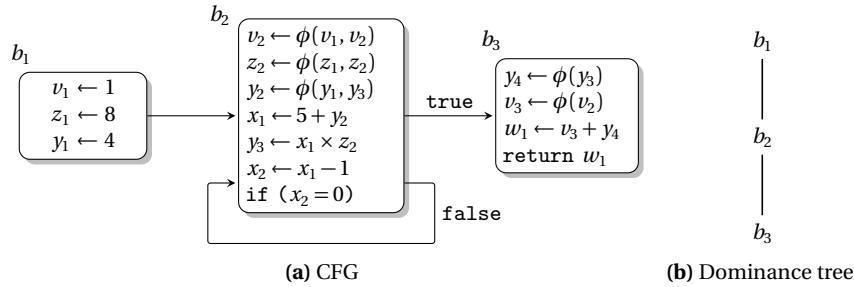
(6.16)


Fig. 6.3 Pruned (non-minimal) SSA form.

Under this perspective, the above construction of parameter lists amounts to equipping each b_i with ϕ -functions for all its live-in variables, with subsequent renaming of the variables. Thus, the above method corresponds to the construction of *pruned* (but not minimal) SSA—see Chapter 2.

While resulting in a legal SSA program, the construction clearly introduces more ϕ -functions than necessary. Each superfluous ϕ -function corresponds to the situation where all call sites to some function f_i pass identical arguments. The technique for eliminating such arguments is called λ -*dropping*, and is the inverse of the more widely known transformation λ -*lifting*.

6.2.2 λ -*dropping*

λ -*dropping*^{*} may be performed before or after variable names are made distinct, but for our purpose, the former option is more instructive. The transformation consists of two phases, *block sinking* and *parameter dropping*.

6.2.2.1 *Block sinking*

Block sinking^{*} analyzes the static call structure to identify which function definitions may be moved inside each other. For example, whenever our set of function declarations contains definitions $f(x_1, \dots, x_n) = e_f$ and $g(y_1, \dots, y_m) = e_g$ where $f \neq g$ and such that all calls to f occur in e_f or e_g , we can move the dec-

laration for f into that of g —note the similarity to the notion of dominance. If applied aggressively, block sinking indeed amounts to making the entire dominance tree structure explicit in the program representation. In particular, algorithms for computing the dominator tree from a CFG discussed elsewhere in this book can be applied to identify block sinking opportunities, where the CFG is given by the call graph of functions.

In our example (6.15), f_3 is only invoked from within f_2 , and f_2 is only called in the bodies of f_2 and f_1 (see the dominator tree in Figure 6.3 (right)). We may thus move the definition of f_3 into that of f_2 , and the latter one into f_1 .

Several options exist as to where f should be placed in its host function. The first option is to place f at the beginning of g , by rewriting to

```
function  $g(y_1, \dots, y_m) =$  function  $f(x_1, \dots, x_n) = e_f$ 
in  $e_g$  end.
```

This transformation does not alter the meaning of the code, as the declaration of f is closed: moving f into the scope of the formal parameters y_1, \dots, y_m (and also into the scope of g itself) does not alter the bindings to which variable uses inside e_f refer.

Applying this transformation to example (6.15) yields the following code:

```
function  $f_1() =$ 
function  $f_2(v, z, y) =$ 
function  $f_3(y, v) =$  let  $w = y + v$  in  $w$  end
in let  $x = 5 + y, y = x \times z, x = x - 1$ 
in if  $x = 0$  then  $f_3(y, v)$  else  $f_2(v, z, y)$  end
end
in let  $v = 1, z = 8, y = 4$  in  $f_2(v, z, y)$  end
end
in  $f_1()$  end
```

(6.17)

An alternative strategy is to insert f near the end of its host function g , in the vicinity of the calls to f . This brings the declaration of f additionally into the scope of all let-bindings in e_g . Again, referential transparency and preservation of semantics are respected as the declaration on f is closed. In our case, the alternative strategy yields the following code:

```
function  $f_1() =$ 
let  $v = 1, z = 8, y = 4$ 
in function  $f_2(v, z, y) =$ 
let  $x = 5 + y, y = x \times z, x = x - 1$ 
in if  $x = 0$ 
then function  $f_3(y, v) =$  let  $w = y + v$  in  $w$  end
in  $f_3(y, v)$  end
else  $f_2(v, z, y)$ 
end
in  $f_2(v, z, y)$  end
end
in  $f_1()$  end
```

(6.18)

*dropping!parameter
parameter dropping*

In general, one would insert f directly prior to its call if g contains only a single call site for f . In case that g contains multiple call sites for f , these are (due to their tail-recursive positioning) in different arms of a conditional, and we would insert f directly prior to this conditional.

Both outlined placement strategies result in code whose nesting structure reflects the dominance relationship of the imperative code. In our example, code (6.17) and (6.18) both nest f_3 inside f_2 inside f_1 , in accordance with the dominator tree of the imperative program show on Figure 6.3.

6.2.2.2 Parameter dropping

The second phase of λ -dropping, *parameter dropping*^{*}, removes superfluous parameters based on the syntactic scope structure. Removing a parameter x from the declaration of some function f has the effect that any use of x inside the body of f will not be bound by the declaration of f any longer, but by the (innermost) binding for x that *contains* the declaration of f . In order to ensure that removing the parameter does not alter the meaning of the program, we thus have to ensure that this f -containing binding for x also contains any call to f , since the binding applicable at the call site determines the value that is passed as the actual argument (before x is deleted from the parameter list).

For example, the two parameters of f_3 in (6.18) can be removed without altering the meaning of the code, as we can statically predict the values they will be instantiated with: parameter y will be instantiated with the result of $x \times z$ (which is bound to y in the first line in the body of f_2), and v will always be bound to the value passed via the parameter v of f_2 . In particular, the bindings for y and v at the *declaration* site of f_3 are identical to those applicable at the *call* to f_3 .

In general, we may drop a parameter x from the declaration of a possibly recursive function f if the following two conditions are met:

1. The tightest scope for x enclosing the declaration of f coincides with the tightest scope for x surrounding each call site to f outside of its declaration;
2. The tightest scope for x enclosing any recursive call to f (i.e., a call to f in the body of f) is the one associated with the formal parameter x in the declaration of f .

Similarly, we may simultaneously drop a parameter x occurring in *all* declarations of a block of mutually recursive functions f_1, \dots, f_n , if the scope for x at the point of declaration of the block coincides with the tightest scope in force at any call site to some f_i outside the block, and if in each call to some f_i inside some f_j , the tightest scope for x is the one associated with the formal parameter x of f_j . In both cases, dropping a parameter means to remove it from the list of formal parameter lists of the function declarations concerned, and also from the argument lists of the corresponding function calls.

In code (6.18), these conditions sanction the removal of both parameters from the nonrecursive function f_3 . The scope applicable for v at the site of dec-

laration of f_3 and also at its call site is the one rooted at the formal parameter v of f_2 . In case of y , the common scope is the one rooted at the let-binding for y in the body of f_2 . We thus obtain the following code:

```
function  $f_1()$  =
  let  $v = 1, z = 8, y = 4$ 
  in function  $f_2(v, z, y) =$ 
    let  $x = 5 + y, y = x \times z, x = x - 1$ 
    in if  $x = 0$ 
      then function  $f_3() = \text{let } w = y + v \text{ in } w \text{ end}$ 
      in  $f_3()$  end
      else  $f_2(v, z, y)$ 
    end
    in  $f_2(v, z, y)$  end
  end
in  $f_1()$  end
```

(6.19)

Considering the recursive function f_2 next we observe that the recursive call is in the scope of the let-binding for y in the body of f_2 , preventing us from removing y . In contrast, neither v nor z have binding occurrences in the body of f_2 . The scopes applicable at the external call site to f_2 coincide with those applicable at its site of declaration and are given by the scopes rooted in the let-bindings for v and z . Thus, parameters v and z may be removed from f_2 :

```
function  $f_1()$  =
  let  $v = 1, z = 8, y = 4$ 
  in function  $f_2(y) =$ 
    let  $x = 5 + y, y = x \times z, x = x - 1$ 
    in if  $x = 0$ 
      then function  $f_3() = \text{let } w = y + v \text{ in } w \text{ end}$ 
      in  $f_3()$  end
      else  $f_2(y)$ 
    end
    in  $f_2(y)$  end
  end
in  $f_1()$  end
```

(6.20)

Interpreting the uniquely-renamed variant of (6.20) back in SSA yields the desired minimal code with a single ϕ -function, for variable y at the beginning of block b_2 , see Figure 6.4. The reason that this ϕ -function can't be eliminated (the redefinition of y in the loop) is precisely the reason why y survives parameter-dropping.

TODO: check minimality

Given this understanding of parameter dropping we can also see why inserting functions near the end of their hosts during block sinking (as in code (6.18)) is in general preferable to inserting them at the beginning of their hosts (as in code (6.17)): the placement of function declarations in the vicinity of their calls potentially enables the dropping of more parameters, namely those that are let-bound in the body of the host function.

An immediate consequence of this strategy is that blocks with a single predecessor indeed do not contain ϕ -functions. Such a block b_f is necessarily dom-

loop-closed|exampleidx

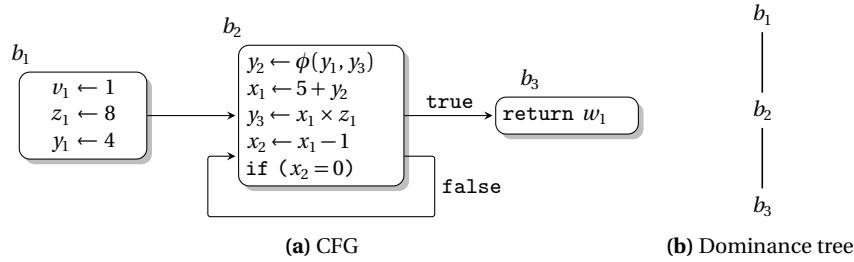


Fig. 6.4 SSA code after λ -dropping.

inated by its predecessor b_g , hence we can always nest f inside g . Inserting f in e_g directly prior to its call site implies that condition (1) is necessarily satisfied for all parameters x of f . Thus, all parameters of f can be dropped and no ϕ -function is generated.

6.2.3 Nesting, dominance, loop-closure

As we observed above, analyzing whether function definitions may be nested inside one another is tantamount to analyzing the imperative dominance structure: function f_i may be moved inside f_j exactly if all non-recursive calls to f_i come from within f_j , i.e., exactly if all paths from the initial program point to block b_i traverse b_j , i.e., exactly if b_j dominates b_i . This observation is merely the extension to function identifiers of our earlier remark that lexical scope coincides with the dominance region, and that points of definition/binding occurrences should dominate the uses. Indeed, functional languages do not distinguish between code and data when aspects of binding and use of variables are concerned, as witnessed by our use of the let-binding construct for binding code-representing expressions to the variables k in our syntax for CPS.

Thus, the dominator tree immediately suggests a function nesting scheme, where all children of a node are represented as a single block of mutually recursive function declarations.²

The choice as to where functions are placed corresponds to variants of SSA. For example, in loop-closed^{*} SSA form (see Chapters 18 and 11), SSA names that are defined in a loop must not be used outside the loop. To this end, special-purpose unary ϕ -nodes are inserted for these variables at the loop exit points. As the loop is unrolled, the arity of these trivial ϕ -nodes grows with the number of unrollings, and the program continuation is always supplied with the value of the variable obtained in the final iteration of the loop. In our example, the only loop-

² Refinements of this representation will be sketched in Section 6.3.

defined variable used in f_3 is y —and we already observed in code (6.17) how we can prevent the dropping of y from the parameter list of f_3 : we insert f_3 at the *beginning* of f_2 , preceding the let-binding for y . Of course, we would still like to drop as many parameters from f_2 as possible, hence we apply the following placement policy during block sinking: functions that are targets of loop-exiting function calls and have live-in variables that are defined in the loop are placed at the beginning of the loop headers. Other functions are placed at the end of their hosts. Applying this policy to our original program (6.15) yields (6.21).

```
function  $f_1()$  =
  let  $v = 1, z = 8, y = 4$ 
  in function  $f_2(v, z, y) =$ 
    function  $f_3(y, v) = \text{let } w = y + v \text{ in } w$  end
    in let  $x = 5 + y, y = x \times z, x = x - 1$ 
      in if  $x = 0$  then  $f_3(y, v)$  else  $f_2(v, z, y)$  end
    end
    in  $f_2(v, z, y)$  end
  end
in  $f_1()$  end
```

(6.21)

We may now drop v (but not y) from the parameter list of f_3 , and v and z from f_2 , to obtain code (6.22).

```
function  $f_1()$  =
  let  $v = 1, z = 8, y = 4$ 
  in function  $f_2(y) =$ 
    function  $f_3(y) = \text{let } w = y + v \text{ in } w$  end
    in let  $x = 5 + y, y = x \times z, x = x - 1$ 
      in if  $x = 0$  then  $f_3(y)$  else  $f_2(y)$  end
    end
    in  $f_2(y)$  end
  end
in  $f_1()$  end
```

(6.22)

The SSA form corresponding to (6.22) contains the desired loop-closing ϕ -node for y at the beginning of b_3 , as shown in Figure 6.5a. The nesting structure of both (6.21) and (6.22) coincides with the dominance structure of the original imperative code and its loop-closed SSA form.

We unroll^{*} the loop by duplicating the body of f_2 , *without duplicating the declaration of f_3* :

unrolling|exampleidx

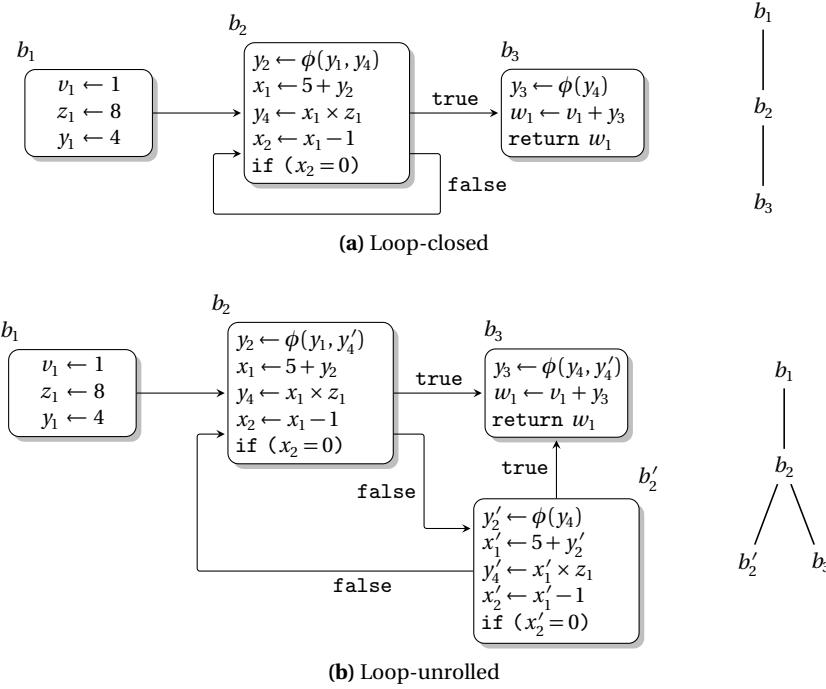


Fig. 6.5 Loop-closed (a) and loop-unrolled (b) forms of running example program, corresponding to codes (6.22) and (6.23), respectively.

```

function f1() =
  let v = 1, z = 8, y = 4
  in function f2(y) =
    function f3(y) = let w = y + v in w end
    in let x = 5 + y, y = x × z, x = x - 1
      in if x = 0 then f3(y)
      else function f2(y) =
        let x = 5 + y, y = x × z, x = x - 1
          in if x = 0 then f3(y) else f2(y) end
        in f2(y) end
      end
      in f2(y) end
    end
  in f1() end

```

(6.23)

Both calls to f_3 are in the scope of the declaration of f_3 and contain the appropriate loop-closing arguments. In the SSA reading of this code—shown in Figure 6.5b—the first instruction in b_3 has turned into a non-trivial ϕ -node. As expected, the parameters of this ϕ -node correspond to the two control-flow arcs leading into b_3 , one for each call site to f_3 in code (6.23). Moreover, the call and

nesting structure of (6.23) is indeed in agreement with the control flow and dominance structure of the loop-unrolled SSA representation.

*destruction!SSA
SSA!destruction
lost-copy problem
swap problem*

6.2.4 Destruction of SSA

The above example code excerpts where variables are not made distinct exhibit a further pattern: the argument list of any call coincides with the list of formal parameters of the invoked function. This discipline is not enjoyed by functional programs in general, and is often destroyed by optimizing program transformations. However, programs that do obey this discipline can be immediately converted to imperative non-SSA form. Thus, the task of SSA destruction^{*} amounts to amounts to converting a functional program with arbitrary argument lists into one where argument lists and formal parameter lists coincide for each function. This can be achieved by introducing additional let-bindings of the form `let $x = y$ in e end`. For example, a call $f(v, z, y)$ where f is declared as `function $f(x, y, z) = e$` may be converted to

```
let  $x = v, a = z, z = y, y = a$  in  $f(x, y, z)$ 
```

in correspondence to the move instructions introduced in imperative in correspondence to the move instructions introduced in imperative formulations of SSA destruction (see Chapters 3 and 22). Appropriate transformations can be formulated as manipulations of the functional representation, although the target format is not immune against α -renaming and thus only syntactically a functional language. For example, we can give a local algorithm that considers each call site individually and avoids the “lost-copy” and “swap” problems (cf. Chapter ??); Instead of introducing let-bindings for all parameter positions of a call, the algorithm scales with the number and size of cycles that span identically named arguments and parameters (like the cycle between y and z above), and employs a single additional variable (called a in the above code) to break all these cycles one by one.

.....

6.3 Refined block sinking and loop nesting forests

As discussed when outlining λ -dropping, block sinking is governed by the dominance relation between basic blocks. Thus, a typical dominance tree with root b and subtrees rooted at b_1, \dots, b_n is most naturally represented as a block of function declarations for the f_i , nested inside the declaration of f :

reducible

```

function  $f(\dots) = \text{let} \dots < \text{body of } f > \dots \text{in}$ 
  function  $f_1(\dots) = e_1$       (* body of  $b_1$ , with calls to  $b, b_i$  *)
  :
  and  $f_n(\dots) = e_n$       (* body of  $b_n$ , with calls to  $b, b_i$  *)
in  $\dots < \text{calls to } b, b_i \text{ from } b > \dots \text{end}$ 

```

(6.24)

By exploiting additional control-flow structure between the b_i , it is possible to obtain refined placements, namely placements that correspond to notions of *loop nesting forests* that have been identified in the SSA literature.

These refinements arise if we enrich the above dominance tree by adding arrows $b_i \rightarrow b_j$ whenever the CFG contains a directed edge from one of the dominance successors of b_i (i.e., the descendants of b_i in the dominance tree) to b_j .

In the case of a *reducible** CFG, the resulting graph contains only trivial loops. Ignoring these self-loops, we perform a post-order DFS (or more generally a reverse topological ordering) amongst the b_i and stagger the function declarations according to the resulting order. As an example, consider the CFG in

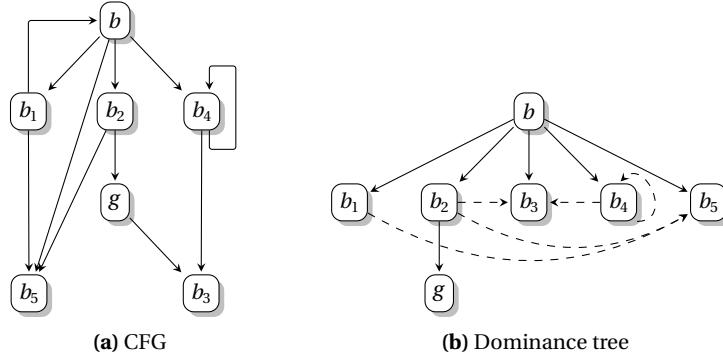


Fig. 6.6 Example placement for reducible flow graphs. (a) Control-flow graph; (b) Overlay of CFG (dashed edges) arcs between dominated blocks onto dominance graph (solid edges).

Figure 6.6a and its enriched dominance tree shown in Figure 6.6b. A possible (but not unique) ordering of the children of b is $[b_5, b_1, b_3, b_2, b_4]$, resulting in the nesting shown in code (6.25).

```

function  $f(\dots) = \text{let} \dots < \text{body of } f > \dots \text{in}$ 
  function  $f_5(\dots) = e_5$ 
  in function  $f_1(\dots) = e_1$           (* contains calls to  $f$  and  $f_5$  *)
    in function  $f_3(\dots) = e_3$ 
      in function  $f_2(\dots) =$ 
        let  $\dots < \text{body of } f_2 > \dots$ 
        in function  $g(\dots) = e_g$       (* contains call to  $f_3$  *)
          in  $\dots < \text{calls to } f_5 \text{ and } g > \dots \text{end}$ 
    in function  $f_4(\dots) = e_4$       (* contains calls to  $f_3$  and  $f_4$  *)
    in  $\dots < \text{calls to } f_1, f_2, f_4, f_5 > \dots \text{end}$ 

```

(6.25)

The code respects the dominance relationship in much the same way as the naive placement, but additionally makes f_1 inaccessible from within e_5 , and makes f_3 inaccessible from within f_1 or f_5 . As the reordering does not move function declarations *inside* each other (in particular: no function declaration is brought into or moved out of the scope of the formal parameters of any other function) the reordering does not affect the potential to subsequently perform parameter dropping.

Declaring functions using λ -abstraction brings further improvements. This enables us not only to syntactically distinguish between loops and non-recursive control-flow structures using the distinction between **let** and **letrec**^{*} present in many functional languages, but also to further restrict the visibility of function names. Indeed, while b_3 is immediately dominated by b in the above example, its only control-flow predecessors are b_2/g and b_4 . We would hence like to make the declaration of f_3 local to the $tuple(f_2, f_4)$, i.e., invisible to f . This can be achieved by combining **let/letrec** bindings with pattern matching, if we insert the shared declaration of f_3 *between* the declaration of the names f_2 and f_4 and the λ -bindings of their formal parameters p_i :

```

letrec  $f = \lambda p. \text{let} \dots < \text{body of } f > \dots \text{in}$ 
  let  $f_5 = \lambda p_5. e_5$ 
  in let  $f_1 = \lambda p_1. e_1$  (* contains calls to  $f$  and  $f_5$  *)
  in letrec ( $f_2, f_4$ ) =
    let  $f_3 = \lambda p_3. e_3$ 
    in ( $\lambda p_2. \text{let} \dots < \text{body of } f_2 > \dots \text{in}$ 
      let  $g = \lambda p_g. e_g$  (* contains call to  $f_3$  *)
      in ... < calls to  $f_5$  and  $g$  > ... end
      ,  $\lambda p_4. e_4$  (* contains calls to  $f_3$  and  $f_4$  *)
    end (* declaration of tuple( $f_2, f_4$ ) ends here *)
    in ... < calls to  $f_1, f_2, f_4, f_5$  > ... end

```

(6.26)

The recursiveness of f_4 is inherited by the function pair (f_2, f_4) but f_3 remains non-recursive. In general, the role of f_3 is played by any merge point b_i that is not directly called from the dominator node b .

In the case of *irreducible*^{*} CFGs, the enriched dominance tree is no longer acyclic (even when ignoring self-loops). In this case, the functional representation not only depends on the chosen DFS order but additionally on the partitioning of the enriched graph into loops. As each loop forms a *strongly connected component* (SCC), different partitioning are possible, corresponding to different notions of *loop nesting forests*^{*}. Of the various loop nesting forest strategies proposed in the SSA-literature, the scheme introduced by Steensgaard is particularly appealing from a functional perspective.

In Steensgaard's notion of loops, the headers H of a loop $L = (B, H)$ are precisely the entry nodes of its body B , i.e., those nodes in B that have a predecessor outside of B . For example, G_0 shown in Figure 6.7 contains the outer loop $L_0 = (\{u, v, w, x\}, \{u, v\})$, whose constituents B_0 are determined as the maximal SCC of G_0 . Removing the back-edges of L_0 from G_0 (i.e., edges from B_0 to H_0) yields G_1 , whose (only) SCC determines a further inner loop $L_1 = (\{w, x\}, \{w, x\})$.

letrec
irreducible
loop nesting forest

Removing the back-edges of L_1 from G_1 results in the acyclic G_2 , terminating the process.

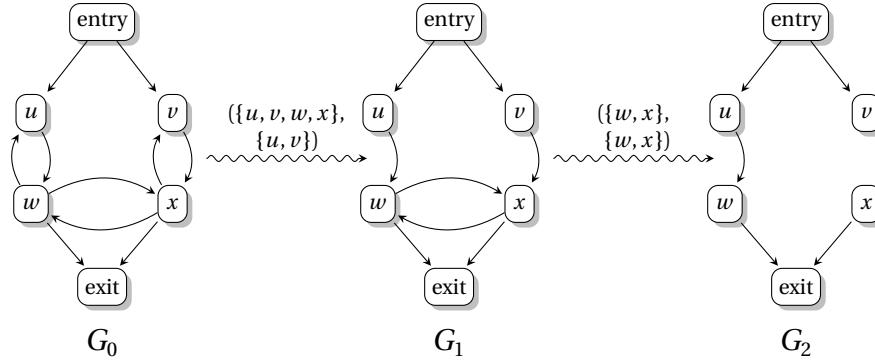


Fig. 6.7 Illustration of Steensgaard's construction of loop nesting forests:

Figure 6.8a shows the CFG-enriched dominance tree of G_0 . The body of loop L_0 is easily identified as the maximal SCC, and likewise the body of L_1 once the cycles (u, w) and (x, v) are broken by the removal of the back-edges $w \rightarrow u$ and $x \rightarrow v$.

The loop nesting forest resulting from Steensgaard's construction is shown in Figure 6.8b. Loops are drawn as ellipsis decorated with the appropriate header nodes, and nested in accordance with the containment relation $B_1 \subset B_0$ between the bodies.

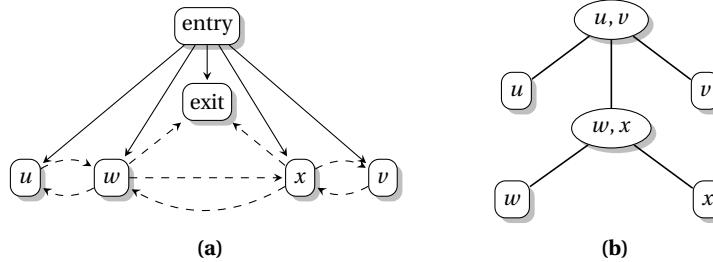


Fig. 6.8 Illustration of Steensgaard's construction of loop nesting forests: (a) CFG-enriched dominance tree; (b) resulting loop nesting forest.

In the functional representation, a loop $L = (B, \{h_1, \dots, h_n\})$ yields a function declaration block for functions h_1, \dots, h_n , with private declarations for the non-headers from $B \setminus H$. In our example, loop L_0 provides entry points for the headers u and v but not for its non-headers w and x . Instead, the loop comprised of

the latter nodes, L_1 , is nested inside the definition of L_0 , in accordance with the loop nesting forest.

```
function entry(...)=  
  let...< body of entry >...  
  in letrec (u, v)= (* define outer loop  $L_0$ , with headers u, v *)  
    letrec (w, x)= (* define inner loop  $L_1$ , with headers w, x *)  
      let exit =  $\lambda p_{exit}.$  ...< body of exit >  
      in ( $\lambda p_w.$  ...< body of w, with calls to u, x, and exit >...  
           , $\lambda p_x.$  ...< body of x, with calls to w, v, and exit >...)  
      end (* end of inner loop *)  
    in ( $\lambda p_u.$  ...< body of u, with call to w >...  
         , $\lambda p_v.$  ...< body of v, with call to x >...)  
    end (* end of outer loop *)  
  in ...< calls from entry to u and v >...
```

(6.27)

By placing L_1 inside L_0 according to the scheme from code (6.26) and making **exit** private to L_1 , we obtain the representation (6.27) which captures all the essential information of Steensgaard's construction. Effectively, the functional reading of the loop nesting forest extends the earlier correspondence between the nesting of individual functions and the dominance relationship to groups of functions and basic blocks: loop L_0 *dominates* L_1 in the sense that any path from **entry** to a node in L_1 passes through L_0 ; more specifically, any path from **entry** to a *header* of L_1 passes through a *header* of L_0 .

In general, each step of Steensgaard's construction may identify several loops, as a CFG may contain several maximal SCCs. As the bodies of these SCCs are necessarily non-overlapping, the construction yields a forest comprised of trees shaped like the loop nesting forest in Figure 6.8b. As the relationship between the trees is necessarily acyclic, the declarations of the function declaration tuples corresponding to the trees can be placed according to the loop-extended notion of dominance.

6.4 Pointers to the literature

Shortly after the introduction of SSA, O'Donnell [211] and Kelsey [161] noted the correspondence between let-bindings and points of variable declaration and its extension to other aspects of program structure using continuation-passing style. Appel [16, 15] popularized the correspondence using a direct-style representation, building on his earlier experience with continuation-based compilation [14].

Continuations and low-level functional languages have been an object of intensive study since their inception about four decades ago [278, 174]. For retrospective accounts of the historical development, see Reynolds [237] and Wadsworth [282]. Early studies of CPS and direct style include work by Reynolds and Plotkin [235, 236, 226]. Two prominent examples of CPS-based compilers are

monadic

those by Sussman et al. [266] and Appel [14]. An active area of research concerns the relative merit of the various functional representations, algorithms for formal conversion between these formats, and their efficient compilation to machine code, in particular with respect to their integration with program analyses and optimizing transformations [93, 234, 162]. A particularly appealing variant is that of Kennedy [162], where the approach to explicitly name control-flow points such as merge points is taken to its logical conclusion. By mandating *all* control-flow points to be explicitly named, a uniform representation is obtained that allows optimizing transformations to be implemented efficiently, avoiding the administrative overhead to which some of the alternative approaches are susceptible.

Occasionally, the term *direct style* refers to the combination of tail-recursive functions and let-normal form, and the conditions on the latter notion are strengthened so that only variables may appear as branch conditions. Variations of this discipline include *administrative normal form* (A-normal form, ANF) [123], B-form [267], and SIL [271].

Closely related to continuations and direct-style representation are *monadic** intermediate languages as used by Benton et al. [29] and Peyton-Jones et al. [221]. These partition expressions into categories of *values* and *computations*, similar to the isolation of primitive terms in let-normal form [236, 226]. This allows one to treat side-effects (memory access, IO, exceptions, etc.) in a uniform way, following Moggi [202], and thus simplifies reasoning about program analyses and the associated transformations in the presence of impure language features.

Lambda-lifting and dropping are well-known transformations in the functional programming community, and are studied in-depth by Johnsson [156] and Danvy et al. [94].

Rideau et al. [238] present an in-depth study of SSA destruction, including a verified implementation in the proof assistant Coq for the “windmills” problem, i.e., the task of correctly introducing ϕ -compensating assignments. The local algorithm to avoid the lost-copy problem and swap problem identified by Briggs et al. [52] was given by Beringer [30]. In this solution, the algorithm to break the cycles is in line with the results of May [197].

We are not aware of previous work that transfers the analysis of loop nesting forests to the functional setting, or of loop analyses in the functional world that correspond to loop nesting forests. Our discussion of Steensgaard’s construction was based on a classification of loop nesting forests by Ramalingam [229], which also served as the source of the example in Figure 6.7. Two alternative constructions discussed by Ramalingam are those by Sreedhar, Gao and Lee [254], and Havlak [142]. To us, it appears that a functional reading of Sreedhar, Gao and Lee’s construction would essentially yield the nesting mentioned at the beginning of Section 6.3. Regarding Havlak’s construction, the fact that entry points of loops are not necessarily classified as headers appears to make an elegant representation in functional form at least challenging.

Extending the syntactic correspondences between SSA and functional languages, similarities may be identified between their characteristic program anal-

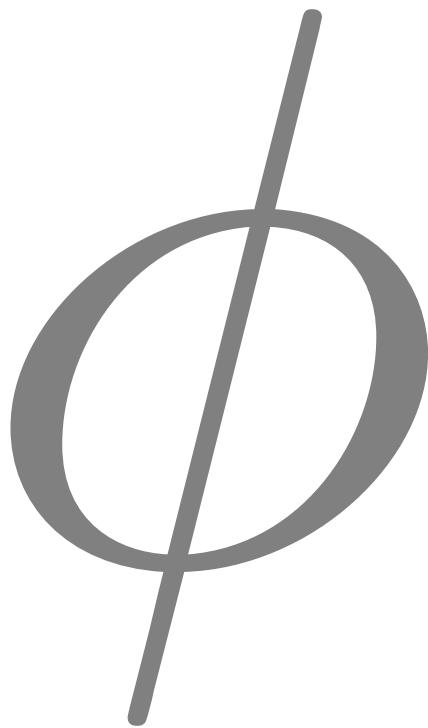
ysis frameworks, *data-flow analyses* and *type systems*. Chakravarty et al. [68] prove the correctness of a functional representation of Wegmann and Zadeck's SSA-based sparse conditional constant propagation algorithm [287]. Beringer et al. [31] consider data-flow equations for liveness and read-once variables, and formally translate their solutions to properties of corresponding typing derivations. Laud et al. [176] present a formal correspondence between data-flow analyses and type systems but consider a simple imperative language rather than SSA. The textbook [209] presents a unifying perspective on various program analysis techniques, including data-flow analysis, abstract interpretation, and type systems. As outlined in this book, the extension of type systems to effect systems is particularly suitable for integrating data with control-flow analyses. Indeed, as functional languages treat code and data in an integrated fashion, type-based analyses provide a natural setting for developing extensions of (SSA-based) program analyses to inter-procedural analyses. For additional pointers on type-based program analysis, see Palsberg's article [217].

6.5 Concluding remarks

In addition to low-level functional languages, alternative representations for SSA have been proposed, but their discussion is beyond the scope of this chapter. Glesner [133] employs an encoding in terms of abstract state machines [136] that disregards the sequential control flow inside basic blocks but retains control flow between basic blocks to prove the correctness of a code generation transformation. Later work by the same author uses a more direct representation of SSA in the theorem prover Isabelle/HOL for studying further SSA-based analyses.

Matsuno and Ohori [196] present a formalism that captures core aspects of SSA in a notion of (non-standard) types while leaving the program text in unstructured non-SSA form. Instead of classifying values or computations, types represent the definition points of variables. Contexts associate program variables at each program point with types in the standard fashion, but the non-standard notion of types means that this association models the reaching-definitions relationship rather than characterizing the values held in the variables at runtime. Noting a correspondence between the types associated with a variable and the sets of def-use paths, the authors admit types to be formulated over type variables whose introduction and use corresponds to the introduction of ϕ -nodes in SSA.

Finally, Pop et al.'s model [227] dispenses with control flow entirely and instead views programs as sets of equations that model the assignment of values to variables in a style reminiscent of partial recursive functions. This model is discussed in more detail in Chapter 11.



Part II

Analysis

Progress: 81%

Progress: 81%

- what structural invariants does ssa provide for analysis (majumdar? ask palsberg to ask him)
- available expressions become available values
- split the analyses between algorithms that need kill calculation and those that dont?
- bidirectional (but dont discuss ssi)

TODO: some throwaway
remarks that never found
a place

CHAPTER 7

Introduction

A. Ertl

Progress: 0%

Material gathering in progress

.....
7.1 benefits of SSA

.....
7.2 properties: suffixed representation can use dense vector, vs DF needs (x,PP) needs hashtable

.....
7.3 phi node construction precomputes DF merge

.....
7.4 dominance properties imply that simple forward analysis works

.....
7.5 Flow insensitive somewhat improved by SSA

.....
7.6 building defuse chains: for each SSA variable can get pointer to next var

.....
7.7 improves SSA DF-propagation

CHAPTER 8

Propagating Information using SSA

*F. Brandner
D. Novillo*

Progress: 90%

 Text and figures formating in progress

8.1 Overview

A central task of compilers is to *optimize* a given input program such that the resulting code is more efficient in terms of execution time, code size, or some other metric of interest. However, in order to perform these optimizations, typically some form of *program analysis* is required to determine if a given program transformation is applicable, to estimate its profitability, and to guarantee its correctness.

Data-flow analysis is a simple yet powerful approach to program analysis that is utilized by many compiler frameworks and program analysis tools today. We will introduce the basic concepts of traditional data-flow analysis in this chapter and will show how the *static single assignment* form (SSA) facilitates the design and implementation of equivalent analyses. We will also show that how the SSA property allows us to reduce the compilation time and memory consumption of the data-flow analyses that this program representation supports.

Traditionally, data-flow analysis is performed on a *control-flow graph* representation (CFG) of the input program. Nodes in the graph represent operations and edges represent the potential flow of program execution. Information on certain *program properties* is propagated among the nodes along the control-flow edges until the computed information stabilizes, i.e., no *new* information can be inferred from the program.

The *propagation engine* presented in the following sections is an extension of the well known approach by Wegman and Zadeck for *sparse conditional constant propagation* (also known as SSA-CCP). Instead of using the CFG they represent the input program as an *SSA graph* as defined in Chapter 18: operations are again represented as nodes in this graph, however, the edges represent *data dependencies* instead of control flow. This representation allows a selective propagation of program properties among data dependent graph nodes only. As before, the processing stops when the information associated with the graph nodes stabilizes. The basic algorithm is not limited to constant propagation and can also be applied to solve a large class of other data-flow problems efficiently. However, not all data-flow analyses can be modeled. In this chapter, we will also investigate the limitations of the SSA-based approach.

The remainder of this chapter is organized as follows. First, the basic concepts of (traditional) data-flow analysis are presented in Section 8.2. This will provide the theoretical foundation and background for the discussion of the SSA-based propagation engine in Section 8. We then provide an example of a data-flow analysis that can be performed efficiently by the aforementioned engine, namely copy propagation in Section 8.4. We conclude and provide links for further reading in Section 8.5.

8.2 Preliminaries

Data-flow analysis is at the heart of many compiler transformations and optimizations, but also finds application in a broad spectrum of analysis and verification tasks in program analysis tools such as program checkers, profiling tools, and timing analysis tools. This section gives a brief introduction to the basics of data-flow analysis. Due to space considerations, we cannot cover this topic in full depth.

As noted before, data-flow analysis derives information from certain interesting program properties that may help to optimize the program. Typical examples of interesting properties are: the set of *live* variables at a given program point, the particular constant value a variable may take, or the set of program points that are reachable at run-time. Liveness information, for example, is critical during register allocation, while the two latter properties help to simplify computations and to identify dead code.

The analysis results are gathered from the input program by propagating information among its operations considering all potential execution paths. The propagation is typically performed iteratively until the computed results stabilize. Formally, a data-flow problem can be specified using a *monotone framework* that consists of:

- a *complete lattice* representing the property space,
- a *flow graph* resembling the control flow of the input program, and

- a set of *transfer functions* modeling the effect of individual operations on the property space.

Property Space: A key concept for data-flow analysis is the representation of the property space via *partially ordered sets* (L, \sqsubseteq) , where L represents some interesting program property and \sqsubseteq represents a reflexive, transitive, and anti-symmetric relation. Using the \sqsubseteq relation, *upper* and *lower bounds*, as well as *least upper* and *greatest lower bounds*, can be defined for subsets of L .

A particularly interesting class of partially ordered sets are *complete lattices*, where all subsets have a least upper bound as well as a greatest lower bound. These bounds are unique and are denoted by \sqcup and \sqcap respectively. In the context of program analysis the former is often referred to as the *join operator*, while the latter is termed the *meet operator*. Complete lattices have two distinguished elements, the *least element* and the *greatest element*, often denoted by \perp and \top respectively.

An *ascending chain* is a totally ordered subset $\{l_1, \dots, l_n\}$ of a complete lattice. A chain is said to *stabilize* if there exists an index m , where $\forall i > m: l_i = l_m$. An analogous definition can be given for *descending chains*.

Program Representation: The functions of the input program are represented as control-flow graphs, where the nodes represent operations, or instructions, and edges denote the potential flow of execution at run-time. Data-flow information is then propagated from one node to another adjacent node along the respective graph edge using *in* and *out* sets associated with every node. If there exists only one edge connecting two nodes, data can be simply copied from one set to the other. However, if a node has multiple incoming edges, the information from those edges has to be combined using the meet or join operator.

Sometimes, it is helpful to reverse the flow graph to propagate information, i.e., reverse the direction of the edges in the control-flow graph. Such analyses are termed *backward analyses*, while those using the regular flow graph are *forward analyses*.

Transfer Functions: Aside from the control flow, the operations of the program need to be accounted for during analysis. Usually these operations change the way data is propagated from one control-flow node to the other. Every operation is thus mapped to a *transfer function*, which transforms the information available from the *in* set of the operation's flow graph node and stores the result in the corresponding *out* set.

8.2.1 Solving Data-flow Problems

Putting all those elements together, i.e., a complete lattice, a flow graph, and a set of transfer functions, yields an instance of a monotone framework. This framework describes a set of *data-flow equations* whose solution will ultimately converge to the solution of the data-flow analysis. A very popular and intuitive

way to solve these equations is to compute the *maximal (minimal) fixed point* (MFP) using an iterative work list algorithm. The work list contains edges of the flow graph that have to be revisited. Visiting an edge consists of first combining the information from the *out* set of the source node with the *in* set of the target node, using the meet or join operator, then applying the transfer function of the target node. The obtained information is then propagated to all successors of the target node by appending the corresponding edges to the work list. The algorithm terminates when the data-flow information stabilizes and the work list becomes empty.

A single flow edge can be appended several times to the work list in the course of the analysis. It may even happen that an infinite feedback loop prevents the algorithm from terminating. We are thus interested in bounding the number of times a flow edge is processed. Recalling the definition of chains from before (see Section 8.2), the *height* of a lattice is defined by the length of its longest chain. We can ensure termination for lattices fulfilling the *ascending chain condition*, which ensures that the lattice has finite height. Given a lattice with finite height h and a flow graph $G = (V, E)$ it is easy to see that the MFP solution can be computed in $O(|E| \cdot h)$ time, where $|E|$ represents the number of edges. Since the number of edges is bounded by the number of graph nodes $|V|$, more precisely, $|E| \leq |V|^2$, this gives an $O(|V|^2 \cdot h)$ general algorithm to solve data-flow analyses. Note that the height of the lattice often depends on properties of the input program, which might ultimately yield bounds worse than cubic in the number of graph nodes. For instance, the lattice for copy propagation consists of the cross product of many smaller lattices, each representing the potential values of a variable occurring in the program. The total height of the lattice thus directly depends on the number of variables in the program.

In terms of memory consumption, we have to propagate data flow to all relevant program points. Nodes are required to hold informations even when these are not directly related to the node, hence, each node must store complete *in* and *out* sets.

8.3 Data-flow Propagation under SSA Form

SSA form allows us to solve a large class of data-flow problems more efficiently than the iterative work list algorithm presented previously. The basic idea is to directly propagate information computed at the unique definition of a variable to all its uses. In this way, intermediate program points that neither define nor use the variable of interest do not have to be taken into consideration, thus reducing memory consumption and compilation time.

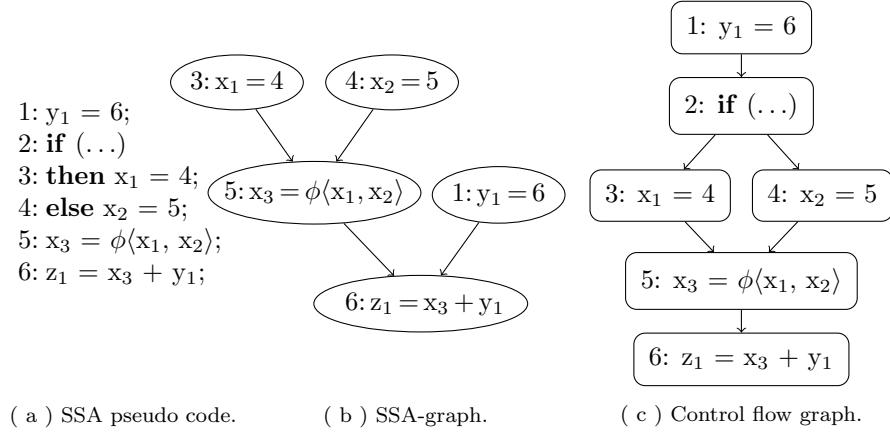


Fig. 8.1 Example program and its SSA graph.

8.3.1 Program Representation

Programs in SSA form exhibit ϕ -operations placed at join points of the original CFG. In the following, we assume that possibly many ϕ -operations are associated with the corresponding CFG nodes at those join points.

Data-flow analyses under SSA form rely on a specialized program representation based on *SSA graphs*, which resemble traditional *def-use chains* and simplify the propagation of data-flow information. The nodes of an SSA graph correspond to the operations of the program, including the ϕ -operations that are represented by dedicated nodes in the graph. The edges of the graph connect the unique definition of a variable with all its uses, i.e., edges represent true dependencies.

Besides the data dependencies, the SSA graph captures the *relevant* join points of the CFG of the program. A join point is relevant for the analysis whenever the value of two or more definitions may reach a use by passing through that join. The SSA form properties ensure that a ϕ -operation is placed at the join point and that any use of the variable that the ϕ -function defines has been properly updated to refer to the correct name.

Consider for example the code excerpt shown in Figure 8.1, along with its corresponding SSA graph and CFG. Assume we are interested in propagating information from the assignment of variable y_1 , at the beginning of the code, down to its unique use at the end. The traditional CFG representation causes the propagation to pass through several intermediate program points. These program points are concerned only with computations of the variables x_1 , x_2 , and x_3 , and are thus irrelevant for y_1 . The SSA graph representation, on the other hand, propagates the desired information directly from definition to use sites, without any intermediate step. At the same time, we also find that the control-

Algorithm 8.1: Sparse Data-flow Propagation

-
1. Initialization
 - Every edge in the CFG is marked not executable.
 - The *CFGWorkList* is seeded with the outgoing edges of the CFG's *start* node.
 - The *SSAWorkList* is empty.
 2. Remove the top element of one of the two work lists.
 3. If the element is a control flow edge, proceed as follows:
 - Mark the edge as executable.
 - Visit every ϕ -operation associated with the edge's target node (Algorithm 8.2).
 - If the target node was reached the first time via the *CFGWorkList*, visit all its operations (Algorithm 8.2).
 - If the target node has a single, non-executable outgoing edge, append that edge to the *CFGWorkList*.
 4. If the element is an edge from the SSA graph, process the target operation as follows:
 - a. When the target operation is a ϕ -operation visit that ϕ -operation.
 - b. For other operations, examine the executable flag of the incoming edges of the respective CFG node. Visit the operation if any of the edges is executable.
 5. Continue with step 2 until both work lists become empty.
-

flow join following the **if** is properly represented by the ϕ -operation defining the variable x_3 in the SSA graph.

Even though the SSA graph captures data dependencies and the relevant join points in the CFG, it lacks information on other *control dependencies*. However, analysis results can often be improved significantly by considering the additional information that is available from the control dependencies in the CFG. As an example consider once more the code of Figure 8.1, and assume that the condition associated with the **if** operation is known to be false for all possible program executions. Consequently, the ϕ -operation will select the value of x_2 in all cases, which is known to be of constant value 5. However, due to the shortcomings of the SSA graph, this information cannot be derived. It is thus important to use both the control-flow graph and the SSA graph during data-flow analysis in order to obtain the best possible results.

8.3.2 Sparse Data-flow Propagation

Similar to monotone frameworks for traditional data-flow analysis, frameworks for *sparse data-flow propagation* under SSA form can be defined using:

- a *complete lattice* representing the property space,
- a set of *transfer functions* for the operations of the program,
- a *control-flow graph* capturing the execution flow of the program, and
- an *SSA graph* representing data dependencies.

Algorithm 8.2: Visiting an Operation

-
1. Propagate data-flow information depending on the operation's kind:
 - a. ϕ -operations:
Combine the data-flow information from the node's operands where the corresponding control-flow edge is executable.
 - b. Conditional branches:
Examine the branch's condition(s) using the data-flow information of its operands.
Determine all outgoing edges of the branch's CFG node whose condition is potentially satisfied. Append the CFG edges that were non-executable to the *CFGWorkList*.
 - c. Other operations:
Update the operation's data-flow information by applying its transfer function.
 2. Whenever the data-flow information of an operation changes, append all outgoing SSA graph edges of the operation to the *SSAWorkList*.
-

We again seek a maximal (minimal) fixed point solution (MFP) using an iterative work list algorithm. However, in contrast to the algorithm described previously, data-flow information is not propagated along the edges of the control-flow graph but along the edges of the SSA graph. For regular uses the propagation is then straightforward due to the fact that every use receives its value from a unique definition. Special care has to be taken only for ϕ -operations, which select a value among their operands depending on the incoming control-flow edges. The data-flow information of the incoming operands has to be combined using the meet or join operator of the lattice. As data-flow information is propagated along SSA edges that have a single source, it is sufficient to store the data-flow information with the SSA graph node. The *in* and *out* sets used by the traditional approach—see Section 8.2—are obsolete, since ϕ -operations already provide the required buffering. In addition, the control-flow graph is used to track which operations are not reachable under any program execution and thus can be safely ignored during the computation of the fixed point solution.

The algorithm is shown in Algorithm 8.1 and processes two work lists, the *CFGWorkList*, which containing edges of the control-flow graph, and the *SSAWorkList*, which contains edges from the SSA graph. It proceeds by removing the top element of either of those lists and processing the respective edge. Throughout the main algorithm, operations of the program are visited to update the work lists and propagate information using Algorithm 8.2. We will highlight the most relevant steps of the algorithms in more detail in the following paragraphs.

The *CFGWorkList* is used to track edges of the CFG that were encountered to be *executable*, i.e., where the data-flow analysis cannot rule out that a program execution traversing the edge exists. Once the algorithm has determined that a CFG edge is executable, it will be processed by Step 3 of the main algorithm. First, all ϕ -operations of its target node need to be reevaluated due to the fact that Algorithm 8.2 discarded the respective operands of the ϕ -operations so far—because the control-flow edge was not yet marked executable. Simi-

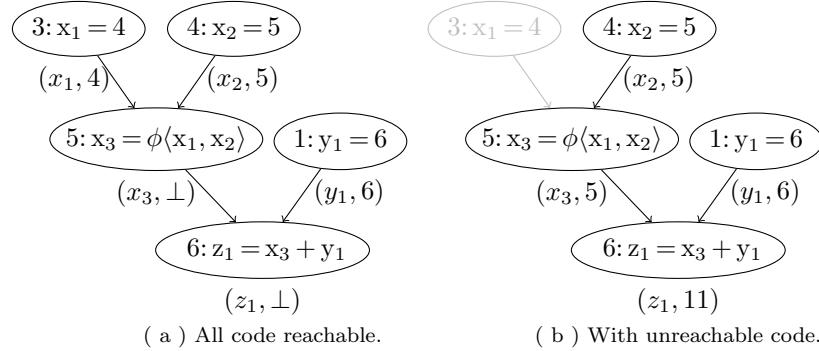


Fig. 8.2 Sparse conditional data-flow propagation using SSA graphs.

larly, the operation of the target node has to be evaluated when the target node is encountered to be executable for the first time, i.e., the currently processed control-flow edge is the first of its incoming edges that is marked executable. Note that this is only required the *first* time the node is encountered to be executable, due to the processing of operations in Step 4b, which thereafter triggers the reevaluation automatically when necessary through the SSA graph.

Regular operations as well as ϕ -operations are visited by Algorithm 8.2 when the corresponding control-flow graph node has become executable, or whenever the data-flow information of one of their predecessors in the SSA graph changed.

At ϕ -operations, the information from multiple control-flow paths is combined using the usual meet or join operator. However, only those operands where the associated control-flow edge is marked executable are considered. Conditional branches are handled by examining its conditions based on the data-flow information computed so far. Depending on whether those conditions are satisfiable or not, control-flow edges are appended to the *CFGWorkList* to ensure that all reachable operations are considered during the analysis. Finally, all regular operations are processed by applying the relevant transfer function and possibly propagating the updated information to all uses by appending the respective SSA graph edges to the *SSAWorkList*.

As an example, consider the program shown in Figure 8.1 and the constant propagation problem. First, assume that the condition of the **if** cannot be statically evaluated, we thus have to assume that all its successors in the CFG are reachable. Consequently, all control-flow edges in the program will eventually be marked executable. This will trigger the evaluation of the constant assignments to the variables x_1 , x_2 , and y_1 . The transfer functions immediately yield that the variables are all constant, holding the values 4, 5, and 6 respectively. This new information will trigger the reevaluation of the ϕ -operation of variable x_3 . As both of its incoming control-flow edges are marked executable, the combined information yields $4 \sqcap 5 = \perp$, i.e., the value is known not to be a particular constant value. Finally, also the assignment to variable z_1 is reevaluated, but

the analysis shows that its value is not a constant as depicted by Figure 8.2a. If, however, the condition of the **if** is known to be false for all possible program executions a more precise result can be computed, as shown in Figure 8.2b. Neither the control-flow edge leading to the assignment of variable x_1 nor its outgoing edge leading to the ϕ -operation of variable x_3 are marked executable. Consequently, the reevaluation of the ϕ -operation considers the data-flow information of its second operand x_2 only, which is known to be constant. This enables the analysis to show that the assignment to variable z_1 is, in fact, constant as well.

8.3.3 Discussion

During the course of the propagation algorithm, every edge of the SSA graph is processed at least once, whenever the operation corresponding to its definition is found to be executable. Afterward, an edge can be revisited several times depending on the height h of the lattice representing the property space of the analysis. On the other hand, edges of the control-flow graph are processed at most once. This leads to an upper bound in execution time of $O(|E_{SSA}| \cdot h + |E_{CFG}|)$, where E_{SSA} and E_{CFG} represent the edges of the SSA graph and the control-flow graph respectively. The size of the SSA graph increases with respect to the original non-SSA program. Measurements indicate that this growth is linear, yielding a bound that is comparable to the bound of traditional data-flow analysis. However, in practice the SSA-based propagation engine outperforms the traditional approach. This is due to the direct propagation from the definition of a variable to its uses, without the costly intermediate steps that have to be performed on the CFG. The overhead is also reduced in terms of memory consumption: instead of storing the *in* and *out* sets capturing the complete property space on every program point, it is sufficient to associate every node in the SSA graph with the data-flow information of the corresponding variable only, leading to considerable savings in practice.

8.3.4 Limitations

Unfortunately, the presented approach also has its limitations, which arise from two sources: First, the exclusive propagation of information between data-dependent operations; Second, the semantics and placement of ϕ -operations. The former issue prohibits the modeling of data-flow problems that propagate information to program points that are not directly related by either a definition or a use of a variable, while the latter prohibits the modeling of general backward problems.

Consider, for example, the problem of available expressions that often occurs in the context of redundancy elimination. An expression is available at a

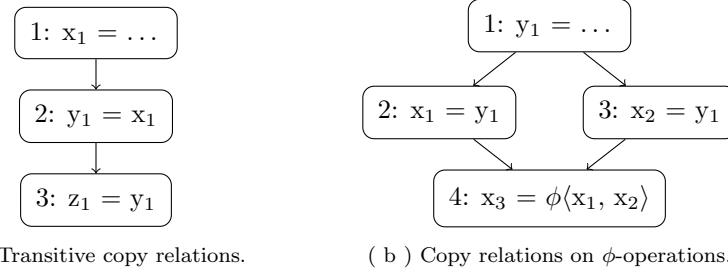


Fig. 8.3 Analysis of copy-related variables.

given program point when the expression is computed and not modified thereafter on all paths leading to that program point. In particular, this might include program points that are independent from the expression and its operands, i.e., neither define nor use any of its operands. The SSA graph does not cover those points, as it propagates information directly from definitions to uses without any intermediate steps.

Furthermore, data-flow analysis using SSA graphs is limited to forward problems. Due to the structure of the SSA graph, it is not possible to simply reverse the edges in the graph as it is done with flow graphs. For one, this would invalidate the nice property of having a single source for incoming edges of a given variable, as variables typically have more than one use. In addition, ϕ -operations are placed at join points with respect to the *forward* control flow and thus do not capture join points in the reversed control-flow graph. SSA graphs are consequently not suited to model backward problems in general. There are, however, program representations akin to the SSA format that can handle backward analyses. Chapter 14 gives an overview of such representations.

8.4 Example – Copy Propagation

Even though data-flow analysis based on SSA graphs has its limitations, it is still a useful and effective solution for interesting problems, as we will show in the following example. Copy propagation under SSA form is, in principle, very simple. Given the assignment $x = y$, all we need to do is to traverse the immediate uses of x and replace them with y , thereby effectively eliminating the original copy operation. However, such an approach will not be able to propagate copies past

ϕ -operations, particularly those in loops. A more powerful approach is to split copy propagation into two phases: Firstly, a data-flow analysis is performed to find copy-related variables throughout the program; Secondly, a rewrite phase eliminates spurious copies and renames variables.

The analysis for copy propagation can be described as the problem of propagating the *copy-of value* of variables. Given a sequence of copies as shown in Figure 8.3a, we say that y_1 is a *copy of* x_1 and z_1 is a *copy of* y_1 . The problem with this representation is that there is no apparent link from z_1 to x_1 . In order to handle transitive copy relations, all transfer functions operate on copy-of values instead of the direct source of the copy. If a variable is not found to be a copy of anything else, its copy-of value is the variable itself. For the above example, this yields that both y_1 and z_1 are copies of x_1 , which in turn is a copy of itself. The lattice of this data-flow problem is thus similar to the lattice used previously for constant propagation. The lattice elements correspond to variables of the program instead of integer numbers. The least element of the lattice represents the fact that a variable is a copy of itself.

Similarly, we would like to obtain the result that x_3 is a copy of y_1 for the example of Figure 8.3b. This is accomplished by choosing the join operator such that a copy relation is propagated whenever the copy-of values of all the operands of the ϕ -operation match. When visiting the ϕ -operation for x_3 , the analysis finds that x_1 and x_2 are both copies of y_1 and consequently propagates that x_3 is also a copy of y_1 .

The next example shows a more complex situation where copy relations are obfuscated by loops—see Figure 8.4. Note that the actual visiting order depends on the shape of the CFG and immediate uses, in other words, the ordering used here is meant for illustration only. Processing starts at the operation labeled 1, with both work lists empty and the data-flow information T associated with all variables:

- Assuming that the value assigned to variable x_1 is not a copy, the data flow information for this variable is lowered to \perp , the SSA edges leading to oper-

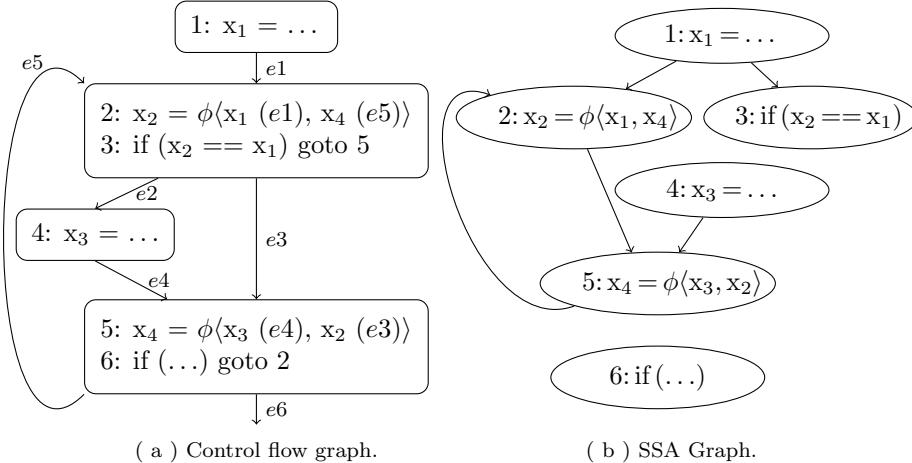


Fig. 8.4 ϕ -operations in loops often obfuscate copy relations.

- ations 2 and 3 are appended to the *SSAWorkList*, and the control-flow graph edge e_1 is appended to the *CFGWorkList*.
2. Processing the control-flow edge e_2 from the work list causes the edge to be marked executable and the operations labeled 2 and 3 to be visited. Since edge e_5 is not yet known to be executable, the processing of the ϕ -operation yields a copy relation between x_2 and x_1 . This information is utilized in order to determine which outgoing control-flow graph edges are executable for the conditional branch. Examining the condition shows that only edge e_3 is executable and thus needs to be added to the work list.
 3. Control-flow edge e_3 is processed next and marked executable for the first time. Furthermore, the ϕ -operation labeled 5 is visited. Due to the fact that edge e_4 is not known to be executable, this yields a copy relation between x_4 and x_1 (via x_2). The condition of the branch labeled 6 cannot be analyzed and thus causes its outgoing control flow edges e_5 and e_6 to be added to the work list.
 4. Now, control-flow edge e_5 is processed and marked executable. Since the target operations are already known to be executable, only the ϕ -operation is revisited. However, variables x_1 and x_4 have the same copy-of value x_1 , which is identical to the previous result computed in Step 2. Thus, neither of the two work lists is modified.
 5. Assuming that the control-flow edge e_6 leads to the exit node of the control-flow graph, the algorithm stops after processing the edge without modifications to the data-flow information computed so far.

The straightforward implementation of copy propagation would have needed multiple passes to discover that x_4 is a copy of x_1 . The iterative nature of the propagation, along with the ability to discover non-executable code, allows to handle even obfuscated copy relations. Moreover, this kind of propagation will only reevaluate the subset of operations affected by newly computed data-flow information instead of the complete control-flow graph once the set of executable operations has been discovered.

8.5 Further Reading

Traditional data-flow analysis is well established and well described in numerous papers. The book by Nielsen, Nielsen, and Hankin [207] gives an excellent introduction to the theoretical foundations and practical applications. For reducible flow graphs the order in which operations are processed by the work list algorithm can be optimized [143, 159, 207], allowing to derive tighter complexity bounds. However, relying on reducibility is problematic because the flow graphs are often *not* reducible even for proper structured languages. For instance, reversed control-flow graphs for backward problems can be—and in fact almost always are—irreducible even for programs with reducible control-flow graphs.

Furthermore, experiments have shown that the tighter bounds not necessarily lead to improved compilation times [83].

Apart from computing a fixed point (MFP) solution, traditional data-flow equations can also be solved using a more powerful approach called the *meet over all paths* (MOP) solution, which computes the *in* data-flow information for a basic block by examining *all* possible paths from the start node of the control-flow graph. Even though more powerful, computing the MOP solution is often harder or even undecidable [207]. Consequently, the MFP solution is preferred in practice.

The sparse propagation engine [210, 285], as presented in the chapter, is based on the underlying properties of the SSA form. Other intermediate representations offer similar properties. *Static Single Information* form (SSI) [249] allows both backward and forward problems to be modeled by introducing σ operations, which are placed at program points where data-flow information for backward problems needs to be merged [248]. Chapter 14 provides additional information on the use of SSI form for static program analysis. Bodík uses an extended SSA form, *e*-SSA, to eliminate array bounds checks [36]. Ruf [242] introduces the *value dependence graph*, which captures both control and data dependencies. He derives a sparse representation of the input program, which is suited for data-flow analysis, using a set of transformations and simplifications.

The *sparse evaluation graph* by Choi et al. [72] is based on the same basic idea as the approach presented in this chapter: intermediate steps are eliminated by by-passing irrelevant CFG nodes and merging the data-flow information only when necessary. Their approach is closely related to the placement of ϕ -operations and similarly relies on the dominance frontier during construction. A similar approach, presented by Johnson and Pingali [153], is based on single-entry/single-exit regions. The resulting graph is usually less sparse, but is also less complex to compute. Ramalingam [231] further extends these ideas and introduces the *compact evaluation graph*, which is constructed from the initial CFG using two basic transformations. The approach is superior to the sparse representations by Choi et al. as well as the approach presented by Johnson and Pingali.

The previous approaches derive a sparse graph suited for data-flow analysis using graph transformations applied to the CFG. Duesterwald et al. [105] instead examine the data-flow equations, eliminate redundancies, and apply simplifications to them.

CHAPTER 9

Liveness

*B. Boissinot
F. Rastello*

Progress: 30%

A horizontal progress bar consisting of a dark red segment followed by a light gray segment.

Material gathering in progress

.....
9.1 TODO

.....
9.2 SSA tells nothing about liveness

Do we really need this? I am not convinced.

Starts with a counter example of the intuition that live-range is “between” the definition and the uses.

Discussion about ϕ -function semantic (this is transverse to the technique).

.....
9.3 Computing liveness sets

Provides Appel’s (In the modern compiler book) technique to compute liveness sets under SSA.

9.4 Liveness check

CHAPTER 10

Alias analysis

???

Progress: 0%

Material gathering in progress

10.1 TODO

CHAPTER 11

Loop tree and induction variables

*S. Pop
A. Cohen*

Progress: 90%

Text and figures formating in progress

This chapter presents an extension of the SSA under which the extraction of the reducible loop tree can be done only on the SSA graph itself. This extension also captures reducible loops in the CFG. This chapter first illustrates this property then shows its usefulness through the problem of induction variable recognition.

11.1 Part of the CFG and Loop Tree can be exposed from the SSA

During the construction of the SSA representation based on a CFG representation, a large part of the CFG information is translated into the SSA representation. As the construction of the SSA has precise rules to place the phi nodes in special points of the CFG (i.e., at the merge of control-flow branches), by identifying patterns of uses and definitions, it is possible to expose a part of the CFG structure from the SSA representation.

Furthermore, it is possible to identify higher level constructs inherent to the CFG representation, such as strongly connected components of basic blocks (or reducible loops), based only on the patterns of the SSA definitions and uses. The induction variable analysis presented in this chapter is based on the detection of self references in the SSA representation and on its characterization.

This first section shows that the classical SSA representation is not enough to represent the semantics of the original program. We will see the minimal

amount of information that has to be added to the classical SSA representation in order to represent the loop information: similar to the η -function used in the Gated SSA presented in Chapter 18, the loop closed SSA form adds an extra variable at the end of a loop for each variable defined in a loop and used after the loop.

11.1.1 An SSA representation without the CFG

In the classic definition of the SSA, the CFG provides the skeleton of the program: basic blocks contain assignment statements defining SSA variable names, and the basic blocks with multiple predecessors contain ϕ -nodes. Let's look at what happens when, starting from a classic SSA representation, we remove the CFG.

In order to remove the CFG, imagine a pretty printer function that dumps only the arithmetic instructions of each basic-blocks and skips the control instructions of an imperative program by traversing the CFG structure in any order. Does the representation, obtained from this pretty printer, contains enough information to enable us to compute the same thing as the original program¹?

Let's see what happens with an example in its CFG based SSA representation:

```
bb_1 (preds = {bb_0}, succs = {bb_2})
{
    a = #some computation independent of b
}
bb_2 (preds = {bb_1}, succs = {bb_3})
{
    b = #some computation independent of a
}
bb_3 (preds = {bb_2}, succs = {bb_4})
{
    c = a + b;
}
bb_4 (preds = {bb_3}, succs = {bb_5})
{
    return c;
}
```

after removing the CFG structure, listing the definitions in an arbitrary order, we could obtain this:

```
return c;
b = #some computation independent of a
```

¹ to simplify the discussion, we consider the original program to be free of side effect instructions

```
c = a + b;
a = #some computation independent of b
```

and this SSA code is enough, in the absence of side effects, to recover an order of computation that leads to the same result as in the original program. For example, the evaluation of this sequence of statements would produce the same result:

```
b = #some computation independent of a
a = #some computation independent of b
c = a + b;
return c;
```

11.1.2 Natural loop structures on the SSA

We will now see how to represent the natural loops in the SSA form by systematically adding extra ϕ -nodes at the end of loops, together with extra information about the loop exit predicate.

Supposing that the original program contains a loop:

```
bb_1 (preds = {bb_0}, succs = {bb_2})
{
    x = 3;
}
bb_2 (preds = {bb_1, bb_3}, succs = {bb_3, bb_4})
{
    i = phi (x, j)
    if (i < N) goto bb_3 else goto bb_4;
}
bb_3 (preds = {bb_2}, succs = {bb_3})
{
    j = i + 1;
}
bb_4 (preds = {bb_2}, succs = {bb_5})
{
    k = phi (i)
}
bb_5 (preds = {bb_4}, succs = {bb_6})
{
    return k;
}
```

Pretty printing, with a random order traversal, we could obtain this SSA code:

```
x = 3;
return k;
```

```
i = phi (x, j)
k = phi (i)
j = i + 1;
```

We can remark that some information is lost in this pretty printing: the exit condition of the loop have been lost. We will have to record this information in the extension of the SSA representation. However, the loop structure still appears through the cyclic definition of the induction variable i . To expose it, we can rewrite this SSA code using simple substitutions, as:

```
i = phi (3, i + 1)
k = phi (i)
return k;
```

Thus, we have the definition of the SSA name “ i ” defined in function of itself. This pattern is characteristic of the existence of a loop. We can remark that there are two kinds of ϕ -nodes used in this example:

- **loop- ϕ nodes** “ $i = \phi(x, j)$ ” have an argument that contains a self reference j and an invariant argument x : here the defining expression “ $j = i + 1$ ” contains a reference to the same loop- ϕ definition i , while x (here 3) is not part of the circuit of dependencies that involves i and j . Note that it is possible to define a canonical SSA form by limiting the number of arguments of loop- ϕ nodes to two.
- **close- ϕ nodes** “ $k = \phi(i)$ ” capture the last value of a name defined in a loop. Names defined in a loop can only be used within that loop or in the arguments of a close- ϕ node (that is “closing” the set of uses of the names defined in that loop). In a canonical SSA form it is possible to limit the number of arguments of close- ϕ nodes to one.

11.1.3 Improving the SSA pretty printer for loops

As we have seen in the above example, the exit condition of the loop disappeared during the basic pretty printing of the SSA. To capture the semantics of the computation of the loop, we have to specify in the close- ϕ node, when we exit the loop so as to be able to derive which value will be available in the end of the loop. With our extension that adds the loop exit condition to the syntax of the close- ϕ , the SSA pretty printing of the above example would be:

```
x = 3;
i = loop-phi (x, j)
j = i + 1;
k = close-phi (i >= N, i)
return k;
```

So k is defined as the “first value” of i satisfying the loop exit condition, “ $i \geq N$ ”. In the case of finite loops, this is well defined as being the first element satisfying

the loop exit condition of the sequence defined by the corresponding loop- ϕ node.

In the following, we will look at an algorithm that translates the SSA representation into a representation of polynomial functions, describing the sequence of values that SSA names take during the execution of a loop. The algorithm is restricted to the loops that are reducible. All such loops are labeled. Note that irreducible control flow is not forbidden: only the loops carrying self-definitions must be reducible.

11.2 Analysis of Induction Variables

The purpose of the induction variables analysis is to provide a characterization of the sequences of values taken by a variable during the execution of a loop. This characterization can be an exact function of the canonical induction variable of the loop (i.e., a loop counter that starts at zero with a step of one for each iteration of the loop) or an approximation of the values taken during the execution of the loop represented by values in an abstract domain. In this section, we will see a possible characterization of induction variables in terms of sequences. The domain of sequences will be represented by *chains of recurrences*: as an example, a canonical induction variable with an initial value 0 and a stride 1 that would occur in the loop with label x will be represented by the chain of recurrence $\{0, +, 1\}_x$.

11.2.1 Stride detection

The first phase of the induction variables analysis is the detection of the strongly connected components of the SSA. This can be performed by traversing the use-def SSA chains and detecting that some definitions are visited twice. For a self referring use-def chain, it is possible to derive the step of the corresponding induction variable as the overall effect of one iteration of the loop on the value of the loop- ϕ node. When the step of an induction variable depends on another cyclic definition, one has to further analyze the inner cycle. The analysis of the induction variable ends when all the inner cyclic definitions used for the computation of the step are analyzed. Note that it is possible to construct SSA graphs with strongly connected components that are impossible to characterize with the chains of recurrences. This is precisely the case of the following example that shows two inter-dependent circuits, the first that involves a and b with step $c + 2$, and the second that involves c and d with step $a + 2$. This leads to an endless loop, which must be detected.

```
a = loop-phi-x (0, b)
```

```
c = loop-phi-x (1, d)
b = c + 2
d = a + 3
```

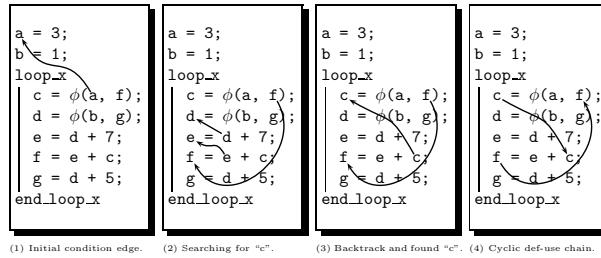


Fig. 11.1 Detection of the cyclic definition using a depth first search traversal of the use-def chains.

Let's now look at an example, presented in Figure 11.1, to see how the stride detection works. The arguments of a ϕ -node are analyzed to determine whether they contain self references or if they are pointing towards the initial value of the induction variable. In this example, (1) represents the use-def edge that points towards the invariant definition. When the argument to be analyzed points towards a longer use-def chain, the full chain is traversed, as shown in (2), until a phi node is reached. In this example, the phi node that is reached in (2) is different to the phi node from which the analysis started, and so in (3) a search starts over the uses that have not yet been analyzed. When the original phi node is found, as in (3), the cyclic def-use chain provides the step of the induction variable: in this example, the step is “+ e”. Knowing the symbolic expression for the step of the induction variable may not be enough, as we will see next, one has to instantiate all the symbols (“e” in the current example) defined in the varying loop to precisely characterize the induction variable.

11.2.2 Translation to chains of recurrences

Once the def-use circuit and its corresponding overall loop update expression have been identified, it is possible to translate the sequence of values of the induction variable to a chain of recurrence. The syntax of a polynomial chain of recurrence is: $\{base, +, step\}_x$, where *base* and *step* may be arbitrary expressions or constants, and *x* is the loop to which the sequence is associated. As a chain of recurrence represents the sequence of values taken by a variable during the execution of a loop, the associated expression of a chain of recurrence is given

by $\{base, +, step\}_x(\ell_x) = base + step \times \ell_x$, that is a function of ℓ_x , the number of times the body of loop x has been executed.

When *base* or *step* translates to sequences varying in outer loops, the resulting sequence is represented by a multivariate chain of recurrences. For example $\{\{0, +, 1\}_x, +, 2\}_y$ defines a multivariate chain of recurrence with a step of 1 in loop x and a step of 2 in loop y , where loop y is enclosed in loop x .

When *step* translates into a sequence varying in the same loop, the chain of recurrence represents a polynomial of a higher degree. For example, $\{3, +, \{8, +, 5\}_x\}_x$ represents a polynomial evolution of degree 2 in loop x . In this case, the chain of recurrence is also written omitting the extra braces: $\{3, +, 8, +, 5\}_x$. The semantics of a chain of recurrences is defined using the binomial coefficient $\binom{n}{p} = \frac{n!}{p!(n-p)!}$, by the equation:

$$\{c_0, +, c_1, +, c_2, +, \dots, +, c_n\}_x(\ell_x) = \sum_{p=0}^n c_p \binom{\ell_x}{p}.$$

with ℓ the iteration domain vector (the iteration loop counters for all the loops in which the chain of recurrence variates), and ℓ_x the iteration counter of loop x . This semantics is very useful in the analysis of induction variables, as it makes it possible to split the analysis into two phases, with a symbolic representation as a partial intermediate result:

1. first, the analysis leads to an expression, where the step part “s” is left in a symbolic form, i.e., $\{c_0, +, s\}_x$;
2. then, by instantiating the step, i.e., $s = \{c_1, +, c_2\}_x$, the chain of recurrence is that of a higher degree polynomial, i.e., $\{c_0, +, \{c_1, +, c_2\}_x\}_x = \{c_0, +, c_1, +, c_2\}_x$.

11.2.3 Instantiation of symbols and region parameters

The last phase of the induction variable analysis consists in the instantiation (or further analysis) of symbolic expressions left from the previous phase. This includes the analysis of induction variables in outer loops, computing the last value of the counter of a preceding loop, and the propagation of closed form expressions for loop invariants defined earlier. In some cases, it becomes necessary to leave in a symbolic form every definition outside a given region, and these symbols are then called parameters of the region.

Let us look again at the example of Figure 11.1 to see how the sequence of values of the induction variable c is characterized with the chains of recurrences notation. The first step, after the cyclic definition is detected, is the translation of this information into a chain of recurrence: in this example, the initial value (or base of the induction variable) is a and the step is e , and so c is represented by a chain of recurrence $\{a, +, e\}_1$ that is varying in loop number 1. The symbols are then instantiated: a is trivially replaced by its definition leading to $\{3, +, e\}_1$. The

analysis of e leads to this chain of recurrence: $\{8, +, 5\}_1$ that is then used in the chain of recurrence of c , $\{3, +, \{8, +, 5\}_1\}_1$ and that is equivalent to $\{3, +, 8, +, 5\}_1$, a polynomial of degree two:

$$\begin{aligned} F(\ell) &= 3\binom{\ell}{0} + 8\binom{\ell}{1} + 5\binom{\ell}{2} \\ &= \frac{5}{2}\ell^2 + \frac{11}{2}\ell + 3. \end{aligned}$$

11.2.4 Number of iterations and computation of the end of loop value

One of the important static analyses for loops is to evaluate their trip count, i.e., the number of times the loop body is executed before the exit condition becomes true. In common cases, the loop exit condition is a comparison of an induction variable against some constant, parameter, or another induction variable. The number of iterations is then computed as the minimum solution of a polynomial inequality with integer solutions, also called a Diophantine inequality. When one or more coefficients of the Diophantine inequality are parameters, the solution is left under a parametric form. The number of iterations can also be an expression varying in an outer loop, in which case, it can be characterized using a chain of recurrence.

Consider a scalar variable varying in an outer loop with strides dependent on the value computed in an inner loop. The expression representing the number of iterations in the inner loop can then be used to express the evolution function of the scalar variable varying in the outer loop.

For example, the following code

```
x = 0;
for (i = 0; i < N; i++)    // loop_1
    for (j = 0; j < M; j++) // loop_2
        x = x + 1;
```

would be written in loop closed SSA form as:

$$\begin{aligned} x_0 &= 0 \\ i &= \text{loop-}\phi_1(0, i+1) \\ x_1 &= \text{loop-}\phi_1(x_0, x_2) \\ x_4 &= \text{close-}\phi_1(i < N, x_1) \\ j &= \text{loop-}\phi_2(0, j+1) \\ x_3 &= \text{loop-}\phi_2(x_1, x_3 + 1) \\ x_2 &= \text{close-}\phi_2(j < M, x_3) \end{aligned}$$

x_3 represents the value of variable x at the end of the original imperative program. The analysis of scalar evolutions for variable x_4 would trigger the analysis

of scalar evolutions for all the other variables defined in the loop closed SSA form as follows:

- first, the analysis of variable x_4 would trigger the analysis of i , N and x_1
 - the analysis of i leads to $i = \{0, +, 1\}_1$ i.e., the canonical loop counter l_1 of $loop_1$.
 - N is a parameter and is left under its symbolic form
 - the analysis of x_1 triggers the analysis of x_0 and x_2
 - the analysis of x_0 leads to $x_0 = 0$
 - analyzing x_2 triggers the analysis of j , M , and x_3
 - $j = \{0, +, 1\}_2$ i.e., the canonical loop counter l_2 of $loop_2$.
 - M is a parameter
 - $x_3 = \text{loop-}\phi_2(x_1, x_3 + 1) = \{x_1, +, 1\}_2$
 - $x_2 = \text{close-}\phi_2(j < M, x_3)$ is then computed as the last value of x_3 after $loop_2$, i.e., it is the chain of recurrence of x_3 applied to the first iteration of $loop_2$ that does not satisfy $j < M$ or equivalently $l_2 < M$. The corresponding Diophantine inequality $l_2 \geq M$ have minimum solution $l_2 = M$. So, to finish the computation of the scalar evolution of x_2 we apply M to the scalar evolution of x_3 , leading to $x_2 = \{x_1, +, 1\}_2(M) = x_1 + M$;
 - the scalar evolution analysis of x_1 then leads to $x_1 = \text{loop-}\phi_1(x_0, x_2) = \text{loop-}\phi_1(x_0, x_1 + M) = \{x_0, +, M\}_1 = \{0, +, M\}_1$
 - and finally the analysis of x_4 ends with $x_4 = \text{close-}\phi_1(i < N, x_1) = \{0, +, M\}_1(N) = M \times N$.

.....

11.3 Further reading

Induction variable detection has been studied extensively in the past because of its central role in loop optimizations. Wolfe [?] designed the first SSA-based induction variable recognition technique. It abstracts the SSA graph and classifies inductions according to a wide spectrum of patterns.

When operating on a low-level intermediate representation with arbitrary gotos, detecting the natural loops is the first step in the analysis of induction variables. In general, and when operating on low-level code in particular, it is preferable to use analyses that are more robust to complex control flow that do not resort to an early classification into predefined patterns. Chains of recurrences [23, 166, 300] have been proposed to characterize the sequence of values taken by a variable during the execution of a loop [?], and it has proven to be more robust to the presence of complex, unstructured control flow, to the characterization of induction variables over modulo-arithmetic such as unsigned wrap-around types in C, and to the implementation in a production compiler [?].

The formalism and presentation of this chapter is derived from the thesis work of Sebastian Pop. The manuscript [?] contains pseudo-code and links to the implementation of scalar evolutions in GCC since version 4.0. The same approach has also influenced the design of LLVM's scalar evolution, but the implementation is different. Parler des types, arith modulo, traitement dans GCC et LLVM, citer comme difficile.

Induction variable analysis is used in dependence tests for scheduling and parallelization [?], and more recently, the extraction of short-vector to SIMD instructions [?].² The Omega test [?] and parametric integer linear programming [?] have typically used to reason about system parametric affine Diophantine inequalities. But in many cases, simplifications and approximations can lead to polynomial decision procedures [?]. Modern parallelizing compilers tend to implement both kinds, depending on the context and aggressiveness of the optimization.

Substituting an induction variable with a closed form expression is also useful to the removal of the cyclic dependences associated with the computation of the induction variable itself [?]. Other applications include enhancements to strength reduction and loop-invariant code motion [?], induction variable canonicalization (reducing induction variables to a single one in a given loop) [?].

The number of iterations of loops can also be computed based on the characterization of induction variables. This information is essential to advanced loop analyses such as value-range propagation propagation [?], and enhance dependence tests for scheduling and parallelization [?, ?]. It also enables more opportunities for scalar optimization when the induction variable is used after its defining loop. Loop transformations also benefit from the replacement of the end of loop value as this removes scalar dependencies between consecutive loops. Another interesting use of the end of loop value is the estimation of the worst case execution time (WCET) where one tries to obtain an upper bound approximation of the time necessary for a program to terminate.

² Note however that the computation of closed form expressions is not required for dependence testing itself [?].

CHAPTER 12

Redundancy Elimination

F. Chow

Progress: 80%

A horizontal progress bar consisting of a dark green segment followed by a light gray segment.

Minor polishing in progress

12.1 Introduction

Redundancy elimination is an important category of optimizations performed by modern optimizing compilers. In the course of program execution, certain computations may be repeated multiple times that yield the same results. Such redundant computations can be eliminated by saving the results of the earlier computations and reusing instead of recomputing them later.

There are two types of redundancies: *full* redundancy and *partial* redundancy. A computation is fully redundant if the computation has occurred earlier regardless of the flow of control. The elimination of full redundancy is also called common subexpression elimination. A computation is partially redundant if the computation has occurred only along certain paths. Full redundancy can be regarded as a special case of partial redundancy where the redundant computation occurs regardless of the path taken.

There are two different views for a computation related to redundancy: how it is computed and the computed value. The former relates to the operator and the operands it operates on, which translates to how it is represented in the program representation. The latter refers to the value generated by the computation in the static sense¹. As a result, algorithms for finding and eliminating redundances can be classified into being syntax-driven or being value-driven. In

¹ All values referred to in this Chapter are static values viewed with respect to the program code. A static value can map to different dynamic values during program execution.

syntax-driven analyses, two computations are the same if they are the same operation applied to the same operands that are program variables or constants. In this case, redundancy can arise only if the variables' values have not changed between the occurrences of the computation. In value-based analyses, redundancy arises whenever two computations yield the same value. For example, $a + b$ and $a + c$ compute the same result if b and c can be determined to hold the same value. In this chapter, we deal mostly with syntax-driven redundancy elimination. The last section will extend our discussion to value-based redundancy elimination.

In our discussion on syntax-driven redundancy elimination, our algorithm will focus on the optimization of a lexically identical expression, like $a + b$, that appears in the program. During compilation, the compiler will repeat the redundancy elimination algorithm on all the other lexically identified expressions in the program.

The style of the program representation can impact the effectiveness of the algorithm applied. We distinguish between statements and expressions. Expressions computes to values without generating any side effect. Statements have side effects by potentially altering memory contents or control flow, and are not candidates for redundancy elimination. In dealing with lexically identified expressions, we advocate a maximal expression tree form of program representation. In this style, a large expression tree like $a + b * c - d$ are represented as is, without having to specify any assignments to temporaries for storing the intermediate values of the computation². We also assume the Conventional SSA Form of program representation, in which each ϕ -congruence class is interference-free and the live ranges of the SSA versions of each variable do not overlap. We further assume the HSSA form that completely models the aliasing in the program [reference to HSSA chapter].

12.2 Why PRE and SSA are related

Figure 12.1 shows the two most basic forms of partial redundancy. In Figure 12.1(a), $a + b$ is redundant when the right path is taken. In Figure 12.1(b), $a + b$ is redundant whenever the branch-back edge of the loop is taken. Both are examples of *strictly* partial redundancies, in which insertions are required to eliminate the redundancies. In contrast, a full redundancy can be deleted without requiring any insertion. Partial redundancy elimination (PRE) is powerful because it subsumes global common subexpressions and loop-invariant code motion.

We can visualize the impact on redundancies of a single computation as shown in Figure 12.2. In the region of the control-flow graph dominated by the occurrence of $a + b$, any further occurrence of $a + b$ is fully redundant, assum-

² The opposite of maximal expression tree form is the triplet form in which each arithmetic operation always defines a temporary.

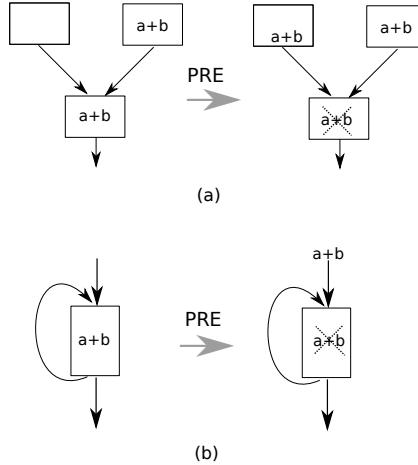


Fig. 12.1 Two basic examples of partial redundancy elimination.

ing a and b are not modified. Following the program flow, once we are past the dominance frontiers, any further occurrence of $a + b$ is partially redundant. In constructing SSA form, dominance frontiers are where ϕ 's are inserted. Since partial redundancies start at dominance frontiers, it must be related to SSA's ϕ 's. In fact, the same sparse approach to modeling the use-def relationships among the occurrences of a program variable can be used to model the redundancy relationships among the different occurrences of $a + b$.

The SSAPRE algorithm that we present performs PRE efficiently by taking advantage of the use-def information inherent in its input conventional SSA form. If an occurrence $a_j + b_j$ is redundant with respect to $a_i + b_i$, SSAPRE builds a redundancy edge that connects $a_i + b_i$ to $a_j + b_j$. To expose potential partial redundancies, we introduce the operator Φ at the dominance frontiers of the occurrences, which has the effect of factoring the redundancy edges at merge points in the control-flow graph.³ The resulting *factored redundancy graph* (FRG) can be regarded as the SSA form for expressions.

To make the expression SSA form more intuitive, we introduce the hypothetical temporary h , which can be thought of as the temporary that will be used to store the value of the expression for reuse in order to suppress redundant computations. The FRG can be viewed as the SSA graph for h which is not precise, because we have not yet determined where h should be defined or used. In referring to the FRG, a *use* node will refer to a node in the FRG that is not a definition.

³ Adhering to SSAPRE's convention, we use lower case ϕ 's in the SSA form of variables and upper case Φ 's in the SSA form for expressions.

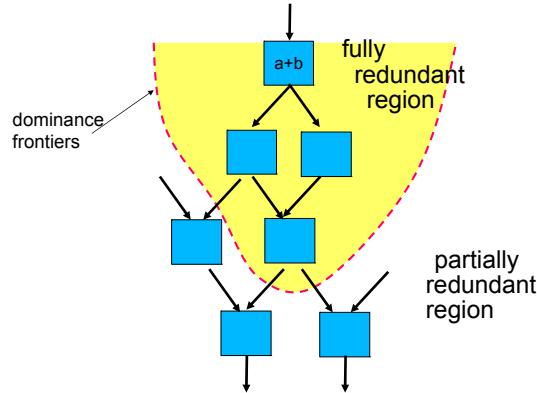


Fig. 12.2 Dominance frontiers are boundaries between fully and partially redundant regions.

The SSA form for h is constructed in two steps similar to ordinary SSA form: the Φ -Insertion step followed by the Renaming step. In the Φ -Insertion step, we insert Φ 's at the dominance frontiers of all the expression occurrences, to ensure that we do not miss any possible placement positions for the purpose of PRE, as in Figure 12.3(a). We also insert Φ 's caused by expression alteration. Such Φ 's are triggered by the occurrence of ϕ 's for any of the operands in the expression. In Figure 12.3(b), the Φ at block 3 is caused by the ϕ for a in the same block, which in turns reflects the assignment to a in block 2.

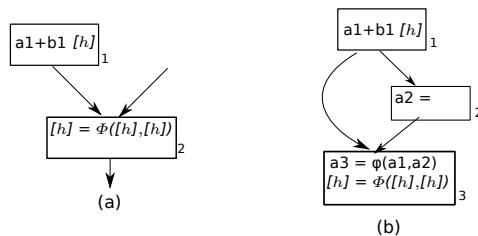


Fig. 12.3 Examples of Φ insertion

The Renaming step assigns SSA versions to h such that occurrences renamed to identical h -versions will compute to the same values. We conduct a pre-order traversal of the dominator tree similar to the renaming step in SSA construction for variables, but with the following modifications. In addition to a renaming stack for each variable, we maintain a renaming stack for the expression. Entries

on the expression stack are popped as our dominator tree traversal backtracks past the blocks where the expression originally received the version. Maintaining the variable and expression stacks together allows us to decide efficiently whether two occurrences of an expression should be given the same h -version.

There are three kinds of occurrences of the expression in the program: (a) the occurrences in the original program, which we call *real* occurrences; (b) the Φ 's inserted in the Φ -Insertion step; and (c) Φ operands, which are regarded as occurring at the ends of the predecessor blocks of their corresponding edges. During the visitation in Renaming, a Φ 's is always given a new version. For a non- Φ , i.e., cases (a) and (c), we check the current version of every variable in the expression (the version on the top of each variable's renaming stack) against the version of the corresponding variable in the occurrence on the top of the expression's renaming stack. If all the variable versions match, we assign it the same version as the top of the expression's renaming stack. If any of the variable versions does not match, for case (a), we assign it a new version, as in the example of Figure 12.4(a); for case (c), we assign the special class \perp to the Φ operand to denote that the value of the expression is unavailable at that point, as in the example of Figure 12.4(b). If a new version is assigned, we push the version on the expression stack.

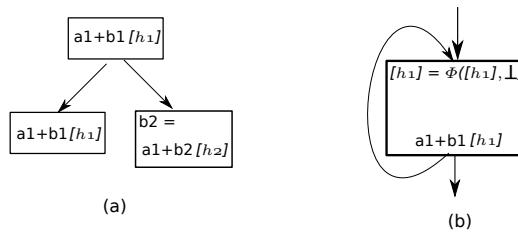


Fig. 12.4 Examples of expression renaming

The FRG captures all the redundancies of $a + b$ in the program. In fact, it contains just the right amount of information for determining the optimal code placement. Because strictly partial redundancies can only occur at the Φ nodes, insertions for PRE only need to be considered at the Φ 's.

12.3 How SSAPRE Works

Referring to the expression being optimized as X , we use the term *placement* to denote the set of points in the *optimized* program where X 's computation occurs. In contrast, *original computation points* refer to the points in the *original* program where X 's computation took place. The *original* program will be transformed to the *optimized* program by performing a set of *insertions* and *deletions*.

The objective of SSAPRE is to find a placement that satisfies the following four criteria in this order:

1. Correctness — X is fully available at all the original computation points.
2. Safety — There is no insertion of X on any path that did not originally contain X .
3. Computational optimality — No other safe and correct placement can result in fewer computations of X on any path from entry to exit in the program.
4. Lifetime optimality — Subject to the computational optimality, the life range of the temporary introduced to store X is minimized.

Each occurrence of X at its original computation point can be qualified with exactly one of the following attributes:

- fully redundant
- strictly partially redundant (SPR)
- non-redundant

As a code placement problem, SSAPRE follows the same two-step process used in all PRE algorithms. The first step determines the best set of insertion points that render as many SPR occurrences fully redundant as possible. The second step deletes fully redundant computations taking into account the effects of the inserted computations. Since the second full redundancy elimination step is trivial and well understood, the challenge lies in the first step for coming up with the best set of insertion points. The first step will tackle the safety, computational optimality and lifetime optimality criteria, while the correctness criterion is delegated to the second step. For the rest of this section, we only focus on the first step for finding the best insertion points, which is driven by the SPR occurrences.

We assume that all *critical edges* in the control-flow graph have been removed by inserting empty basic blocks at such edges⁴. In the SSAPRE approach, insertions are only performed at Φ operands. When we say a Φ is a candidate for insertion, it means we will consider inserting at its operands to render X available at the entry to the basic block containing that Φ . An insertion at a Φ operand means inserting X at the incoming edge corresponding to that Φ operand. In reality, the actual insertion is done at the end of the predecessor block.

12.3.1 The Safety Criterion

As we have pointed out at the end of Section 12.2, insertions only need to be considered at the Φ 's. The safety criterion implies that we should only insert at Φ 's where X is *downsafe* (fully anticipated). Thus, we perform data-flow analysis on the FRG to determine the *downsafe* attribute for Φ 's. Data-flow analysis can

⁴ A critical edge is one whose tail block has multiple successors and whose head block has multiple predecessors.

be performed with linear complexity on SSA graphs, which we illustrate with the Downsafety computation.

A Φ is not *downsafe* if there is a control-flow path from that Φ along which the expression is not computed before program exit or before being altered by redefinition of one of its variables. Except for loops with no exit, this can happen only due to one of the following cases: (a) there is a path to exit or an alteration of the expression along which the Φ result version is not used; or (b) the Φ result version appears as the operand of another Φ that is not *downsafe*. Case (a) represents the initialization for our backward propagation of $\neg \text{downsafe}$; all other Φ 's are initially marked *downsafe*. The Downsafety propagation is based on case (b). Since a real occurrence of the expression blocks the case (b) propagation, we define a *has_real_use* flag attached to each Φ operand and set this flag to true when the Φ operand is defined by another Φ and the path from its defining Φ to its appearance as a Φ operand crosses a real occurrence. The propagation of $\neg \text{downsafe}$ is blocked whenever the *has_real_use* flag is true. Figure 12.5 gives the DownSafety propagation algorithm. The initialization of the *has_real_use* flags is performed in the earlier Renaming phase.

```

Procedure Reset_downsafe( $X$ ) {
1:   if  $\text{def}(X)$  is not a  $\Phi$ 
2:     return
3:    $f \leftarrow \text{def}(X)$ 
4:   if (not  $\text{downsafe}(f)$ )
5:     return
6:    $\text{downsafe}(f) \leftarrow \text{false}$ 
7:   for each operand  $\omega$  of  $f$  do
8:     if (not  $\text{has\_real\_use}(\omega)$ )
9:       Reset_downsafe( $\omega$ )
}
Procedure DownSafety {
10:  for each  $f \in \{\Phi\text{'s in the program}\}$  do
11:    if  $\exists$  path  $P$  to program exit or alteration of expression along which  $f$  is not used
12:       $\text{downsafe}(f) \leftarrow \text{false}$ 
13:    for each  $f \in \{\Phi\text{'s in the program}\}$  do
14:      if (not  $\text{downsafe}(f)$ )
15:        for each operand  $\omega$  of  $f$  do
16:          if (not  $\text{has\_real\_use}(\omega)$ )
17:            Reset_downsafe( $\omega$ )
}

```

Fig. 12.5 Algorithm for DownSafety

12.3.2 The Computational Optimality Criterion

At this point, we have eliminated the unsafe Φ 's based on the safety criterion. Next, we want to identify all the Φ 's that are possible candidates for insertion by disqualifying Φ 's that cannot be insertion candidates in any computationally optimal placement. An unsafe Φ can still be an insertion candidate if the expression is fully available there, though the inserted computation will itself be fully redundant. We define the *can_be_avail* attribute for the current step, whose purpose is to identify the region where, after appropriate insertions, the computation can become fully available. A Φ is $\neg\text{can_be_avail}$ if and only if inserting there violates computational optimality. The *can_be_avail* attribute can be viewed as:

$$\text{can_be_avail}(\Phi) = \text{downsafe}(\Phi) \cup \text{avail}(\Phi)$$

We could compute the *avail* attribute separately using the full availability analysis, which involves propagation in the forward direction with respect to the control-flow graph. But this would have performed some useless computation because we do not need to know its values within the region where the Φ 's are *downsafe*. Thus, we choose to compute *can_be_avail* directly by initializing a Φ to be $\neg\text{can_be_avail}$ if the Φ is not *downsafe* and one of its operands is \perp . In the propagation phase, we propagate $\neg\text{can_be_avail}$ forward when a $\neg\text{downsafe}$ Φ has an operand that is defined by a $\neg\text{can_be_avail}$ Φ and that operand is not marked *has_real_use*.

After *can_be_avail* has been computed, it is possible to perform insertions at all the *can_be_avail* Φ 's. There would be full redundancies created among the insertions themselves, but they would not affect computational optimality because the subsequent full redundancy elimination step will remove any fully redundant inserted or non-inserted computation, leaving the earliest computations as the optimal code placement.

12.3.3 The Lifetime Optimality Criterion

To fulfill lifetime optimality, we perform a second forward propagation called *Later* that is derived from the well-understood partial availability analysis. The purpose is to disqualify *can_be_avail* Φ 's where the computation is partially available based on the original occurrences of X . A Φ is marked *later* if it is not necessary to insert there because a later insertion is possible. In other words, there exists a computationally optimal placement under which X is not available immediately after the Φ . We optimistically regard all the *can_be_avail* Φ 's to be *later*, except the following cases: (a) the Φ has an operand defined by a real computation; or (b) the Φ has an operand that is *can_be_avail* Φ marked not *later*. Case (a) represents the initialization for our forward propagation of

not *later*; all other *can_be_avail* Φ 's are marked *later*. The Later propagation is based on case (b).

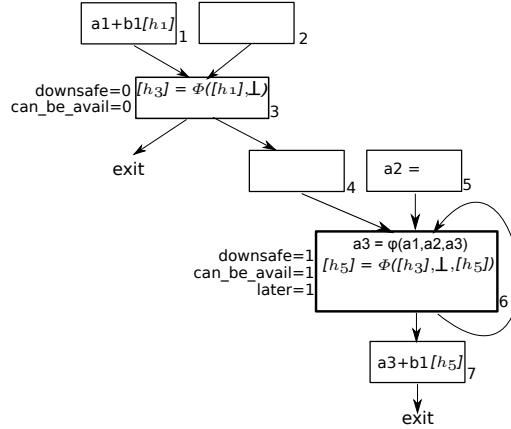


Fig. 12.6 Example to show the need of the *Later* attribute

The final criterion for performing insertion is to insert at the Φ 's where *can_be_avail* and $\neg \text{later}$ hold. We call such Φ 's *will_be_avail*. At these Φ 's, insertion is performed at each operand that satisfies either of the following condition:

1. it is \perp ; or
2. *has_real_use* is false and it is defined by a $\neg \text{will_be_avail}$ Φ .

We illustrate our discussion in this section with the example of Figure 12.6, where the program exhibits partial redundancy that cannot be removed by safe code motion. The two Φ 's with their computed data-flow attributes are as shown. If insertions were based on *can_be_avail*, $a + b$ would have been inserted at the exits of blocks 4 and 5 due to the Φ in block 6, which would have resulted in unnecessary code motion that increases register pressure. By considering *later*, no insertion is performed, which is optimal under safe PRE for this example.

12.4 Speculative PRE

If we ignore the safety requirement of PRE discussed in Section 12.3, the resulting code motion will involve speculation. Speculative code motion suppresses redundancy in some path at the expense of another path where the computation is added but result is unused. As long as the paths that are burdened with more computations are executed less frequently than the paths where the redundant computations are avoided, a net gain in program performance can be

achieved. Thus, speculative code motion should only be performed when there are clues about the relative execution frequencies of the paths involved.

Without profile data, speculative PRE can be conservatively performed by restricting it to loop-invariant computations. Figure 12.7 shows a loop-invariant computation $a + b$ that occurs in a branch inside the loop. This loop-invariant code motion is speculative because, depending on the branch condition inside the loop, it may be executed zero time, while moving it to the loop header causes it to execute one time. This speculative loop-invariant code motion is profitable unless the path inside the loop containing the expression is never taken, which is usually not the case. When performing SSAPRE, marking Φ 's located at the start of loop bodies downsafe will effect speculative loop invariant code motion.

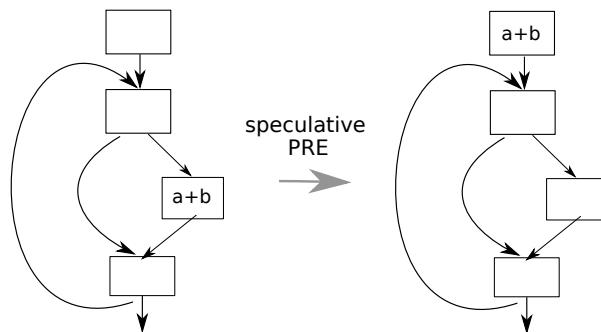


Fig. 12.7 Speculative loop-invariant code motion

Computations like indirect loads and divides are called *dangerous* computations because they may fault. Dangerous computations in general should not be speculated. As an example, if we replace the expression $a + b$ in Figure 12.7 by a/b and the speculative code motion is performed, it may cause a run-time divide-by-zero fault after the speculation because b can be 0 at the loop header while it is never 0 in the branch that contains a/b inside the loop body.

Dangerous computations are sometimes protected by tests (or guards) placed in the code by the programmers or automatically generated by language compilers like those for Java. When such a test occurs in the program, we say the dangerous computation is *safety-dependent* on the control-flow point that establishes its safety. At the points in the program where its safety dependence is satisfied, the dangerous instruction is *fault-safe* and can still be speculated.

We can represent safety dependences as value dependences in the form of abstract *tau* variables. Each run-time test that succeeds defines a tau variable on its fall-through path. During SSAPRE, we attach these tau variables as additional operands to the dangerous computations related to the test. The tau variables are also put into SSA form, so their definitions can be found by following

the use-def edges. The definitions of the tau variables have abstract right-hand-side values that are not allowed to be involved in any optimization. Because they are abstract, they are also omitted in the generated code after the SSAPRE phase. A dangerous computation can be defined to have more than one tau operands, depending on its semantics. When all its tau operands have definitions, it means the computation is fault-safe; otherwise, it is unsafe to speculate. By including the tau operands into consideration, speculative PRE automatically honors the fault-safety of dangerous computations when it performs speculative code motion.

In Figure 12.8, the program contains a non-zero test for b . We define an additional tau operand for the divide operation in a/b in SSAPRE to provide the information whether a non-zero test for b is available. At the start of the region guarded by the non-zero test for b , the compiler inserts the definition of τau_1 with the abstract right-hand-side value \tauauedge . Any appearance of a/b in the region guarded by the non-zero test for b will have τau_1 as its tau operand. Having a defined tau operand allows a/b to be freely speculated in the region guarded by the non-zero test, while the definition of τau_1 prevents any hoisting of a/b past the non-zero test.

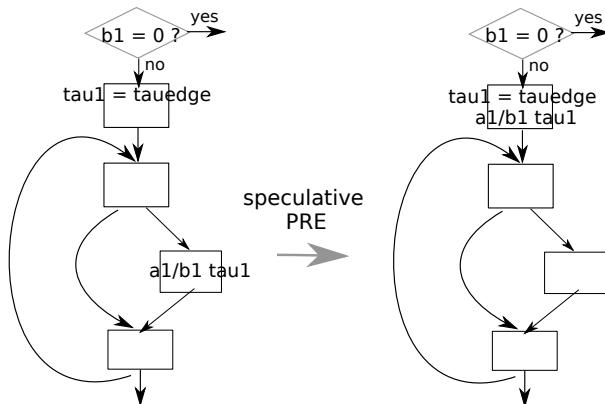


Fig. 12.8 Speculative and fault-safe loop-invariant code motion

12.5 Register Promotion via PRE

Variables and most data in programs normally start out residing in memory. It is the compiler's job to promote these memory contents to registers as much as

possible to speed up program execution. But when these memory contents are implicitly accessed through aliasing, it is necessary that their values are up-to-date with respect to their assigned registers. Load and store instructions have to be generated to transfer contents between memory locations and registers. The compiler also has to deal with the limited number of physical registers and find an allocation that makes the best use of them. Instead of solving these problems all at once, we can tackle them as two smaller problems separately:

1. Register promotion — We assume there is an unlimited number of registers, called *pseudo-registers*, *symbolic registers* or *virtual registers*. Register promotion will allocate variables to pseudo-registers whenever possible and optimize the placement of the loads and stores that transfer their values between memory and registers.
2. Register allocation — This phase will fit the unlimited number of pseudo-registers to the limited number of *real* or *physical* registers. Because pseudo-registers have no alias, assigning them to registers involves only renaming them. When it runs out of registers due to high register pressure, it will generate code to spill contents from registers back to memory.

In this Chapter, we only address the register promotion problem because it can be cast as a redundancy elimination problem.

12.5.1 Register Promotion as Placement Optimization

Variables with no aliases are trivial register promotion candidates. They include the temporaries generated during PRE to hold the values of redundant computations. Variables in the program can also be determined via compiler analysis or by language rules to be alias-free. For these trivial candidates, we can just rename them to unique pseudo-registers, and no load or store needs to be generated.

Our register promotion is mainly concerned with scalar variables that have aliases, indirectly accessed memory locations and program constants. A scalar variable can have aliases whenever its address is taken, or if it is a global variable, since it can be accessed by function calls. A program constant is a register promotion candidate whenever some operation is involved to put its value in a register and its register residence is required for its value to be used during execution.

Since the goal of register promotion is to obtain the most efficient placements for the loads and stores, register promotion can be modeled as two separate problems: PRE of loads, followed by PRE of stores. In the case of program constants, our use of the term *load* will extend to refer to the computation operation performed to put the constant value in a register. The PRE of stores does not apply to program constants.

From the point of view of redundancy, loads are like expressions because the later occurrences are the ones to be deleted. For stores, the reverse is true: the earlier stores are the ones to be deleted, as is evident in the examples of Figure 12.9(a) and (b). The PRE of stores, also called *partial dead code elimination*, can thus be treated as the dual of the PRE of loads. Performing PRE of stores thus has the effects of moving stores forward while inserting them as early as possible. Combining the effects of the PRE of loads and stores results in optimal placements of loads and stores while minimizing the live ranges of the pseudo-registers, by virtue of the computational and lifetime optimalities of our PRE algorithm.

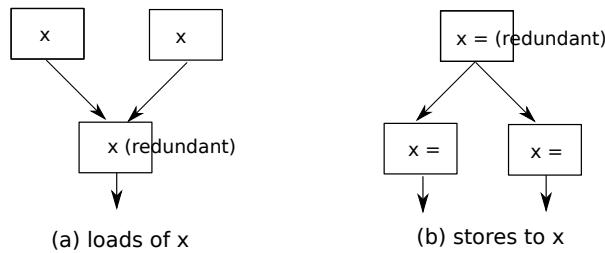


Fig. 12.9 Duality between load and store redundancies

12.5.2 Load Placement Optimization

PRE applies to any computation including loads from memory locations and creation of constants. In program representations, loads can either be indirect through a pointer or direct. Indirect loads are automatically covered by the PRE for expressions. Direct loads correspond to scalar variables in the program, and since our input program representation is in HSSA form, the aliasing that affects the scalar variables are completely modeled by the χ and μ operations. In our tree representation, both direct loads and constants are leaves in the expression trees. When we apply SSAPRE to direct loads, since the hypothetical temporary h can be regarded as the candidate variable itself, the FRG corresponds somewhat to the variable's SSA graph, so the Φ -insertion step and Rename step can be streamlined.

When working on the PRE of memory loads, it is important to also take into account the stores, which we call *l-value* occurrences. A store of the form $x \leftarrow \langle \text{expr} \rangle$ can be regarded as being made up of the sequence:

```

r ← <expr>
x ← r

```

Because the pseudo-register r contains the current value of x , any subsequent occurrences of the load of x can reuse the value from r , and thus can be regarded as redundant. Figure 12.10 gives examples of loads made redundant by stores.

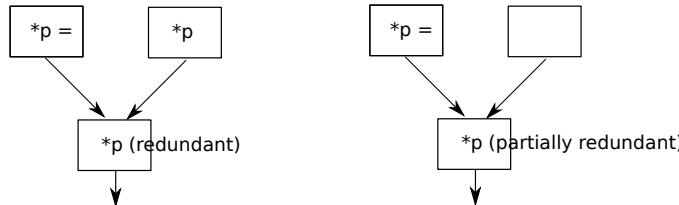


Fig. 12.10 Redundant loads after stores

When we perform the PRE of loads, we thus include the store occurrences into consideration. The Φ -insertion step will insert Φ 's at the iterated dominance frontiers of store occurrences. In the Rename step, a store occurrence is always given a new h -version, because a store is a definition. Any subsequent load renamed to the same h -version is redundant with respect to the store.

We apply the PRE of loads first, followed by the PRE of stores. This ordering is based on the fact that the PRE of loads is not affected by the results of the PRE of stores, but the PRE of loads creates more opportunities for the PRE of stores by deleting loads that would otherwise have blocked the movement of stores. In addition, speculation is required for the PRE of loads and stores in order for register promotion to do a decent job in loops.

The example in Figure 12.11 illustrates what we discuss in this section. During the PRE of loads (LPRE), $a =$ is regarded as a store occurrence. The hoisting of the load of a to the loop header does not involve speculation. The occurrence of $a =$ causes r to be updated by splitting the store into the two statements $r =$ followed by $a = r$. In the PRE of stores (SPRE), speculation is needed to sink $a =$ to outside the loop because the store occurs in a branch inside the loop. Without performing LPRE first, the load of a inside the loop would have blocked the sinking of $a =$.

12.5.3 Store Placement Optimization

As mentioned earlier, SPRE is the dual of LPRE. In the presence of store redundancies, the earlier occurrences are redundant. Code motion in SPRE will have the effect of moving stores forward with respect to the control-flow graph. Any presence of (aliased) loads have the effect of blocking the movement of stores or rendering the earlier stores non-redundant.

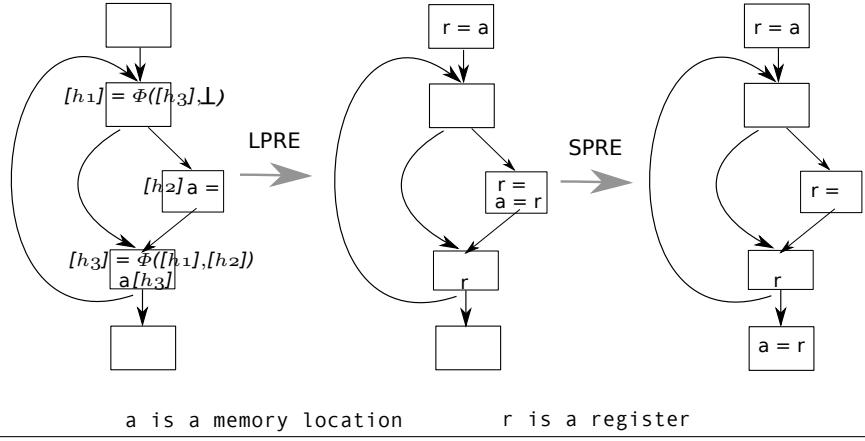


Fig. 12.11 Register promotion via load PRE followed by store PRE

To apply the dual of the SSAPRE algorithm, it is necessary to compute a program representation that is the dual of the SSA form, and we call this the *static single use* form. In SSU, use-def edges are factored at divergence points in the control-flow graph. We call this factoring operator Σ . Each use of a variable establishes a new version (we say the load *uses* the version), and every store reaches exactly one load. The Σ is regarded as a use of a new version of the variable. The use post-dominates all the stores of its version. The left-hand-side of the Σ is the multiple definitions, each of which is post-dominated by their single uses [reference to SSI chapter].

We call our store PRE algorithm SSUPRE, which is made up of the corresponding steps in SSAPRE. Σ -Insertion and SSU-Rename, construct the SSU form for the variable whose store is being optimized. The data-flow analyses consist of UpSafety to compute the *upsafe* (fully available) attribute, CanBeAnt to compute the *can_be_ant* attribute and Earlier to compute the *earlier* attribute. Though store elimination itself does not require the introduction of temporaries, lifetime optimality still needs to be considered for the temporaries introduced in the LPRE phase which hold the values to the point where the stores are placed. It is desirable not to sink the stores too far down.

Figure 12.12 gives an example program with (b) being the SSU representation for the program in (a). (c) shows the result of applying SSUPRE to the code. The store can be sunk to outside the loop only if it is also inserted in the branch inside the loop that contains the load. The optimized code no longer exhibits any store redundancy.

value numbering

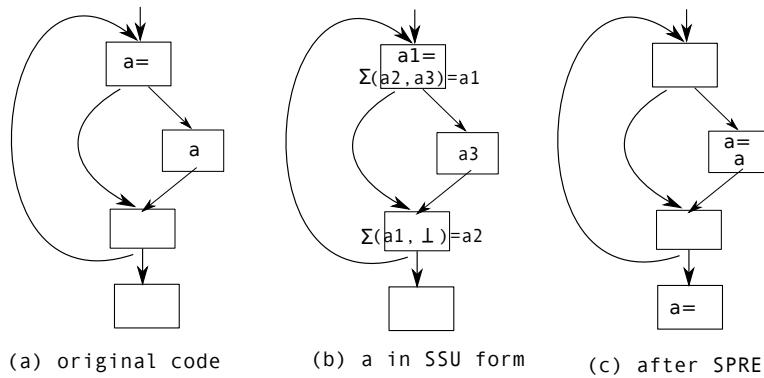


Fig. 12.12 Example of program in SSU form and the result of applying SSUPRE

12.6 Value-based Redundancy Elimination

The redundant computations stored into a temporary introduced by PRE may be of different static values because the same lexically identified expression may yield different results at different code paths in the program. The PRE algorithm we have described so far is not capable of recognizing redundant computations among lexically different expressions that yield the same value. In this Section, we discuss redundancy elimination based on value analysis.

12.6.1 Value Numbering

*

The term *value number* originated from the hash-based method developed by Cocke and Schwartz for recognizing when two expressions evaluate to the same value within a basic block. The value number of an expression tree can be regarded as the index of its hashed entry in the hash table. An expression tree is hashed recursively bottom up starting with the leaf nodes. Each internal node is hashed based on its operator and the value numbers of its operands. The local algorithm for value numbering will conduct a scan down the instructions in a basic block, assigning value numbers to the expressions. At an assignment, the assigned variable will be assigned the value number of the right hand side expression. The assignment will also cause any value number that refers to that variable to be killed. For example, the program code in Figure 12.13(a) will result in the value numbers v_1 , v_2 and v_3 shown in Figure 12.13(b). Note that variable c is involved with both value numbers v_2 and v_3 because it has been re-defined.

SSA form enables value numbering to be easily extended to the global scope, called global value numbering (GVN), because each SSA version of a variable

$b \leftarrow c$	$v_1: a$
$c \leftarrow a + c$	$v_2: b, c$
$d \leftarrow a + b$	$v_3: v_1+v_2, c, d$
(a) Code in a basic block	(b) Value numbers

Fig. 12.13 Value numbering in a local scope

$b_1 \leftarrow c_1$	$v_1: a_1$
$c_2 \leftarrow a_1 + c_1$	$v_2: b_1, c_1$
$d_1 \leftarrow a_1 + b_1$	$v_3: v_1+v_2, c_2, d_1$
(a) Statements processed	(b) Value numbers

Fig. 12.14 Global value numbering on SSA form

corresponds to at most one static value for the variable. In the example of 12.14, the same value numbers can be arrived at regardless of the order of processing the program code. One subtlety is regarding the ϕ 's. When we value number a ϕ , we would like the value numbers for its operands to have been determined. One strategy is to perform the global value numbering by visiting the nodes in the control-flow graph in a reverse postorder traversal of the dominitor tree. This traversal strategy can minimize the instances when a ϕ operand has unknown value number when the ϕ itself needs to be value-numbered, except in the case of back edges arising from loops. When this arises, we have no choice but to assign a new value number to the ϕ . For example, in this loop:

```

 $i_1 \leftarrow 0;$ 
 $j_1 \leftarrow 0;$ 
do {
     $i_2 \leftarrow \phi(i_3, i_1);$ 
     $j_2 \leftarrow \phi(j_3, j_1);$ 
     $i_3 \leftarrow i_2 + 4;$ 
     $j_3 \leftarrow j_2 + 4;$ 
while <cond>;

```

when we try to hash a value number for either of the two ϕ 's, the value numbers for i_3 and j_3 are not yet determined. As a result, we create different value numbers for i_2 and j_2 . This makes the above algorithm unable to recognize that i_2 and j_2 can be given the same value number, or i_3 and j_3 can be given the same value number.

The above hash-based value numbering algorithm can be regarded as pessimistic, because it will not assign the same value number to two different expressions unless it can prove they compute the same value. There is a different approach to performing value numbering that is not hash-based and is optimistic. It does not depend on any traversal over the program's flow of control, and so is not affected by the presence of back edges. The algorithm partitions all the expressions in the program into *congruence classes*. Expressions in the same congruence class are considered *equivalent* because they evaluate to the

```

1: Place expressions with the same operator in the same congruence class
2:  $N \leftarrow$  maximum number of operands among all operators in the program
3:  $worklist \leftarrow$  set of all congruence classes
4: while  $worklist \neq \{\}$ 
5:   Select and delete an arbitrary congruence class  $c$  from  $worklist$ 
6:   for ( $p = 0; p < N; p++$ )
7:      $touched \leftarrow$  set of expressions in program that has a member of  $c$  in position  $p$ 
8:     for each class  $s$  such that  $(s \cap touched)$  is not empty and  $(s \cap touched) \subset s$ 
9:       Create a new class  $n \leftarrow s \cap touched$ 
10:       $s \leftarrow s - n$ 
11:      if  $s \in worklist$ 
12:        Add  $n$  to  $worklist$ 
13:      else
14:        Add smaller of  $n$  and  $s$  to  $worklist$ 

```

Fig. 12.15 The partitioning algorithm

same static value. The algorithm is optimistic because when it starts, it assumes all expressions that have the same operator to be in the same congruence class. Given two expressions within the same congruence class, if their operands at the same operand position belong to different congruence classes, the two expressions may compute to different values, and thus should not be in the same congruence class. This is the subdivision criterion. As the algorithm iterates, the congruence classes are subdivided into smaller ones while the total number of congruence classes increases. The algorithm terminates when no more subdivision can occur. At this point, the set of congruence classes in this final partition will represent all the values in the program that we care about, and each congruence class is assigned a unique value number.

The details of this partitioning algorithm is shown in Figure 12.15. At line 8, $(s \cap touched)$ being not empty implies s uses a member of c as its p th operand. If $(s \cap touched)$ is a proper subset of s , it implies some member of s does not have a member of c as its p th operand. This is the criterion for subdividing class s because those members of s that do not have a member of c at its p th operand potentially compute to a different value. After the partition into n and the new s , if s is not in the worklist (i.e., processed already), the partitioning was already stable with respect to the old s , and we can add either n or the new s to the worklist to re-stabilize with respect to that split of s . Choosing the smaller one results in less overhead.

While this partition-based algorithm is not obstructed by the presence of back edges, it does have its own deficiencies. Because it has to consider one operand position at a time, it is not able to apply commutativity to detect more equivalences. Since it is not applied bottom-up with respect to the expression tree, it is not able to apply algebraic simplifications while value numbering. To get the best of both the hash-based and the partition-based algorithms, it is possible to apply both algorithms independently and then combine their results together to shrink the final set of value numbers.

12.6.2 Redundancy Elimination under Value Numbering

So far, we have discussed finding computations that compute to the same values, but have not addressed eliminating the redundancies among them. Two computations that compute to the same value exhibit redundancy only if there is a control-flow path that leads from one to the other.

An obvious approach is to consider PRE for each value number separately. A temporary t will be introduced to store the redundant computations for each value number, but in this case, its value will stay the same throughout the lifetime of t . If there are ϕ 's introduced for the t , they will be merging identical values, and we know from experience that such ϕ 's are rare. A subset of such ϕ 's is expected to come from PRE's insertions, and that implies that insertions introduced by value-number-based PRE are also rare.

Value-number-based PRE also has to deal with the additional issue of *how* to generate an insertion. Because the same value can come from different forms of expressions at different points in the program, it is necessary to determine which form to use at each insertion point. If the insertion point is outside the live range of any variable version that can compute that value, then the insertion point has to be disqualified. Due to this complexity, and the expectation that strictly partial redundancy is rare among computations that yield the same value, it is sufficient to perform only full redundancy elimination among computations that have the same value number.

But it is possible to broaden the scope and consider PRE among lexically identical expressions and value numbers at the same time. In this hybrid approach, it is best to relax our restriction on the style of program representation described in Section 12.1. By not requiring Conventional SSA Form, we can more effectively represent the flow of values among the program variables⁵. By regarding the live range of each SSA version to extend from its definition to program exit, we allow its value to be used whenever convenient. The program representation can even be in the form of triplets, in which the result of every operation is immediately stored in a temporary. It will just assign the value number of the right hand sides to the left hand side variables.

This hybrid approach can be implemented based on an adaptation of the SSAPRE framework. Since each ϕ in the input can be viewed as merging different value numbers from the predecessor blocks to form a new value number, the Φ -Insertion step will be driven by the presence of ϕ 's for the program variables. The FRGs can be formed from some traversal of the program code. Each FRG can be regarded as a representation of the flow and merging of computed values based on which PRE can be performed by applying the remaining steps of the SSAPRE algorithm.

⁵ We would remove the restriction that the result and operands of each ϕ have to be SSA versions derived from the same original program variable

12.7 Additional Reading

The concept of partial redundancy was first introduced by Morel and Renvoise. In their seminal work [203], Morel and Renvoise showed that global common subexpressions and loop-invariant computations are special cases of partial redundancy, and they formulated PRE as a code placement problem. The PRE algorithm developed by Morel and Renvoise involves bi-directional data-flow analysis, which incurs more overhead than uni-directional data-flow analysis. In addition, their algorithm does not yield optimal results in certain situations. A better placement strategy, called lazy code motion (LCM), was later developed by Knoop *et al* [169][171]. It improved on Morel and Renvoise's results by avoiding unnecessary code movements, by removing the bi-directional nature of the original PRE data-flow analysis and by proving the optimality of their algorithm. After lazy code motion was introduced, there have been alternative formulations of PRE algorithms that achieve the same optimal results, but differ in the formulation approach and implementation details [104][103][216][293].

The above approaches to PRE are all based on encoding program properties in bit vector forms and the iterative solution of data-flow equations. Since the bit vector representation uses basic blocks as its granularity, a separate algorithm is needed to detect and suppress local common subexpressions. Chow *et al* [74][163] came up with the first SSA-based approach to perform PRE. Their SSAPRE algorithm is an adaptation of LCM that take advantage of the use-def information inherent in SSA. It avoids having to encode data-flow information in bit vector form, and eliminates the need for a separate algorithm to suppress local common subexpressions. Their algorithm was first to make use of SSA to solve data-flow problems for expressions in the program, taking advantage of SSA's sparse representation so that fewer number of steps are needed to propagate data-flow information. The SSAPRE algorithm thus brings the many desirable characteristics of SSA-based solution techniques to PRE.

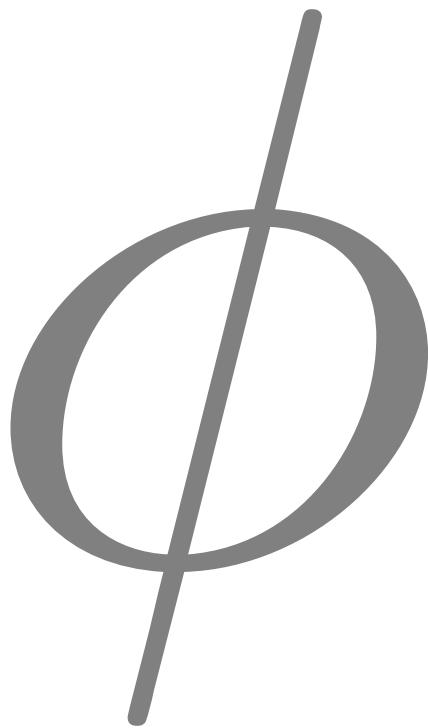
In the area of speculative PRE, Murphy *et al.* introduced the concept of fault-safety and use it in the SSAPRE framework for the speculation of dangerous computations [205]. When execution profile data are available, it is possible to tailor the use of speculation to maximize run-time performance for the execution that matches the profile. Xue and Cai presented a computationally and lifetime optimal algorithm for speculative PRE based on profile data [292]. Their algorithm uses bit-vector-based data-flow analysis and applies minimum cut to flow networks formed out of the control-flow graph to find the optimal code placement. Zhou et al. applies the minimum cut approach to flow networks formed out of the FRG in the SSAPRE framework to achieve the same computational and lifetime optimal code motion [298]. They showed their sparse approach based on SSA results in smaller flow networks, enabling the optimal code placements to be computed more efficiently.

Lo *et al.* showed that register promotion can be achieved by load placement optimization followed by store placement optimization [184]. Other optimiza-

tions can potentially be implemented using the SSAPRE framework, like code hoisting, register shrink-wrapping [76] and live range shrinking. Moreover, PRE has traditionally provided the context for integrating additional optimizations into its framework. They include operator strength reduction [170] and linear function test replacement [164].

Hashed-based value numbering originated from Cocke and Schwartz [81], and Rosen *et al.* extended it to global value number based on SSA [240]. The partition-based algorithm was developed by Alpern *et al.* [10]. Briggs *et al.* presented refinements to both the hash-based and partition-based algorithms [49], including applying the hash-based method in a postorder traversal of the dominator tree.

VanDrunen and Hosking proposed their anticipation-based SSAPRE (A-SSAPRE) that removes the requirement of Convention SSA Form and is best for program representations in the form of triplets [279]. Their algorithm determines optimization candidates and construct FRGs via a depth-first, preorder traversal over the basic blocks of the program. Within each FRG, non-lexically identical expressions are allowed, as long as there are potential redundancies among them. VanDrunen and Hosking subsequently presented their Value-based Partial Redundancy Elimination algorithm (GVN-PRE) that they claim to subsume both PRE and GVN [280]. But in this algorithm, the PRE part is not SSA-based because it uses a system of data-flow equations at the basic block level to solve for insertion points.



Part III

Extensions

Progress: 48%

Progress: 48%



TODO: Aimed at experts,
but with a track for be-
ginners

TODO: Section on
Sigma-SSA?

TODO: Section on Dy-
namic SSA?

CHAPTER 13

Introduction

V. Sarkar
F. Rastello

Progress: 50%



Review in progress

Thus far, we have introduced the foundations of SSA form and its use in different program analyses. We now motivate the need for extensions to SSA form to enable a larger class of program analyses. The extensions arise from the fact that many analyses need to make finer-grained distinctions between program points and data accesses than is achieved by vanilla SSA form. However, these richer flavors of extended SSA-based analyses still retain many of the benefits of SSA form (e.g., sparse data-flow propagation) which distinguish them from classical data-flow frameworks for the same analysis problems.

Static Single Information form

The sparseness in vanilla SSA form arises from the observation that information for an unaliased scalar variable can be safely propagated from its (unique) definition to all its reaching uses without examining any intervening program points. As an example, SSA-based constant propagation aims to compute for each single assignment variable, the (usually over-approximated) set of possible values carried by the definition for that variable. At an instruction that defines variable a and uses variables b and c (say), the information for b and c is propagated directly to the instruction from the definitions of those variables. However, there are many analyses for which definition points are not the only source of new information. For example, consider an if-then-else statement with uses of a in both the then and else parts. If the branch condition involves a , say $a == 0$, we now have additional information that can distinguish the value of a in the then and else parts, even though both uses have the same reaching definition.

Likewise, the use of a reference variable p in certain contexts can be the source of new information for subsequent uses e.g., the fact that p is non-null because no exception was thrown at the first use.

The goal of Chapter 14 is to present a systematic approach to dealing with these additional sources of information, while still retaining the space and time benefits of vanilla SSA form. This approach is called Static Single Information (SSI) form, and it involves additional renaming of variables to capture new sources of information. For example, SSI form can provide distinct names for the two uses of a in the then and else parts of the if-then-else statement mentioned above. This additional renaming is also referred to as *live range splitting*, akin to the idea behind splitting live ranges in optimized register allocation. The sparseness of SSI form follows from formalization of the Partitioned Variable Lattice (PVL) problem and the Partitioned Variable Problem (PVP), both of which establish orthogonality among transfer functions for renamed variables. The ϕ functions inserted at join nodes in vanilla SSA form are complemented by σ functions in SSI form that can perform additional renaming at branch nodes and interior (instruction) nodes. Information can be propagated in the forward and backward directions using SSI form, enabling analyses such as range analysis (leveraging information from branch conditions) and null pointer analysis (leveraging information from prior uses) to be performed more precisely than with vanilla SSA form.

Hashed SSA form

The motivation for SSI form arose from the need to perform additional renaming to distinguish among different uses of an SSA variable. Hashed SSA (HSSA) form introduced in Chapter 15 addresses another important requirement viz., the need to model aliasing among variables. For example, a static use or definition of indirect memory access $*p$ in the C language could represent the use or definition of multiple local variables whose addresses can be taken and may potentially be assigned to p along different control-flow paths. To represent aliasing of local variables, HSSA form extends vanilla SSA form with $MayUse(\mu)$ and $MayDef(\chi)$ functions to capture the fact that a single static use or definition could potentially impact multiple variables. Note that MayDef functions can result in the creation of new names (versions) of variables, compared to vanilla SSA form. HSSA form does not take a position on the accuracy of alias analysis that it represents. It is capable of representing the output of any alias analysis performed as a pre-pass to HSSA construction. As summarized above, a major concern with HSSA form is that its size could be quadratic in the size of the vanilla SSA form, since each use or definition can now be augmented by a set of MayUse's and MayDef's respectively. A heuristic approach to dealing with this problem is to group together all variable versions that have no “real” occurrence in the program i.e., do not appear in a real instruction outside of a ϕ , μ or χ

function. These versions are grouped together into a single version called the *zero version* of the variable.

In addition to aliasing of locals, it is important to handle the possibility of aliasing among heap-allocated variables. For example, $*p$ and $*q$ may refer to the same location in the heap, even if no aliasing occurs among local variables. HSSA form addresses this possibility by introducing a *virtual variable* for each address expression used in an indirect memory operation, and renaming virtual variables with ϕ functions as in SSA form. Further, the alias analysis pre-pass is expected to provide information on which virtual variables may potentially be aliased, thereby leading to the insertion of μ or χ functions for virtual variables as well. Global value numbering is used to increase the effectiveness of the virtual variable approach, since all indirect memory accesses with the same address expression can be merged into a single virtual variable (with SSA renaming as usual). In fact, the Hashed SSA name in HSSA form comes from the use of hashing in most value-numbering algorithms.

Array SSA form

In contrast to HSSA form, Array SSA form (Chapter 16) takes an alternate approach to modeling aliasing of indirect memory operations by focusing on aliasing in arrays as its foundation. The aliasing problem for arrays is manifest in the fact that accesses to elements $A[i]$ and $A[j]$ of array A refer to the same location when $i = j$. This aliasing can occur with just local array variables, even in the absence of pointers and heap-allocated data structures. Consider a program with a definition of $A[i]$ followed by a definition of $A[j]$. The vanilla SSA approach can be used to rename these two definitions to (say) $A_1[i]$ and $A_2[j]$. The challenge with arrays arises when there is a subsequent use of $A[k]$. For scalar variables, the reaching definition for this use can be uniquely identified in vanilla SSA form. However, for array variables, the reaching definition depends on the subscript values. In this example, the reaching definition for $A[k]$ will be A_2 or A_1 if $k == j$ or $k == i$ (or a prior definition A_0 if $k \neq j$ and $k \neq i$). To provide $A[k]$ with a single reaching definition, Array SSA form introduces a definition- Φ ($d\Phi$) operator that represents the merge of $A_2[j]$ with the prevailing value of array A prior to A_2 . The result of this $d\Phi$ operator is given a new name, A_3 (say), which serves as the single definition that reaches use $A[k]$ (which can then be renamed to $A_3[k]$). This extension enables sparse data-flow propagation algorithms developed for vanilla SSA form to be applied to array variables, as illustrated by the algorithm for *sparse constant propagation of array elements* presented in this chapter. The accuracy of analyses for Array SSA form depends on the accuracy with which pairs of array subscripts can be recognized as being *definitely-same* (\mathcal{DS}) or *definitely-different* (\mathcal{DD}).

To model heap-allocated objects, Array SSA form builds on the observation that all indirect memory operations can be modeled as accesses to elements of abstract arrays that represent disjoint subsets of the heap. For modern object-

oriented languages like Java, type information can be used to obtain a partitioning of the heap into disjoint subsets e.g., instances of field x are guaranteed to be disjoint from instances of field y . In such cases, the set of instances of field x can be modeled as a logical array (map) \mathcal{H}^x that is indexed by the object reference (key). The problem of resolving aliases among field accesses $p.x$ and $q.x$ then becomes equivalent to the problem of resolving aliases among array accesses $\mathcal{H}^x[p]$ and $\mathcal{H}^x[q]$, thereby enabling Array SSA form to be used for analysis of objects as in the algorithm for *redundant load elimination* among object fields presented in this chapter. For weakly-typed languages such as C, the entire heap can be modeled as a single heap array. As in HSSA form, an alias analysis pre-pass can be performed to improve the accuracy of *definitely-same* and *definitely-different* information. In particular, global value numbering is used to establish *definitely-same* relationships, analogous to its use in establishing equality of address expressions in HSSA form. In contrast to HSSA form, the size of Array SSA form is guaranteed to be linearly proportional to the size of a comparable scalar SSA form for the original program (if all array variables were treated as scalars for the purpose of the size comparison).

Psi-SSA form

Psi-SSA form (Chapter 17) addresses the need for modeling static single assignment form in predicated operations. A predicated operation is an alternate representation of a fine-grained control-flow structure, often obtained by using the well known *if conversion* transformation. A key advantage of using predicated operations in a compiler's intermediate representation is that it can often enable more optimizations by creating larger basic blocks compared to approaches in which predicated operations are modeled as explicit control-flow graphs. From an SSA form perspective, the challenge is that a predicated operation may or may not update its definition operand, depending on the value of the predicate guarding that assignment. This challenge is addressed in Psi-SSA form by introducing ψ functions that perform merge functions for predicated operations, analogous to the merge performed by ϕ functions at join points in vanilla SSA form.

In general, a ψ function has the form, $a0 = \text{Psi}(p1?a1, \dots, pi?ai, \dots, pn?an)$, where each input argument a_i is associated with a predicate p_i as in a nested if-then-else expression for which the value returned is the rightmost argument whose predicate is true. If all predicates are false, then the assignment is a no-op and the result is undefined. (*TODO: check if this is correct.*) A number of algebraic transformations can be performed on ψ functions, including ψ -inlining, ψ -reduction, ψ -projection, ψ -permutation and ψ -promotion. Chapter 17 also includes an algorithm for transforming a program out of Psi-SSA form that extends the standard algorithm for destruction of vanilla SSA form. So as to expose parallelism / locality one need to get rid of the CFG at some points. For loop transformations, software pipelining, one need to manipulate a higher degree of

abstraction to represent the iteration space of nested loops and to extend data-flow information to this abstraction. One can expose even more parallelism (at the level of instructions) by replacing control flow by control dependences: the goal is either to express a predicate expression under which a given basic block is to be executed, or select afterward (using similar predicate expressions) the correct value among a set of eagerly computed ones.

1. technically, we say that SSA provides data flow (data dependences). The goal is to enrich it with control dependences. Program dependence graph (PDG) constitute the basis of such IR extensions. Gated-SSA (GSA) provides an interpretable (data or demand driven) IR that uses this concept. PDG and GSA are described in Chapter 18. Psi-SSA is a very similar IR but more appropriated to code generation for architectures with predication. Psi-SSA is described in Chapter 17.
2. Note that such extensions sometimes face difficulties to handle loops correctly (need to avoid deadlock between the loop predicate and the computation of the loop body, replicate the behavior of infinite loops, etc.). However, we believe that, as we will illustrate further, loop carried control dependences complicate the recognition of possible loop transformations: it is usually better to represent loops and there corresponding iteration space using a dedicated abstraction.

Memory based data flow

SSA provides data flow / dependences between *scalar* variables. Execution order of side effect instructions must also be respected. Indirect memory access can be considered very conservatively as such and lead to (sometimes called state, see Chapter 18) dependence edges. Too conservative dependences annihilate the potential of optimizations. Alias analysis is the first step toward more precise dependence information. Representing this information efficiently in the IR is important. One could simply add a dependence (or a flow arc) between two "consecutive" instructions that may or must alias. Then in the same spirit than SSA phi-nodes aim at combining the information as early as possible (as opposed to standard def-use chains, see the discussion in Chapter 2), similar nodes can be used for memory dependences. Consider the following sequential code

```
*p=...; *q=...; ...=*p; ...=*p
```

where p, and q may alias, to illustrate this point. Without the use of phi-nodes, the amount of def-use chains required to link the assignments to their uses would be quadratic (4 here). Hence the usefulness of generalizing SSA and its phi-node for scalars to handle memory access for sparse analyzes. HSSA (see Chapter 15) and Array-SSA (see the Chapter 16) are two different implementations of this idea. One have to admit that if this early combination is well suited for analysis or interpretation, the introduction of a phi-function might add a

control dependence to an instruction that would not exist otherwise. In other words only simple loop carried dependences can be expressed this way. Let us illustrate this point using a simple example:

```
for i { A[i]=f(A[i-2]) }
```

where the computation at iteration indexed by $i + 2$ accesses the value computed at iteration indexed by i . Suppose we know that $f(x) > x$ and that prior to entering the loop $A[*] \geq 0$. Then a SSA like representation

```
for i {
    A2=phi(A0,A1);
    A1=PHI(j==i:A1[j]=f(A2[j-2]), j!=i:A1[j]=A2[j])
}
```

would easily allow for the propagation of information that $A2 \geq 0$. On the other hand, by adding this phi-node, it becomes difficult to devise that iteration i and $i + 1$ can be executed in parallel: the phi node adds a loop carried dependence. If one is interested in performing more sophisticated loop transformations than just exposing fully parallel loops (such as loop interchange, loop tiling or multidimensional software pipelining), then (Dynamic) Single Assignment forms should be his friend. There exists many formalisms including Kahn Process Networks (KPN), or Fuzzy Data-flow Analysis (FADA) that implement this idea. But each time restrictions apply. This is part of the huge research area of automatic parallelization outside of the scope of this book. For further details we refer to the corresponding Encyclopedia of Parallel Computing [215].

CHAPTER 14

Static Single Information Form

F. Pereira

F. Rastello

Progress: 90%



Text and figures formating in progress

14.1 Introduction

The objective of a data-flow analysis is to discover facts that are true about a program. We call such facts *information*. Using the notation introduced in Section 8, an information is an element in the data-flow lattice. For example, the information that concerns liveness analysis is the set of variables alive at a certain program point. Similarly to liveness analysis, many other classic data-flow approaches bind information to pairs formed by a variable and a program point. However, if an invariant occurs for a variable v at any program point where v is alive, then we can associate this invariant directly to v . If a program's intermediate representation guarantees this correspondence between information and variable for every variable, then we say that the program representation provides the *Single Static Information* (SSI) property.

In Chapter 8 we have shown how SSA form allows us to solve sparse forward data-flow problems such as constant propagation. In the particular case of constant propagation, the SSA format let us assign to each variable the invariant – or information – of being constant or not. The SSA intermediate representation gives us this invariant because it splits the live ranges of variables, in such a way that each variable name is defined only once. Now we will show that live range splitting can also provide the SSI property not only to forward, but also to backward data-flow analyses.

Different data-flow analysis might extract information from different program facts. Therefore, a program representation may afford the SSI property to some data-flow analysis, but not to all of them. For instance, the SSA form naturally provides the SSI property to the reaching definition analysis. Indeed, the SSA form provides the static single information property to any data-flow analysis that obtains information at the definition sites of variables. These analyses and transformations include copy and constant propagation as illustrated in Chapter 8. However, for a data-flow analysis that derives information from the use sites of variables, such as the class inference analysis described in Section 14.2.6, the information associated with v might not be unique along its entire live-range: in that case the SSA form does not provide the SSI property.

There exists extensions of the SSA form that provide the SSI property to more data-flow analyses than the original SSA does. Two classic examples, that we will detail further, are the *Extended-SSA* (e-SSA) form, and the *Static Single Information* (SSI) form. The e-SSA form provides the SSI property to analyses that take information from the definition site of variables, and also from conditional tests where these variables are used. The SSI form provides the static single information property to data-flow analyses that extract information from the definition sites of variables, and from last use sites (which we define later). These different intermediate representations rely on a common strategy to achieve the SSI property: *live range splitting*. In this chapter we show how to use live range splitting to build program representations that provide the static single information property to different types of data-flow analysis.

14.2 Static Single Information

The goal of this section is to define the notion of Static Single Information, and to explain how it supports the sparse data-flow analyses discussed in Chapter 8. With this purpose, we revisit the concept of sparse analysis in Section 14.2.1. There exists a special class of data-flow problems, which we call *Partitioned Lattice per Variable* (PLV), that fits in this chapter's sparse data-flow framework very well. We will look more carefully into these problems in Section 14.2.2. The intermediate program representations that we discuss in this chapter lets us provide the Static Single Information property – formalized in Section 14.2.3 – to any PLV problem. And the algorithms that we give in Section 14.2.5 lets us solve sparsely any data-flow problem that contains the SSI property. This sparse framework is very broad: many well-known data-flow problems are partitioned lattice, as we will see in the examples from Section 14.2.6.

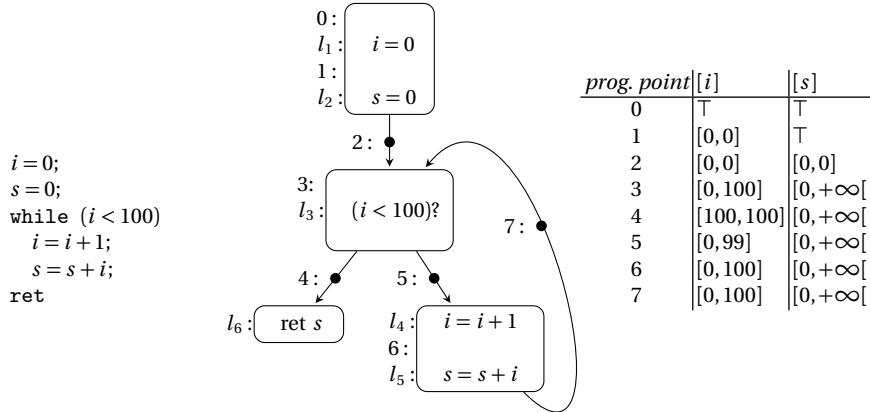


Fig. 14.1 An example of a dense data-flow analysis that finds the range of possible values associated with each variable at each program point .

14.2.1 Sparse Analysis

Traditionally, non relational data-flow analyses bind information to pairs formed by a variable and a program point . Let's consider, as an example, the problem of estimating the interval of values that any integer variable may assume throughout the execution of a program. An algorithm that solves this problem is called a *range analysis*. A traditional implementation of this analysis would find, for each pair formed by a variable v plus a program point p , the interval of possible values that v might assume at p . Figure 14.1 illustrates this analysis with an example. In this case we call a program point any region between two consecutive instructions and we let $[v]$ denote the abstract information that is associated with variable v . Because this approach keeps information at each program point , we call it *dense*, in contrast with the sparse analyses seen in Section 8.3.

The dense approach might keep more redundant information during the data-flow analysis. For instance, if we let $[v]^p$ denote the abstract state of variable v at program point p , then we have that, in our example, $[i]^1 = [i]^2$, $[s]^5 = [s]^6$ and $[i]^6 = [i]^7$. This redundancy happens because some transfer functions are identities. For instance, in range analysis, an instruction that neither defines nor uses any variable is associated with an identity transfer function. The transfer function that updates the abstract state of i at program point 2 is an identity. Because the instruction immediately before 2 does not add any new information to the abstract state of i , $[i]^2$ is updated with the information that flows directly from the predecessor point 1.

The goal of sparse data-flow analysis is to shortcut the identity transfer functions, a task that we accomplish by grouping contiguous program points bound to identities into larger regions. Solving a data-flow analysis sparsely has many advantages over doing it densely: because we do not need to keep bitvectors

associated with each program point, the sparse solution tends to be more economical in terms of space and time. Going back to our example, a given variable v may be mapped to the same interval along many consecutive program points. Furthermore, if the information associated with a variable is invariant along its entire live range, then we can bind this information to the variable itself. In other words, we can replace all the constraint variables $[v]^p$ by a single constraint variable $[v]$, for each variable v and every $i \in \text{live}(v)$. Not every data-flow problem can be easily solved sparsely; however, many of them can, as they fit into the family of PLV problems that we describe in the next section.

14.2.2 Partitioned Lattice per Variable (PLV) data-flow Problems

The class of non-relational data-flow analysis problems we are interested in are the ones that bind information to pairs of program variables and program points. We design this class of problems as *Partitioned Lattice per Variable* problems that we formally describe as follow.

Definition 1 (PLV). Let $\mathcal{V} = \{v_1, \dots, v_n\}$ be the set of program variables. Let us consider, without loss of generality, a forward data-flow analysis. This data-flow analysis can be written as an equation system that associates, with each program point p , an element of a lattice \mathcal{L} , given by the equation $x^p = \bigwedge_{s \in \text{preds}(p)} F^{s,p}(x^s)$, where: x^p denotes the abstract state associated with program point p ; for s a control-flow predecessor of p , $F^{s,p}$ is the transfer function from s to p . The analysis can alternatively be written¹ as a constraint system that binds to each program point p and each $s \in \text{preds}(p)$ the equation $x^p = x^p \wedge F^{s,p}(x^s)$ or, equivalently, the inequation $x^p \sqsubseteq F^{s,p}(x^s)$. The corresponding Maximum Fixed Point (MFP) problem is said to be a *Partitioned Lattice per Variable Problem* iff \mathcal{L} can be decomposed into the product of $\mathcal{L}_{v_1} \times \dots \times \mathcal{L}_{v_n}$ where each \mathcal{L}_{v_i} is the lattice associated with program variable v_i . In other words x^s can be written as $([v_1]^s, \dots, [v_n]^s)$ where $[v]^s$ denotes the abstract state associated to variable v and program point s . $F^{s,p}$ can thus be decomposed into the product of $F_{v_1}^{s,p} \times \dots \times F_{v_n}^{s,p}$ and the constraint system decomposed into the inequalities $[v_i]^p \sqsubseteq F_{v_i}^{s,p}([v_1]^s, \dots, [v_n]^s)$.

Going back to range analysis, if we denote by \mathcal{I} the lattice of integer intervals, then the overall lattice can be written as $\mathcal{L} = \mathcal{I}^n$, where n is the number of variables. Note that the class of PLV problems include a smaller class of problems called *Partitioned Variable Problems* (PVP). These analyses, which include live variables, reaching definitions and forward/backward printing, can be decomposed into a set of sparse data-flow problems – usually one per variable – each independent on the other.

¹ As far as we are concerned with finding its maximum solution. See for example Section 1.3.2 of [208].

14.2.3 The Static Single Information Property

If the information associated with a variable is invariant along its entire live range, then we can bind this information to the variable itself. In other words, we can replace all the constraint variables $[v]^p$ by a single constraint variable $[v]$, for each variable v and every $p \in \text{live}(v)$. Consider the problem of range analysis again. There are two types of control-flow points associated with non-identity transfer functions: definitions and conditionals: (1) at the definition point of variable v , F_v simplifies to a function that depends only on some $[u]$ where each u is an argument of the instruction defining v ; (2) at the conditional tests that use a variable v , F_v can be simplified to a function that uses $[v]$ and possibly other variables that appear in the test. The other program points are associated with an identity transfer function and can thus be ignored: $[v]^p = [v]^p \wedge F_v^{s,p}([v_1]^s, \dots, [v_n]^s)$ simplifies to $[v]^p = [v]^p \wedge [v]^p$ i.e., $[v]^p = [v]^p$. This gives the intuition on why a propagation engine along the def-use chains of a SSA-form program can be used to solve the constant propagation problem in an equivalent, yet “sparser”, manner. This also paves the way toward a formal definition of the Static Single Information property.

Property 1 (SSI). STATIC SINGLE INFORMATION: Consider a forward (resp. backward) monotone PLV problem E_{dense} stated as a set of constraints $[v]^p \subseteq F_v^{s,p}([v_1]^s, \dots, [v_n]^s)$ for every variable v , each program point p , and each $s \in \text{preds}(p)$ (resp. $s \in \text{succs}(p)$). A program representation fulfills the Static Single Information property iff:

[SPLIT]: let s be the unique predecessor (resp. successor) of a program point where a variable v is live and such that $F_v^{s,p} \neq \lambda x. \perp$ is non-trivial, i.e., is not the simple projection on \mathcal{L}_v , then s should contain a definition (resp. last use) of v ; Let $(Y_v^p)_{(v,p) \in \text{variables} \times \text{progPoints}}$ be a maximum solution to E_{dense} . Each node p that has several predecessors (resp. successors), and for which $F_v^{s,p}(Y_{v_1}^s, \dots, Y_{v_n}^s)$ has different values on its incoming edges (s, p) (resp. outgoing edges (p, s)) should have a ϕ -function at the entry of p (resp. σ -function at the exit of p) for v as defined in Section 14.2.4.

[INFO]: each program point i such that $v \notin \text{live-out}(i)$ (resp. $v \notin \text{live-in}(i)$) should be bound to an undefined transfer function, i.e., $F_v^i = \lambda x. \perp$.

[LINK]: each instruction $inst$ for which F_v^{inst} depends on some $[u]^s$ should contain an use (resp. def) of u live-in (resp. live-out) at $inst$.

[VERSION]: for each variable v , $\text{live}(v)$ is a connected component of the CFG.

These properties allow us to attach the information to variables, instead of program points. The SPLIT property forces the information related to a variable to be invariant along its entire live-range. INFO forces this information to be irrelevant outside the live range of the variable. The LINK property forces the def-use chains to reach the points where information is available for a transfer function to be evaluated. The VERSION property provides an one-to-one mapping between variable names and live ranges.

We must split live ranges to provide the SSI properties. If we split them between each pair of consecutive instructions, then we would automatically provide these properties, as the newly created variables would be live at only one program point. However, this strategy would lead to the creation of many trivial program regions, and we would lose sparsity. In Section 14.3 we provide a sparser way to split live ranges that fit Property 1. Possibly, we may have to extend the live-range of a variable to cover every program point where the information is relevant. We accomplish this last task by inserting into the program pseudo-uses and pseudo-definitions of this variable.

14.2.4 Special instructions used to split live ranges

We perform live range splitting via special instructions: the σ -functions and parallel copies that, together with ϕ -functions, create new definitions of variables. These notations are important elements of the propagation engine described in the section that follows. In short, a σ -function (for a branch point) is the dual of a ϕ -function (for a join point), and a parallel copy is a copy that *must* be done in parallel with another instruction. Each of these special instructions, ϕ -function, σ -functions, and parallel copies split live ranges at different kinds of program points: interior nodes, branches and joins.

Interior nodes are program points that have a unique predecessor and a unique successor. At these points we perform live range splitting via copies. If the program point already contains another instruction, then this copy *must* be done *in parallel* with the existing instruction. The notation,

$$\textit{inst} \parallel v_1 = v'_1 \parallel \dots \parallel v_m = v'_m$$

denotes m copies $v_i = v'_i$ performed in parallel with instruction \textit{inst} . This means that all the uses of \textit{inst} plus all right-hand variables v'_i are read simultaneously, then \textit{inst} is computed, then all definitions of \textit{inst} plus all left-hand variables v_i are written simultaneously. The null pointer analysis example of Figure 14.4(d), that we will explain later, uses parallel copies at nodes l_4 and l_6 .

We call *joins* the program points that have one successor and multiple predecessors. For instance, two different definitions of the same variable v might be associated with two different constants; hence, providing two different pieces of information about v . To avoid that these definitions reach the same use of v we merge them at the earliest program point where they meet. We do it via our well-known ϕ -functions.

In backward analyses the information that emerges from different uses of a variable may reach the same *branch point*, which is a program point with a unique predecessor and multiple successors. To ensure Property 1, the use that reaches the definition of a variable must be unique, in the same way that in a SSA-form program the definition that reaches a use is unique. We ensure this

property via special instructions called σ -functions. The σ -functions are the dual of ϕ -functions, performing a parallel assignment depending on the execution path taken. The assignment

$$(l^1 : v_1^1, \dots, l^q : v_1^q) = \sigma(v_1) \parallel \dots \parallel (l^1 : v_m^1, \dots, l^q : v_m^q) = \sigma(v_m)$$

represents m σ -functions that assign to each variable v_i^j the value in v_i if control flows into block l^j . These assignments happen in parallel, i.e., the m σ -functions encapsulate m parallel copies. Also, notice that variables live in different branch targets are given different names by the σ -function that ends that basic block.

14.2.5 Propagating Information Forwardly and Backwardly

Let us consider a unidirectional forward (resp. backward) PLV problem E_{dense}^{ssi} stated as a set of equations $[v]^p \sqsubseteq F_v^{s,p}([v_1]^s, \dots, [v_n]^s)$ (or equivalently $[v]^p = [v]^p \wedge F_v^{s,p}([v_1]^s, \dots, [v_n]^s)$) for every variable v , each program point p , and each $s \in preds(p)$ (resp. $s \in succs(p)$). To simplify the discussion, any ϕ -function (resp. σ -function) is seen as a set of copies, one per predecessor (resp. successor), which leads to as many constraints. In other words, a ϕ -function such as $p : a = \phi(a_1 : l^1, \dots, a_m : l^m)$, gives us n constraints such as $[a]^p \sqsubseteq F_v^{l^j,p}([a_1]^{l^j}, \dots, [a_n]^{l^j})$, which usually simplifies into $[a]^p \sqsubseteq [a_j]^{l^j}$. This last can be written equivalently into the classical meet $[a]^p \sqsubseteq \bigwedge_{l^j \in preds(p)} [a_j]^{l^j}$ used in Chapter 8. Similarly, a σ -function $(l^1 : a_1, \dots, l^m : a_m) = \sigma(p : a)$ after program point p yields n constraints such as $[a_j]^{l^j} \sqsubseteq F_v^{p,l^j}([a_1]^p, \dots, [a_n]^p)$, which usually simplifies into $[a_j]^{l^j} \sqsubseteq [a]^p$. Given a program that fulfills the SSI property for E_{dense}^{ssi} and the set of transfer functions F_v^s , we show here how to build an equivalent sparse constrained system.

Definition 2 (SSI constrained system). Consider that a program in SSI form gives us a constraint system that associates with each variable v the constraints $[v]^p \sqsubseteq F_v^{s,p}([v_1]^s, \dots, [v_n]^s)$. We define a system of sparse equations E_{sparse}^{ssi} as follows:

- For each instruction at a program point i that defines (resp. uses) a variable v , we let $a \dots z$ be its set of used (resp. defined) variables. Because of the LINK property, $F_v^{s,p}$ (that we will denote F_v^i from now) depends only on some $[a]^s \dots [z]^s$. Thus, there exists a function G_v^i defined as the restriction of F_v^i on $\mathcal{L}_a \times \dots \times \mathcal{L}_z$, i.e., informally “ $G_v^i([a], \dots, [z]) = F_v^i([v_1], \dots, [v_n])$ ”.
- The sparse constrained system associates with each variable v , and each definition (resp. use) point i of v , the corresponding constraint $[v] \sqsubseteq G_v^i([a], \dots, [z])$ where a, \dots, z are used (resp. defined) at i .

Algorithm 14.1: Backward propagation engine under SSI

```

1 function back_propagate(transfer_functions  $\mathcal{G}$ )
2   worklist =  $\emptyset$ 
3   foreach  $v \in \text{vars}$ :  $[v] = \top$ 
4   foreach  $i \in \text{insts}$ :  $\text{worklist} += i$ 
5   while  $\text{worklist} \neq \emptyset$ :
6     let  $i \in \text{worklist}$ ;  $\text{worklist} -= i$ 
7     foreach  $v \in i.\text{uses}()$ :
8        $[v]_{\text{new}} = [v] \wedge G_v^i([i.\text{def}()])$ 
9       if  $[v] \neq [v]_{\text{new}}$ :
10         $\text{worklist} += v.\text{def}()$ 
11         $[v] = [v]_{\text{new}}$ 
```

Algorithm 14.2: Forward propagation engine under SSI

```

1 function forward_propagate(transfer_functions  $\mathcal{G}$ )
2   worklist =  $\emptyset$ 
3   foreach  $v \in \text{vars}$ :  $[v] = \top$ 
4   foreach  $i \in \text{insts}$ :  $\text{worklist} += i$ 
5   while  $\text{worklist} \neq \emptyset$ :
6     let  $i \in \text{worklist}$ ;  $\text{worklist} -= i$ 
7     foreach  $v \in i.\text{def}()$ :
8        $[v]_{\text{new}} = G_v^i([i.\text{uses}()])$ 
9       if  $[v] \neq [v]_{\text{new}}$ :
10         $\text{worklist} += v.\text{uses}()$ 
11         $[v] = [v]_{\text{new}}$ 
```

The propagation engine discussed in Chapter 8 sends information forwardly along the def-use chains naturally formed by the SSA-form program. If a given program fulfills the SSI property for a backward analysis, we can use a very similar propagation algorithm to communicate information backwardly. Figure 14.1 shows a worklist algorithm that propagates information backwardly. A slightly modified version of this algorithm, seen in Figure 14.2 propagates information forwardly. If necessary, these algorithms can be made control-flow sensitive, like Algorithm 8.1.

Still we should outline a quite important subtlety that appears in line 8 of algorithms 14.1 and 14.2. $[v]$ appears (and has to) on the right hand side of the assignment for Algorithm 14.1 while it does not for Algorithm 14.2. This comes from the asymmetry of our SSI form that ensures (for practical purpose only as we will explain soon) the Static Single Assignment property but not the Static Single Use (SSU) property. If we have several uses of the same variable, then the sparse backward constraint system will have several inequations – one per variable use – with the same left-hand-side. Technically this is the reason why we manipulate a constraint system (system with inequations) and not an equation system as in Chapter 8. Both systems can be solved ² using a scheme known as

² in the ideal world with monotone framework and lattice of finite height

chaotic iteration such as the worklist algorithm we provide here. The slight and important difference for a constraint system as opposed to an equation system, is that one needs to meet $G_v^i(\dots)$ with the old value of $[v]$ to ensure the monotonicity of the consecutive values taken by $[v]$. It would be still possible to enforce the SSU property, in addition to the SSA property, of our intermediate representation at the expenses of adding more ϕ -functions and σ -functions. However, this guarantee is not necessary to every sparse analysis. The dead-code elimination problem illustrates well this point: for a program under SSA form, replacing G_v^i in Algorithm 14.1 by “*i is a useful instruction or one of its definitions is marked as useful*” leads to the standard SSA-based dead-code elimination algorithm. The sparse constraint system does have several equations (one per variable use) for the same left-hand-side (one for each variable). It is not necessary to enforce the SSU property in this instance of dead-code elimination, and doing so would lead to a less efficient solution, in terms of compilation time and memory consumption. In other words, a code under SSA form fulfills the SSI property for dead-code elimination.

14.2.6 Examples of sparse data-flow analyses

As we have mentioned before, many data-flow analyses can be classified as PLV problems. In this section we present some meaningful examples.

Range analysis revisited

We start this section revising this chapter’s initial example of data-flow analysis, given in Figure 14.1. A range analysis acquires information from either the points where variables are defined, or from the points where variables are tested. In Figure 14.2 we know that i must be bound to the interval $[0, 0]$ immediately after program point l_1 . Similarly, we know that this variable is upper bounded by 100 at entry of program point l_4 , due to the conditional test that happens before. Therefore, in order to achieve the SSI property, we should split the live ranges of variables at their definition points, or at the conditionals where they are used. Figure 14.2(a) shows the original example after live range splitting. In order to ensure the SSI property in this example, the live range of variable i must be split at its definition, and at the conditional test. The live range of s , on the other hand, must be split only at its definition point, as it is not used in the conditional. Splitting at conditionals is done via σ -functions. The representation that we obtain by splitting live ranges at definitions and conditionals is called the Extended Static Single Assignment (e-SSA) form. Figure 14.2(b) shows the result that we obtain by running the range analysis on this intermediate representation. This solution assigns to each variable a unique range interval.

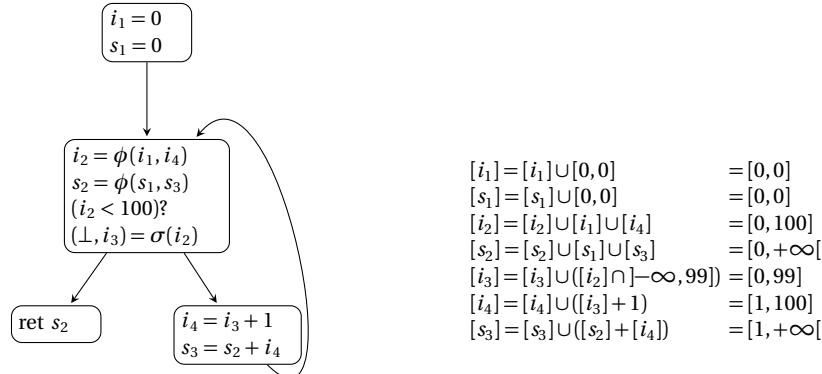


Fig. 14.2 (a) The example from Figure 14.1 after live range splitting. (b) A solution to this instance of the range analysis problem.

Class Inference

Some dynamically typed languages, such as Python, JavaScript, Ruby or Lua, represent objects as tables containing methods and fields. It is possible to improve the execution of programs written in these languages if we can replace these simple tables by actual classes with virtual tables. A class inference engine tries to assign a class to a variable v based on the ways that v is defined and used. The Python program in Figure 14.3(a) illustrates this optimization. Our objective is to infer the correct suite of methods for each object bound to variable v . Figure 14.3(b) shows the control-flow graph of the program, and Figure 14.3(c) shows the results of a dense implementation of this analysis. Because type inference is a backward analysis that extracts information from use sites, we split live ranges at these program points, and rely on σ -functions to merge them back, as we see in Figure 14.3(e). The use-def chains that we derive from the program representation lead naturally to a constraint system, which we show in Figure 14.3(f), where $[v_j]$ denotes the set of methods associated with variable v_j . A fix-point to this constraint system is a solution to our data-flow problem. This instance of class inference is a Partitioned Variable Problem (PVP)³, because the data-flow information associated with a variable v can be computed independently from the other variables.

Null pointer analysis

The objective of null pointer analysis is to determine which references may hold null values. This analysis allows compilers to remove redundant null-exception

³ Actually class inference is no more a PVP as soon as we want to propagate the information through copies.

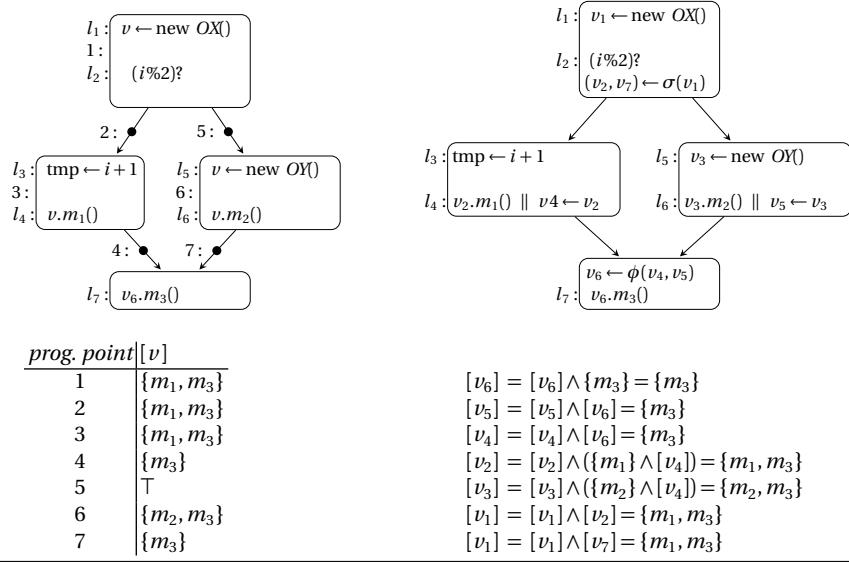


Fig. 14.3 Class inference analysis as an example of backward data-flow analysis that takes information from the uses of variables.

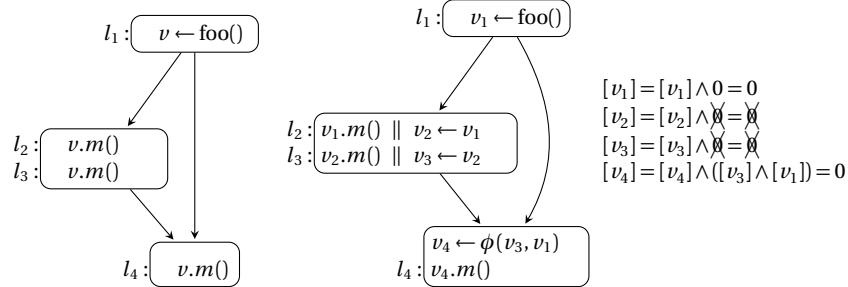


Fig. 14.4 Null pointer analysis as an example of forward data-flow analysis that takes information from the definitions and uses of variables (0 represents the fact that the pointer is possibly null, \emptyset if it cannot be).

tests and helps developers to find null pointer dereferences. Figure 14.4 illustrates this analysis. Because information is both produced at definition but also at use sites, we split live ranges after each variable is used, as we show in Figure 14.4(b). For instance, we know that v_2 cannot be null and hence the call $v_2.m()$ cannot result in a null pointer dereference exception, otherwise an exception would have been thrown during the invocation $v_1.m()$. On the other hand, in Figure 14.4(c) we notice that the state of v_4 is the meet of the state of v_3 , definitely not-null, and the state of v_1 , possibly null, and we must conservatively assume that v_4 may be null.

14.3 Construction and Destruction of the Intermediate Program Representation

In the previous section we have seen how the static single information property gives the compiler the opportunity to solve a data-flow problem sparsely. However, we have not yet seen how to convert a program to a format that provides the SSI property. This is a task that we address in this section, via the three-steps algorithm from Section 14.3.2.

14.3.1 Splitting strategy

A *live range splitting strategy* $\mathcal{P}_v = I_\uparrow \cup I_\downarrow$ over a variable v consists of a set of “oriented” program points. We let I_\downarrow denote a set of points i with forward direction. Similarly, we let I_\uparrow denote a set of points i with backward direction. The live-range of v must be split at least at every point in \mathcal{P}_v . Going back to the examples from Section 14.2.6, we have the live range splitting strategies enumerated below. The list in Figure 14.5 gives further examples of live range splitting strategies. Corresponding references are given in the last section of this chapter.

- Range analysis, is a forward analysis that takes information from points where variables are defined, and conditional tests that use these variables. For instance, in Figure 14.2, we have that $\mathcal{P}_i = \{l_1, \text{Out}(l_3), l_4\}_\downarrow$ where $\text{Out}(l_i)$ denotes the exit of l_i (i.e., the program point immediately after l_i), and that $\mathcal{P}_s = \{l_2, l_5\}_\downarrow$.
- Class inference is a backward analysis that takes information from the uses of variables; thus, for each variable, the live-range splitting strategy is characterized by the set Uses_\uparrow where Uses is the set of use points. For instance, in Figure 14.3(d), we have that $\mathcal{P}_v = \{l_4, l_6, l_7\}_\uparrow$.
- Null pointer analysis takes information from definitions and uses and propagates this information forwardly. For instance, in Figure 14.4, we have that $\mathcal{P}_v = \{l_1, l_2, l_3, l_4\}_\downarrow$.

The algorithm SSIfy in Figure 14.6 implements a live range splitting strategy in three steps. Firstly, it splits live ranges, inserting new definitions of variables into the program code. Secondly, it renames these newly created definitions; hence, ensuring that the live ranges of two different re-definitions of the same variable do not overlap. Finally, it removes dead and non-initialized definitions from the program code. We describe each of these phases in the rest of this section.

Client	Splitting strategy \mathcal{P}
Alias analysis, reaching definitions cond. constant propagation	$Defs_{\downarrow}$
Partial Redundancy Elimination	$Defs_{\downarrow} \cup LastUses_{\uparrow}$
ABCD, taint analysis, range analysis	$Defs_{\downarrow} \cup Out(Conds)_{\downarrow}$
Stephenson's bitwidth analysis	$Defs_{\downarrow} \cup Out(Conds)_{\downarrow} \cup Uses_{\uparrow}$
Mahlke's bitwidth analysis	$Defs_{\downarrow} \cup Uses_{\uparrow}$
An's type inference, Class inference	$Uses_{\uparrow}$
Hochstadt's type inference	$Uses_{\uparrow} \cup Out(Conds)_{\uparrow}$
Null-pointer analysis	$Defs_{\downarrow} \cup Uses_{\downarrow}$

Fig. 14.5 Live range splitting strategies for different data-flow analyses. $Defs$ (resp. $Uses$) denotes the set of instructions that define (resp. use) the variable. $Conds$ denotes the set of instructions that apply a conditional test on a variable; $Out(Conds)$ denotes the exits of the corresponding basic blocks; $LastUses$ denotes the set of instructions where a variable is used, and after which it is no longer live.

```

1 function SSIify(var  $v$ , Splitting_Strategy  $\mathcal{P}_v$ )
2   split( $v$ ,  $\mathcal{P}_v$ )
3   rename( $v$ )
4   clean( $v$ )

```

Fig. 14.6 Split the live ranges of v to convert it to SSI form

14.3.2 Splitting live ranges

In order to implement \mathcal{P}_v we must split the live ranges of v at each program point listed by \mathcal{P}_v . However, these points are not the only ones where splitting might be necessary. As we have pointed out in Section 14.2.4, we might have, for the same original variable, many different sources of information reaching a common program point. For instance, in Figure 14.1, there exist two definitions of variable i , e.g., l_1 and l_4 , that reach the use of i at l_3 . The information that flows forward from l_1 and l_4 collides at l_3 , the loop entry. Hence the live-range of i has to be split immediately before l_3 – at $In(l_3) -$, leading, in our example, to a new definition i_1 . In general, the set of program points where information collides can be easily characterized by the notion of join sets and iterated dominance frontier (DF^+) seen in Chapter 4. Similarly, split sets created by the backward propagation of information can be over-approximated by the notion of *iterated post-dominance frontier* (pDF^+), which is the dual of DF^+ . That is, the post-dominance frontier is the dominance frontier in a CFG where direction of edges have been reversed. Note that, just as the notion of dominance requires the existence of a unique entry node that can reach every CFG node, the notion of post dominance requires the existence of a unique exit node reachable by any

```

1 function split(var  $v$ , Splitting_Strategy  $\mathcal{P}_v = I_{\downarrow} \cup I_{\uparrow}$ )
2     "compute the set of split nodes"
3      $S_{\downarrow} = \emptyset$ 
4     foreach  $i \in I_{\downarrow}$ :
5         if  $i$ .is_join:
6             foreach  $e \in \text{incoming\_edges}(i)$ :
7                  $S_{\downarrow} = S_{\downarrow} \cup \text{Out}(pDF^+(e))$ 
8             else:
9                  $S_{\downarrow} = S_{\downarrow} \cup \text{Out}(pDF^+(i))$ 
10             $S_{\downarrow} = \emptyset$ 
11            foreach  $i \in S_{\downarrow} \cup \text{Defs}(v) \cup I_{\downarrow}$ :
12                if  $i$ .is_branch:
13                    foreach  $e \in \text{outgoing\_edges}(i)$ 
14                         $S_{\downarrow} = S_{\downarrow} \cup \text{In}(DF^+(e))$ 
15                else:
16                     $S_{\downarrow} = S_{\downarrow} \cup \text{In}(DF^+(i))$ 
17             $S = \mathcal{P}_v \cup S_{\downarrow} \cup S_{\uparrow}$ 
18            "Split live range of  $v$  by inserting  $\phi$ ,  $\sigma$ , and copies"
19            foreach  $i \in S$ :
20                if  $i$  does not already contain any definition of  $v$ :
21                    if  $i$ .is_join: insert " $v = \phi(v, \dots, v)$ " at  $i$ 
22                    elseif  $i$ .is_branch: insert " $(v, \dots, v) = \sigma(v)$ " at  $i$ 
23                    else: insert a copy " $v = v$ " at  $i$ 

```

Fig. 14.7 Live range splitting. We use $\text{In}(l)$ to denote a program point immediately before l , and $\text{Out}(l)$ to denote a program point immediately after l .

CFG node. For control-flow graphs that contain several exit nodes or loops with no exit, we can ensure the single-exit property by creating a dummy common exit node and inserting some never-taken exit edges into the program.

Figure 14.7 shows the algorithm that we use to create new definitions of variables. This algorithm has three main phases. First, in lines 3-9 we create new definitions to split the live ranges of variables due to backward collisions of information. These new definitions are created at the iterated post-dominance frontier of points that originate information. If a program point is a join node, then each of its predecessors will contain the live range of a different definition of v , as we ensure in lines 6-7 of our algorithm. Notice that these new definitions are not placed parallel to an instruction, but in the region immediately after it, which we denote by $\text{Out}(\dots)$. In lines 10-16 we perform the inverse operation: we create new definitions of variables due to the forward collision of information. Our starting points S_{\downarrow} , in this case, include also the original definitions of v , as we see in line 11, because we want to stay in SSA form in order to have access to a fast liveness check as described in Chapter 9. Finally, in lines 17-23 we actually insert the new definitions of v . These new definitions might be created by σ functions (due to \mathcal{P}_v or to the splitting in lines 3-9); by ϕ -functions (due to \mathcal{P}_v or to the splitting in lines 10-16); or by parallel copies.

```

1 function rename(var  $v$ )
2   "Compute use-def & def-use chains"
3   "We consider here that stack.peek() =  $\perp$  if stack.isEmpty(),"
4   "and that  $def(\perp) = entry$ "
5   stack =  $\emptyset$ 
6   foreach CFG node  $n$  in dominance order:
7     if exists  $v = \phi(v : l^1, \dots, v : l^q)$  in  $In(n)$ :
8       stack.set_def( $v = \phi(v : l^1, \dots, v : l^q)$ )
9     foreach instruction  $u$  in  $n$  that uses  $v$ :
10      stack.set_use( $u$ )
11      if exists instruction  $d$  in  $n$  that defines  $v$ :
12        stack.set_def( $d$ )
13        foreach instruction (...) =  $\sigma(v)$  in  $Out(n)$ :
14          stack.set_use(...) =  $\sigma(v)$ )
15      if exists  $(v : l^1, \dots, v : l^q) = \sigma(v)$  in  $Out(n)$ :
16        foreach  $v : l^i = v$  in  $(v : l^1, \dots, v : l^q) = \sigma(v)$ :
17          stack.set_def( $v : l^i = v$ )
18        foreach  $m$  in successors( $n$ ):
19          if exists  $v = \phi(\dots, v : l^n, \dots)$  in  $In(m)$ :
20            stack.set_use( $v = v : l^n$ )
21 function stack.set_use(instruction  $inst$ ):
22   while  $def(stack.peek())$  does not dominate  $inst$ : stack.pop()
23    $v_i = stack.peek()$ 
24   replace the uses of  $v$  by  $v_i$  in  $inst$ 
25   if  $v_i \neq \perp$ : set  $Uses(v_i) = Uses(v_i) \cup inst$ 
26   function stack.set_def(instruction  $inst$ ):
27     let  $v_i$  be a fresh version of  $v$ 
28     replace the defs of  $v$  by  $v_i$  in  $inst$ 
29     set  $Def(v_i) = inst$ 
30     stack.push( $v_i$ )

```

Fig. 14.8 Versioning

14.3.3 Variable Renaming

The algorithm in Figure 14.8 builds def-use and use-def chains for a program after live range splitting. This algorithm is similar to the classic algorithm used to rename variables during the SSA construction that we saw in Chapter 3. To rename a variable v we traverse the program's dominance tree, from top to bottom, stacking each new definition of v that we find. The definition currently on the top of the stack is used to replace all the uses of v that we find during the traversal. If the stack is empty, this means that the variable is not defined at this point. The renaming process replaces the uses of undefined variables by \perp (line 3). We have two methods, `stack.set_use` and `stack.set_def` that build the chains of relations between variables. Notice that sometimes we must rename a single use inside a ϕ -function, as in lines 19-20 of the algorithm. For simplicity we consider this single use as a simple assignment when calling `stack.set_use`, as one can see in line 20. Similarly, if we must rename a single definition inside

```

1 clean(var  $v$ )
2   let web = { $v_i$  |  $v_i$  is a version of  $v$ }
3   let defined =  $\emptyset$ 
4   let active = { $inst$  |  $inst$  actual instruction and  $web \cap inst.defs \neq \emptyset$ }
5   while  $\exists inst \in active$  s.t.  $web \cap inst.defs \setminus defined \neq \emptyset$ :
6     foreach  $v_i \in web \cap inst.defs \setminus defined$ :
7       active = active  $\cup$  Uses( $v_i$ )
8       defined = defined  $\cup$  { $v_i$ }
9   let used =  $\emptyset$ 
10  let active = { $inst$  |  $inst$  actual instruction and  $web \cap inst.uses \neq \emptyset$ }
11  while  $\exists inst \in active$  s.t.  $inst.uses \setminus used \neq \emptyset$ :
12    foreach  $v_i \in web \cap inst.uses \setminus used$ :
13      active = active  $\cup$  Def( $v_i$ )
14      used = used  $\cup$  { $v_i$ }
15  let live = defined  $\cap$  used
16  foreach non actual  $inst \in Def(web)$ :
17    foreach  $v_i$  operand of  $inst$  s.t.  $v_i \notin live$ :
18      replace  $v_i$  by  $\perp$ 
19    if  $inst.defs = \{\perp\}$  or  $inst.uses = \{\perp\}$ 
20      remove  $inst$ 

```

Fig. 14.9 Dead and undefined code elimination. Original instructions not inserted by split are called *actual* instruction. We let $inst.defs$ denote the set of variable(s) defined by $inst$, and $inst.uses$ denote the set of variables used by $inst$.

a σ -function, then we treat it as a simple assignment, like we do in lines 15-16 of the algorithm.

14.3.4 Dead and Undefined Code Elimination

Just as Algorithm 3.7, the algorithm in Figure 14.9 eliminates ϕ -functions and parallel copies that define variables not actually used in the code. By symmetry, it also eliminates σ -functions and parallel copies that use variables not actually defined in the code. We mean by “actual” instructions, those instructions that already existed in the program before we transformed it with split. In line 3 we let “web” be the set of versions of v , so as to restrict the cleaning process to variable v , as we see in lines 4-6 and lines 10-12. The set “active” is initialized to actual instructions in line 4. Then, during the loop in lines 5-8 we add to active ϕ -functions, σ -functions, and copies that can reach actual definitions through use-def chains. The corresponding version of v is then marked as *defined* (line 8). The next loop, in lines 11-14 performs a similar process, this time to add to the active set instructions that can reach actual uses through def-use chains. The corresponding version of v is then marked as *used* (line 14). Each non live variable (see line 15), i.e., either undefined or dead (non used) is replaced by \perp in all ϕ , σ , or copy functions where it appears in. This is done by

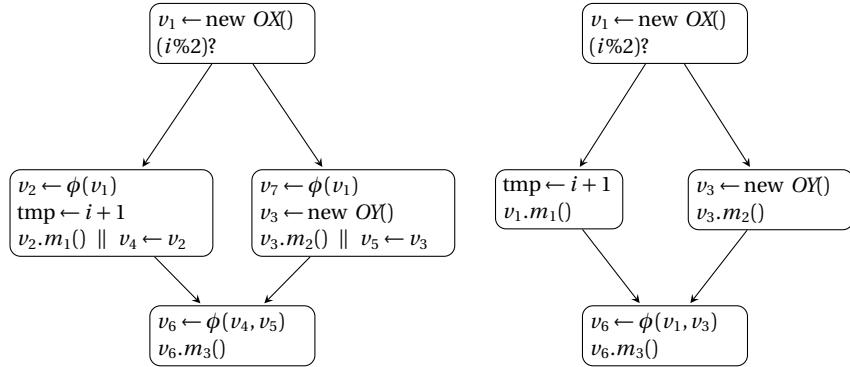


Fig. 14.10 (a) implementing σ -functions via single arity ϕ -functions; (b) getting rid of copies and σ -functions.

lines 15-18. Finally every useless ϕ , σ , or copy functions are removed by lines 19-20.

14.3.5 Implementation Details

Implementing σ -functions:

The most straightforward way to implement σ -functions, in a compiler that already supports the SSA form, is to represent them by ϕ -functions. In this case, the σ -functions can be implemented as single arity ϕ -functions. As an example, Figure 14.10(a) shows how we would represent the σ -functions of Figure 14.3(d). If l is a branch point with n successors that would contain a σ -function ($l^1 : v_1, \dots, l^n : v_n = \sigma(v)$), then, for each successor l^j of l , we insert at the beginning of l^j an instruction $v_j = \phi(l^j : v)$. Notice that it is possible that l^j already contains a ϕ -function for v . This case happens when the control-flow edge $l \rightarrow l^j$ is *critical*. A critical edge links a basic block with several successors to a basic block with several predecessors. If l^j already contains a ϕ -function $v' = \phi(\dots, v_j, \dots)$, then we rename v_j to v .

SSI Destruction:

Traditional instruction sets do not provide ϕ -functions nor σ -functions. Thus, before producing an executable program, the compiler must implement these instructions. We have already seen in Chapter 3 how to replace ϕ -functions with actual assembly instructions; however, now we must also replace σ -functions and parallel copies. A simple way to eliminate all the σ -functions and parallel

copies is via copy-propagation. In this case, we copy-propagate the variables that these special instructions define. As an example, Figure 14.10(b) shows the result of copy folding applied on Figure 14.10(a).

14.4 Further Reading

The monotone data-flow framework is an old ally of compiler writers. Since the work of pionners like Prosser [?], Allen [?, ?], Kildall [?] and Hecht [?], data-flow analyses such as reaching definitions, available expressions and liveness analysis have made their way into the implementation of virtually every important compiler. Many compiler textbooks describes the theoretic basis of the notions of lattice, monotone data-flow framework and fixed points. For a comprehensive overview of these concepts, including algorithms and formal proofs, we refer the interested reader to Nielson *et al.*'s book [208] on static program analysis.

The original description of the intermediate program representation known as Static Single Information form was given by Ananian in his Master's thesis [13]. The notation for σ -functions that we use in this chapter was borrowed from Ananian's work. The SSI program representation was subsequently revisited by Jeremy Singer in his PhD thesis [250]. Singer proposed new algorithms to convert programs to SSI form, and also showed how this program representation could be used to handle truly bidirectional data-flow analyses. We did not discuss bidirectional data-flow problems, but the interested reader can find examples of such analyses in Khedker *et al.*'s work [165]. Working on top of Ananian's and Singer's work, Boissinot *et al.* [39] have proposed a new algorithm to convert a program to SSI form. Boissinot *et al.* have also separated the SSI program representation in two flavors, which they call *weak* and *strong*. Tavares *et al.* [268] have extended the literature on SSI representations, defining building algorithms and giving formal proofs that these algorithms are correct. The presentation that we use in this chapter is mostly based on Tavares *et al.*'s work.

There exist other intermediate program representations that, like the SSI form, make it possible to solve some data-flow problems sparsely. Well-known among these representations is the *Extended Static Single Assignment* form, introduced by Bodik *et al.* to provide a fast algorithm to eliminate array bound checks in the context of a JIT compiler [37]. Another important representation, which supports data-flow analyses that acquire information at use sites, is the *Static Single Use* form (SSU). As uses and definitions are not fully symmetric (the live-range can “traverse” a use while it cannot traverse a definition) there exists different variants of SSU (eg. [225, 130, 185]). For instance, the “strict” SSU form enforces that each definition reaches a single use, whereas SSI and other variations of SSU allow two consecutive uses of a variable on the same path. All these program representations are very effective, having seen use in a number

of implementations of flow analyses; however, they only fit specific data-flow problems.

The notion of *Partitioned Variable Problem* (PVP) was introduced by Zadeck, in his PhD dissertation [297]. Zadeck proposed fast ways to build data-structures that allow one to solve these problems efficiently. He also discussed a number of data-flow analyses that are partitioned variable problems. There are data-flow analyses that do not meet the Partitioned Lattice per Variable property. Noticeable examples include abstract interpretation problems on relational domains, such as Polyhedrons [85], Octagons [201] and Pentagons [186].

In terms of data-structures, the first, and best known method proposed to support sparse data-flow analyses is Choi *et al.*'s *Sparse Evaluation Graph* (SEG) [73]. The nodes of this graph represent program regions where information produced by the data-flow analysis might change. Choi *et al.*'s ideas have been further expanded, for example, by Johnson *et al.*'s *Quick Propagation Graphs* [154], or Ramalingam's *Compact Evaluation Graphs* [231]. Nowadays we have efficient algorithms that build such data-structures [222, 223, 152]. These data-structures work best when applied on partitioned variable problems.

As opposed to those approaches, the solution promoted by this chapter consists in an intermediate representation (IR) based evaluation graph, and has advantages and disadvantages when compared to the data-structure approach. The intermediate representation based approach has two disadvantages, which we have already discussed in the context of the standard SSA form. First it has to be maintained and at some point destructured. Second, because it increases the number of variables, it might add some overhead to analyses and transformations that do not require it. On the other hand, IR based solutions to sparse data-flow analyses have many advantages over data-structure based approaches. For instance, an IR allows concrete or abstract interpretation. Solving any coupled data-flow analysis problem along with a SEG was mentioned in [73] as an opened problem. However, as illustrated by the conditional constant propagation problem described in Chapter 8, coupled data-flow analysis can be solved naturally in IR based evaluation graphs. Last, SSI is compatible with SSA extensions such as gated-SSA described in Chapter 18 which allows demand driven interpretation.

The data-flow analyses discussed in this chapter are well-known in the literature. Class inference was used by Chambers *et al.* in order to compile Self programs more efficiently [69]. Nanda and Sinha have used a variant of null-pointer analysis to find which method dereferences may throw exceptions, and which may not [206]. Ananian [13], and later Singer [250], have showed how to use the SSI representation to do partial redundancy elimination sparsely. In addition to being used to eliminate redundant array bound checks [37], the e-SSA form has been used to solve Taint Analysis [239], and range analysis [265, 127]. Stephenson *et al.* [259] described a bit-width analysis that is both forward, and backwards, taking information from definitions, uses and conditional tests. For another example of bidirectional bitwidth analysis, see Mahlke *et al.*'s algo-

rithm [190]. The type inference analysis that we mentioned in Figure 14.5 was taken from Hochstadt *et al.*'s work [270].

CHAPTER 15

Hashed SSA form: HSSA

M. Mantione
F. Chow

Progress: 85%



Minor polishing in progress

★
★
★

15.1 Introduction

Hashed SSA (or in short HSSA), is an SSA extension that can effectively represent how aliasing relations affect a program in SSA form. It works equally well for aliasing among scalar variables and, more generally, for indirect load and store operations on arbitrary memory locations. This allows the application of all common SSA based optimizations to perform uniformly both on local variables and on external memory areas.

It should be noted, however, that HSSA is a technique useful for representing aliasing effects, but not for detecting aliasing. For this purpose, a separate alias analysis pass must be performed, and the effectiveness of HSSA will be influenced by the accuracy of this analysis.

The following sections explain how HSSA works. Initially, *given* aliasing information, we will see how to represent them in SSA form for scalar variables. Then we will introduce a technique that reduces the overhead of the above representation, avoiding an explosion in the number of SSA versions for aliased variables. Subsequently we will represent indirect memory operations on external memory areas as operations on "virtual variables" in SSA form, which will be handled uniformly with scalar (local) variables. Finally we will apply global

value numbering (GVN) to all of the above, obtaining the so called Hashed SSA form ¹.

15.2 SSA and aliasing: μ and χ -functions

★
★

Aliasing occurs inside a compilation unit when a single one single storage location (that contains a value) can be potentially accessed through different program "variables". This can happen in one of the four following ways:

- First, when two or more storage locations partially overlap. This, for instance, happens with the C "union" construct, where different parts of a program can access the same storage location under different names.
- Second, when a local variable is referred by a pointer used in an indirect memory operation. In this case the variable can be accessed in two ways: *directly*, through the variable name, and *indirectly*, through the pointer that holds its address.
- Third, when the address of a local variable is passed to a function, which in turn can then access the variable indirectly.
- Finally, storage locations with global scope can obviously be accessed by different functions. In this case every function call can potentially access every global location, unless the compiler uses global optimizations techniques where every function is analyzed before the actual compilation takes place.

The real problem with aliasing is that these different accesses to the same program variable are difficult to predict. Only in the first case (explicitly overlapping locations) the compiler has full knowledge of when each access takes place. In all the other cases (indirect accesses through the address of the variable) the situation becomes more complex, because the access depends on the address that is effectively stored in the variable used in the indirect memory operation. This is a problem because every optimization pass is concerned with the actual value that is stored in every variable, and when those values are used. If variables can be accessed in unpredictable program points, the only safe option for the compiler is to handle them as "volatile" and avoid performing optimizations on them, which is not desirable.

Intuitively, in the presence of aliasing the compiler could try to track the values of variable addresses inside other variables (and this is exactly what HSSA does), but the formalization of this process is not trivial. The first thing that is needed is a way to model the *effects* of aliasing on a program in SSA form. To do this, assuming that we have already performed alias analysis, we must formally define the effects of indirect definitions and uses of variables. Particularly,

¹ The name *Hashed* SSA comes from the use of hashing in value-numbering

each definition can be a "*MustDef*" operand in the direct case, or a "*MayDef*" operand in the indirect case. We will represent MayDef through the use of χ -functions. Similarly, uses can be "*MustUse*" or "*MayUse*" operands (respectively in the direct and indirect case), and we will represent MayUse through the use of μ -functions. The semantic of μ and χ operators can be illustrated through the C like example of Figure 15.1 where $*p$ represents an indirect access with address p . Obviously the argument of the μ -operator is the potentially used variable. Less obviously, the argument to the χ -operator is the assigned variable itself. This expresses the fact that the χ -operator only *potentially* modifies the variable, so the original value could "flow through" it.

The use of μ and χ -operators does not alter the complexity of transforming a program in SSA form. All that is necessary is a pre-pass that inserts them in the program. Ideally, a μ and a χ should be placed *in parallel*^{*} to the instruction that led to its insertion. Parallel instructions are represented in Figure 15.1 using the notation introduced in Section 14.2.4. Still, practical implementations may choose to insert μ and χ s before or after the instructions that involve aliasing. Particularly, μ -functions could be inserted immediately *before* the involved statement or expression, and χ -operators immediately *after* it. This distinction allows us to model call effects correctly: the called function appears to potentially use the values of variables before the call, and the potentially modified values appear after the call.

Thanks to the systematic insertion of μ -functions and χ -functions, an assignment of any scalar variable can be safely considered dead if it is not marked live by a standard SSA based dead-code elimination. In our running example of Figure 15.1(c), the potential side effect of any assignment to i , represented through the μ -function at the return, allows detecting that the assignment to i_4 is not dead. The assignment of value 2 to i would have been considered as dead in the absence of the function call to $f()$, which potentially uses it (detected through its corresponding μ -function).

<pre> 1 i = 2; 2 if j then 3 ... 4 else f(); 5 *p = 3; 6 i = 4; 7 return; </pre>	<pre> 1 i = 2; 2 if j then 3 ... 4 else f() $\mu(i)$; 5 *p = 3 i = $\chi(i)$; 6 i = 4; 7 return $\mu(i)$; </pre>	<pre> 1 i₁ = 2; 2 if j₁ then 3 ... 4 else f() $\mu(i_1)$; 5 *p = 3 i₂ = $\chi(i_1)$; 6 i₃ = $\phi(i_1, i_2)$; 7 i₄ = 4; 8 return $\mu(i_4)$; </pre>
(a) Initial C code	(b) After μ and χ insertion	(c) After ϕ insertion and versioning

Fig. 15.1 A program example where $*p$ might alias i , and function f might indirectly use i but not alter it

parallel instruction

*zero version
variable version
data flow
use-def chains*

15.3 Introducing “zero versions” to limit the explosion of the number of variables

*

While it is true that μ and χ insertion does not alter the complexity of SSA construction, applying it to a production compiler as described in the previous section would make working with code in SSA form terribly inefficient. This is because χ -operators cause an explosion in the number of variable values, inducing the insertion of new ϕ -functions which in turn create new variable versions. In practice the resulting IR, and especially the number of distinct variable versions, would be needlessly large. The biggest issue is that the SSA versions introduced by χ -operators are useless for most optimizations that deal with variable values. χ definitions adds uncertainty to the analysis of variables values: the actual value of a variable after a χ definition could be its original value, or it could be the one indirectly assigned by the χ .

Intuitively, the solution to this problem is to factor all variable versions that are considered *useless* together, so that no space is wasted to distinguish among them. We assign number 0 to this special variable version, and call it "*zero version*".

Our notion of useless versions relies on the concept of "*real occurrence of a variable*", which is an actual definition or use of a variable in the original program. Therefore, in SSA form, variable occurrences in μ , χ and ϕ -functions are not "real occurrences". In our example of Figure 15.1, i_2 have no real occurrence while i_1 and i_3 have. The idea is that variable versions that have no real occurrence do not influence the program output. Once the program is converted back from SSA form these variables are removed from the code. Since they do not directly appear in the code, and their value is usually unknown, distinguishing among them is almost pointless. For those reasons, we consider zero versions, versions of variables that have no real occurrence, and whose value comes from at least one χ -function (optionally through ϕ -functions). An equivalent, recursive definition is the following:

- The result of a χ has zero version if it has no real occurrence.
- If the operand of a ϕ has zero version, the ϕ result has zero version if it has no real occurrence.

Algorithm 15.1^{*} performs zero-version detection if only use-def chains^{*}, and not def-use chains, are available. A "HasRealOcc" flag is associated to each variable version, setting it to true whenever a real occurrence is met in the code. This can be done while constructing the SSA form. A list "NonZeroPhiList", initially empty, is also associated to each original program variable.

The time spent in the first iteration grows linearly with the number of variable versions, which in turn is proportional to the number of definitions and therefore to the code size. On the other hand, the while loop may, in the worst case, iterate as many times as the longest chain of contiguous ϕ assignments

Algorithm 15.1: Zero-version detection based on SSA use-def chains

```

1 foreach variable  $v$  do
2   foreach version  $v_i$  of  $v$  do
3     if  $\neg v_i.\text{HasRealOcc} \wedge v_i.\text{def.operator} = \chi$  then
4       |  $v_i.\text{version} = 0;$ 
5       if  $\neg v_i.\text{HasRealOcc} \wedge v_i.\text{def.operator} = \phi$  let  $V = v_i.\text{def.operands};$ 
6       if  $\forall v_j \in V, v_j.\text{HasRealOcc}$  then
7         |  $v_i.\text{HasRealOcc} = \text{true};$ 
8       else if  $\exists v_j \in V, v_j.\text{version} = 0$  then
9         |  $v_i.\text{version} = 0;$ 
10        else  $v.\text{NonZeroPhiList.add}(v_i)$ 
11      end
12       $changes = \text{true};$ 
13      while  $changes$   $changes = \text{false};$ 
14      foreach  $v_i \in v.\text{NonZeroPhiList}$  do
15        let  $V = v_i.\text{def.operands};$ 
16        if  $\forall v_j \in V, v_j.\text{HasRealOcc}$  then
17          |  $v_i.\text{HasRealOcc} = \text{true};$ 
18           $v.\text{NonZeroPhiList.remove}(v_i);$ 
19           $changes = \text{true};$ 
20        else if  $\exists v_j \in V, v_j.\text{version} = 0$  then
21          |  $v_i.\text{version} = 0;$ 
22           $v.\text{NonZeroPhiList.remove}(v_i);$ 
23           $changes = \text{true};$ 
24        end
25      end
26 end

```

*forward control-flow
graph
dead code elimination*

in the program. This bound can easily be reduced to the largest loop depth of the program by traversing the versions using a topological order of the forward control-flow graph.^{*} All in all, zero version detection in the presence of μ and χ -functions does not change the complexity of SSA construction in a significant way, while the corresponding reduction in the number of variable versions is definitely desirable.

This loss of information has almost no consequences on the effectiveness of subsequent optimization passes. Since variables with zero versions have uncertain values, not being able to distinguish them usually only slightly affects the quality of optimizations that operate on values. On the other hand, when performing sparse dead-code elimination^{*} along use-def chains, zero versions for which use-def chains have been broken must be assumed live. However this has no practical effect. Zero versions have no real occurrence, so there is no real statement that would be eliminated by the dead-code elimination pass if we could detect that a zero version occurrence is dead. There is only one case in dead-code elimination where the information loss is evident. A zero version can be used in a μ and defined in a χ , which in turn can have a real occurrence as argument. If the μ is eliminated by some optimization, the real occurrence used in the χ becomes dead but will not be detected as it, since the χ

is conservatively marked as non dead. This case is sufficiently rare in real world code that using zero versions is anyway convenient.

15.4 SSA and indirect memory operations: virtual variables

The technique described in the previous sections only apply to "regular" variables in a compilation unit, and not to arbitrary memory locations accessed indirectly. As an example, in Figure 15.1, μ , χ , and ϕ -functions have been introduced to keep track of i 's live-range, but $*p$ is not considered as an SSA variable. In other words, up to now, HSSA allows to apply SSA based optimizations to variables also when they are affected by aliasing, but memory access operations are still excluded from the SSA form.

This situation is far from ideal, because code written in current mainstream imperative languages (like C++, Java or C#) typically contains many operations on data stored in global memory areas. For instance, in C we can imagine to define bidimensional vector as a struct: "typedef struct {double x; double y;} point;", and then to have a piece of code that computes the modulus of a vector of address "p": "m = (p->x * p->x) + (p->y * p->y);". Looking at the code it is obvious that "x" is accessed twice but both accesses give the same value so the second access could be optimized away (the same for "y"). The problem is that "x" and "y" are not "regular" variables: "p" is a variable while "p->x" and "p->y" are indirect memory operations. Putting that code snipped in SSA form tells us that the value of "p" never changes, but it reveals nothing about the values stored in the locations "p->x" and "p->y". It is worth noting that operations on array elements suffer from the same problem.

The purpose of HSSA is to handle indirect memory operations just like accesses to scalar variables, and be able to apply all SSA based optimizations uniformly on both of them. To do this, in the general case, we assume that the code intermediate representation supports a *dereference* operator, which performs an indirect load (memory read) from the given address. This operator can be placed in expressions on both the left and right side of the assignment operator, and we will represent it with the usual "*" C language operator. We will then, for simplicity, represent indirect stores with the usual combination of dereference and assignment operators, like in C. Some examples of this notation are:

- $*p$: access memory at address p .
- $*(p+1)$: access memory at address $p+1$ (like to access an object field at offset 1).
- $**p$: double indirection.
- $*p = \text{expression}$: indirect store.

As we noted above indirect memory operations cannot be handled by the SSA construction algorithm because they are *operations* while SSA construction works renaming *variables*. What HSSA does is to represent the "target" locations of indirect memory operations with *virtual variables*. A virtual variable is an abstraction of a memory area and appears under HSSA thanks to the insertion of μ and χ -functions: indeed as any other variable, it may alias with an operand. The example of Figure 15.2 shows two possible forms after μ -function and χ -function insertion for the same code. In Figure 15.2(b), two virtual variables v^* and w^* are respectively associated to the memory expressions $*p$ and $*q$; v^* and w^* alias with all indirect memory operands (of lines 3, 5, 6, and 8). In Figure 15.2(c) v is associated to the base pointer of b and w to $b + 1$ respectively. As one can see, there exists many possible ways to chose virtual variables. The only discipline imposed by HSSA is that each indirect memory operand must be associated to a single virtual variable. So on one extreme, there would be one virtual variable for each indirect memory operation. *Assignment factoring* corresponds to make each virtual variable represents more than one single indirect memory operand. On the other extreme, the most factored HSSA form would have only one single virtual variable on the overall. In the example of Figure 15.2(c) we considered as given by alias analysis the non aliasing between b and $b + 1$, and choose two virtual variables to represent the two corresponding distinct memory locations. In the general case, virtual variables can obviously alias with one another as in Figure 15.2(b).

(a) Initial C code	(b) v and w alias with ops 3,5,7, and 8	(c) x alias with op 3; y with 5,7, and 8
<pre> 1 p = b; 2 q = b; 3 *p = ...; 4 p = p + 1; 5 ... = *p; 6 q = q + 1; 7 *p = ...; 8 ... = *q; </pre>	<pre> 1 p = b; 2 q = b; 3 *p = ... v* = $\chi(v^*)$ w* = $\chi(w^*)$; 4 p = p + 1; 5 ... = *p $\mu(v^*)$ $\mu(w^*)$; 6 q = q + 1; 7 *p = ... v* = $\chi(v^*)$ w* = $\chi(w^*)$; 8 ... = *q $\mu(v^*)$ $\mu(w^*)$; </pre>	<pre> 1 p = b; 2 q = b; 3 *p = ... x* = $\chi(x^*)$; 4 p = p + 1; 5 ... = *p $\mu(y^*)$; 6 q = q + 1; 7 *p = ... y* = $\chi(y^*)$; 8 ... = *q $\mu(y^*)$; </pre>

Fig. 15.2 Some virtual variables and their insertion depending on how they alias with operands.

To summarise, alias analysis, considered as a given, drives the choice of virtual variables². The insertion of μ -function and χ -function for virtual variables is performed in the same pass than for scalar variables. Also, when the program is converted into SSA form, virtual variables can be renamed just like scalar variables. Use-def relationships of virtual variables now represent use-def relationships of the memory locations accessed indirectly. Note that the zero version-

² I'm not sure what this means. Rephrase? (dnovillo)

ing technique can be applied unaltered to virtual variables. At this point, we can complete the construction of HSSA form by applying global value numbering.

15.5 GVN and indirect memory operations: HSSA

In the previous sections we sketched the foundations of a framework for dealing with aliasing and indirect memory operations in SSA form: we identified the effects of aliasing on local variables, introduced μ and χ -operators to handle them, applied zero versioning to keep the number of SSA versions acceptable, and defined virtual variables as a way to apply SSA also to memory locations accessed indirectly. However, HSSA is complete only once *Global Value Numbering*^{*} is applied to scalar and virtual variables, handling all of them uniformly (GVN). See Chapter 12).

GVN is normally used as a way to perform *redundancy elimination* which means removing redundant expression computations from a program, typically storing the result of a computation in a temporary variable and reusing it later instead of performing the computation again. As the name suggests, GVN works assigning a unique number to every expression in the program with the idea that expression identified by the same number are guaranteed to give the same result. This value number is obtained using a hash function represented here as $H(key)$. To identify identical expressions, each expression tree is hashed bottom up: as an example, for $p_1 + 1$, p_1 and 1 are replaced by their respective value numbers, then the expression is put into canonical form, and finally hashed. Which ends up to $H(+H(b), H(1))$, as $H(p_1) = H(b)$ in our example. If the program is in SSA form computing value numbers that satisfy the above property is straightforward: variables are defined only once, therefore two expressions that apply the same operator to the same SSA variables are guaranteed to give the same result and so can have the same value number. While computing global value numbers for code that manipulates scalar variables is beneficial because it can be used to implement redundancy elimination, applying GVN to virtual variables has additional benefits.

First of all, it can be used to determine when two address expressions compute the same address: this is guaranteed if they have the same global value number. In our example, p_2 and q_2 will have the same value number h_7 while p_1 will not. This allows indirect memory operands that have the same GVN for their address expressions and the same virtual variable version to become a single entity in the representation. This puts them in the same rank as scalar variables and allows the transparent application of SSA based optimizations on indirect memory operations. For instance, in the vector modulus computation described above, every occurrence of the expression "p->x" will always have the same GVN and therefore will be guaranteed to return the same value, allowing the compiler to emit code that stores it in a temporary register instead of per-

(a) Initial C code	(b) with one virtual variable	(c) HSSA statements
<pre> 1 p = b; 2 q = b; 3 *p = 3; 4 p = p + 1; 5 ...=*p; 6 q = q + 1; 7 *p = 4; 8 ...=*q; </pre>	<pre> 1 p1 = b; 2 q1 = b; 3 *p1 = 3 v1* = χ(v0*); 4 p2 = p1 + 1; 5 ...=*p2 μ(v1*); 6 q2 = q1 + 1; 7 *p2 = 4 v2* = χ(v1*); 8 ...=*q2 μ(v2*); </pre>	<pre> 1 h1 = h0; 2 h2 = h0; 3 h5 = h3 h4 = χ(H(v0*)); 4 h8 = h7; 5 ...=h10 μ(h4); 6 h11 = h7; 7 h14 = h12 h13 = χ(h4); 8 ...=h14 μ(h13); </pre>
(d) Hash Table		

key	hash	value
<i>b</i>	<i>h</i> ₀	<i>b</i>
<i>p</i> ₁	<i>h</i> ₁	<i>h</i> ₀
<i>q</i> ₁	<i>h</i> ₂	<i>h</i> ₀
3	<i>h</i> ₃	<i>const</i> (3)
<i>v</i> ₁ *	<i>h</i> ₄	<i>v</i> ₁ *
<i>ivar</i> (* <i>p</i> ₁ , <i>v</i> ₁ *) ~> <i>ivar</i> (* <i>h</i> ₀ , <i>v</i> ₁ *)	<i>h</i> ₅	<i>h</i> ₃
1	<i>h</i> ₆	<i>const</i> (1)
<i>p</i> ₁ + 1 ~> +(<i>h</i> ₀ , <i>h</i> ₆)	<i>h</i> ₇	+(<i>h</i> ₀ , <i>h</i> ₆)
<i>p</i> ₂	<i>h</i> ₈	<i>h</i> ₇
<i>ivar</i> (* <i>p</i> ₂ , <i>v</i> ₁ *) ~> <i>ivar</i> (* <i>h</i> ₇ , <i>v</i> ₁ *)	<i>h</i> ₁₀	<i>ivar</i> (* <i>h</i> ₇ , <i>v</i> ₁ *)
<i>q</i> ₂	<i>h</i> ₁₁	<i>h</i> ₇
4	<i>h</i> ₁₂	<i>const</i> (4)
<i>v</i> ₂ *	<i>h</i> ₁₃	<i>v</i> ₂ *
<i>ivar</i> (* <i>p</i> ₂ , <i>v</i> ₂ *) ~> <i>ivar</i> (* <i>h</i> ₇ , <i>v</i> ₂ *)	<i>h</i> ₁₄	<i>h</i> ₁₂

Fig. 15.3 Some code after variables versioning, its corresponding HSSA form along with its hash table entries. $q_1 + 1$ that simplifies into $+(h_0, h_6)$ will be hashed to h_7 , and $ivar(*q_2, v_2^*)$ that simplifies into $ivar(*h_7, v_2^*)$ will be hashed to h_{14} .

forming the redundant memory reads (the same holds for "p->y"). Similarly, consider the example of Figure 15.3. Loads of lines 5 and 8 can not be considered as redundant because the versions for v (v_1^* then v_2^*) are different. On the other hand the load of line 8 can be safely avoided using a rematerialization of the value computed in line 7 as both the version for v (v_2^*) and the value number for the memory operands are identical. As a last example, if all the associated virtual variable versions for an indirect memory store (defined in parallel by χ -functions) are found to be dead, then it can be safely eliminated³. In other words HSSA transparently extends DEADCE to work also on indirect memory stores.

Another advantage of GVN in this context is that it enables uniform treatment of indirect memory operations regardless of the levels of indirections (like

³ Note that any virtual variable that aliases with a memory region live-out of the compiled procedure is considered to alias with the return instruction of the procedure, and as a consequence will lead to a live μ -function

in the "`**p`" expression which represents a double indirection). This happens naturally because each node of the expression is identified by its value number, and the fact that it is used as an address in another expression does not raise any additional complication.

Having explained why HSSA uses GVN we are ready to explain how the HSSA intermediate representation is structured. A program in HSSA keeps its CFG structure, with basic-blocks made of a sequence of statements (assignments and procedure calls), but the left and right hand-side of each statement are replaced by their corresponding entry in the hash table. Each constant, address and variable (both scalar and virtual) finds its entry in the hash table. As already explained, expressions are also hashed using the operator and the global value number of its operands. Each operand that corresponds to an indirect memory operation gets a special treatment on two aspects. First of all, it is considered semantically both as an expression (that has to be computed) and as a memory location (just as other variables). Second, its entry in the hash table is made up of both the address expression and its corresponding virtual variables version. This allows two entries with identical value numbers to "carry" the same value. This is illustrated in lines 3, 7, and 8 of Example 15.3. Such entries are referred as *ivar* node, for *indirect variables* so as to recall its operational semantic and distinguish them from the original program scalar variables.

It should be noted that ivar nodes, in terms of code generation, are operations and not variables: the compiler back end, when processing them, will emit indirect memory accesses to the address passed as their operand. While each ivar node is *related* to a virtual variable (the one corresponding to its address expression), and its GVN is determined taking into account the SSA version of its virtual variable; the virtual variable has no real counterpart in the program code. In this sense virtual variables are just an aid to apply aliasing effects to indirect memory operations, be aware of liveness and perform SSA renaming (with any of the regular SSA renaming algorithms). However, even though virtual variables do not contribute to the emission of machine code, their use-def links can be examined by optimization passes to determine the program data dependency paths (just like use-def links of scalar variables).

Therefore in HSSA value numbers have the same role of SSA versions in "plain" SSA (while in HSSA variable versions can effectively be discarded: values are identified by GVN and data dependencies by use-def links between IR nodes. It is in this way that HSSA extends all SSA based optimizations to indirect memory operations, even if they were originally designed to be applied to "plain" program variables.

15.6 Building HSSA

We now present the HSSA construction algorithm. It is straightforward, because it is a simple composition of μ and χ insertion, zero versioning and virtual variable introduction (described in previous sections), together with regular SSA renaming and GVN application.

Algorithm 15.2: SSA form construction

1. Perform alias analysis and assign a virtual variables to each indirect memory operand
 2. Insert μ -functions and χ -functions for scalar and virtual variables
 3. Insert ϕ -functions (considering both regular and χ assignments) as for standard SSA construction
 4. Perform SSA renaming on all scalar and virtual variables as for standard SSA construction
-

At the end of Algorithm 15.2 we have code in plain SSA form. The use of μ and χ operands guarantees that SSA versions are correct also in case of aliasing. Moreover, indirect memory operations are "annotated" with virtual variables, and also virtual variables have SSA version numbers. However, note how virtual variables are sort of "artificial" in the code and will not contribute to the final code generation pass, because what really matters are the indirect memory operations themselves.

The next steps corresponds to Algorithm 15.3 where steps 5 and 6 can be done using a single traversal of the program. At the end of this phase the code has exactly the same structure as before, but the number of unique SSA versions had diminished because of the application of zero versions.

Algorithm 15.3: Detecting zero versions

5. perform DEADCE (also on χ and ϕ stores)
 6. initialize *HasRealOcc* and *NonZeroPhiList* as for Algorithm 15.1, then run Algorithm 15.1 (Zero-version detection)
-

The application of global value numbering is straightforward. Some basics about GVN are recalled in Chapter 12.

Algorithm 15.4: Applying GVN

7. Perform a pre-order traversal of the dominator tree, applying GVN to the whole code (generate a unique hash table *var* node for each scalar and virtual variable version that is still live)
 - a. expressions are processed bottom up, reusing existing hash table expression nodes and using var nodes of the appropriate SSA variable version (the current one in the dominator tree traversal)
 - b. two ivar nodes have the same value number if these conditions are both true:
 - their address expressions have the same value number, and
 - their virtual variables have the same versions, or are separated by definitions that do not alias the ivar (possible to verify because of the dominator tree traversal order)
-

register promotion

As a result of this, each node in the code representation has a proper value number, and nodes with the same number are guaranteed to produce the same value (or hold it in the case of variables). The crucial issue is that the code must be traversed following the dominator tree in pre-order. This is important because when generating value numbers we must be sure that all the involved definitions already have a value number. Since the SSA form is a freshly created one, it is strict (i.e., all definitions dominate their use). As a consequence, a dominator tree pre-order traversal satisfies this requirement.

Note that after this step virtual variables are not needed anymore, and can be discarded from the code representation: the information they convey about aliasing of indirect variables has already been used to generate correct value numbers for ivar nodes.

Algorithm 15.5: Linking definitions

8. The left-hand side of each assignment (direct and indirect, real, ϕ and χ) is updated to point to its var or ivar node (which will point back to the defining statement)
9. Also all ϕ , μ and χ operands are updated to point to the corresponding GVN table entry

At the end of the last steps listed in Algorithm 15.5, HSSA form is complete, and every value in the program code is represented by a reference to a node in the HSSA value table.

.....

15.7 Using HSSA

As seen in the previous sections, HSSA is an internal code representation that applies SSA to indirect memory operations and builds a global value number table, valid also for values computed by indirect memory operations.

This representation, once built, is particularly memory efficient because expressions are shared among use points (they are represented as HSSA table nodes). In fact the original program code can be discarded, keeping only a list of statements pointing to the shared expressions nodes.

All optimizations conceived to be applied on scalar variables work "out of the box" on indirect locations: of course the implementation must be adapted to use the HSSA table, but their algorithm (and computational complexity) is the same even when dealing with values accessed indirectly.

Indirect memory operation (ivar) nodes are both variables and expressions, and benefit from the optimizations applied to both kinds of nodes. Particularly, it is relatively easy to implement register promotion.*

Of course the effectiveness of optimizations applied to indirect memory operations depends on the quality of alias analysis: if the analysis is poor the compiler will be forced to "play safe", and in practice the values will have "zero version" most of the time, so few or no optimizations will be applied to them. On the other hand, for all indirect memory operations where the alias analyzer de-

termines that there are no interferences caused by aliasing all optimizations can happen naturally, like in the scalar case.

15.8 Further readings

Cite the original paper from Zhou et al. [299] Cite the original paper from Chow. Cite work done in the GCC compiler (which was later scrapped due to compile time and memory consumption problems, but the experience is valuable): <http://www.airs.com/dnovillo/Papers/mem-ssa.pdf> Memory SSA - A Unified Approach for Sparsely Representing Memory Operations, D. Novillo, 2007 GCC Developers' Summit, Ottawa, Canada, July 2007. Talk about the possible differences (in terms of notations) that might exist between this chapter and the paper. Cite Alpern, Wegman, Zadech paper for GVN. Discuss the differences with Array SSA. Cite the paper mentioned in Fred's paper about factoring. Give some pointers to computation of alias analysis, but also representation of alias information (eg points-to). Add a reference on the paper about register promotion and mention the authros call it indirect removal

CHAPTER 16

Array SSA Form

*Vivek Sarkar
Kathleen Knobe
Stephen Fink*

Progress: 80%



Minor polishing in progress

16.1 Introduction

In this chapter, we introduce an Array SSA form that captures element-level data-flow information for array variables, and coincides with standard SSA form when applied to scalar variables. Any program with arbitrary control-flow structures and arbitrary array subscript expressions can be automatically converted to this Array SSA form, thereby making it applicable to structures, heap objects and any other data structure that can be modeled as a logical array. A key extension over standard SSA form is the introduction of a *definition-Φ* function that is capable of merging values from distinct array definitions on an element-by-element basis. There are several potential applications of Array SSA form in compiler analysis and optimization of sequential and parallel programs. In this chapter, we focus on sequential programs and use *constant propagation* as an exemplar of a program analysis that can be extended to array variables using Array SSA form, and *redundant load elimination* as an exemplar of a program optimization that can be extended to heap objects using Array SSA form.

The rest of the chapter is organized as follows. Section 16.2 introduces *full Array SSA form* for run-time evaluation and *partial Array SSA form* for static analysis. Section 16.3 extends the scalar SSA constant propagation algorithm to enable constant propagation through array elements. This includes an extension

to the constant propagation lattice to efficiently record information about array elements and an extension to the work-list algorithm to support *definition- Φ* functions (section 16.3.1), and a further extension to support non-constant (symbolic) array subscripts (section 16.3.2). Section 16.4 shows how Array SSA form can be extended to support elimination of redundant loads of object fields and array elements in strongly typed languages, and section 16.5 contains suggestions for further reading.

16.2 Array SSA Form

To introduce full Array SSA form with runtime evaluation of Φ functions, we use the concept of an *iteration vector* to differentiate among multiple dynamic instances of a static definition, S_k , that occur in the same dynamic instance of S_k 's enclosing procedure, $f()$. Let n be the number of loops that enclose S_k in procedure $f()$. These loops could be for-loops, while-loops, or even loops constructed out of goto statements. For convenience, we treat the outermost region of acyclic control flow in a procedure as a dummy outermost loop with a single iteration, thereby ensuring that $n \geq 1$.

A single point in the iteration space is specified by the iteration vector $\mathbf{i} = (i_1, \dots, i_n)$, which is an n -tuple of iteration numbers, <one for each enclosing loop. For convenience, this definition of iteration vectors assumes that all loops are single-entry, or equivalently, that the control-flow graph is *reducible*. (This assumption is not necessary for partial Array SSA form.) For single-entry loops, we know that each def executes at most once in a given iteration of its surrounding loops, hence the iteration vector serves the purpose of a “timestamp”. The key extensions in Array SSA form relative to standard SSA form are as follows.

1. **Renamed array variables:** All array variables are renamed so as to satisfy the static single assignment property. Analogous to standard SSA form, control Φ operators are introduced to generate new names for merging two or more prior definitions at control-flow join points, and to ensure that each use refers to precisely one definition.
2. **Array-valued @ variables:** For each static definition A_j , we introduce an *@ variable* (pronounced “at variable”) $@A_j$ that identifies the most recent *iteration vector* at which definition A_j was executed. We assume that all *@* variables are initialized to the empty vector, $()$, at the start of program execution. Each update of a single array element, $A_j[k] := \dots$, is followed by the statement, $@A_j[k] := \mathbf{i}$ where \mathbf{i} is the iteration vector for the loops surrounding the definition of A_j .
3. **Definition Φ 's:** A *definition- Φ* operator is introduced in Array SSA form to deal with preserving (“non-killing”) definitions of arrays. Consider A_0 and A_1 , two renamed arrays that originated from the same array variable in the source program such that $A_1[k] := \dots$ is an update of a single array element

and A_0 is the prevailing definition at the program point just prior to the definition of A_1 . A definition Φ , $A_2 := d\Phi(A_1, @A_1, A_0, @A_0)$, is inserted immediately after the definitions for A_1 and $@A_1$. Since definition A_1 only updates one element of A_0 , A_2 represents an element-level merge of arrays A_1 and A_0 . Definition Φ 's did not need to be inserted in standard SSA form because a scalar definition completely kills the old value of the variable.

4. **Array-valued Φ operators:** Another consequence of renaming arrays is that a Φ operator for array variables must also return an array value. Consider a (control or definition) Φ operator of the form, $A_2 := \Phi(A_1, @A_1, A_0, @A_0)$. Its semantics can be specified precisely by the following conditional expression for each element, $A_2[j]$, in the result array A_2 :

$$A_2[j] = \begin{array}{ll} \textbf{if} & @A_1[j] \succeq @A_0[j] \textbf{then } A_1[j] \\ & \textbf{else } A_0[j] \\ & \textbf{end if} \end{array} \quad (16.1)$$

The key extension over the scalar case is that the conditional expression specifies an element-level merge of arrays A_1 and A_0 .

Figures 16.1 and 16.2 show an example program with an array variable, and the conversion of the program to full Array SSA form as defined above.

```

 $A[*]:=$  initial value of  $A$ 
 $i:=1$ 
 $C:=i < 2$ 
 $\text{if } C \text{ then}$ 
     $k:=2 * i$ 
     $A[k]:=i$ 
     $\text{print } A[k]$ 
 $\text{endif}$ 
 $\text{print } A[2]$ 

```

Fig. 16.1 Example program with array variables

We now introduce a *partial Array SSA form* for static analysis, that serves as an approximation of full Array SSA form. Consider a (control or definition) Φ statement, $A_2 := \Phi(A_1, @A_1, A_0, @A_0)$. A static analysis will need to approximate the computation of this Φ operator by some data-flow transfer function, \mathcal{L}_Φ . The inputs and output of \mathcal{L}_Φ will be *lattice elements* for scalar/array variables that are compile-time approximations of their run-time values. We use the notation $\mathcal{L}(V)$ to denote the lattice element for a scalar or array variable V . Therefore, the statement, $A_2 := \Phi(A_1, @A_1, A_0, @A_0)$, will in general be modeled by the data-flow equation, $\mathcal{L}(A_2) = \mathcal{L}_\Phi(\mathcal{L}(A_1), \mathcal{L}(@A_1), \mathcal{L}(A_0), \mathcal{L}(@A_0))$.

While the *runtime* semantics of Φ functions for array variables critically depends on $@$ variables (Equation 16.1), many *compile-time* analyses do not need

```

@A0[*]:=() ; @A1[*]:=()

A0[*]:=initial value of A
@A0[*]:=(1)
i := 1
C := i < n
if C then
    k := 2 * i
    A1[k]:=i
    @A1[k]:=(1)
    A2:=dΦ(A1,@A1,A0,@A0)
    @A2:=max(@A1,@A0)
    print A2[k]
endif
A3:=Φ(A2,@A2,A0,@A0)
@A3:=max(@A2,@A0)
fig: print A3[2]

```

Fig. 16.2 Conversion of program in figure 16.1 to Full Array SSA Form

the full generality of @ variables. For analyses that do not distinguish among iteration instances, it is sufficient to model $A_2 := \Phi(A_1, @A_1, A_0, @A_0)$ by a data-flow equation, $\mathcal{L}(A_2) = \mathcal{L}_\phi(\mathcal{L}(A_1), \mathcal{L}(A_0))$, that does not use lattice variables $\mathcal{L}(@A_1)$ and $\mathcal{L}(@A_0)$. For such cases, a *partial* Array SSA form can be obtained by dropping @ variables, and using the ϕ operator, $A_2 := \phi(A_1, A_0)$ instead of $A_2 := \Phi(A_1, @A_1, A_0, @A_0)$. A consequence of dropping @ variables is that partial Array SSA form does not need to deal with iteration vectors, and therefore does not require the control-flow graph to be *reducible* as in full Array SSA form. For scalar variables, the resulting ϕ operator obtained by dropping @ variables exactly coincides with standard SSA form.

16.3 Sparse Constant Propagation of Array Elements

16.3.1 *Array Lattice for Sparse Constant Propagation*

In this section, we describe the lattice representation used to model array values for constant propagation. Let \mathcal{U}_{ind}^A and \mathcal{U}_{elem}^A be the universal set of *index values* and the universal set of array *element values* respectively for an array variable A . For an array variable, the set denoted by lattice element $\mathcal{L}(A)$ is a subset of index-element pairs in $\mathcal{U}_{ind}^A \times \mathcal{U}_{elem}^A$. There are three kinds of lattice elements for array variables that are of interest in our framework:

1. $\mathcal{L}(A) = \top \Rightarrow \text{SET}(\mathcal{L}(A)) = \{\}$

This “top” case indicates that the set of possible index-element pairs that have been identified thus far for A is the empty set, $\{\}$.

2. $\mathcal{L}(A) = \langle (i_1, e_1), (i_2, e_2), \dots \rangle$

$$\Rightarrow \text{SET}(\mathcal{L}(A)) = \{(i_1, e_1), (i_2, e_2), \dots\} \cup (\mathcal{U}_{\text{ind}}^A - \{i_1, i_2, \dots\}) \times \mathcal{U}_{\text{elem}}^A$$

The lattice element for this “constant” case is represented by a finite list of constant index-element pairs, $\langle (i_1, e_1), (i_2, e_2), \dots \rangle$. The constant indices, i_1, i_2, \dots , must represent distinct (non-equal) index values. The meaning of this lattice element is that the current stage of analysis has identified some finite number of constant index-element pairs for array variable A , such that $A[i_1] = e_1, A[i_2] = e_2$, etc. All other elements of A are assumed to be non-constant. (Extensions to handle non-constant indices are described in section 16.3.2.)

3. $\mathcal{L}(A) = \perp \Rightarrow \text{SET}(\mathcal{L}(A)) = \mathcal{U}_{\text{ind}}^A \times \mathcal{U}_{\text{elem}}^A$

This “bottom” case indicates that, according to the approximation in the current stage of analysis, array A may take on any value from the universal set of index-element pairs. Note that $\mathcal{L}(A) = \perp$ is equivalent to an empty list, $\mathcal{L}(A) = \langle \rangle$, in case (2) above; they both denote the universal set of index-element pairs.

$\mathcal{L}(A_1[k])$	$\mathcal{L}(k) = \top$	$\mathcal{L}(k) = \text{Constant}$	$\mathcal{L}(k) = \perp$
$\mathcal{L}(A_1) = \top$	\top	\top	\perp
$\mathcal{L}(A_1) = \langle (i_1, e_1), \dots \rangle$	\top	$e_j, \text{ if } \exists (i_j, e_j) \in \mathcal{L}(A_1) \text{ with } \mathcal{D}\mathcal{S}(i_j, \mathcal{L}(k)) = \text{true}$ $\perp, \text{ otherwise}$	\perp
$\mathcal{L}(A_1) = \perp$	\perp	\perp	\perp

Fig. 16.3 Lattice computation for $\mathcal{L}(A_1[k]) = \mathcal{L}_{[]}(\mathcal{L}(A_1), \mathcal{L}(k))$, where $A_1[k]$ is an array element read operator

$\mathcal{L}(A_1)$	$\mathcal{L}(i) = \top$	$\mathcal{L}(i) = \text{Constant}$	$\mathcal{L}(i) = \perp$
$\mathcal{L}(k) = \top$	\top	\top	\perp
$\mathcal{L}(k) = \text{Constant}$	\top	$\langle \mathcal{L}(k), \mathcal{L}(i) \rangle$	\perp
$\mathcal{L}(k) = \perp$	\perp	\perp	\perp

Fig. 16.4 Lattice computation for $\mathcal{L}(A_1) = \mathcal{L}_{d[]}(\mathcal{L}(k), \mathcal{L}(i))$, where $A_1[k] := i$ is an array element write operator

We now describe how array lattice elements are computed for various operations that appear in Array SSA form. We start with the simplest operation *viz.*, a read access to an array element. Figure 16.3 shows how $\mathcal{L}(A_1[k])$, the lattice element for array reference $A_1[k]$, is computed as a function of $\mathcal{L}(A_1)$ and $\mathcal{L}(k)$, the lattice elements for A_1 and k . We denote this function by $\mathcal{L}_{[]}$

i.e., $\mathcal{L}(A_1[k]) = \mathcal{L}_{[]}(\mathcal{L}(A_1), \mathcal{L}(k))$. The interesting case in figure 16.3 occurs in the middle cell when neither $\mathcal{L}(A_1)$ nor $\mathcal{L}(k)$ is \top or \perp . The notation $\mathcal{D}\mathcal{S}$ in the middle cell in figure 16.3 represents a “definitely-same” binary relation i.e., $\mathcal{D}\mathcal{S}(a, b) = \text{true}$ if and only if a and b are known to have exactly the same value.

Next, consider a write access of an array element, which in general has the form $A_1[k] := i$. Figure 16.4 shows how $\mathcal{L}(A_1)$, the lattice element for the array being written into, is computed as a function of $\mathcal{L}(k)$ and $\mathcal{L}(i)$, the lattice elements for k and i . We denote this function by $\mathcal{L}_{d[]}$ i.e., $\mathcal{L}(A_1) = \mathcal{L}_{d[]}(\mathcal{L}(k), \mathcal{L}(i))$. As before, the interesting case in figure 16.4 occurs in the middle cell when both $\mathcal{L}(k)$ and $\mathcal{L}(i)$ are constant. For this case, the value returned for $\mathcal{L}(A_1)$ is simply the singleton list, $\langle (\mathcal{L}(k), \mathcal{L}(i)) \rangle$, which contains exactly one constant index-element pair.

Now, we turn our attention to the ϕ functions. Consider a definition ϕ operation of the form, $A_2 := d\phi(A_1, A_0)$. The lattice computation for $\mathcal{L}(A_2) = \mathcal{L}_{d\phi}(\mathcal{L}(A_1), \mathcal{L}(A_0))$ is shown in figure 16.5. Since A_1 corresponds to a definition of a single array element, the list for $\mathcal{L}(A_1)$ can contain at most one pair (see figure 16.4). Therefore, the three cases considered for $\mathcal{L}(A_1)$ in figure 16.5 are $\mathcal{L}(A_1) = \top$, $\mathcal{L}(A_1) = \langle (i', e') \rangle$, and $\mathcal{L}(A_1) = \perp$.

The notation $\text{UPDATE}((i', e'), \langle (i_1, e_1), \dots \rangle)$ used in the middle cell in figure 16.5 denotes a special update of the list $\mathcal{L}(A_0) = \langle (i_1, e_1), \dots \rangle$ with respect to the constant index-element pair (i', e') . UPDATE involves four steps:

1. Compute the list $T = \{ (i_j, e_j) \mid (i_j, e_j) \in \mathcal{L}(A_0) \text{ and } \mathcal{D}\mathcal{D}(i', i_j) = \text{true} \}$.
Analogous to $\mathcal{D}\mathcal{S}$, $\mathcal{D}\mathcal{D}$ denotes a “definitely-different” binary relation i.e., $\mathcal{D}\mathcal{D}(a, b) = \text{true}$ if and only if a and b are known to have distinct (non-equal) values.
2. Insert the pair (i', e') into T to obtain a new list, I .
3. (Optional) If there is a desire to bound the height of the lattice due to compile-time considerations, and the size of list I exceeds a threshold size Z , then one of the pairs in I can be dropped from the output list so as to satisfy the size constraint.
4. Return I as the value of $\text{UPDATE}((i', e'), \langle (i_1, e_1), \dots \rangle)$.

$\mathcal{L}(A_2)$	$\mathcal{L}(A_0) = \top$	$\mathcal{L}(A_0) = \langle (i_1, e_1), \dots \rangle$	$\mathcal{L}(A_0) = \perp$
$\mathcal{L}(A_1) = \top$	\top	\top	\top
$\mathcal{L}(A_1) = \langle (i', e') \rangle$	\top	$\text{UPDATE}((i', e'), \langle (i_1, e_1), \dots \rangle)$	$\langle (i', e') \rangle$
$\mathcal{L}(A_1) = \perp$	\perp	\perp	\perp

Fig. 16.5 Lattice computation for $\mathcal{L}(A_2) = \mathcal{L}_{d\phi}(\mathcal{L}(A_1), \mathcal{L}(A_0))$ where $A_2 := d\phi(A_1, A_0)$ is a definition ϕ operation

Finally, consider a control ϕ operation that merges two array values, $A_2 := \phi(A_1, A_0)$. The join operator (\sqcap) is used to compute $\mathcal{L}(A_2)$, the lattice element for A_2 , as a function of $\mathcal{L}(A_1)$ and $\mathcal{L}(A_0)$, the lattice elements for A_1 and A_0 i.e., $\mathcal{L}(A_2) = \mathcal{L}_\phi(\mathcal{L}(A_1), \mathcal{L}(A_0)) = \mathcal{L}(A_1) \sqcap \mathcal{L}(A_0)$. The rules for computing this join

$\mathcal{L}(A_2) = \mathcal{L}(A_1) \sqcap \mathcal{L}(A_0)$	$\mathcal{L}(A_0) = \top$	$\mathcal{L}(A_0) = \langle (i_1, e_1), \dots \rangle$	$\mathcal{L}(A_0) = \perp$
$\mathcal{L}(A_1) = \top$	\top	$\mathcal{L}(A_0)$	\perp
$\mathcal{L}(A_1) = \langle (i'_1, e'_1), \dots \rangle$	$\mathcal{L}(A_1)$	$\mathcal{L}(A_1) \sqcap \mathcal{L}(A_0)$	\perp
$\mathcal{L}(A_1) = \perp$	\perp	\perp	\perp

Fig. 16.6 Lattice computation for $\mathcal{L}(A_2) = \mathcal{L}_\phi(\mathcal{L}(A_1), \mathcal{L}(A_0)) = \mathcal{L}(A_1) \sqcap \mathcal{L}(A_0)$, where $A_2 := \phi(A_1, A_0)$ is a control ϕ operation

operator are shown in figure 16.6, depending on different cases for $\mathcal{L}(A_1)$ and $\mathcal{L}(A_0)$. The notation $\mathcal{L}(A_1) \sqcap \mathcal{L}(A_0)$ used in the middle cell in figure 16.6 denotes a simple intersection of lists $\mathcal{L}(A_1)$ and $\mathcal{L}(A_0)$ — the result is a list of pairs that appear in both $\mathcal{L}(A_1)$ and $\mathcal{L}(A_0)$.

We conclude this section by discussing the example program in figure 16.7. The partial Array SSA form for this example is shown in figure 16.8, and the data-flow equations for this example are shown in figure 16.9. Each assignment statement in the partial Array SSA form (in figure 16.8) results in one data-flow equation (in figure 16.9); the numbering S1 through S8 indicates the correspondence. Any solver can be used for these data-flow equations, including the standard worklist-based algorithm for constant propagation using scalar SSA form. The fixpoint solution is shown in figure 16.10. This solution was obtained assuming $\mathcal{L}(I) = \perp$. If, instead, variable I is known to equal 3 i.e., $\mathcal{L}(I) = 3$, then the lattice variables that would be obtained after the fixpoint iteration step has completed are shown in figure 16.11. In either case ($\mathcal{L}(I) = \perp$ or $\mathcal{L}(I) = 3$), the resulting array element constants revealed by the algorithm can be used in whatever analyses or transformations the compiler considers to be profitable to perform.

```

Y[3]:=99
if C then
    D[1]:=Y[3]*2
else
    D[1]:=Y[I]*2
endif
Z := D[1]

```

Fig. 16.7 Sparse Constant Propagation Example

16.3.2 Beyond Constant Indices

In this section we address constant propagation through *non-constant array subscripts*, as a generalization of the algorithm for constant subscripts described in section 16.3.1. As an example, consider the program fragment in figure 16.12.

```

 $Y_0$  and  $D_0$  in effect here.
...
S1:  $Y_1[3]:=99$ 
S2:  $Y_2:=d\phi(Y_1, Y_0)$ 
if C then
S3:  $D_1[1]:=Y_2[3]*2$ 
S4:  $D_2:=d\phi(D_1, D_0)$ 
else
S5:  $D_3[1]:=Y_2[I]*2$ 
S6:  $D_4:=d\phi(D_3, D_0)$ 
endif
S7:  $D_5:=\phi(D_2, D_4)$ 
S8:  $Z:=D_5[1]$ 

```

Fig. 16.8 Array SSA form for the Sparse Constant Propagation Example

```

S1:  $\mathcal{L}(Y_1) = \langle(3, 99)\rangle$ 
S2:  $\mathcal{L}(Y_2) = \mathcal{L}_{d\phi}(\mathcal{L}(Y_1), \mathcal{L}(Y_0))$ 
S3:  $\mathcal{L}(D_1) = \mathcal{L}_{d[\ ]}\mathcal{L}_*(\mathcal{L}_{[\ ]}(\mathcal{L}(Y_2), 3)), 2))$ 
S4:  $\mathcal{L}(D_2) = \mathcal{L}_{d\phi}(\mathcal{L}(D_1), \mathcal{L}(D_0))$ 
S5:  $\mathcal{L}(D_3) = \mathcal{L}_{d[\ ]}\mathcal{L}_*(\mathcal{L}_{[\ ]}(\mathcal{L}(Y_2), \mathcal{L}(I))), 2))$ 
S6:  $\mathcal{L}(D_4) = \mathcal{L}_{d\phi}(\mathcal{L}(D_3), \mathcal{L}(D_0))$ 
S7:  $\mathcal{L}(D_5) = \mathcal{L}_\phi(\mathcal{L}(D_2), \mathcal{L}(D_4))$ 
S8:  $\mathcal{L}(Z) = \mathcal{L}_{[\ ]}(\mathcal{L}(D_5), 1)$ 

```

Fig. 16.9 Data-flow Equations for the Sparse Constant Propagation Example

```

S1:  $\mathcal{L}(Y_1) = \langle(3, 99)\rangle$ 
S2:  $\mathcal{L}(Y_2) = \langle(3, 99)\rangle$ 
S3:  $\mathcal{L}(D_1) = \langle(1, 198)\rangle$ 
S4:  $\mathcal{L}(D_2) = \langle(1, 198)\rangle$ 
S5:  $\mathcal{L}(D_3) = \perp$ 
S6:  $\mathcal{L}(D_4) = \perp$ 
S7:  $\mathcal{L}(D_5) = \perp$ 
S8:  $\mathcal{L}(Z) = \perp$ 

```

Fig. 16.10 Solution to data-flow equations from figure 16.9, assuming I is unknown

```

S1:  $\mathcal{L}(Y_1) = \langle(3, 99)\rangle$ 
S2:  $\mathcal{L}(Y_2) = \langle(3, 99)\rangle$ 
S3:  $\mathcal{L}(D_1) = \langle(1, 198)\rangle$ 
S4:  $\mathcal{L}(D_2) = \langle(1, 198)\rangle$ 
S5:  $\mathcal{L}(D_3) = \langle(1, 198)\rangle$ 
S6:  $\mathcal{L}(D_4) = \langle(1, 198)\rangle$ 
S7:  $\mathcal{L}(D_5) = \langle(1, 198)\rangle$ 
S8:  $\mathcal{L}(Z) = 198$ 

```

Fig. 16.11 Solution to data-flow equations from figure 16.9, assuming I is known to be = 3

```

k := 2
do i := ...
...
a[i]:=k*5
...:=a[i]
enddo

```

Fig. 16.12 Example of Constant Propagation through Non-constant Index

In the loop in figure 16.12, we see that the read access of $a[i]$ will have a constant value ($k * 5 = 10$), even though the index/subscript value i is not a constant. We would like to extend the framework from section 16.3.1 to be able to recognize the read of $a[i]$ as constant in such programs. There are two key extensions that need to be considered for non-constant (symbolic) subscript values:

- For constants, C_1 and C_2 , $\mathcal{DS}(C_1, C_2) \neq \mathcal{DD}(C_1, C_2)$. However, for two symbols, S_1 and S_2 , it is possible that both $\mathcal{DS}(S_1, S_2)$ and $\mathcal{DD}(S_1, S_2)$ are FALSE, that is, we don't know if they are the same or different.
- For constants, C_1 and C_2 , the values for $\mathcal{DS}(C_1, C_2)$ and $\mathcal{DD}(C_1, C_2)$ can be computed by inspection. For symbolic indices, however, some program analysis is necessary to compute the \mathcal{DS} and \mathcal{DD} relations.

We now discuss the compile-time computation of \mathcal{DS} and \mathcal{DD} for symbolic indices. Observe that, given index values I_1 and I_2 , only one of the following three cases is possible:

Case 1: $\mathcal{DS}(I_1, I_2) = \text{FALSE}$; $\mathcal{DD}(I_1, I_2) = \text{FALSE}$
 Case 2: $\mathcal{DS}(I_1, I_2) = \text{TRUE}$; $\mathcal{DD}(I_1, I_2) = \text{FALSE}$
 Case 3: $\mathcal{DS}(I_1, I_2) = \text{FALSE}$; $\mathcal{DD}(I_1, I_2) = \text{TRUE}$

The first case is the most conservative solution. In the absence of any other knowledge, it is always correct to state that $\mathcal{DS}(I_1, I_2) = \text{false}$ and $\mathcal{DD}(I_1, I_2) = \text{false}$.

The problem of determining if two symbolic index values are the same is equivalent to the classical problem of *global value numbering*. If two indices i and j have the same value number, then $\mathcal{DS}(i, j)$ must = *true*. The problem of computing \mathcal{DD} is more complex. Note that \mathcal{DD} , unlike \mathcal{DS} , is not an equivalence relation because \mathcal{DD} is not transitive. If $\mathcal{DD}(A, B) = \text{true}$ and $\mathcal{DD}(B, C) = \text{true}$, it does not imply that $\mathcal{DD}(A, C) = \text{true}$. However, we can leverage past work on array dependence analysis to identify cases for which \mathcal{DD} can be evaluated to *true*. For example, it is clear that $\mathcal{DD}(i, i+1) = \text{true}$, and that $\mathcal{DD}(i, 0) = \text{true}$ if i is a loop index variable that is known to be ≥ 1 .

Let us consider how the \mathcal{DS} and \mathcal{DD} relations for symbolic index values are used by our constant propagation algorithms. Note that the specification of how \mathcal{DS} and \mathcal{DD} are used is a separate issue from the precision of the \mathcal{DS} and \mathcal{DD} values. We now describe how the lattice and the lattice operations presented in section 16.3.1 can be extended to deal with non-constant subscripts.

First, consider the lattice itself. The \top and \perp lattice elements retain the same meaning as in section 16.3.1 *viz.*, $\text{SET}(\top) = \{\}$ and $\text{SET}(\perp) = \mathcal{U}_{\text{ind}}^A \times \mathcal{U}_{\text{elem}}^A$. Each element in the lattice is a list of index-value pairs where the value is still required to be constant but the index may be symbolic — the index is represented by its value number.

We now revisit the processing of an array element read of $A_1[k]$ and the processing of an array element write of $A_1[k]$. These operations were presented in section 16.3.1 (figures 16.3 and 16.4) for constant indices. The versions for non-constant indices appear in figure 16.13 and figure 16.14. For the read operation in figure 16.13, if there exists a pair (i_j, e_j) such that $\mathcal{DS}(i_j, \text{VALNUM}(k)) = \text{true}$ (*i.e.*, i_j and k have the same value number), then the result is e_j . Otherwise, the result is \top or \perp as specified in figure 16.13. For the write operation in figure 16.14, if the value of the right-hand-side, i , is a constant, the result is the singleton list $\langle (\text{VALNUM}(k), \mathcal{L}(i)) \rangle$. Otherwise, the result is \top or \perp as specified in figure 16.14.

$\mathcal{L}(A_1[k])$	$\mathcal{L}(k) = \top$	$\mathcal{L}(k) = \text{VALNUM}(k)$	$\mathcal{L}(k) = \perp$
$\mathcal{L}(A_1) = \top$	\top	\top	\perp
$\mathcal{L}(A_1) = \langle (i_1, e_1), \dots \rangle$	\top	$e_j, \text{ if } \exists (i_j, e_j) \in \mathcal{L}(A_1) \text{ with } \mathcal{DS}(i_j, \text{VALNUM}(k)) = \text{true}$ $\perp, \text{ otherwise}$	\perp
$\mathcal{L}(A_1) = \perp$	\perp	\perp	\perp

Fig. 16.13 Lattice computation for $\mathcal{L}(A_1[k]) = \mathcal{L}_{[]}(\mathcal{L}(A_1), \mathcal{L}(k))$, where $A_1[k]$ is an array element read operator. If $\mathcal{L}(k) = \text{VALNUM}(k)$, the lattice value of index k is a value number that represents a constant or a symbolic value.

$\mathcal{L}(A_1)$	$\mathcal{L}(i) = \top$	$\mathcal{L}(i) = \text{Constant}$	$\mathcal{L}(i) = \perp$
$\mathcal{L}(k) = \top$	\top	\top	\perp
$\mathcal{L}(k) = \text{VALNUM}(k)$	\top	$\langle (\text{VALNUM}(k), \mathcal{L}(i)) \rangle$	\perp
$\mathcal{L}(k) = \perp$	\perp	\perp	\perp

Fig. 16.14 Lattice computation for $\mathcal{L}(A_1) = \mathcal{L}_{d[]}(\mathcal{L}(k), \mathcal{L}(i))$, where $A_1[k] := i$ is an array element write operator. If $\mathcal{L}(k) = \text{VALNUM}(k)$, the lattice value of index k is a value number that represents a constant or a symbolic value.

Let us now consider the propagation of lattice values through $d\phi$ operators. The only extension required relative to figure 16.5 is that the \mathcal{DD} relation used in performing the UPDATE operation should be able to determine when $\mathcal{DD}(i', i_j) = \text{true}$ if i' and i_j are symbolic value numbers rather than constants.

(If no symbolic information is available for i' and i_j , then it is always safe to return $\mathcal{D}(i', i_j) = \text{false}$.)

16.4 Extension to Objects: Redundant Load Elimination

In this section, we introduce redundant load elimination as an exemplar of a program optimization that can be extended to heap objects in strongly typed languages by using Array SSA form. This extension models object references (pointers) as indices into hypothetical *heap arrays* (Section 16.4.1). We then describe how definitely-same and definitely-different analyses can be extended to heap array indices (Section 16.4.2), followed by a *scalar replacement* transformation that uses the analysis results to perform load elimination (Section 16.4.3).

16.4.1 Analysis Framework

We introduce a formalism called *heap arrays* which allows us to model object references as associative arrays. An *extended Array SSA* form is constructed on heap arrays by adding *use- ϕ* functions. For each field x , we introduce a hypothetical one-dimensional heap array, \mathcal{H}^x . Heap array \mathcal{H}^x consolidates all instances of field x present in the heap. Heap arrays are indexed by object references. Thus, a *GETFIELD* of $p.x$ is modeled as a read of element $\mathcal{H}^x[p]$, and a *PUTFIELD* of $q.x$ is modeled as a write of element $\mathcal{H}^x[q]$. The use of distinct heap arrays for distinct fields leverages the fact that accesses to distinct fields must be directed to distinct memory locations in a strongly typed language. Note that field x is considered to be the same field for objects of types C_1 and C_2 , if C_2 is a subtype of C_1 . Accesses to one-dimensional array objects with the same element type are modeled as accesses to a single *two-dimensional* heap array for that element type, with one dimension indexed by the object reference as in heap arrays for fields, and the second dimension indexed by the integer subscript.

Heap arrays are renamed in accordance with an *extended Array SSA* form that contains three kinds of ϕ functions:

1. A *control ϕ* from scalar SSA form.
2. A *definition ϕ* ($d\phi$) from Array SSA form.
3. A *use ϕ* ($u\phi$) function creates a new name whenever a statement reads a heap array element. $u\phi$ functions represent the extension in “extended” Array SSA form.

The main purpose of the $u\phi$ function is to link together load instructions for the same heap array in control-flow order. While $u\phi$ functions are used by the re-

dundant load elimination optimization presented in this chapter, it is not necessary for analysis algorithms (e.g., constant propagation) that do not require the creation of a new name at each use.

16.4.2 Definitely-Same and Definitely-Different Analyses for Heap Array Indices

In this section, we show how global value numbering and allocation site information can be used to efficiently compute *definitely-same* (\mathcal{DS}) and *definitely-different* (\mathcal{DD}) information for heap array indices, thereby reducing pointer analysis queries to array index queries. If more sophisticated pointer analyses are available in a compiler, they can be used to further refine the \mathcal{DS} and \mathcal{DD} information.

As an example, Figure 16.15 illustrates two different cases of scalar replacement (load elimination) for object fields. The notation $Hx[p]$ refers to a read/write access of heap array element, $\mathcal{H}^x[p]$. For the original program in figure 16.15(a), introducing a scalar temporary $T1$ for the store (def) of $p.x$ can enable the load (use) of $r.x$ to be eliminated *i.e.*, to be replaced by a use of $T1$. Figure 16.15(b) contains an example in which a scalar temporary ($T2$) is introduced for the first load of $p.x$, thus enabling the second load of $r.x$ to be eliminated *i.e.*, replaced by $T2$. In both cases, the goal of our analysis is to determine that the load of $r.x$ is redundant, thereby enabling the compiler to replace it by a use of scalar temporary that captures the value in $p.x$. We need to establish two facts to perform this transformation: 1) object references p and r are identical (definitely same) in all program executions, and 2) object references q and r are distinct (definitely different) in all program executions.

As before, we use the notation $\mathcal{V}(i)$ to denote the value number of SSA variable i . Therefore, if $\mathcal{V}(i) = \mathcal{V}(j)$, then $\mathcal{DS}(i, j) = \text{true}$. For the code fragment above, the statement, $p := r$, ensures that p and r are given the same value number, $\mathcal{V}(p) = \mathcal{V}(r)$, so that $\mathcal{DS}(p, r) = \text{true}$. The problem of computing \mathcal{DD} for object references is more complex than value numbering, and relates to pointer alias analysis. We outline a simple and sound approach below, which can be replaced by more sophisticated techniques as needed. It relies on two observations related to allocation-sites:

1. Object references that contain the results of distinct allocation-sites must be different.
2. An object reference containing the result of an allocation-site must be different from any object reference that occurs at a program point that dominates the allocation site in the control-flow graph. For example, in Figure 16.15, the presence of the allocation site in $q := \text{new Type1}$ ensures that $\mathcal{DD}(p, q) = \text{true}$.

Original program:

```
r := p
q := new Type1
...
p.x := ... // Hx[p] := ...
q.x := ... // Hx[q] := ...
... := r.x // ... := Hx[r]
```

Original program:

```
r := p
q := new Type1
...
... := p.x // ... := Hx[p]
q.x := ... // Hx[q] := ...
... := r.x // ... := Hx[r]
```

After redundant load elimination:

```
r := p
q := new Type1
...
T1 := ...
p.x := T1
q.x := ...
... := T1
```

After redundant load elimination:

```
r := p
q := new Type1
...
T2 := p.x
... := T2
q.x := ...
... := T2
```

(a)

(b)

Fig. 16.15 Examples of scalar replacement

16.4.3 Scalar Replacement Algorithm

$\mathcal{L}(A_2)$	$\mathcal{L}(A_0) = \top$	$\mathcal{L}(A_0) = \langle i_1, \dots \rangle$	$\mathcal{L}(A_0) = \perp$
$\mathcal{L}(A_1) = \top$	\top	\top	\top
$\mathcal{L}(A_1) = \langle i' \rangle$	\top	UPDATE(i' , $\langle i_1, \dots \rangle$)	$\langle i' \rangle$
$\mathcal{L}(A_1) = \perp$	\perp	\perp	\perp

Fig. 16.16 Lattice computation for $\mathcal{L}(A_2) = \mathcal{L}_{d\phi}(\mathcal{L}(A_1), \mathcal{L}(A_0))$ where $A_2 := d\phi(A_1, A_0)$ is a definition ϕ operation

$\mathcal{L}(A_2)$	$\mathcal{L}(A_0) = \top$	$\mathcal{L}(A_0) = \langle (i_1), \dots \rangle$	$\mathcal{L}(A_0) = \perp$
$\mathcal{L}(A_1) = \top$	\top	\top	\top
$\mathcal{L}(A_1) = \langle i' \rangle$	\top	$\mathcal{L}(A_1) \cup \mathcal{L}(A_0)$	$\mathcal{L}(A_1)$
$\mathcal{L}(A_1) = \perp$	\perp	\perp	\perp

Fig. 16.17 Lattice computation for $\mathcal{L}(A_2) = \mathcal{L}_{u\phi}(\mathcal{L}(A_1), \mathcal{L}(A_0))$ where $A_2 := u\phi(A_1, A_0)$ is a use ϕ operation

The main program analysis needed to enable redundant load elimination is *index propagation*, which identifies the set of indices that are *available* at a spe-

$\mathcal{L}(A_2) = \mathcal{L}(A_1) \sqcap \mathcal{L}(A_0)$	$\mathcal{L}(A_0) = \top$	$\mathcal{L}(A_0) = \langle \mathbf{i}_1, \dots \rangle$	$\mathcal{L}(A_0) = \perp$
$\mathcal{L}(A_1) = \top$	\top	$\mathcal{L}(A_0)$	\perp
$\mathcal{L}(A_1) = \langle \mathbf{i}', \dots \rangle$	$\mathcal{L}(A_1)$	$\mathcal{L}(A_1) \sqcap \mathcal{L}(A_0)$	\perp
$\mathcal{L}(A_1) = \perp$	\perp	\perp	\perp

Fig. 16.18 Lattice computation for $\mathcal{L}(A_2) = \mathcal{L}_\phi(\mathcal{L}(A_1), \mathcal{L}(A_0)) = \mathcal{L}(A_1) \sqcap \mathcal{L}(A_0)$, where $A_2 := \phi(A_1, A_0)$ is a control ϕ operation

(a) Extended Array SSA form:	(b) After index propagation:	(c) After transformation:
$r := p$ $q := \text{new Type1}$ \dots $\mathcal{H}_1^x[p] := \dots$ $\mathcal{H}_2^x := d\phi(\mathcal{H}_1^x, \mathcal{H}_0^x)$ $\mathcal{H}_3^x[q] := \dots$ $\mathcal{H}_4^x := d\phi(\mathcal{H}_3^x, \mathcal{H}_2^x)$ $\dots := \mathcal{H}_4^x[r]$	$\mathcal{L}(\mathcal{H}_0^x) = \{\}$ $\mathcal{L}(\mathcal{H}_1^x) = \{\mathcal{V}(p) = \mathcal{V}(r)\}$ $\mathcal{L}(\mathcal{H}_2^x) = \{\mathcal{V}(p) = \mathcal{V}(r)\}$ $\mathcal{L}(\mathcal{H}_3^x) = \{\mathcal{V}(q)\}$ $\mathcal{L}(\mathcal{H}_4^x) = \{\mathcal{V}(p) = \mathcal{V}(r), \mathcal{V}(q)\}$	$r := p$ $q := \text{new Type1}$ \dots $T1 := \dots$ $p.x := T1$ $q.x := \dots$ $\dots := T1$

Fig. 16.19 Trace of index propagation and load elimination transformation for program in figure 16.15(a)

cific def/use A_i of heap array A . Index propagation is a data-flow problem, the goal of which is to compute a lattice value $\mathcal{L}(\mathcal{H})$ for each renamed heap variable \mathcal{H} in the Array SSA form such that a load of $\mathcal{H}[\mathbf{i}]$ is *available* if $\mathcal{V}(\mathbf{i}) \in \mathcal{L}(\mathcal{H})$. Note that the lattice element does not include the value of $\mathcal{H}[\mathbf{i}]$ (as in constant propagation), just the fact that it is available. Figures 16.16, 16.17 and 16.18 give the lattice computations which define the index propagation solution. The notation $\text{UPDATE}(\mathbf{i}', \langle \mathbf{i}_1, \dots \rangle)$ used in the middle cell in figure 16.16 denotes a special update of the list $\mathcal{L}(A_0) = \langle \mathbf{i}_1, \dots \rangle$ with respect to index \mathbf{i}' . UPDATE involves four steps:

1. Compute the list $T = \{ \mathbf{i}_j \mid \mathbf{i}_j \in \mathcal{L}(A_0) \text{ and } \mathcal{D}\mathcal{D}(\mathbf{i}', \mathbf{i}_j) = \text{true} \}$. List T contains only those indices from $\mathcal{L}(A_0)$ that are *definitely different* from \mathbf{i}' .
2. Insert \mathbf{i}' into T to obtain a new list, I .
3. (Optional) As before, if there is a desire to bound the height of the lattice due to compile-time considerations, and the size of list I exceeds a threshold size Z , then any one of the indices in I can be dropped from the output list.
4. Return I as the value of $\text{UPDATE}(\mathbf{i}', \langle \mathbf{i}_1, \dots \rangle)$.

After index propagation, the algorithm selects a load, $A_j[x]$, for scalar replacement if and only if index propagation determines that an index with value number $\mathcal{V}(x)$ is available at the def of A_j . Figure 16.19 illustrates a trace of this load elimination algorithm for the example program in figure 16.15(a). Figure 16.19(a) shows the extended Array SSA form computed for this example program. The results of index propagation are shown in figure 16.19(b). These results depend on definitely-different analysis establishing that $\mathcal{V}(p) \neq \mathcal{V}(q)$ and definitely-same analysis establishing that $\mathcal{V}(p) = \mathcal{V}(r)$. Figure 16.19(c) shows the transformed code after performing the scalar replacement actions. The load

of $p.x$ has thus been eliminated in the transformed code, and replaced by a use of the scalar temporary, $T1$.

16.5 Further Reading

In this chapter, we introduced an Array SSA form that captures element-level data-flow information for array variables, illustrated how it can be used to extend program analyses for scalars to array variables using constant propagation as an exemplar, and illustrated how it can be used to extend optimizations for scalars to array variables using load elimination in heap objects as an exemplar. In addition to reading the other chapters in this book for related topics, the interested reader can consult [167] for details on full Array SSA form, [168] for details on constant propagation using Array SSA form, [120] for details on load and store elimination for pointer-based objects using Array SSA form, [243] for efficient dependence analysis of pointer-based array objects, and [24] for extensions of this load elimination algorithm to parallel programs.

There are many possible directions for future research based on this work. The definitely-same and definitely-different analyses outlined in this chapter are sound, but conservative. In restricted cases, they can be made more precise using array subscript analysis from polyhedral compilation frameworks. Achieving a robust integration of Array SSA and polyhedral approaches is an interesting goal for future research. Past work on Fuzzy Array Data-flow Analysis (FADA) [25] may provide a useful foundation for exploring such an integration. Another interesting direction is to extend the value numbering and definitely-different analyses mentioned in section 16.3.2 so that they can be combined with constant propagation rather than performed as a pre-pass. An ultimate goal is to combine conditional constant, type propagation, value numbering, partial redundancy elimination, and scalar replacement analyses within a single framework that can be used to analyze scalar variables, array variables, and pointer objects with a unified approach.

*single definition property
predicated execution
guard
predicate*

CHAPTER 17

Psi-SSA Form

F. de Ferrière

Progress: 95%



Final reading in progress

17.1 Overview

In the SSA representation, each definition of a variable is given a unique name,^{*} and new pseudo definitions are introduced on ϕ -functions to merge values coming from different control-flow paths. An example is given figure 17.1(b). Each definition is an unconditional definition, and the value of a variable is the value of the expression on the unique assignment to this variable. This essential property of the SSA representation does not any longer hold when definitions may be conditionally executed.^{*} When a variable is defined by a predicated operation, the value of the variable will or will not be modified depending on the value of a guard register. As a result, the value of the variable after the predicated operation is either the value of the expression on the assignment if the predicate is true, or the value the variable had before this operation if the predicate is false. This is represented in figure 17.1(c) where we use the notation $p ? a = \text{op}$ to indicate that an operation $a = \text{op}$ is executed only if predicate p is true,^{*} and is ignored otherwise. We will also use the notation \bar{p} to refer to the complement of predicate p . The goal of the ψ -SSA form advocated in this chapter is to express these conditional definitions while keeping the static single assignment property.

*Gated-SSA form
dependence, flow-
 ψ -function
back-end, compiler
guard
gate
predicate*

$a = \text{op1};$	$a_1 = \text{op1};$	$a_1 = \text{op1};$	$a_1 = \text{op1};$
$\text{if } (p)$	$\text{if } (p)$	then	then
then			
$a = \text{op2};$	$a_2 = \text{op2};$	$p? a_2 = \text{op2};$	$p? a_2 = \text{op2};$
	$a_3 = \phi(a_1, a_2)$		$a_3 = \psi(a_1, p?a_2)$
$= a$	$= a_3$	$= a_{??}$	$= a_3$
(a) non-SSA	(b) SSA	(c) With predication	(d) ψ -SSA

Fig. 17.1 SSA representation

17.2 Definition and Construction

Predicated operations are used to convert control-flow regions into straight-line code. Predicated operations may be used by the intermediate representation in an early stage of the compilation process as a result of inlining intrinsic functions. Later on, the compiler may also generate predicated operations through if-conversion optimizations as described in Chapter 21.

In figure 17.1(c), the use of a on the last instruction refers to the variable a_1 if p is false, or to the variable a_2 if p is true. These multiple reaching definitions on the use of a cannot be represented by the standard SSA representation. One possible representation would be to use the Gated-SSA form*, presented in Chapter 18. In such a representation, the ϕ -function would be augmented with the predicate p to tell which value between a_1 and a_2 is to be considered. However, Gated-SSA is a completely different intermediate representation where the control flow is no longer represented. This representation is more suited for program interpretation than for optimizations at code generation level as addressed in this chapter. Another possible representation would be to add a reference to a_1 on the definition of a_2 . $p? a_2 = \text{op2} | a_1$ would have the following semantic: a_2 takes the value computed by op2 if p is true, or holds the value of a_1 if p is false. The use of a on the last instruction of Figure 17.1(c) would now refer to the variable a_2 , which holds the correct value. The drawback of this representation is that it adds dependencies* between operations (here a flow dependence from op1 to op2), which would prevent code reordering for scheduling.

Our solution is presented in figure 17.1(d). The ϕ -function of the SSA code with control flow is “replaced” by a ψ -function* on the corresponding predicated code, with information on the predicate associated with each argument. This representation is adapted to code optimization and code generation on a low-level intermediate representation.* A ψ -function $a_0 = \psi(p_1?a_1, \dots, p_i?a_i, \dots, p_n?a_n)$ defines one variable, a_0 , and takes a variable number of arguments a_i ; each argument a_i is associated with a predicate p_i . In the notation, the predicate p_i will be omitted if $p_i \equiv \text{true}$.*

A ψ -function has the following properties:

dependence, control-

- *It is an operation* : A ψ -function is a regular operation. It can occur at any location in a basic block where a regular operation is valid. Each argument a_i , and each predicate p_i , must be dominated by its definition.
- *It is predicated* : A ψ -function is a predicated operation, under the predicate $\bigcup_{k=1}^n p_k$, although this predicate is not explicit in the representation.
- *It has an ordered list of arguments* : The order of the arguments in a ψ -function is significant. A ψ -function is evaluated from left to right. The value of a ψ -function is the value of the right most argument whose predicate evaluates to true.
- *Rule on predicates* : The predicate p_i associated with the argument a_i in a ψ -function must be included in or equal to the predicate on the definition of the variable a_i . In other words, for the code $q ? a_i = \text{op}; a_0 = \psi(\dots, p_i ? a_i, \dots)$, we must have $p_i \subseteq q$ (or $p_i \Rightarrow q$).

<pre> if (p) then a₁ = 1; else a₂ = -1; x₁ = φ(a₁, a₂) if (q) then a₃ = 0; x₂ = φ(x₁, a₃) </pre> <p>(a) control-flow code</p>	<pre> p? a₁ = 1; p? a₂ = -1; x₁ = ψ(p?a₁, p?a₂) q? a₃ = 0; x₂ = ψ(p?a₁, p?a₂, q?a₃) </pre> <p>(b) Predicated code</p>
--	---

Fig. 17.2 ψ -SSA with non-disjoint predicates

A ψ -function can represent cases where variables are defined on arbitrary independent predicates such as p and q in the example of Figure 17.2 : For this example, during the SSA construction a unique variable a was renamed into the variables a_1 , a_2 and a_3 and the variables x_1 and x_2 were introduced to merge values coming from different control-flow paths. In the control-flow version of the code, there is a control-dependence* between the basic-block that defines x_1 and the operation that defines a_3 , which means the definition of a_3 must be executed after the value for x_1 has been computed. In the predicated form of this example, there is no longer any control dependencies between the definitions of a_1 , a_2 and a_3 . A compiler transformation can now freely move these definitions independently of each other, which may allow more optimizations to be performed on this code. However, the semantics of the original code requires that the definition of a_3 occurs after the definitions of a_1 and a_2 . The order of the arguments in a ψ -function gives information on the original order of the definitions. We take the convention that the order of the arguments in a ψ -function

is, from left to right, equal to the original order of their definitions, from top to bottom, in the control-flow dominance tree of the program in a non-SSA representation. This information is needed to maintain the correct semantics of the code during transformations of the ψ -SSA representation and to revert the code back to a non ψ -SSA representation.

The construction of the ψ -SSA representation is a small^{*} modification on the standard algorithm to built an SSA representation (see Section 3.1). The insertion of ψ -functions is performed during the SSA renaming^{*} phase. During the SSA renaming phase, basic blocks are processed in their dominance order, and operations in each basic block are scanned from top to bottom. On an operation, for each predicated definition of a variable, a new ψ -function will be inserted just after the operation : Consider the definition of a variable x under predicate p_2 ($p_2 ? x = \text{op}$); suppose x_1 is the current version of x before to proceeding op , and that x_1 is defined through predicate p_1 (possibly true); after renaming x into a freshly created version, say x_2 , a ψ -function of the form $x = \psi(p_1?x_1, p_2?x)$, is inserted right after op . Then renaming of this new operation proceeds. The first argument of the ψ -function is already renamed and thus is not modified. The second argument is renamed into the current version of x which is x_2 . On the definition of the ψ -function, the variable x is given a new name, x_3 , which becomes the current version for further references to the x variable. This insertion and renaming of a ψ -function is shown on Figure 17.3.

$p_2? x = \text{op}$	$p_2? x = \text{op}$ $x = \psi(p_1?x_1, p_2?x)$	$p_2? x_2 = \text{op}$ $x = \psi(p_1?x_1, p_2?x)$	$p_2? x_2 = \text{op}$ $x_3 = \psi(p_1?x_1, p_2?x_2)$
(a) Initial	(b) ψ -insertion	(c) op-renaming	(d) ψ -renaming

Fig. 17.3 Construction and renaming of ψ -SSA

ψ -functions can also be introduced in an SSA representation by applying an if-conversion transformation, such as the one that is described in Chapter 21^{*}. Local transformations on control-flow patterns can also require to replace ϕ -functions by ψ -functions.

17.3 SSA algorithms

With this definition of the ψ -SSA representation, implicit data-flow links^{*} to predicated operations are now explicitly expressed through ψ -functions. Usual algorithms that perform optimizations or transformations on the SSA representation can now be easily adapted to the ψ -SSA representation, without compromising the efficiency of the transformations performed. Actually, within the ψ -SSA representation, predicated definitions behave exactly the same as non pred-

icated ones for optimizations on the SSA representation. Only the ψ -functions have to be treated in a specific way. As an example, the classical constant propagation algorithm under SSA can be easily adapted to the ψ -SSA representation. In this algorithm, the only modification is that ψ -functions have to be handled with the same rules as the ϕ -functions. Other algorithms such as dead code elimination^{*} (see Chapter 3), global value numbering, partial redundancy elimination^{*} (see Chapter 12), and induction variable analysis^{*} (see Chapter 11), are examples of algorithm that can easily be adapted to this representation with minor efforts.

constant propagation
dead code elimination
global value numbering
partial redundancy elimination
induction variable recognition
 ψ -inlining
 ψ -reduction
 ψ -projection
 ψ -predication

17.4 Psi-SSA algorithms

In addition to standard algorithms that can be applied to ψ -functions and predicated code, a number of specific transformations can be performed on the ψ -functions, namely ψ -Inlining, ψ -reduction, ψ -projection, ψ -permutation and ψ -promotion. For a ψ -function $a_0 = \psi(p_1?a_1, \dots, p_i?a_i, \dots, p_n?a_n)$, those transformations are defined as follows:

ψ -Inlining^{*} recursively replaces in a ψ function an argument a_i that is defined on another ψ function by the arguments of this other ψ -function. The predicate p_i associated with argument a_i will be distributed with an \wedge operation over the predicates associated with the inlined arguments. This is shown in figure 17.4.

$$\begin{array}{ll}
 \begin{array}{l}
 a_1 = \text{op1} \\
 p_2? a_2 = \text{op2} \\
 x_1 = \psi(a_1, p_2?a_2) \\
 p_3? a_3 = \text{op3} \\
 x_2 = \psi(p_1?x_1, p_3?a_3)
 \end{array} &
 \begin{array}{l}
 a_1 = \text{op1} \\
 p_2? a_2 = \text{op2} \\
 x_1 = \psi(a_1, p_2?a_2) // \text{dead} \\
 p_3? a_3 = \text{op3} \\
 x_2 = \psi(p_1?a_1, p_1 \wedge p_2?a_2, p_3?a_3)
 \end{array}
 \end{array}$$

Fig. 17.4 ψ -Inlining of the definition of x_1

ψ -Reduction^{*} removes from a ψ -function an argument a_i whose value will always be overridden by arguments on its right in the argument list. An argument a_i associated with predicate p_i can be removed if $p_i \subseteq \bigcup_{k=i+1}^n p_k$. This can be illustrated by the example of Figure 17.5.

ψ -Projection^{*} creates from a ψ -function a new ψ -function on a restricted predicate say p . In this new ψ -function, an argument a_i initially guarded by p_i shall be guarded by the conjunction $p_i \wedge p$. If p_i is known to be disjoint with p , a_i actually contributes no value to the ψ -function and thus can be removed. ψ -projection on predicate p is usually performed when the result of a ψ -function is used in an operation predicated by p . This is illustrated in Figure 17.6.

ψ -permutation
 ψ -promotion
speculation
predication, partial

$$\begin{array}{l} a_1 = \text{op1} \\ p_2? a_2 = \text{op2} \\ \overline{p}_2? a_3 = \text{op3} \\ x_2 = \psi(a_1, p_2? a_2, \overline{p}_2? a_3) \end{array}$$

$$\begin{array}{l} a_1 = \text{op1} \\ p_2? a_2 = \text{op2} \\ \overline{p}_2? a_3 = \text{op3} \\ x_2 = \psi(p_2? a_2, \overline{p}_2? a_3) \end{array}$$

Fig. 17.5 ψ -reduction. The first argument a_1 of the ψ -function can safely be removed

$$\begin{array}{ll} \begin{array}{l} p_2? a_2 = \text{op2} \\ \overline{p}_2? a_3 = \text{op3} \\ x_2 = \psi(p_2? a_2, \overline{p}_2? a_3) \\ p_2? y_1 = x_2 \end{array} & \begin{array}{l} p_2? a_2 = \text{op2} \\ \overline{p}_2? a_3 = \text{op3} \\ x_2 = \psi(p_2? a_2, \overline{p}_2? a_3) \\ x_3 = \psi(p_2? a_2) \\ p_2? y_1 = x_3 \end{array} \end{array}$$

Fig. 17.6 ψ -projection of x_2 on p_2 . Second argument a_3 can be removed.

ψ -permutation* changes the order of the arguments in a ψ -function. In a ψ -function the order of the arguments is significant. Two arguments in a ψ -function can be permuted if the intersection of their associated predicate in the ψ -function is empty. An example of such a permutation is shown on Figure 17.7.

$$\begin{array}{ll} \begin{array}{l} \overline{p}_2? a_3 = \text{op3} \\ p_2? a_2 = \text{op2} \\ x_2 = \psi(p_2? a_2, \overline{p}_2? a_3) \end{array} & \begin{array}{l} \overline{p}_2? a_3 = \text{op3} \\ p_2? a_2 = \text{op2} \\ x_2 = \psi(\overline{p}_2? a_3, p_2? a_2) \end{array} \end{array}$$

Fig. 17.7 ψ -permutation of arguments a_2 and a_3

ψ -promotion* changes one of the predicates used in a ψ -function by a larger predicate. Promotion must obey the following condition so that the semantics of the ψ -function is not altered by the transformation : consider an operation $a_0 = \psi(p_1?x_1, \dots, p_i?x_i, \dots, p_n?x_n)$ promoted into $a_0 = \psi(p_1?x_1, \dots, p'_i?x_i, \dots, p_n?x_n)$ with $p_i \subseteq p'_i$, then p'_i must fulfill

$$\left(p'_i \setminus \bigcup_{k=i}^n p_k \right) \cap \bigcup_{k=1}^{i-1} p_k = \emptyset \quad (17.1)$$

where $p'_i \setminus \bigcup_{k=i}^n p_k$ corresponds to the possible increase of the predicate of the ψ -function, $\bigcup_{k=1}^n p_k$. This promotion must also satisfy the properties of ψ -functions, and in particular, that the predicate associated with a variable in a ψ -function must be included in or equal to the predicate on the definition of that variable (which itself can be a ψ -function). A simple ψ -promotion is illustrated in Figure 17.8(c).

The ψ -SSA representation can be used on a partially predicated architecture*, where only a subset of the instructions supports a predicate operand. Figure 17.8 shows an example where some code with control-flow edges was transformed into a linear sequence of instructions. Taking the example of an architecture

			<i>speculation</i> <i>SSA!destruction</i> <i>phiweb</i> <i>C-SSA</i>
if (p)			
then	$a_1 = \text{ADD } i_1, 1;$	$a_1 = \text{ADD } i_1, 1;$	
else	$a_2 = \text{ADD } i_1, 2;$	$a_2 = \text{ADD } i_1, 2;$	
	$x = \phi(a_1, a_2)$	$x = \psi(p?a_1, \bar{p}?a_2)$	
(a) Control flow	(b) ψ -SSA	(c) after ψ -promotion	

Fig. 17.8 ψ -SSA for partial predication. ψ -promotion of argument a_1

where the ADD operation cannot be predicated, the ADD operation must be speculated^{*} under the true predicate. On an architecture where the ADD operation can be predicated, it may also be profitable to perform speculation in order to reduce the number of predicates on predicated code and to reduce the number of operations to compute these predicates. Once speculation has been performed on the definition of a variable used in a ψ -function, the predicate associated with this argument can be promoted, provided that the semantic of the ψ -function is maintained (Equation 17.1).

Usually, the first argument of a ψ -function can be promoted under the true predicate. Also, when disjoint conditions are computed, one of them can be promoted to include the other conditions, usually reducing the number of predicates. A side effect of this transformation is that it may increase the number of copy instructions to be generated during the ψ -SSA destruction phase, as will be explained in the following section.

17.5 Psi-SSA destruction

The SSA destruction^{*} phase reverts an SSA representation into a non-SSA representation. This phase must be adapted to the ψ -SSA representation. This algorithm uses ψ - ϕ -webs to create a conventional ψ -SSA representation. The notion of ϕ -webs is^{*} extended to ϕ and ψ operations so as to derive the notion of conventional ψ -SSA (ψ -C-SSA) form. A ψ - ϕ -web is a non empty, minimal, set of variables such that if two variables are referenced on the same ϕ or ψ -function then they are in the same ψ - ϕ -web. The property of the ψ -C-SSA form is that the renaming into a single variable of all variables that belong to the same ψ - ϕ -web, and the removal of the ψ and ϕ functions, results in a program with the same semantics as the original program.

Now, consider Figure 17.9 to illustrate the transformations that must be performed to convert a program from a ψ -SSA form into a program in ψ -C-SSA form.

$p? b = \dots$	$p? b = \dots$	$p? b = \dots$
$a = \dots$	$a = \dots$	$x = \dots$
$x = \psi(a, p?b)$	$p? c = b$	$p? x = b$
	$x = \psi(a, p?c)$	
(a) ψ-T-SSA form	(b) ψ-C-SSA form	(c) non-SSA form
$a = \dots$	$a = \dots$	$x = \dots$
$b = \dots$	$b = \dots$	$b = \dots$
$x = \psi(a, p?b)$	$p? c = b$	$p? x = b$
	$x = \psi(a, p?c)$	
(d) ψ-T-SSA form	(e) ψ-C-SSA form	(f) non-SSA form
$a = \dots$	$a = \dots$	$x = \dots$
$p? b = \dots$	$d = a$	$y = x$
$q? c = \dots$	$p? b = \dots$	$p? x = \dots$
$x = \psi(a, p?b)$	$q? c = \dots$	$q? y = \dots$
$y = \psi(a, q?c)$	$x = \psi(a, p?b)$	$y = \psi(d, q?c)$
(g) ψ-T-SSA form	(h) ψ-C-SSA form	(i) non-SSA form

Fig. 17.9 Non-conventional ψ -SSA (ψ -T-SSA)* form, ψ -C-SSA forms and non-SSA form after destruction

Looking at the first example (Figure 17.9(a)), the dominance order of the definitions for the variables a and b differs from their order from left to right in the ψ -function. Such code may appear after a code motion algorithm has moved the definitions for a and b relatively to each other. Here, the renaming of the variables a , b and x into a single variable will not restore the semantics of the original program. The order in which the definitions of the variables a , b and x occur must be corrected. This is done through the introduction of the variable c that is defined as a predicated copy of the variable b , after the definition of a . Now, the renaming of the variables a , c and x into a single variable will result in the correct behavior.

In Figure 17.9(d) the definition of the variable b has been speculated. However, the semantics of the ψ -function is that the variable x will only be assigned the value of b when p is true. A new variable c must be defined as a predicated copy of the variable b , after the definition of b and p ; in the ψ -function, variable b is then replaced by variable c . The renaming of variables a , c and x into a single variable will now follow the correct behavior.

In Figure 17.9(g), the renaming of the variables a , b , c , x and y into a single variable will not give the correct semantics. In fact, the value of a used in the second ψ -function would be overridden by the definition of b before the definition of the variable c . Such code will occur after copy folding has been applied

on a ψ -SSA representation. We see that the value of a has to be preserved before the definition of b . This is done through the definition of a new variable (d here), resulting in the code given in Figure 17.9(h). Now, the variables a , b and x can be renamed into a single variable, and the variables d , c and y will be renamed into another variable, resulting in a program in a non-SSA form with the correct behavior.

We will now present an algorithm that will transform a program from a ψ -SSA form into its ψ -C-SSA form. This algorithm is made of three parts.

- *Psi-normalize*: This phase puts all ψ -functions^{*} in what we call a *normalized* form.
- *Psi-web*^{*}: This phase grows ψ -webs from ψ -functions, and introduces repair code where needed such that each ψ -web is interference free.
- *Phi-web* : This phase is the standard SSA-destruction algorithm (e.g., see Chapter 22) with the additional constraint that all variables in a ψ -web must be coalesced together. This can be done using the pinning^{*} mechanism presented in Chapter 22.

We detail now the implementation of each of the first two parts.

normalized- ψ
 ψ -web
pining
coalescing
normalized- ψ
 ψ -T-SSA

17.5.1 Psi-normalize

We define the notion of *normalized- ψ* ^{*}. The normalized form of a ψ -function has two characteristics:

- The order of the arguments in a normalized- ψ -function is, from left to right, equal to the order of their definitions, from top to bottom, in the control-flow dominance tree.
- The predicate associated with each argument in a normalized- ψ -function is equal to the predicate used on the unique definition of this argument.

These two characteristics correspond respectively to the two cases presented in Figure 17.9(a) and Figure 17.9(d). When some arguments of a ψ -function are also defined by ψ -functions, the normalized- ψ characteristics must hold on a virtual ψ -function where ψ -Inlining has been performed on these arguments.

When ψ -functions are created during the construction of the ψ -SSA representation, they are naturally built in their normalized form. Later, transformations are applied to the ψ -SSA representation. Predicated definitions may be moved relatively to each others. Also, operation speculation and copy folding may enlarge the domain of the predicate used on the definition of a variable. These transformations may cause some ψ -functions to be in a non-normalized form.

PSI-normalize implementation

Each ψ -function is processed independently. An analysis of the ψ -functions in a top down traversal of the dominator tree reduces the amount of repair code that is inserted during this pass. We only detail the algorithm for such a traversal.

For a ψ -function $a_0 = \psi(p_1?a_1, \dots, p_i?a_i, \dots, p_n?a_n)$, the argument list is processed from left to right. For each argument a_i , the predicate p_i associated with this argument in the ψ -function and the predicate used on the definition of this argument are compared. If they are not equal, a new variable a'_i is introduced and is initialized at the highest point in the dominator tree after the definition of a_i and p_i . a'_i is defined by the operation $p_i ? a'_i = a_i$. Then, a_i is replaced by a'_i in the ψ -function.

Then, we consider the dominance order of the definition for a_i , with the definition for a_{i-1} . When a_i is defined on a ψ -function, we recursively look for the definition of the first argument of this ψ -function, until a definition on a non- ψ -function is found. If the definition we found for a_i dominates the definition for a_{i-1} , some correction is needed. If the predicates p_{i-1} and p_i are disjoint, a ψ -permutation can be applied between a_{i-1} and a_i , so as to reflect into the ψ -function the actual dominance order of the definitions of a_{i-1} and a_i . If ψ -permutation cannot be applied, a new variable a'_i is created for repair. a'_i is defined by the operation $p_i ? a'_i = a_i$. This copy operation is inserted at the highest point that is dominated by the definitions of a_{i-1} and a_i .¹ Then, a_i is replaced in the ψ -function by a'_i .

The algorithm continues with the argument a_{i+1} , until all arguments of the ψ -function are processed. When all arguments are processed, the ψ is in its normalized form. When all ψ functions are processed, the function will contain only normalized- ψ -functions.

17.5.2 Psi-web

*

The role of the psi-web phase is to repair the ψ -functions that are part of a non interference-free ψ -web.^{*} This case corresponds to the example presented in Figure 17.9(g). In the same way as there is a specific point of use for arguments on ϕ -functions for liveness analysis (e.g., see Section 22.3), we give a definition of the actual point of use of arguments on normalized ψ -functions for liveness analysis. With this definition, liveness^{*} analysis is computed accurately and an interference graph can be built. The cases where repair code is needed can be easily and accurately detected by observing that variables in a ψ -function interfere.

¹ When a_i is defined by a ψ -function, its definition may appear after the definition for a_{i-1} , although the non- ψ definition for a_i appears before the definition for a_{i-1} .

Liveness and interferences in Psi-SSA.

*uses, of ψ -functions
interference
 ψ -web*

Consider the code in Figure 17.10 (b). Instead of using a representation with ψ -functions, predicated definitions have been modified to make a reference to the value the predicated definition will have in case the predicate evaluates to false. We use in this example the notation of the select operator $x = cond? exp1 : exp2$ that assigns $exp1$ to x if $cond$ is true and $exp2$ otherwise. Each of the predicated definitions make an explicit use of the variable immediately to its left in the argument list of the original ψ -function from Figure 17.10 (a). We can see that a renaming of the variables a, b, c and x into a single representative name will still compute the same value for the variable x . Note that this transformation can only be performed on normalized ψ -functions, since the definition of an argument must be dominated by the definition of the argument immediately at its left in the argument list of the ψ -function, and the same predicate must be used on the definition of an argument and with this argument in the ψ operation. Using this equivalence for the representation of a ψ -function, we now give a definition of the point of use for the arguments of a ψ -function.

Definition 3 (use points).* Let $a_0 = \psi(p_1?a_1, \dots, p_i?a_i, \dots, p_n?a_n)$ be a normalized ψ -function. For $i < n$, the point of use of argument a_i occurs at the operation that defines a_{i+1} . The point of use for the last argument a_n occurs at the ψ -function itself.

$\begin{aligned} a &= op1 \\ p? b &= op2 \\ q? c &= op3 \\ x &= \psi(a, p?b, q?c) \end{aligned}$	$\begin{aligned} a &= op1 \\ b &= p ? op2 : a \\ c &= q ? op3 : b \\ x &= c \end{aligned}$
(a) ψ-SSA form	(b) conditional form

Fig. 17.10 ψ -functions and C conditional operations equivalence

Given this definition of point of use of ψ -function arguments, and using the usual point of use of ϕ -function arguments, a traditional liveness analysis can be run. Then an interference graph can be built to collect the interferences between variables involved in ψ or ϕ -functions. For the construction of the interference graph, an interference between two variables that are defined on disjoint predicates can be ignored.

Repairing interferences on ψ -functions.

We now present an algorithm that resolves the interferences as it builds the ψ -webs. A pseudo-code of this algorithm is given in Figure 17.1. First, the ψ -

*live-range
liveness check*

webs are initialized with a single variable per ψ -web. Then, ψ -functions are processed one at a time, in no specific order, merging when non-interfering the ψ -webs of its operands together. Two ψ -webs interfere if at least one variable in the first ψ -web interferes with at least one variable in the other one. The arguments of the ψ -function, say $a_0 = \psi(p_1?a_1, \dots, p_i?a_i, \dots, p_n?a_n)$, are processed from right (a_n) to left (a_1). If the ψ -web that contains a_i does not interfere with the ψ -web that contains a_0 , they are merged together. Otherwise, repair code is needed. A new variable, a'_i , is created and is initialized with a predicated copy $p_i ? a'_i = a_i$, inserted just above the definition for a_{i+1} , or just above the ψ -function in case of the last argument. The current argument a_i in the ψ -function is replaced by the new variable a'_i . The interference graph is updated. This can be done by considering the set of variables, say U , a_i interferes with. For each $u \in U$, if u is in the merged ψ -web, it should not interfere with a'_i ; if the definition of u dominates the definition of a_i , it is live-through the definition of a_i , thus is should be made interfering with a'_i ; last, if the definition of a_i dominates the definition of b , it should be made interfering only if this definition is within the live-range^{*} of a'_i (see Chapter 9).

Algorithm 17.1: ψ -webs merging during the processing of a ψ -function
 $a_0 = \psi(p_1?a_1, \dots, p_i?a_i, \dots, p_n?a_n)$

```

1 begin
2   let psiWeb be the web containing  $a_0$  ;
3   foreach  $a_i$  in  $[a_n, a_{n-1}, \dots, a_1]$  do
4     let opndWeb be the web containing  $a_i$  ;
5     if opndWeb ≠ psiWeb if IGraph.interfere(psiWeb, opndWeb) let  $a'_i$  be a freshly
6     created variable ;
7     let Ci be a new predicated copy  $p_i ? a'_i \leftarrow a_i$  ;
8     let op be the operation that defines  $a_{i+1}$ ,  $a_i.def.op$ , or the psi operation ;
9     while op is a  $\psi$ -function replace op by the operation that defines its first
10    argument ;
11    append Ci right before op ;
12    replace  $a_i$  by  $a'_i$  in the  $\psi$ -function;
13    opndWeb ←  $\{a'_i\}$  ;
14    foreach  $u$  in IGraph.interferenceSet( $a_i$ ) do
15      if  $u \notin \psi$  if  $u.def.op$  dominates  $a_i.def.op \vee a'_i \in \text{livein}(u.def.op)$ 
16        IGraph.addInterference( $a'_i, a$ ) ;
17    end
18  psiWeb ← psiWeb ∪ opndWeb ;
19 end
20 end
21 end

```

Consider the code in Figure 17.11 to see how this algorithm works. The liveness on the ψ -function creates a live-range for variable a that extends down to the definition of b , but not further down. Thus, the variable a does not interfere with the variables b , c or x . The live-range for variable b extends down to its use in the definition of variable d . This live-range interferes with the variables

c and x . The live-range for variable c extends down to its use in the ψ -function that defines the variable x . At the beginning of the processing on the ψ -function $x = \psi(p?a, q?b, r?c)$, ψ -webs are singletons $\{a\}$, $\{b\}$, $\{c\}$, $\{x\}$, $\{d\}$. The argument list is processed from right to left i.e., starting with variable c . $\{c\}$ does not interfere with $\{x\}$, they can be merged together, resulting in $psiWeb = \{x, c\}$. Then, variable b is processed. Since it interferes with both x and c , repairing code is needed. A variable b' is created, and is initialized just below the definition for b , as a predicated copy of b . The interference graph is updated conservatively, with no changes. $psiWeb$ now becomes $\{x, b', c\}$. Then variable a is processed, and as no interference is encountered, $\{a\}$ is merged to $psiWeb$. The final code after SSA destruction is shown in Figure 17.11(c).

$p? a = ...$ $q? b = ...$ $r? c = ...$ $x = \psi(p?a, q?b, r?c)$ $s? d = b + 1$	$p? a = ...$ $q? b = ...$ $q? b' = b$ $r? c = ...$ $x = \psi(p?a, q?b', r?c)$ $s? d = b + 1$	$p? x = ...$ $q? b = ...$ $q? x = b$ $r? x = ...$ $s? d = b + 1$
(a) before processing the ψ -function	(b) after processing the ψ -function	(c) after actual coalescing

Fig. 17.11 Elimination of ψ live-interference

In this chapter we mainly described the ψ -SSA representation and we detailed specific transformations that can be performed thanks to this representation. More details on the implementation of the ψ -SSA algorithms, and figures on the benefits of this representation, can be found in [262] and [98].

We mentioned in this chapter that a number of classical SSA-based algorithm can be easily adapted to the ψ -SSA representation, usually by just adapting the rules on the ϕ -functions to the ψ -functions. Among these algorithm, we can mention the constant propagation algorithm described in [286], dead code elimination [204], global value numbering [79], partial redundancy elimination [75] and induction variable analysis [289] which have already been implemented into a ψ -SSA framework.

There are also other SSA representations that can handle predicated instruction, of which is the Predicated SSA representation [64]. This representation is targeted at very low level optimization to improve operation scheduling in pres-

17.6 Additional reading

ence of predicated instructions. Another representation is the Gated SSA form, presented in Chapter 18.

The ψ -SSA destruction algorithm presented in this chapter is inspired from the SSA destruction algorithm of Sreedhar et al. [255] that introduces repair code when needed as it grows ϕ -webs from ϕ -functions. The phi-web phase mentioned in this chapter to complete the ψ -SSA destruction algorithm can use exactly the same approach by simply initializing ψ - ϕ -webs by ψ -webs.

CHAPTER 18

Graphs and gating functions

J. Stanier

Progress: 90%



Text and figures formating in progress

18.1 Introduction

Many compilers represent the input program as some form of graph in order to aid analysis and transformation. A cornucopia of program graphs have been presented in the literature and implemented in real compilers. Therefore it comes as no surprise that a number of program graphs use SSA concepts as the core principle of their representation. These range from very literal translations of SSA into graph form to more abstract graphs which are implicitly SSA. We aim to introduce a selection of program graphs which use SSA concepts, and examine how they may be useful to a compiler writer.

One of the seminal graph representations is the Control-Flow Graph (CFG), which was introduced by Allen to explicitly represent possible control paths in a program. Traditionally, the CFG is used to convert a program into SSA form. Additionally, representing a program in this way makes a number of operations simpler to perform, such as identifying loops, discovering irreducibility and performing interval analysis techniques.

The CFG models control flow, but many graphs model *data flow*. This is useful as a large number of compiler optimizations are based on data flow. The graphs we consider in this chapter are all data-flow graphs, representing the data dependencies in a program. We will look at a number of different SSA-based graph representations. These range from those which are a very literal translation of SSA into a graph form to those which are more abstract in nature. An introduction to each graph will be given, along with diagrams to show how

sample programs look when translated into that particular graph. Additionally, we will touch on the literature describing a usage of a given graph with the application that it was used for.

18.2 Data-flow Graphs

The data-flow graph (DFG) is a directed graph $G = (V, E)$ where the edges E represent the flow of data from the result of one operation to the input of another. An instruction executes once all of its input data values have been consumed. When an instruction executes it produces a new data value which is propagated to other connected instructions.

Whereas the CFG imposes a total ordering on instructions, the DFG has no such concept, nor does the DFG contain whole program information. Thus, target code cannot be generated directly from the DFG. The DFG can be seen as a companion to the CFG, and they can be generated alongside each other. With access to both graphs, optimisations such as dead code elimination, constant folding and common subexpression elimination can be performed effectively. However, keeping both graphs updated during optimisation can be costly and complicated.

18.3 The SSA Graph

We begin our exploration with a graph that is very similar to SSA: the SSA Graph. Notice that many different variations exist in the literature. For us, an SSA Graph consists of vertices which represent operations (such as add and load) or ϕ -functions, and directed edges connect uses to definitions of values. The outgoing edges from a vertex represent the arguments required for that operation, and the incoming edge(s) to a vertex represents the propagation of that operation's result(s) after it has been computed. This graph is therefore a *demand-based* representation. In order to compute a vertex, we must first *demand* the results of the operands and then perform the operation indicated on that vertex. Reversing the direction of each edge would provide the widely used *data-based* representation. The SSA Graph can be constructed from a program in SSA form by *explicitly* adding use-definition chains. We present some sample code in Figure 18.1 which is then translated into an SSA Graph. Note that each node mixes up the operation and the variable(s) it defines, as actual data structures might be able to find one from the other.

The textual representation of SSA is much easier for a human to read. However, the primary benefit of representing the input program in this form is that the compiler writer is able to apply a wide array of graph-based optimizations by

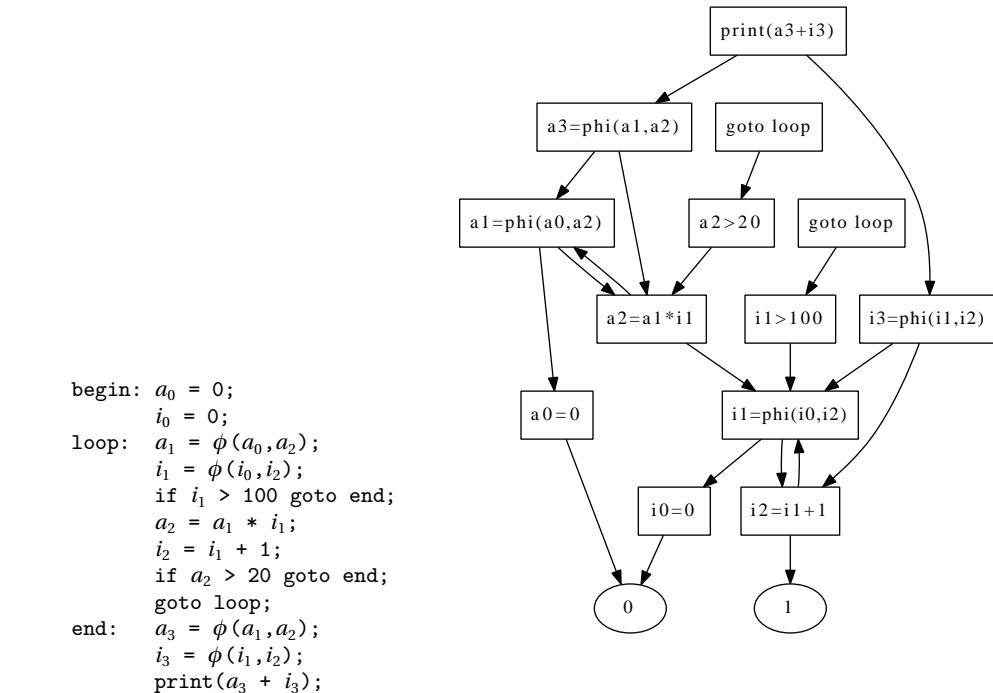


Fig. 18.1 Some SSA code translated into a (demand-based) SSA graph.

using standard graph traversal and transformation techniques. It is possible to augment the SSA Graph to model memory dependencies. This can be achieved by adding *state edges* that enforce an order of interpretation. These edges are extensively used in the Value State Dependence Graph, which we will look at later, after we touch on the concept of gating functions.

In the literature, the SSA Graph has been used to detect a variety of induction variables in loops, also for performing instruction selection techniques, operator strength reduction, rematerialization, and has been combined with

an extended SSA language to aid compilation in a parallelizing compiler. The reader should note that the exact specification of what constitutes an SSA Graph changes from paper to paper. The essence of the intermediate representation (IR) has been presented here, as each author tends to make small modifications for their particular implementation.

18.3.1 Finding induction variables with the SSA Graph

We illustrate the usefulness of the SSA Graph through a basic induction variable (IV) recognition technique. A more sophisticated technique is developed in Chapter 11. Given that a program is represented as an SSA Graph, the task of finding induction variables is simplified. A *basic linear induction variable* i is a variable that appears only in the form:

```
i = 10
loop
...
i = i + k
...
endloop
```

where k is a constant or loop invariant. A simple IV recognition algorithm is based on the observation that each basic linear induction variable will belong to a non-trivial strongly connected component (SCC) in the SSA graph. SCCs can be easily discovered in linear time using any depth first search traversal. Each such SCC must conform to the following constraints:

- The SCC contains only one ϕ -function at the header of the loop.
- The SCC contains only addition and subtraction operators, and the right operand of the subtraction is not part of the SCC (no $i=n-i$ assignments).
- The other operand of each addition or subtraction is loop invariant.

This technique can be expanded to detect a variety of other classes of induction variables, such as wrap-around variables, non-linear induction variables and nested induction variables. Scans and reductions also show a similar SSA Graph pattern and can be detected using the same approach.

18.4 Program Dependence Graph

The Program Dependence Graph (PDG) represents both control and data dependencies together in one graph. The PDG was developed to aid optimizations

requiring reordering of instructions and graph rewriting for parallelism, as the strict ordering of the CFG is relaxed and complemented by the presence of data dependence information. The PDG is a directed graph $G = (V, E)$ where nodes V are statements, predicate expressions or region nodes, and edges E represent either control or data dependencies. Thus, the set of all edges E has two distinct subsets: the control dependence subgraph E_C and the data dependence subgraph E_D . Similar to the CFG, a PDG also has two nodes ENTRY and EXIT, through which control flow enters and exits the program respectively. For this purpose, it is assumed that every node of the CFG is reachable from the entry node and can reach the exit node.

Statement nodes represent instructions in the program. Predicate nodes test a conditional statement and have true and false edges to represent the choice taken on evaluation of the predicate. Region nodes group control dependencies with identical source and label together. If the control dependence for a region node is satisfied, then it follows that all of its children can be executed. Thus, if a region node has three different control-independent statements as immediate children, then these could potentially be executed in parallel. Diagrammatically, rectangular nodes represent statements, diamond nodes predicates, and circular nodes are region nodes. Dashed edges represent control dependence, and solid edges represent data dependence. Loops in the PDG are represented by back edges in the control dependence subgraph. We show an example code translated into a PDG in Figure 18.2.

Construction of the PDG is tackled in two steps from the CFG: construction of the control dependence subgraph and construction of the data dependence subgraph. The construction of the control dependence subgraph falls itself into two steps. First, control dependence between statements and predicate nodes are computed; then region nodes are added. A node w is said to be *control dependent on node u along CFG edge (u, v)* or simply *control dependent on edge (u, v)* if w post-dominates v and w does not strictly post-dominate u . Control dependence between nodes is nothing else than post-dominance frontier, i.e., dominance frontier on the reverse CFG. A slight difference in its construction process is that the corresponding control dependence edges from u to w is labelled by the boolean value taken by the predicate computed in u when branching on edge (u, v) . To compute the control dependence subgraph, the postdominator tree is constructed for the procedure. Then, the ENTRY node is added with one edge labelled true pointing to the CFG entry node, and another labelled false going to the CFG exit node. Then, let S consist of all edges (A, B) in the CFG such that B is not an ancestor of A in the postdominator tree. Each of these edges has an associated label true or false. Then, each edge in S is considered in turn. Given (A, B) , the postdominator tree is traversed backwards from B until we reach A 's parent, marking all nodes visited (including B) as control dependent on A with the label of S .

Next, region nodes are added to the PDG. Each region node summarizes a set of control conditions and “groups” all nodes with the same set of control conditions together. Region nodes are also inserted so that predicate nodes will only

```

begin: i = 1;
loop: if i > 100 goto end;
      a = 2 * B[i];
      A[i] = a;
      i = i + 1;
      if a > 20 goto end;
      goto loop;
end:   return a;
  
```

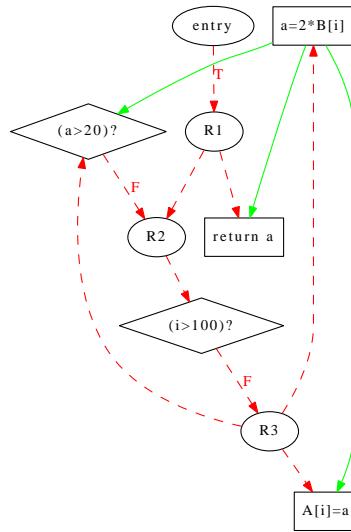


Fig. 18.2 Some code translated into a PDG. Nodes associated with the evolution of the induction variable i are omitted.

have two successors. To begin with, an unpruned PDG is created by checking, for each node of the CFG, which control region it depends on. This is done by traversing the postdominator tree in postorder, and using a hash table to map sets of control dependencies to region nodes. For each node N visited in the postdominator tree, the hash table is checked for an existing region node with the same set CD of control dependencies. If none exists, a new region node R is created with these control dependencies and entered into the hash table. R is made to be the only control dependence predecessor of N . Next, the intersection INT of CD is computed for each immediate child of N in the postdominator tree. If $INT = CD$ then the corresponding dependencies are deleted from the child and replaced with a single dependence on the child's control predecessor. Then, a pass over the graph is made to make sure that each predicate node has a unique successor for each truth value. If more than one exists, the corre-

sponding edges are replaced by a single edge to a freshly created region node that itself points to the successor nodes.

A statement B that is to be executed after a statement A in the original sequential ordering of the program depends on A in the following situations: (flow) B reads to a storage location that was lastly accessed by A through a write; (anti) B writes to a storage location previously accessed through a read by A ; (output) A and B both have a write access to the same storage location. Side effects can also dictate the insertion of a dependence between A and B to force the sequential ordering of the final schedule. Memory accesses can not always be analyzed with enough precision. In the presence of a may-alias between two consecutive accesses, a conservative dependence is to be inserted also. Memory access locations often vary with the iterations of the enclosing loops. Dependence analysis can take advantage of some abstract representations of the access function such as when the memory address can be represented as an affine function of the induction variables. Not only does this enables a refinement of the alias information, but also the dependence can be labelled with a distance vector as a function itself of the loop indices. As an example, a loop indexed by i , that would access array $A[i]$ twice, first as a write at iteration i , then as a read at iteration $2i + 1$, would lead to a flow dependence of distance $i + 1$.

The PDG can also be constructed during parsing. Sequential code can be derived from a PDG, but generating the *minimal size* CFG from a PDG turns out to be an NP-Complete problem. The PDG's structure has been exploited for generating code for vectorisation, and has also been used in order to perform accurate program slicing and testing.

18.4.1 Detecting parallelism with the PDG

The structure of the PDG allows for parallelism to be detected easily. On a regular CFG representation, a scheduler based on data dependencies will generally restrict its scope to hyper-blocks. In this context, code transformations such as loop unrolling or if-conversion (see Chapter 21) that effectively change control dependencies into data dependencies can expose instruction level parallelism. However, the PDG can *directly* be used to detect parallelism. As an example, any node of a CFG loop, that is not contained in an SCC of the PDG (considering *both* control and data dependence edges) can be vectorized. In the example in Figure 18.2, since the statement $A[i] = a$ in the loop do not form an SCC in the PDG, it can be vectorized provided array expansion of variable a . On the other hand, because of the circuit involving the test on a , the statement $a=2*B[i]$ cannot.

18.5 Gating functions and GSA

In SSA form, ϕ -functions are used to identify points where variable definitions converge. However, they cannot be directly *interpreted*, as they do not specify the condition which determines which of the variable definitions to choose. By this logic, we cannot directly interpret the SSA Graph. Being able to interpret our IR is a useful property as it gives the compiler writer more information when implementing optimizations, and also reduces the complexity of performing code generation. Gated Single Assignment form (GSA; sometimes called Gated SSA) is an extension of SSA with *gating functions*. These gating functions are directly interpretable versions of ϕ -nodes, and replace ϕ -nodes in the representation. We usually distinguish the three following forms of gating functions:

- The ϕ_{if} function explicitly represents the condition which determines which ϕ value to select. A ϕ_{if} function of the form $\phi_{if}(P, V_1, V_2)$ has P as a predicate, and V_1 and V_2 as the values to be selected if the predicate evaluates to true or false respectively. This can be read simply as *if-then-else*.
- The ϕ_{entry} function is inserted at loop headers to select the initial and loop carried values. A ϕ_{entry} function of the form $\phi_{entry}(V_{init}, V_{iter})$, has V_{init} as the initial input value for the loop, and V_{iter} as the iterative input. We replace ϕ -functions at loop headers with ϕ_{entry} functions.
- The ϕ_{exit} function determines the value of a variable when a loop terminates. A ϕ_{exit} function of the form $\phi_{exit}(P, V_{final})$ has P as predicate and V_{final} as the definition reaching beyond the loop.

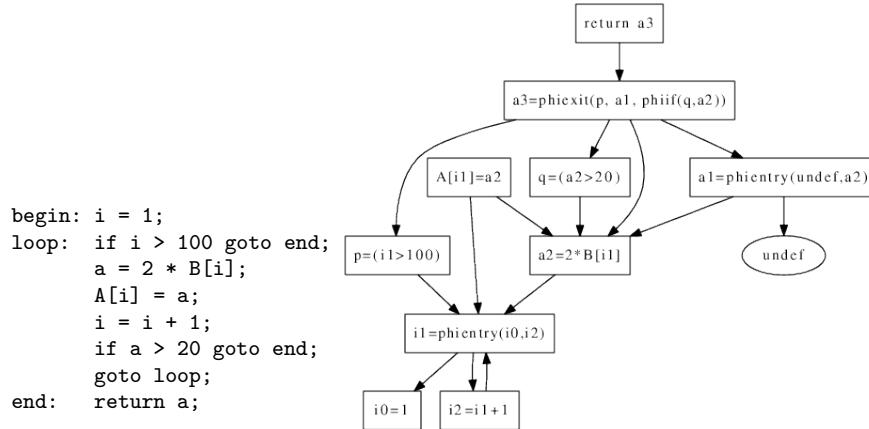


Fig. 18.3 A graph representation of our sample code in GSA form.

It is easiest to understand these gating functions by means of an example. Figure 18.3 shows how our earlier code in Figure 18.2 translates into GSA form. Here, we can see the use of both ϕ_{entry} and ϕ_{exit} gating functions. At the header of our sample loop, the ϕ -function has been replaced by a ϕ_{entry} function which determine between the initial and iterative value of i . After the loop has finished executing, the nested ϕ_{exit} functions selects the correct live-out version of a .

This example shows several interesting points. First, the semantic of both the ϕ_{exit} and ϕ_{if} are strict in their gate: here a_1 or $\phi_{exit}(q, a_2)$ are not evaluated before p is known¹. Similarly, a ϕ_{if} function that results from the nested if-then-else code of Figure 18.4 would be itself nested as $a = \phi_{if}(p, \phi_{if}(q, a_2, a_3), a_1)$. Second, this representation of the program does not allow for an interpreter to decide whether an instruction with a side effect (such as $A[i_1] = a_2$ in our running example) has to be executed or not. Finally, computing the values of gates is highly related to the simplification of path expressions: in our running example a_2 should be selected when the path $\neg p$ followed by q (denoted $\neg p.q$) is taken while a_1 should be selected when the path p is taken; for our nested if-then-else example, a_1 should be selected either when the path $\neg p.r$ is taken or when the path $\neg p.\neg r$ is taken which simplifies to $\neg p$. Diverse approaches can be used to generate the correct nested ϕ_{if} or ϕ_{exit} gating functions.

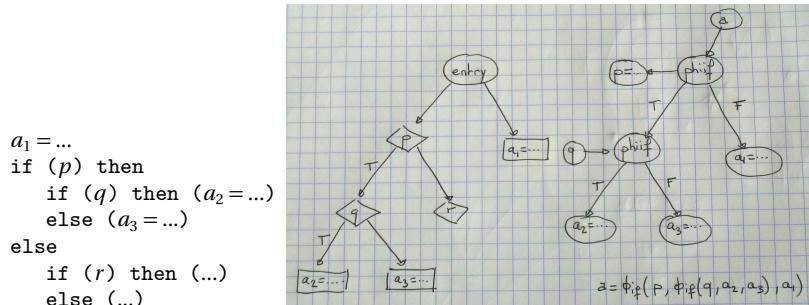


Fig. 18.4 (a) A structured code; (b) the CDG (with region nodes omitted); (c) the DAG representation of the nested gated ϕ_{if}

The more natural way uses a data-flow analysis that computes, for each program points and each variable, its unique reaching definition and the associated set of reaching paths. This set of paths is abstracted using a *path expression*. If the code is not already under SSA, and if at a merge point of the CFG its predecessor basic blocks are reached by different variables, a ϕ -function is inserted. The gates of each operand is set to the path expression of its corresponding incoming edge. If a unique variable reaches all the predecessor basic blocks, the corresponding path expressions are merged. Of course, a classical path com-

¹ As opposed to the ψ -function described in Chapter 17 that would use a syntax such as $a_3 = \phi((p \wedge \neg q)?a_1, (\neg p \wedge q)?a_2)$ instead.

pression technique can be used to minimize the number of visited edges. One can observe the similarities with the ϕ -function placement algorithm described in Section 4.5.

There also exists a relationship between the control dependencies and the gates: from a code already under strict and conventional SSA form, one can derive the gates of a ϕ_{if} function from the control dependencies of its operands. This relationship is illustrated by Figure 18.4 in the simple case of a structured code.

These gating functions are important as the concept will form components of the Value State Dependence Graph later. GSA has seen a number of uses in the literature including analysis and transformations based on data flow. With the diversity of applications (see chapters 24 and 11), many variants of GSA have been proposed. Those variations concern the correct handling of loops in addition to the computation and representation of gates.

By using gating functions it becomes possible to construct IRs based solely on data dependencies. These IRs are sparse in nature compared to the CFG, making them good for analysis and transformation. This is also a more attractive proposition than generating and maintaining both a CFG and DFG, which can be complex and prone to human error. One approach has been to combine both of these into one representation, as is done in the PDG. Alternatively, we can utilize gating functions along with a data-flow graph for an effective way of representing whole program information using data-flow information.

18.5.1 Backwards symbolic analysis with GSA

GSA is useful for performing symbolic analysis. Traditionally, symbolic analysis is performed by forward propagation of expressions through a program. However, complete forward substitution is expensive and can result in a large quantity of unused information and complicated expressions. Instead, *backward*, demand-driven substitution can be performed using GSA which only substitutes *needed* information. Consider the following program:

```
R: JMAX = Expr
S: if(P) then J = JMAX - 1
   else J = JMAX
T: assert(J ≤ JMAX)
```

Fig. 18.5 A program on which to perform symbolic analysis.

If forwards substitution were to be used in order to determine whether the assertion is correct, then the symbolic value of J must be discovered, starting at the top of the program in statement R. Forward propagation through this program results in statement T being $assert(if P then Expr - 1 else Expr) \leq Expr$,

thus the *assert* statement evaluates to true. In real, non-trivial programs, these expressions can get unnecessarily long and complicated.

Using GSA instead allows for backwards, demand-driven substitution. The program above has the following GSA form:

```
R: JMAX1 = Expr
S: if(P) then J1 = JMAX1 - 1
   else J2 = JMAX1
   J3 = φif(P, J1, J2)
T: assert(J3 ≤ JMAX1)
```

Fig. 18.6 Figure 18.5 in GSA form.

Using this backwards substitution technique, we start at statement T, and follow the SSA links of the variables from J_3 . This allows for skipping of any intermediate statements that do not affect variables in T. Thus the substitution steps are:

$$\begin{aligned} J_3 &= \phi_{if}(P, J_1, J_2) \\ &= \phi_{if}(P, JMAX_1 - 1, JMAX_1) \end{aligned}$$

Fig. 18.7 Substitution steps in backwards symbolic analysis.

The backwards substitution then stops because enough information has been found, avoiding the redundant substitution of $JMAX_1$ by $Expr$. In non-trivial programs this can greatly reduce the number of redundant substitutions, making symbolic analysis significantly cheaper.

.....

18.6 Value State Dependence Graph

The gating functions defined in the previous section were used in the development of a sparse data-flow graph IR called the Value State Dependence Graph (VSDG). The VSDG is a directed graph consisting of operation nodes, loop and merge nodes together with value and state dependency edges. Cycles are permitted but must satisfy various restrictions. A VSDG represents a single procedure: this matches the classical CFG.

An example VSDG is shown in Figure 18.8. In (a) we have the original C source for a recursive factorial function. The corresponding VSDG (b) shows both value and state edges and a selection of nodes.

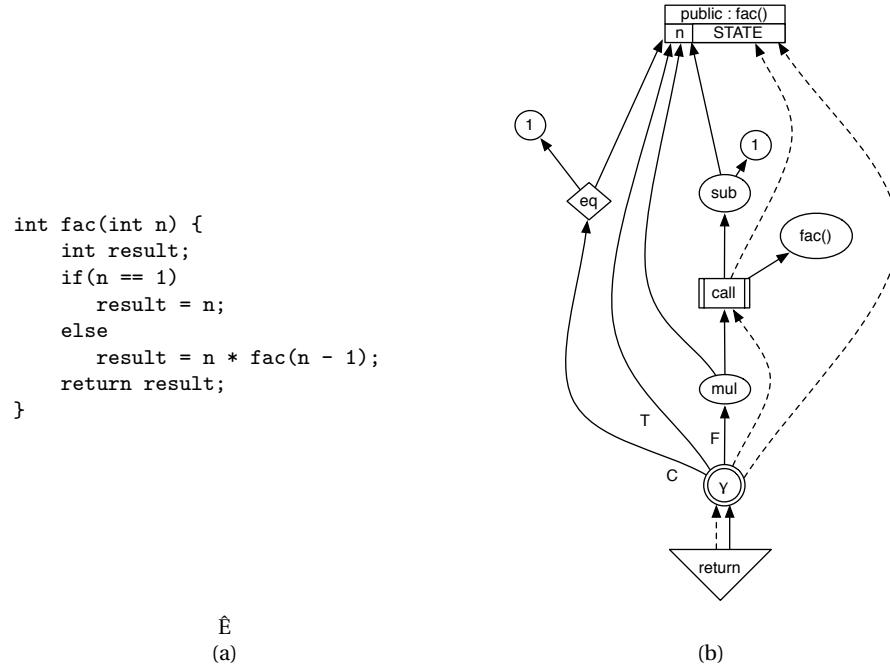


Fig. 18.8 A recursive factorial function, whose VSDG illustrates the key graph components—value dependency edges (solid lines), state dependency edges (dashed lines), a `const` node, a `call` node, a γ -node, a conditional node, and the function entry and exit nodes.

18.6.1 Definition of the VSDG

A VSDG is a labelled directed graph $G = (N, E_V, E_S, \ell, N_0, N_\infty)$ consisting of nodes N (with unique entry node N_0 and exit node N_∞), value dependency edges $E_V \subseteq N \times N$, state dependency edges $E_S \subseteq N \times N$. The labelling function ℓ associates each node with an operator.

The VSDG corresponds to a reducible program, e.g. there are no cycles in the VSDG except those mediated by θ (loop) nodes.

Value dependency (E_V) indicates the flow of values between nodes. State dependency (E_S) represents two things; the first is essential sequential dependency required by the original program, e.g. a given `load` instruction may be required to follow a given `store` instruction without being re-ordered, and a `return` node in general must wait for an earlier loop to terminate even though there might be no value-dependency between the loop and the `return` node. The second purpose is that state dependency edges can be added incrementally until the VSDG corresponds to a unique CFG. Such state dependency edges are called *serializing* edges.

The VSDG is implicitly represented in SSA form: a given operator node, n , will have zero or more E_V -consumers using its value. Note that, in implementation terms, a single register can hold the produced value for consumption at all consumers; it is therefore useful to talk about the idea of an output *port* for n being allocated a specific register, r , to abbreviate the idea of r being used for each edge (n_1, n_2) where $n_2 \in succ(n_1)$.

18.6.2 Nodes

There are four main classes of VSDG nodes: value nodes (representing pure arithmetic), γ -nodes (conditionals), $\hat{\theta}$ -nodes (loops), and state nodes (side-effects). The majority of nodes in a VSDG generate a value based on some computation (add, subtract, etc) applied to their dependent values (constant nodes, which have no dependent nodes, are a special case).

18.6.3 γ -Nodes

The γ -node is similar to the ϕ_{if} gating function in being dependent on a control predicate, rather than the control-independent nature of SSA ϕ -functions.

A γ -node $\gamma(C, T, F)$ evaluates the condition dependency C , and returns the value of T if C is true, otherwise F .

We generally treat γ -nodes as single-valued nodes (contrast θ -nodes, which are treated as tuples), with the effect that two separate γ -nodes with the same condition can be later combined into a tuple using a single test. Figure 18.9 illustrates two γ -nodes that can be combined in this way. Here, we see how two syntactically different programs can map to the same structure in the VSDG.

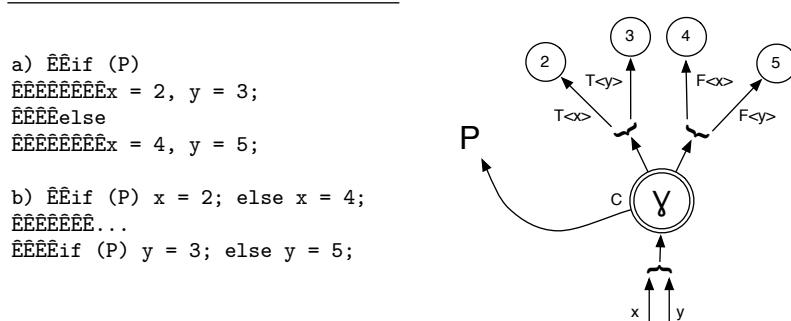


Fig. 18.9 Two different code schemes (a) & (b) map to the same γ -node structure.

18.6.4 θ -Nodes

The θ -node models the iterative behaviour of loops, modelling loop state with the notion of an *internal value* which may be updated on each iteration of the loop. It has five specific ports which represent dependencies at various stages of computation. The θ -node corresponds to a merge of the ϕ_{entry} and ϕ_{exit} nodes in Gated SSA.

A θ -node $\theta(C, I, R, L, X)$ sets its internal value to initial value I then, while condition value C holds true, sets L to the current internal value and updates the internal value with the repeat value R . When C evaluates to false computation ceases and the last internal value is returned through the X port.

A loop which updates k variables will have: a single condition port C , initial-value ports I_1, \dots, I_k , loop iteration ports L_1, \dots, L_k , loop return ports R_1, \dots, R_k , and loop exit ports X_1, \dots, X_k . The example in Figure 18.10 also shows a pair (2-tuple) of values being used for I, R, L, X , one for each loop-variant value.

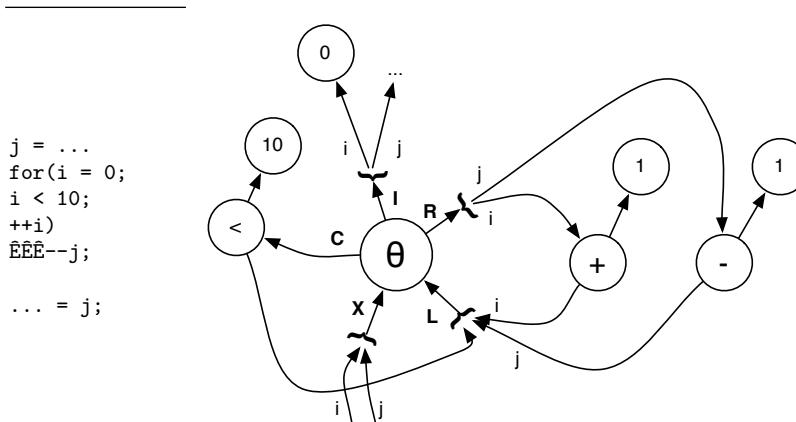


Fig. 18.10 An example showing a `for` loop. Evaluating the X port triggers it to evaluate the I value (outputting the value on the L port). While C evaluates to true, it evaluates the R value (which in this case also uses the θ -node's L value). When C is false, it returns the final internal value through the X port. As i is not used after the loop there is no dependency on the i port of X .

The θ -node directly implements pre-test loops (`while`, `for`); post-test loops (`do...while`, `repeat...until`) are synthesised from a pre-test loop preceded by a duplicate of the loop body. At first this may seem to cause unnecessary duplication of code, but it has two important benefits: (i) it exposes the first loop body iteration to optimization in post-test loops (cf. loop-peeling), and (ii) it normalizes all loops to one loop structure, which both reduces the cost of optimization, and increases the likelihood of two schematically-dissimilar loops being isomorphic in the VSDG.

18.6.5 State Nodes

Loads and stores compute a value and state. The `call` node takes both the name of the function to call and a list of arguments, and returns a list of results; it is treated as a state node as the function body may read or update state.

We maintain the simplicity of the VSDG by imposing the restriction that *all* functions have *one* return node (the exit node N_∞), which returns at least one result (which will be a state value in the case of `void` functions). To ensure that function calls and definitions are able to be allocated registers easily, we suppose that the number of arguments to, and results from, a function is smaller than the number of physical registers—further arguments can be passed via a stack as usual.

Note also that the VSDG neither forces loop invariant code into nor out-of loop bodies, but rather allows later phases to determine, by adding serializing edges, such placement of loop invariant nodes for later phases.

18.6.6 Dead node elimination with the VSDG

By representing a program as a VSDG, many optimisations become trivial. For example, consider dead node elimination (Figure 18.11). This combines both dead code elimination and unreachable code elimination. Dead code generates VSDG nodes for which there is no value or state dependency path from the return node, i.e., the result of the function does not in any way depend on the results of the dead nodes. Unreachable code generates VSDG nodes that are either dead, or become dead after some other optimisation. Thus, a *dead node* is a node that is not postdominated by the exit node N_∞ . To perform dead node elimination, only two passes are required over the VSDG resulting in linear runtime complexity: one pass to identify all of the live nodes, and a second pass to delete the unmarked (i.e., dead) nodes. It is safe because all nodes which are deleted are guaranteed never to be reachable from the `return` node.

18.7 Further readings

A compiler’s intermediate representation can be a graph, and many different graphs exist in the literature. We can represent the control flow of a program as a Control-Flow Graph (CFG) [7], where straight-line instructions are contained within basic blocks and edges show where the flow of control may be transferred to once leaving that block. A CFG is traditionally used to convert a program to SSA form [91]. We can also represent programs as a type of Data-flow Graph (DFG) [101, 102], and SSA can be represented in this way as an SSA Graph [84].

Input: A VSDG $G(N, E_V, E_S, N_\infty)$ with zero or more dead nodes.

Output: A VSDG with no dead nodes.

```

Procedure DNE( $G$ ) {
1:   WalkAndMark( $N_\infty, G$ );
2:   DeleteMarked( $G$ );
}
Procedure WalkAndMark( $n, G$ ) {
1:   if  $n$  is marked then finish;
2:   mark  $n$ ;
3:   foreach (node  $m \in N \wedge (n, m) \in (E_V \cup E_S)$ ) do
4:     WalkAndMark( $m$ );
}
Procedure DeleteMarked( $G$ ) {
1:   foreach (node  $n \in N$ ) do
2:     if  $n$  is unmarked then delete( $n$ );
}
```

Fig. 18.11 Dead node elimination on the VSDG.

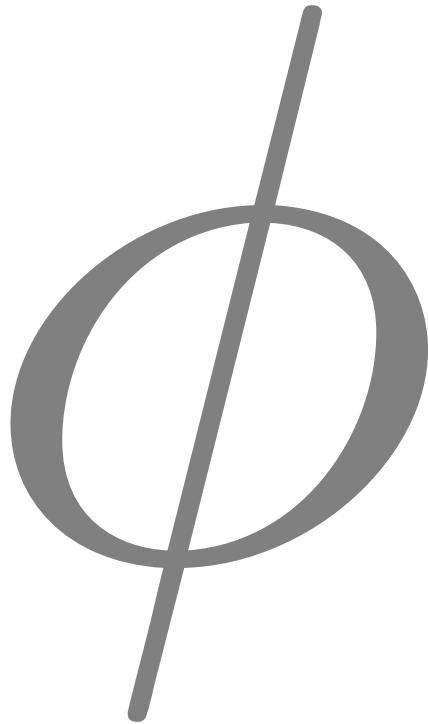
An example was given that used the SSA Graph to detect a variety of induction variables in loops [290, 131]. It has also been used for performing instruction selection techniques [107, 244], operator strength reduction [84], rematerialization [54], and has been combined with an extended SSA language to aid compilation in a parallelizing compiler [260].

The Program Dependence Graph (PDG) as defined by Ferrante et al. [119] represents control and data dependencies in one graph. Their definition of control dependencies that turns out to be equivalent to post-dominance frontier leads to confusions at it uses a non standard definition of post-dominance. We choose to report the definition of Bilardi and Pingali [33]. Section 18.4 mentions possible abstractions to represent data dependencies for dynamically allocated objects. Among others, the book of Darte et al. [95] provides a good overview of such representations. The PDG has been used for program slicing [214], testing [26], and widely for parallelization [117, 115, 247, 27]. We showed an example of how the PDG directly exposes parallel code.

Gating functions can be used to create directly interpretable ϕ -functions. These are used in Gated Single Assignment Form. Alpern et al. [12] presented a precursor of GSA for structured code, to detect equality of variables. This chapter adopts their notations, i.e., a ϕ_{if} for a if-then-else construction, a ϕ_{entry} for the entry of a loop, and a ϕ_{exit} for its exit. The original usage of GSA was by Ballance et al. [213] as an intermediate stage in the construction of the Program Dependence Web IR. Further GSA papers replaced ϕ_{if} by γ , ϕ_{entry} by μ , and ϕ_{exit} by η . Havlak [141] presented an algorithm for construction of a simpler version of GSA—Thinned GSA—which is constructed from a CFG in SSA form. The construction technique sketched in this chapter is developed in more detail in [274]. GSA has been used for a number of analyses and transformations based on data flow. The example given of how to perform backwards demand-driven symbolic analysis using GSA has been borrowed from [275]. If conversion (see Chapter 21), converts control dependencies into data dependencies. To avoid

the potential loss of information related to the lowering of ϕ -functions into conditional moves or select instructions, gating ψ -functions (see Chapter 17) can be used.

We then described the Value State Dependence Graph (VSDG) [151], which is an improvement on a previous, unmentioned graph, the Value Dependence Graph [288]. It uses the concept of gating functions, data dependencies and state to model a program. We gave an example of how to perform dead node elimination on the VSDG. Detailed semantics of the VSDG are available [151], as well as semantics of a related IR: the Gated Data Dependence Graph [277]. Further study has taken place on the problem of generating code from the VSDG [276, 178, 257], and it has also been used to perform a combined register allocation and code motion algorithm [150].



Part IV

Machine code generation and optimization

Progress: 70%

Progress: 70%



TODO: Style: instructive,
no performance results,
try for max pages 10 per
ssainput

CHAPTER 19

SSA form and code generation

B. Dupont de Dinechin

Progress: 55%

Review in progress

In a compiler for imperative languages like C, C++, or FORTRAN, the code generator covers the set of code transformations and optimizations that operate on a program representation close to the target processor ISA, and produce an assembly source or relocatable file with debugging information as result.

The main duties of code generation are: lowering the program intermediate representation to the target processor instructions and calling conventions; laying out data objects in sections and composing the stack frames; allocating variable live ranges to architectural registers; scheduling instructions to exploit micro-architecture; and producing assembly source or object code.

Precisely, the 1986 edition of the “Compilers Principles, Techniques, and Tools” Dragon Book by Aho et al. lists the tasks of code generation as:

- Instruction selection and calling conventions lowering.
- Control flow (dominators, loops) and data flow (variable liveness) analyses.
- Register allocation and stack frame building.
- Peephole optimizations.

Ten years later, the 1997 textbook “Advanced Compiler Design & Implementation” by Muchnick extends code generation with the following:

- Loop unrolling and basic block replication.
- Instruction scheduling and software pipelining.
- Branch optimizations and basic block alignment.

In current releases of compilers such as the Open64 or GCC, code generation techniques have significantly evolved, as they are mainly responsible for exploiting the performance-oriented features of processor architectures and micro-architectures. In these compilers, code generator optimizations include:

- If-conversion using SELECT, conditional move, or predicated, instructions.
- Use of specialized addressing modes such as auto-modified and modulo.
- Exploitation of hardware looping or static branch prediction hints.
- Matching fixed-point arithmetic and SIMD idioms by special instructions.
- Memory hierarchy optimizations, including pre-fetching and pre-loading.
- VLIW instruction bundling, that may interfere with instruction scheduling.

This increasing sophistication of compiler code generation motivates the introduction of the SSA form in order to simplify analyses and optimizations. However, engineering the SSA form on machine code raises a number of issues, some of which we present in Section 19.1. The suitability of the SSA form for main code generation phases is discussed in Section 19.2.

19.1 SSA form engineering issues

19.1.1 *About instructions, operands, operations, and operators*

In this chapter, we adopt the following terminology. An *instruction* is a member of the processor Instruction Set Architecture (ISA). Instructions access values and modify the processor state through *operands*. We distinguish *explicit operands*, which are associated with a specific bit-field in the instruction encoding, from *implicit operands*, without any encoding bits. Explicit operands correspond to allocatable architectural registers, immediate values, or instruction modifiers. Implicit operands correspond to single instance architectural registers and to registers implicitly used by some instructions, such as the processor status register, the procedure link register, or even the stack pointer.

An *operation* is an instance of an instruction that compose a program. It is seen by the compiler as an *operator* applied to a list of operands (explicit & implicit), along with operand naming constraints, and has a set of clobbered registers. The compiler view of operations also involves *indirect operands*, which are not apparent in the instruction behavior, but are required to connect the flow of values between operations. Implicit operands correspond to the registers used for passing arguments and returning results at function call sites, and may also be used for the registers encoded in register mask immediates.

19.1.2 *Representation of instruction semantics*

Unlike IR operators, there is no straightforward mapping between machine instruction and their operational semantics. For instance, a subtract with operands (a, b, c) may either compute $c \leftarrow a - b$ or $c \leftarrow b - a$ or any such expression with

permuted operands. Yet basic SSA form code cleanups such as constant propagation and sign extension removal need to know what is actually computed by machine instructions. Machine instructions may also have multiple target operands, such as memory accesses with auto-modified addressing, or combined division-modulus instructions. There are at least two ways to address this issue.

- Add properties to the instruction operator and to its operands, a technique used by the Open64 compiler. Typical operation properties include 'isAdd', 'isLoad', etc. Typical operand properties include 'isLeft', 'isRight', 'isBase', 'isOffset', 'isPredicated', etc. Extended properties that involve the instruction operator and some of its operands include 'isAssociative', 'isCommutative', etc.
- Associate a 'semantic combinator', that is, a tree of IR-like operators, to each target operand of a machine instruction. This more ambitious alternative was implemented in the SML/NJ [181] compiler and the LAO compiler [96].

An issue related to the representation of instruction semantics is how to factor it. Most information can be statically tabulated by the instruction operator, yet properties such as safety for control speculation, or being equivalent to a simple IR instruction, can be refined by the context where the instruction appears. For instance, range propagation may ensure that an addition cannot overflow, that a division by zero is impossible, or that a memory access is safe for control speculation. Alternate semantic combinators, or modifiers of the instruction operator semantic combinator, need to be associated with each machine instruction of the code generator internal representation.

Finally, code generation for some instruction set architectures require that pseudo-instructions with known semantics be available, besides variants of ϕ -functions and parallel COPY instructions.

- Machine instructions that operate on register pairs, such as the long multiplies on the ARM, or more generally on register tuples, are common. In such cases there is a need for pseudo-instructions to compose wide operands in register tuples, and to extract independently register allocatable operands from wide operands.
- Embedded processor architectures such as the Tensilica Xtensa provide hardware loops, where an implicit conditional branch back to the loop header is taken whenever the program counter matches some address. The implied loop-back branch is also conveniently materialized by a pseudo-instruction.
- Register allocation for predicated architectures need that the live-ranges of pseudo-registers or SSA variables with predicated definitions be contained by kill pseudo-instructions [132].

19.1.3 Operand naming constraints

Implicit operands and indirect operands are constrained to specific architectural registers either by the instruction set architecture (ISA constraints), or by application binary interface (ABI constraints). An effective way to deal with these constraints in the SSA form is by inserting parallel COPY instruction that write to the constrained source operands or read from the constrained target operands of the instructions. The new SSA variables thus created are pre-colored with the required architectural register. The parallel COPY operations are coalesced away or sequentialized when going out of SSA form [42].

Explicit instruction operands may also be constrained to use the same register between a source and a target operand, or even to use different registers between two source operands (MUL instructions on the ARM). Operand constraints between one source and the target operand are the general case on popular instruction set architectures such as x68 and the ARM Thumb. In the setting of [42], these constraints are represented by inserting a COPY between the constrained source operand and a new variable, then using this new variable as the constrained source operand. The COPY is a parallel copy in case of multiple constrained source operands. Again, the COPY operations are processed when going out of SSA form.

A difficult case of ISA or ABI operand constraint is when a variable must be bound to a specific register at all points in the program. This is the case of the stack pointer, as interrupt handling may reuse the program run-time stack. One possibility is to inhibit the promotion of the stack pointer to a SSA variable. Stack pointer definitions including memory allocations through `alloca()`, activation frame creation/destruction, are then encapsulated in a specific pseudo-instruction. Instructions that use the stack pointer must be treated as special cases as far as the SSA form analyses and optimizations are concerned.

19.1.4 Non-kill target operands

The SSA form assumes that variable definitions are kills. This is not the case for target operands such as a processor status register that contains several independent bit-fields. Moreover, some instruction effects on bit-field may be 'sticky', that is, with an implied OR with the previous value. Typical sticky bits include exception flags of the IEEE 754 arithmetic, or the integer overflow flag on DSPs with fixed-point arithmetic. When mapping a processor status register to a SSA variable, any operation that partially reads or modifies the register bit-fields should appear as reading and writing the corresponding variable.

Predicated execution and conditional execution are the other main source of definitions that do not kill the target register. The execution of predicated instructions is guarded by the evaluation of a single bit operand. The execution of conditional instructions is guarded by the evaluation of a condition on a multi-

bit operand. We extend the classification of [192] and distinguish four classes of ISA:

Partial predicated execution support SELECT instructions like those of the Multiflow TRACE architecture [82] are provided. The Multiflow TRACE 500 architecture was to include predicated store and floating-point instructions [187].

Full predicated execution support Most instructions accept a Boolean predicate operand which nullifies the instruction effects if the predicate evaluates to false. EPIC-style architectures also provide predicate define instructions (PDIs) to efficiently evaluate predicates corresponding to nested conditions: Unconditional, Conditional, parallel-OR, parallel-AND [132].

Partial conditional execution support Conditional move (CMOV) instructions similar to those introduced in the Alpha AXP architecture [35] are provided. CMOV instructions are available in the ia32 ISA since the Pentium Pro.

Full conditional execution support Most instructions are conditionally executed depending on the evaluation of a condition of a source operand. On the ARM architecture, the implicit source operand is a bit-field in the processor status register and the condition is encoded on 4 bits. On the VelociTI™ TMS230C6xxx architecture, the source operand is a general register encoded on 3 bits and the condition is encoded on 1 bit.

Extensions of the SSA form such as the ψ -SSA [262] presented in Chapter 17 specifically address the handling operations that do not kill the target operand because of predicated or conditional execution.

Observe however that under the SSA form, a CMOV is equivalent to a SELECT with a 'must be same register' naming constraint between one source and the target operand. Unlike other predicated or conditional instructions, the SELECT instructions kill the target register. Generalizing this observation provides a simple way to handle predicated or conditional execution in vanilla SSA form:

- For each target operand of the predicated or conditional instruction, add a corresponding source operand in the instruction signature.
- For each added source operand, add a 'must be same same register' naming constraint with the corresponding target operand.

This simple transformation enables SSA form analyses and optimizations to remain oblivious to predicated and conditional code. The drawback of this solution is that non-kill definitions of a given variable (before SSA renaming) remain in dominance order across transformations, as opposed to ψ -SSA where predicate value analysis may enable to relax this order.

19.1.5 Program representation invariants

Engineering a code generator requires decisions about what information is transient, or belongs to the invariants of the program representation. By invariant we

mean a property which is ensured before and after each phase. Transient information is recomputed as needed by some phases from the program representation invariants. The applicability of the SSA form only spans the early phases of the code generation process: from instruction selection, down to register allocation. After register allocation, program variables are mapped to architectural registers or to memory locations, so the SSA form analyses and optimizations no longer apply. In addition, a program may be only partially converted to the SSA form. This motivates the engineering of the SSA form as extensions to a baseline code generator program representation.

Some extensions to the program representation required by the SSA form are better engineered as invariants, in particular for operands, operations, basic blocks, and control-flow graph. Operands which are SSA variables need to record the unique operation that defines them as target operand, and possibly to maintain the list of where they appear as source operands. Operations such as ϕ -functions, σ -functions of the SSI form [39], and parallel copies, may appear as regular operations constrained to specific places in the basic blocks. The incoming arcs of basic blocks need also be kept in the same order as the source operands of each of its ϕ -functions.

A program representation invariant that impacts SSA form engineering is the structure of loops. The modern way of identifying loops in a CFG is the construction of a loop nesting forest as defined by Ramalingam [228]. Non-reducible control flow allow for different loop nesting forest for a given CFG, yet high-level information such as loop-carried memory dependences, or user-level loop annotations, are provided to the code generator. This information is attached to a loop structure, which thus becomes an invariant. The impact on the SSA form is that some loop nesting forests, such as the Havlak [142] loop structure, are better than others for key analyses such as SSA variable liveness [38].

Up-to-date live-in and live-out sets at basic block boundaries are also candidates for being program representation invariants. However, when using and updating liveness information under the SSA form, it appears convenient to distinguish the ϕ -function contributions from the results of data-flow fixpoint computation. In particular, Sreedhar et al. [255] introduced the ϕ -function semantics that became later known as *multiplexing mode* (see Chapter 22), where a ϕ -function $B_0 : a_0 = \phi(B_1 : a_1, \dots, B_n : a_n)$ makes a_0 live-in of basic block B_0 , and a_1, \dots, a_n live-out of basic blocks B_1, \dots, B_n . The classic basic block invariants $\text{LiveIn}(B)$ and $\text{LiveOut}(B)$ are then complemented with $\text{PhiDefs}(B)$ and $\text{PhiUses}(B)$ [38].

Finally, we know one compiler who adopted the invariant that the SSA form be conventional across the code generation phases. This approach was motivated in part by the fact that some classic optimizations such as SSAPRE [75] require that the 'live ranges of different versions of h [an arbitrary SSA variable] do not overlap', implying the SSA form is conventional. Other compilers that use SSA numbers and omit the ϕ -functions from the program representation [175] are similarly constrained. Work by Sreedhar et al. [255] makes it clear how

to make the transformed SSA form conventional wherever required, so there is no reason nowadays for this property to be an invariant.

19.2 Code generation phases and the SSA form

19.2.1 *Instruction selection*

Instruction selection (Chapter 20). Unlike classic techniques that match one IR tree or one DAG at a time, using the SSA form as input extends the scope of pattern matching to more complex IR graphs, in particular those resulting from control-flow constructs.

19.2.2 *If-conversion*

If-conversion refers to optimization techniques that remove conditional branches from a program region. The basic idea is to replace conditional branches by straight-line computations that directly use the condition as a source operand. The scope and effectiveness of if-conversion depends on ISA support for predicated execution [192]. In principle, any if-conversion technique targeted to full predicated or conditional execution support may be adapted to partial predicated or conditional execution support. For instance, non-predicated instructions with side-effects such as memory accesses can be used in combination with SELECT to provide a harmless effective address in case the operation must be nullified [192].

Besides predicated execution, architectural support for if-conversion can be improved by supporting speculative execution. Speculative execution (control speculation) refers to executing an operation before knowing that its execution is required, such as when moving code above a branch [187] or promoting predicates [192]. Speculative execution assumes instructions have reversible side effects, so speculating potentially excepting instructions requires architectural support. On the Multiflow TRACE 300 architecture and later on the Lx VLIW architecture [113], non-trapping memory loads known as 'dismissible' are provided. The IMPACT EPIC architecture speculative execution [19] is generalized from the 'sentinel' model [191].

If-conversion is primarily motivated by instruction scheduling on instruction-level parallel processors [192], as removing conditional branches:

- eliminates branch resolution stalls in the instruction pipeline,
- reduces uses of the branch unit, which is often single-issue,
- increases the size of the instruction scheduling regions.

In case of inner loop bodies, if-conversion further enables vectorization [9] and software pipelining (modulo scheduling) [219]. Consequently, control-flow regions selected for if-conversion are acyclic, even though seminal techniques [9, 219] consider more general control flow. The result of an if-converted region is typically a hyperblock, that is, a set of predicated basic blocks in which control may only enter from the top, but may exit from one or more locations [194].

Classic contributions to if-conversion do not consider the SSA form:

- Park & Schlansker [219] propose the RK algorithm based the control dependences. They assume a fully predicated architecture with only Conditional PDIs. The R function assigns a minimal set of Boolean predicates to basic blocks, and the K function express the way these predicates are computed. The algorithm is general enough to process cyclic and irreducible rooted flow graphs, but in practice it is applied to single entry acyclic regions.
- Blickstein et al. [35] pioneer the use of CMOV instructions to replace conditional branches in the GEM compilers for the Alpha AXP architecture.
- Lowney et al. [187] match the innermost if-then constructs in the Multiflow compiler in order to generate the SELECT and the predicated memory store operations.
- Fang [111] assumes a fully predicated architecture with Conditional PDIs. The proposed algorithm is tailored to acyclic regions with single entry and multiple exits, and as such is able to compute R and K functions without relying on explicit control dependences. The main improvement of this algorithm over [219] is that it also speculates instructions up the dominance tree through predicate promotion, except for stores and PDIs. This work further proposes a pre-optimization pass to hoist or sink common sub-expressions before predication and speculation.
- Leupers [183] focuses on if-conversion of nested if-then-else (ITE) statements on architectures with full conditional execution support. A dynamic programming technique appropriately selects either a conditional jump or a conditional instruction based implementation scheme for each ITE statement, and the objective is the reduction of worst-case execution time (WCET).

A few contributions to if-conversion only use the SSA form internally:

- Jacome et al. [147] propose the Static Single Assignment – Predicated Switching (SSA-PS) transformation aimed at clustered VLIW architectures, with predicated move instructions that operate inside clusters (internal moves) or between clusters (external moves). The first idea of the SSA-PS transformation is to realize the conditional assignments corresponding to ϕ -functions via predicated switching operations, in particular predicated move operations. The second idea is that the predicated external moves leverage the penalties associated with inter-cluster data transfers. The SSA-PS transformation predicates non-move operations and is apparently restricted to innermost if-then-else statements.

- Chuang et al. [77] introduce a predicated execution support aimed at removing non-kill register writes from the processor micro-architecture. They propose SELECT instructions called 'phi-ops', predicated memory accesses, Unconditional PDIs, and ORP instructions for OR-ing multiple predicates. A restriction of the RK algorithm to single-entry single-exit regions is proposed, adapted to the Unconditional PDIs and the ORP instructions. The other contribution from this work is the generation of 'phi-ops', whose insertion points are computed like the SSA form placement of the ϕ -functions. The ϕ -functions source operands are replaced by ϕ -lists, where each operand is associated with the predicate of its source basic block. The ϕ -lists are processed by topological order of the operand predicates to generate the 'phi-ops'.

Recent contributions to if-conversion input SSA form, and output ψ -SSA form:

- Stoutchinin & Ferrière [262] introduce ψ -functions in order to bring predicated code under the SSA form for a fully predicated architecture. The ψ -functions arguments are paired with predicates and are ordered in dominance order, a correctness condition re-discovered later by Chuang et al. [77] for their phi-ops.
- Stoutchinin & Gao [263] propose an if-conversion technique based on the predication of Fang [111] and the replacement of ϕ -functions by ψ -functions. They prove the conversion is correct provided the SSA form is conventional. The technique is implemented in the Open64 for the ia64 architecture.
- Bruel [56] targets VLIW architectures with SELECT and dismissible load instructions. The proposed framework reduces acyclic control-flow constructs from innermost to outermost, and the monitoring of the if-conversion benefits provides the stopping criterion. The core technique control speculates operations, reduces height of predicate computations, and performs tail duplication. It can also generate ψ -functions instead of SELECT operations. A generalization of this framework, which also accepts ψ -SSA form as input, is described in Chapter 21.
- Ferrière [98] extends the ψ -SSA form algorithms of [262] to the partial predicated execution support, by formulating simple correctness conditions for the predicate promotion of operations that do not have side-effects. This work also details how to transform the ψ -SSA form to conventional ψ -SSA by generating CMOV operations. A self-contained explanation of these techniques appears in Chapter 17.

Thanks to these two latter contributions, virtually all if-conversion techniques formulated without the SSA form can be adapted to the ψ -SSA form, with the added benefit that already predicated code can be part of input. The main issues that remain for the selection of an if-conversion technique are the match with the architectural support, and the approach to if-conversion region formation.

19.2.3 Inner loop optimizations

Non-inner loop transformations, such as unroll-and-jam, are usually left to the IR loop optimizer. Loop unrolling replicates the loop body and removes all but one of the replicated loop exit branches. Loop unwinding replicates the loop body while keeping the loop exit branches. This style of loop unrolling necessarily applies to counted loops, and requires that pre-conditioning or post-conditioning code be inserted [187]. Inner loop unrolling or unwinding is facilitated by using the loop-closed SSA form (Chapter 11), the SSA version of the self assignment technique also pioneered by the Multiflow Trace Scheduling compiler [187]. A main motivation for inner loop unrolling is to unlock modulo scheduling benefits [177].

Induction variable classification (Chapter 11),

Memory addressing optimizations include selection of auto-modified addressing modes. Memory packing optimizations select wider memory access instructions whenever the effective addresses are provably adjacent, and no side effects such as possible address misalignment traps are introduced.

19.2.4 Code cleanups and instruction re-selection

Constant propagation, copy folding, and dead code elimination.

Bit-width analysis,

GVN or LICM or PRE

Instruction re-selection. Uses bit width analysis. SIMD idioms. Required because analyses and code specialization

Reassociation

19.2.5 Pre-pass instruction scheduling

Further down the code generator, the next major phase is pre-pass instruction scheduling. Innermost loops with a single basic block, super-block or hyper-block body are candidates for software pipelining techniques such as modulo scheduling. For innermost loops that are not software pipelined, and for other program regions, acyclic instruction scheduling techniques apply: basic block scheduling [134]; super-block scheduling [146]; hyper-block scheduling [194]; tree region scheduling [140]; or trace scheduling [187].

By definition, pre-pass instruction scheduling operates before register allocation. On a classic code generator, instruction operands are mostly virtual registers, except for instructions with ISA or ABI constraints that binds them to specific architectural registers. Moreover, preparation to pre-pass instruction scheduling include virtual register renaming, also known as register web con-

struction, in order to reduce the number of anti dependences and output dependences in the instruction scheduling problem. Other reasons why it seems there is little to gain to schedule instructions on a SSA form program representation include:

- Except in case of trace scheduling which pre-dates the use of SSA form in production compilers, the classic scheduling regions are single-entry and do not have control-flow merge. So there are no ϕ -functions in case of acyclic scheduling, and only ϕ -functions in the loop header in case of software pipelining. Keeping those ϕ -functions in the scheduling problem has no benefits and raises engineering issues, due to their parallel execution semantics and the constraint to keep them first in basic blocks.
- Instruction scheduling must account for all the instruction issue slots required to execute a code region. If the only ordering constraints between instructions, besides control dependences and memory dependences, are limited to true data dependences on operands, code motion will create interferences that must later be resolved by inserting COPY operations in the scheduled code region. (Except for interferences created by the overlapping of live ranges that results from modulo scheduling, as these are resolved by modulo renaming.) So scheduling instructions with SSA variables as operands is not effective unless extra dependences are added to the scheduling problem to prevent such code motion.
- Some machine instructions have side effects on special resources such as the processor status register. Representing these resources as SSA variables even though they are not operated like the processor architectural registers requires coarsening the instruction effects to the whole resource, as discussed in Section 19.1.4. In turn this implies def-use variable ordering that prevents aggressive instruction scheduling. For instance, all 'sticky' bit-field definitions can be reordered with regards to the next use, and an instruction scheduler is expected to do so. Scheduling OR-type predicate define operations [245] raises the same issues. An instruction scheduler is also expected to precisely track accesses to unrelated or partially overlapping bit-fields in a processor status register.
- Aggressive instruction scheduling relaxes some flow data dependences that are normally implied by SSA variable def-use ordering. A first example is 'move renaming' [296], the dynamic switching of the definition of a source operand defined by a COPY operation when the consumer operations ends up being scheduled at the same cycle or earlier. Another example is 'inductive relaxation' [97], where the dependence between additive induction variables and their use as base in base+offset addressing modes is relaxed to the extent permitted by the induction step and the range of the offset. These techniques apply to acyclic scheduling and to modulo scheduling.

To summarize, trying to keep the SSA form inside the pre-pass instruction scheduling currently appears more complex than operating on the program representation with classic compiler temporary variables. This representation is

obtained from the SSA form after aggressive coalescing and SSA destruction. If required by the register allocation, the SSA form should be re-constructed.

19.2.6 SSA form destruction

The destruction of the SSA form in a code generator is required before the pre-pass instruction scheduling and software pipelining, at least with the established techniques, and also before classic register allocation such as George & Appel [128]. A weaker form is the conversion of transformed SSA form to conventional, which is required by some register allocators operating on the SSA form [220]. For all these uses cases, the main objective besides removing the SSA form extensions from the program representation is to ensure that the ISA and ABI operand naming constraints are satisfied. These objectives are met with aggressive coalescing.

A summary of the SSA form destruction techniques is that correctness issues of the original out-of-SSA translation of Cytron et al. [86] were addressed first. The focus then moved to the handling of ISA and ABI operand naming constraints. Simultaneously, the time and space constraints of JIT compilation motivated new approaches to SSA variable interference checks. State-of-the-art algorithms such as those presented in Chapter 22 fix the remaining correctness issues, handle all ISA and ABI operand naming constraints, support aggressive coalescing, and can be adapted to JIT compilation constraints.

- Cytron et al. [86] describe the process of 'translating out of SSA' as 'naive replacement preceded by dead code elimination and followed by coloring' (graph coloring coalescing). The naive replacement replaces each ϕ -function $B_0 : a_0 = \phi(B_1 : a_1, \dots, B_n : a_n)$ by n copies $a_0 = a_i$, one per basic block B_i .
- Briggs et al. [50] identifies correctness issues in Cytron et al. translation out of (transformed) SSA form and illustrate them by the 'lost-copy' and the 'swap' problems. These issues are related to the fact that critical edges must be split, and that a sequence of ϕ -functions at the start of a basic block has parallel assignment semantics. Two SSA form destruction algorithms are proposed, depending on the presence of critical edges in the control-flow graph.
- Sreedhar et al. [255]
- Budimlić et al. "Fast Copy Coalescing and Live-Range Identification" [PLDI'02]
 - lightweight SSA destruction motivated by JIT compilation
 - use the SSA form dominance of definitions over uses to avoid explicit interference graph
 - construct SSA-webs with early pruning of interfering variables, then partition into non-interfering classes

- introduce the “dominance forest” data-structure to avoid quadratic number of interference tests
- critical edge splitting is required
- Sreedhar et al. “Translating Out of Static Single Assignment Form” [SAS’99] (US patent 6182284):
 - Method 1 inserts COPY for the arguments of Φ -functions in the predecessors *and* for the Φ -functions targets in the current block, then applies a new SSA-based coalescing algorithm
 - Method 3 maintains liveness and interference graph to insert COPY that will not be removed by the new SSA-based coalescing algorithm
 - the new SSA-based coalescing algorithm is more effective than register allocation coalescing
- Leung & George “Static Single Assignment Form for Machine Code” [PLDI’99]
 - handles the operand constraints of machine-level SSA form
 - builds on the algorithm by Briggs et al. [SPE 28(8) 1998]
 - Errors in missing parallel copies in case of W shaped CFG, the copies in the pointy edge need to be in parallel
- Rastello et al. “Optimizing Translation Out of SSA using Renaming Constraints” [CGO’04] (STMicroelectronics)
 - fix bugs and generalize Leung & George [PLDI’99]
 - generalize Sreedhar et al. [SAS’99] (and avoid patent)

Sreedhar et al. [SAS 1999] Liveness and Congruence Example

- variables $x_1, x_3, x_3, y_1, y_2, y_3$ are in the same congruence class
- in this example, several interferences inside the congruence class

Insights of Sreedhar et al. [SAS 1999]

- a Φ -congruence class is the closure of the Φ -connected relation
- liveness under SSA form: Φ arguments are live-out of predecessor blocks and Φ targets are live-in of Φ block
- SSA form is *conventional* if no two members of a Φ -congruence class interfere under this liveness
- correct SSA destruction is the removal of Φ -functions from a conventional SSA form
- after SSA construction (without COPY propagation), the SSA form is conventional
- Methods 1 – 3 restore a conventional SSA form
- the new SSA-based coalescing is able to coalesce interfering variables, as long as the SSA form remains conventional

19.2.7 Register allocation

The last major phase of code generation where SSA form has demonstrated benefits is register allocation and its three sub-problems: variable spilling, variable coloring, and variable coalescing (Chapter 23).

Issue with predicated or conditional instructions. [109] [155]

Issue with aliased registers.

Pre conditioning to reduce MaxLive, saturation, sufficiency

CHAPTER 20

Instruction Code Selection

D. Ebner
A. Krall
B. Scholz

Progress: 95%

Final reading in progress

20.1 Introduction

*

Instruction code selection is a transformation step in a compiler that translates a machine-independent intermediate code representation into a low-level intermediate representation or to machine code for a specific target architecture. Instead of hand-crafting an instruction selector for each target architecture, generator tools have been designed and implemented that generate the instruction code selector based on a specification of the machine description of the target. This approach is used in large compiler infrastructures such as GCC or LLVM that target a range of architectures. A possible scenario of a code generator in a compiler is depicted in Figure 20.1. The Intermediate Representation (IR) of an input program is passed on to an optional lowering phase that breaks down instructions and performs other machine dependent transformations. Thereafter, the instruction selection performs the mapping to machine code or lowered IR based on the machine description of the target architecture.

One of the widely used techniques in code generation is tree pattern matching.* The unit of translation for tree pattern matching is expressed as a tree structure that is called a data-flow tree* (DFT). The basic idea is to describe the target instruction set using an *ambiguous* cost-annotated graph grammar.* The instruc-

*symbol, non-terminal
symbol, terminal
rule, chain*

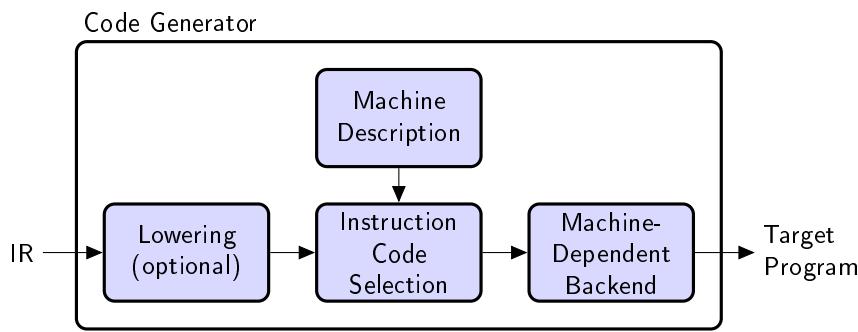


Fig. 20.1 Scenario: An instruction code selector translates a compiler's IR to a low-level machine-dependent representation.

tion code selector seeks for a cost-minimal *cover* of the DFT. Each of the selected rules have an associated semantic action that is used to emit the corresponding machine instructions, either by constructing a new intermediate representation or by rewriting the DFT bottom-up.

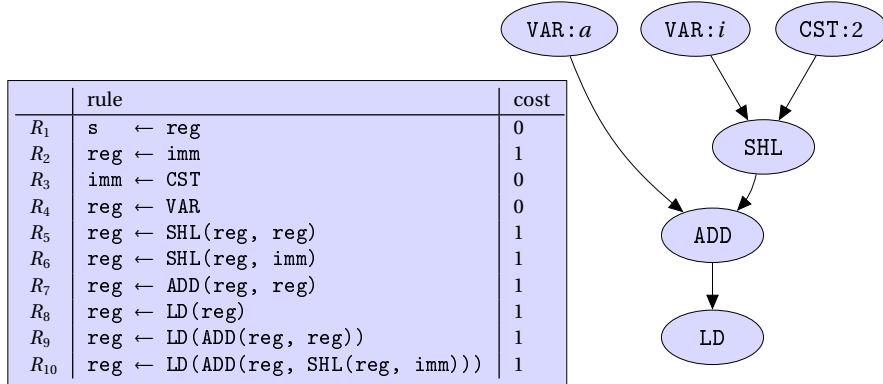


Fig. 20.2 Example of a data-flow tree and a rule fragment with associated costs.

An example of a DFT along with a set of rules representing valid ARM instructions is shown in Figure 20.2. Each rule consists of non-terminals^{*} (shown in lower-case), and terminal symbols^{*} (shown in upper-case). Non-terminals are used to chain individual rules together. Non-terminal s denotes a distinguished start symbol for the root node of the DFT. Terminal symbols match the corresponding labels of nodes in the data-flow trees. The terminals of the grammar are VAR, CST, SHL, ADD, and LD. Rules that translate from one non-terminal to another are called *chain rules*^{*}, e.g., $\text{reg} \leftarrow \text{imm}$ that translates an immediate value to a register. Note that there are multiple possibilities to obtain a cover of

the data-flow tree for the example shown in Figure 20.2. Each rule has associated costs. The cost of a tree cover is the sum of the costs of the selected rules. For example, the DFT could be covered by rules R_3 , R_4 , and R_{10} which would give a total cost for the cover of one cost unit. Alternatively, the DFT could be covered by rule R_2 , R_3 , R_5 , R_7 , and R_8 which yields four cost units for the cover for issuing four assembly instructions. A dynamic programming algorithm selects a cost optimal cover for the DFT.

Tree pattern matching on a DFT is limited to the scope of tree structures. To overcome this limitation, we can extend the scope of the matching algorithm to the computational flow of a whole procedure. The use of the SSA form as an intermediate representation improves the code generation by making def-use relationships explicit. Hence, SSA exposes the data flow of a translation unit and utilizes the code generation process. Instead of using a textual SSA representations, we employ a graph representation of SSA called the *SSA graph*¹ that is an extension of DFTs and represents the data flow for scalar variables of a procedure in SSA form. SSA graphs are a suitable representation for code generation: First, SSA graphs capture acyclic and cyclic information flow beyond basic block boundaries. Second, SSA graphs often arise naturally in modern compilers as the intermediate code representation usually already is in SSA form. Third, output or anti-dependencies in SSA graph do not exist.

As even acyclic SSA graphs are in the general case not restricted to a tree, no dynamic programming approach can be employed for instruction code selection. To get a handle on instruction code selection for SSA graphs, we will discuss in the following an approach based on a reduction to a quadratic mathematical programming problem (PBQP). Consider the code fragment of a dot-product routine and the corresponding SSA graph shown in Figure 20.3. The code implements a simple vector dot-product using fixed-point arithmetic. Nodes in the SSA graph represent a single operation while edges describe the flow of data that is produced at the source node and consumed at the target node. Incoming edges have an order which reflects the argument order of the operation. In the figure the color of the nodes indicates to which basic block the operations belong to.

The example in Figure 20.3 has fixed-point computations that need to be modeled in the grammar. For fixed-point values most arithmetic and bit-wise operations are identical to their integer equivalents. However, some operations have different semantics, e.g., multiplying two fixed-point values in format $m.i$ results in a value with $2i$ fractional digits. The result of the multiplication has to be adjusted by a shift operation to the right (LSR). To accommodate for fixed-point values, we add the following rules to the grammar introduced in Figure 20.2:

¹ We consider its data-based representation here.^{*} See Chapter 18

rule	cost	instruction
<code>reg ← VAR</code>	is_fixed_point ? ∞ otherwise 0	
<code>fp ← VAR</code>	is_fixed_point ? 0 otherwise ∞	
<code>fp2 ← MUL(fp, fp)</code>	1	MUL Rd, Rm, Rs
<code>fp ← fp2</code>	1	LSR Rd, Rm, i
<code>fp ← ADD(fp, fp)</code>	1	ADD Rd, Rm, Rs
<code>fp2 ← ADD(fp2, fp2)</code>	1	ADD Rd, Rm, Rs
<code>fp ← PHI(fp, ...)</code>	0	
<code>fp2 ← PHI(fp2, ...)</code>	0	

In the example the accumulation for double-precision fixed point values (`fp2`) can be performed at the same cost as for the single-precision format (`fp`). Thus, it would be beneficial to move the necessary shift from the inner loop to the return block, performing the intermediate calculations in the extended format. However, as a tree-pattern matcher generates code at statement-level, the information of having values as double-precision cannot be hoisted across basic block boundaries. An instruction code selector that is operating on the SSA graph, is able to propagate non-terminal `fp2` across the ϕ node prior to the return and emits the code for the shift to the right in the return block.

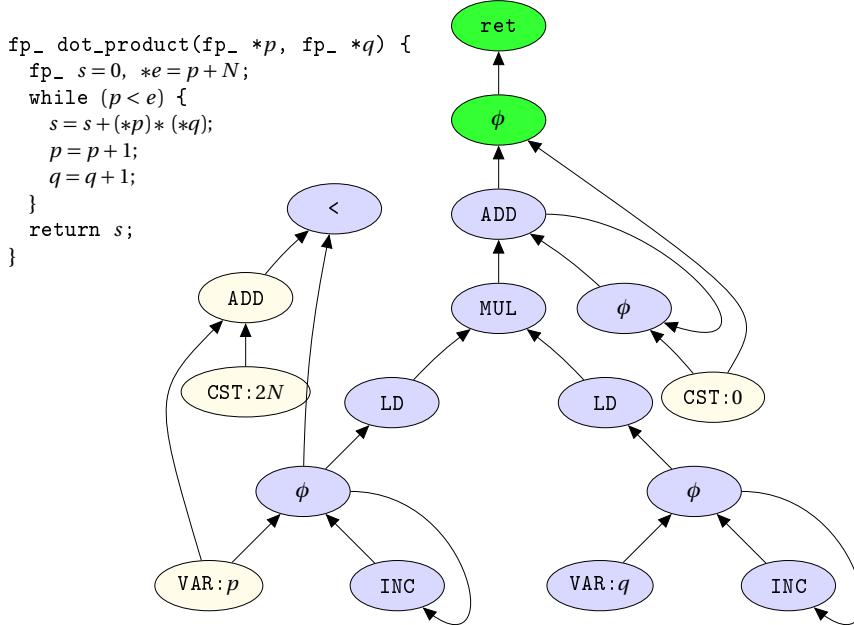


Fig. 20.3 Instruction code selection SSA Graph for a vector dot-product in fixed-point arithmetic. `fp_` stands for unsigned short fixed point type.

In the following, we will explain how to perform instruction code selection on SSA graphs with the means of a specialized quadratic assignment problem (PBQP). First, we discuss the instruction code selection problem by employing a discrete optimization problem called partitioned boolean quadratic problem. An extension of *patterns* to arbitrary acyclic graph structures, which we refer to as DAG grammars, is discussed in Sub-Section 20.3.1.

20.2 Instruction Code Selection for Tree Patterns on SSA-Graphs

The matching problem for SSA graphs reduces to a discrete optimization problem called Partitioned Boolean Quadratic Problem (PBQP). First, we will introduce the PBQP problem and then we will describe the mapping of the instruction code selection problem to PBQP.

20.2.1 Partitioned Boolean Quadratic Problem

*

Partitioned Boolean Quadratic Programming (PBQP) is a generalized quadratic assignment problem that has proven to be effective for a wide range of applications in embedded code generation, e.g., register assignment, address mode selection, or bank selection for architectures with partitioned memory. Instead of problem-specific algorithms, these problems can be modeled in terms of generic PBQPs that are solved using a common solver library. PBQP is flexible enough to model irregularities of embedded architectures that are hard to cope with using traditional heuristic approaches.

Consider a set of discrete variables $X = \{x_1, \dots, x_n\}$ and their finite domains $\{\mathbb{D}_1, \dots, \mathbb{D}_n\}$. A solution of PBQP is a mapping h of each variable to an element in its domain, i.e., an element of \mathbb{D}_i needs to be chosen for variable x_i . The chosen element imposes *local costs* and *related costs* with neighboring variables. Hence, the quality of a solution is based on the contribution of two sets of terms.

1. For assigning variable x_i to the element d_i in \mathbb{D}_i . The quality of the assignment is measured by a *local cost function* $c(x_i, d_i)$.
2. For assigning two related variables x_i and x_j to the elements $d_i \in \mathbb{D}_i$ and $d_j \in \mathbb{D}_j$. We measure the quality of the assignment with a *related cost function* $C(x_i, x_j, d_i, d_j)$.

The total cost of a solution h is given as,

$$f = \sum_{1 \leq i \leq n} c(x_i, h(x_i)) + \sum_{1 \leq i < j \leq n} C(x_i, x_j, h(x_i), h(x_j)). \quad (20.1)$$

The PBQP problem seeks for an assignment of variables x_i with minimum total costs.

In the following we represent both the local cost function and the related cost function in matrix form, i.e., the related cost function $C(x_i, x_j, d_i, d_j)$ is decomposed for each pair (x_i, x_j) . The costs for the pair are represented as $|\mathbb{D}_i|$ -by- $|\mathbb{D}_j|$ matrix/table $\mathcal{C}_{i,j}$. A matrix element corresponds to an assignment (d_i, d_j) . Similarly, the local cost function $c(x_i, d_i)$ is represented by a cost vector \vec{c}_i enumerating the costs of the elements. A PBQP problem has an underlying graph structure graph $G = (V, E, C, c)$, which we refer to as a PBQP graph. For each decision variable x_i we have a corresponding node $v_i \in V$ in the graph, and for each cost matrix $\mathcal{C}_{i,j}$ that is not the zero matrix, we introduce an edge $e = (v_i, v_j)$. The cost functions c and C map nodes and edges to the original cost vectors and matrices respectively. We will present an example later in this chapter in the context of instruction code selection.

In general, finding a solution to this minimization problem is NP hard. However, for many practical cases, the PBQP instances are sparse, *i.e.*, many of the cost matrices $\mathcal{C}_{i,j}$ are zero matrices and do not contribute to the overall solution. Thus, optimal or near-optimal solutions can often be found within reasonable time limits. Currently, there are two algorithmic approaches for PBQP that have been proven to be efficient in practice for instruction code selection problems, *i.e.*, a polynomial-time heuristic algorithm and a branch-&-bound based algorithm with exponential worst case complexity. For a certain subclass of PBQP, the algorithm produces provably optimal solutions in time $\mathcal{O}(nm^3)$, where n is the number of discrete variables and m is the maximal number of elements in their domains, *i.e.*, $m = \max(|\mathbb{D}_1|, \dots, |\mathbb{D}_n|)$. For general PBQPs, however, the solution may not be optimal. To obtain still an optimal solution outside the subclass, branch-&-bound techniques can be applied.

20.2.2 Instruction Code Selection with PBQP

In the following, we describe the modeling of instruction code selection for SSA graphs as a PBQP problem. In the basic modeling, SSA and PBQP graphs coincide. The variables x_i of the PBQP are decision variables reflecting the choices of applicable rules (represented by \mathbb{D}_i) for the corresponding node of x_i . The local costs reflect the costs of the rules and the related costs reflect the costs of chain rules making rules compatible with each other. This means that the number of decision vectors and the number of cost matrices in the PBQP are determined by the number of nodes and edges in the SSA graph respectively. The sizes of \mathbb{D}_i depend on the number of rules in the grammar. A solution for the PBQP instance induces a complete cost minimal cover of the SSA graph.

As in traditional tree pattern matching, an ambiguous graph grammar consisting of tree patterns with associated costs and semantic actions is used. Input grammars have to be *normalized*. This means that each rule is either a

so-called *base rule* or a *chain rule*. A base rule^{*} is a production of the form $nt_0 \leftarrow OP(nt_1, \dots, nt_{k_p})$ where nt_i are non-terminals and OP is a terminal symbol, i.e., an operation represented by a node in the SSA graph. A chain-rule is a production of the form $nt_0 \leftarrow nt_1$, where nt_0 and nt_1 are non-terminals. A production rule $nt \leftarrow OP_1(\alpha, OP_2(\beta), \gamma)$ can be normalized by rewriting the rule into two production rules $nt \leftarrow OP_1(\alpha, nt', \gamma)$ and $nt' \leftarrow OP_2(\beta)$ where nt' is a new non-terminal symbol and α, β and γ denote arbitrary pattern fragments. This transformation can be iteratively applied until all production rules are either chain rules or base rules. To illustrate this transformation, consider the grammar in Figure 20.4, which is a normalized version of the tree grammar introduced in Figure 20.2. Temporary non-terminal symbols $t1, t2$, and $t3$ are used to decompose larger tree patterns into simple base rules. Each base rule spans across a single node in the SSA graph.

The instruction code selection problem for SSA graphs is modeled in PBQP as follows. For each node u in the SSA graph, a PBQP variable x_u is introduced. The domain of variable x_u is determined by the subset of base rules whose terminal symbol matches the operation of the SSA node, e.g., there are three rules (R_4, R_5, R_6) that can be used to cover the shift operation SHL in our example. The last rule is the result of automatic normalization of a more complex tree pattern. The cost vector $\vec{c}_u = w_u \cdot \langle c(R_1), \dots, c(R_{k_u}) \rangle$ of variable x_u encodes the local costs for a particular assignment where $c(R_i)$ denotes the associated cost of base rule R_i . Weight w_u is used as a parameter to optimize for various objectives including speed (e.g., w_u is the expected execution frequency of the operation at node u) and space (e.g., the w_u is set to one). In our example, both R_4 and R_5 have associated costs of one. Rule R_6 contributes no local costs as we account for the full costs of a complex tree pattern at the root node. All nodes have the same weight of one, thus the cost vector for the SHL node is $\langle 1, 1, 0 \rangle$.

An edge in the SSA graph represents data transfer between the result of an operation u , which is the source of the edge, and the operand v which is the tail of the edge. To ensure consistency among base rules and to account for the costs of chain rules, we impose costs dependent on the selection of variable x_u and variable x_v in the form of a cost matrix \mathcal{C}_{uv} . An element in the matrix corresponds to the costs of selecting a specific base rule $r_u \in R_u$ of the result and a specific base rule $r_v \in R_v$ of the operand node. Assume that r_u is $nt \leftarrow OP(\dots)$ and r_v is $\dots \leftarrow OP(\alpha, nt', \beta)$ where nt' is the non-terminal of operand v whose value is obtained from the result of node u . There are three possible cases:

1. If the non-terminal nt and nt' are identical, the corresponding element in matrix \mathcal{C}_{uv} is zero, since the result of u is compatible with the operand of node v .
2. If the non-terminals nt and nt' differ and there exists a rule $r : nt' \leftarrow nt$ in the transitive closure of all chain rules, the corresponding element in \mathcal{C}_{uv} has the costs of the chain rule, i.e., $w_v \cdot c(r)$.
3. Otherwise, the corresponding element in \mathcal{C}_{uv} has infinite costs prohibiting the selection of incompatible base rules.

*rule, base
rule, production*

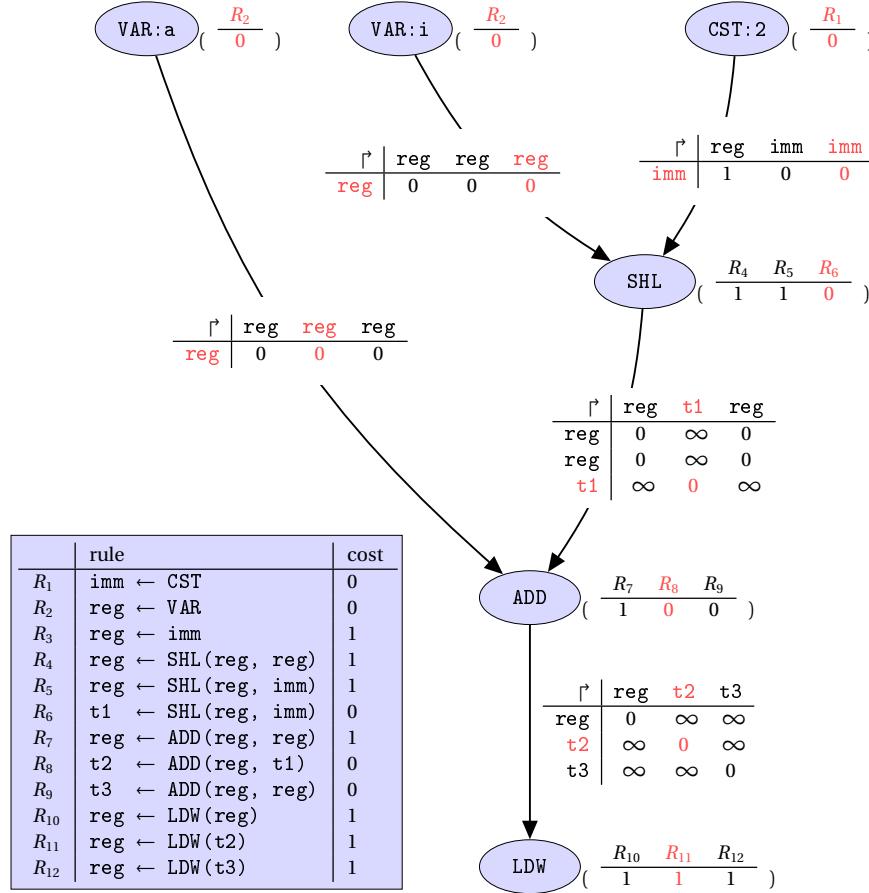


Fig. 20.4 PBQP instance derived from the example shown in Figure 20.2. The grammar has been normalized by introducing additional non-terminals.

As an example, consider the edge from **CST:2** to node **SHL** in Figure 20.4. There is a single base rule R_1 with local costs 0 and result non-terminal **imm** for the constant. Base rules R_4 , R_5 , and R_6 are applicable for the shift, of which the first one expects non-terminal **reg** as its second argument, rules R_5 and R_6 both expect **imm**. Consequently, the corresponding cost matrix accounts for the costs of converting from **reg** to **imm** at index (1, 1) and is zero otherwise.

Highlighted elements in Figure 20.4 show a cost-minimal solution of the PBQP with costs one. A solution of the PBQP directly induces a selection of base and chain rules for the SSA graph. The execution of the semantic action rules inside a basic block follow the order of basic blocks. Special care is necessary for chain rules that link data flow across basic blocks. Such chain rules may be

placed inefficiently and a placement algorithm [244] is required for some grammars.

20.3 Extensions and Generalizations

20.3.1 Instruction Code Selection for DAG Patterns

In the previous section we have introduced an approach based on code patterns that resemble simple tree fragments. This restriction often complicates code generators for modern CPUs with specialized instructions and SIMD extensions, e.g., there is no support for machine instructions with multiple results.

Consider the introductory example shown in Figure 20.3. Many architectures have some form of auto-increment addressing modes. On such a machine, the load and the increment of both p and q can be done in a single instruction benefiting both code size and performance. However, post-increment loads cannot be modeled using a single tree-shaped pattern. Instead, it produces multiple results and spans across two non-adjacent nodes in the SSA graph, with the only restriction that their arguments have to be the same.

Similar examples can be found in most architectures, e.g., the DIVU instruction in the Motorola 68K architecture performs the division and the modulo operation for the same pair of inputs. Other examples are the RMS (read-modify-store) instructions on the IA32/AMD64 architecture, autoincrement- and decrement addressing modes of several embedded systems architectures, the IRC instruction of the HPPA architecture, or fsincos instructions of various math libraries. Compiler writers are forced to pre- or post-process these patterns heuristically often missing much of the optimization potential. These architecture-specific tweaks also complicate re-targeting, especially in situations where patterns are automatically derived from generic architecture descriptions.

We will now outline, through the example in Figure 20.5, a possible problem formulation for these generalized patterns in the PBQP framework discussed so far. The code fragment contains three feasible instances of a post-increment store pattern. Assuming that p , q , and r point to mutually distinct memory locations, there are no further dependencies apart from the edges shown in the SSA graph. If we select *all* three instances of the post-increment store pattern concurrently, the graph induced by SSA edges becomes acyclic, and the code cannot be emitted. To overcome this difficulty, the idea is to express in the modeling of the problem, a numbering of chosen nodes, that reflects the existence of a topological order.

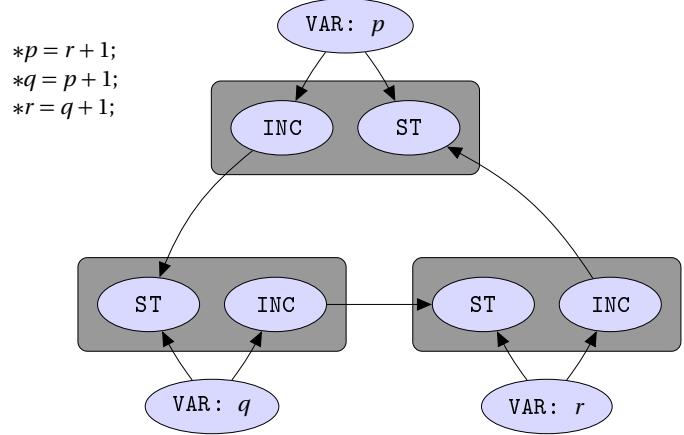


Fig. 20.5 DAG patterns may introduce cyclic data dependencies.

Modeling

The first step is to explicitly enumerate *instances* of complex patterns, *i.e.*, concrete tuples of nodes that match the terminal symbols specified in a particular production. There are three instances of the post-increment store pattern (surrounded by boxes) in the example shown in Figure 20.5. As for tree patterns, DAG patterns are decomposed into simple base rules for the purpose of modeling, *e.g.*, the post-increment store pattern

$$P_1: \text{tmt} \leftarrow \text{ST}(x:\text{reg}, \text{reg}), \text{reg} \leftarrow \text{INC}(x) : 3$$

is decomposed into the individual pattern fragments

$$\begin{aligned} P_{1,1}: \text{stmt} &\leftarrow \text{ST}(\text{reg}, \text{reg}) \\ P_{1,2}: \text{reg} &\leftarrow \text{INC}(\text{reg}) \end{aligned}$$

For our modeling, new variables are created for each enumerated instance of a complex production. They encode whether a particular instance is chosen or not, *i.e.*, the domain basically consists of the elements on and off. The local costs are set to the combined costs for the particular pattern for the on state and to 0 for the off state. Furthermore, the domain of existing nodes is augmented with the base rule fragments obtained from the decomposition of complex patterns. We can safely squash all identical base rules obtained from this process into a single state. Thus, each of these new states can be seen as a proxy for the whole set of instances of (possibly different) complex productions including the node. The local costs for these proxy states are set to 0.

Continuing our example, the PBQP for the SSA graph introduced in Figure 20.5 is shown in Figure 20.6. In addition to the post-increment store pattern with costs three, we assume regular tree patterns for the store and the increment nodes with costs two denoted by P_2 and P_3 respectively. Rules for the VAR nodes are omitted for simplicity.

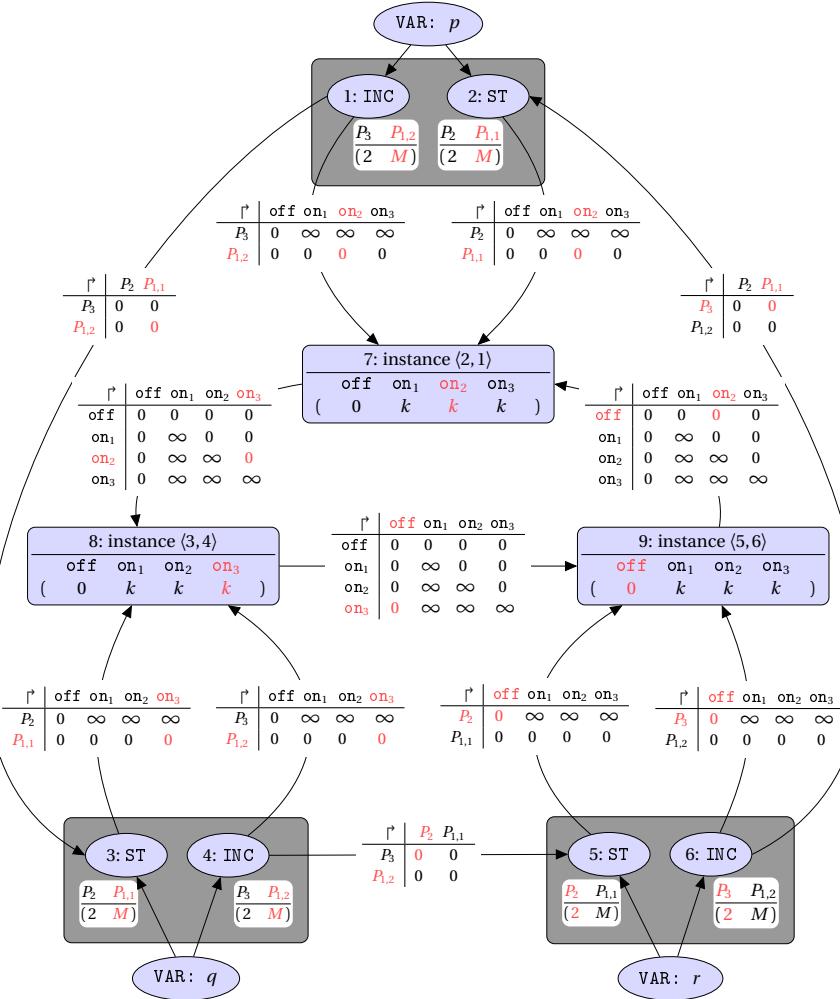


Fig. 20.6 PBQP Graph for the Example shown in Figure 20.5. M is a large integer value. We use k as a shorthand for the term $3 - 2M$.

Nodes 1 to 6 correspond to the nodes in the SSA graph. Their domain is defined by the simple base rule with costs two and the proxy state obtained from the decomposition of the post-increment store pattern. Nodes 7, 8, and 9 correspond to the three instances identified for the post-increment store pattern. As noted before, we have to guarantee the existence of a topological order among the chosen nodes. To this end, we refine the state on such that it reflects a particular index in a concrete topological order. Matrices among these nodes account for data dependencies, e.g., consider the matrix established among nodes 7 and 8. Assuming instance 7 is on at index 2 (i.e., mapped to on₂), the only remaining choices for instance 8 are not to use the pattern (i.e., mapped

to `off`) or to enable it at index 3 (i.e., mapped to on_3), as node 7 has to precede node 8.

Additional cost matrices are required to ensure that the corresponding proxy state is selected on all the variables forming a particular pattern instance (which can be modeled with combined costs of 0 or ∞ respectively). However, this formulation allows for the trivial solution where all of the related variables encoding the selection of a complex pattern are set to `off` (accounting for 0 costs) even though the artificial proxy state has been selected. We can overcome this problem by adding a large integer value M to the costs for all proxy states. In exchange, we subtract these costs from the cost vector of instances. Thus, the penalties for the proxy states are effectively eliminated unless an invalid solution is selected.

Cost matrices among nodes 1 to 6 do not differ from the basic approach discussed before and reflect the costs of converting the non-terminal symbols involved. It should be noted that for general grammars and irreducible graphs, that the heuristic solver of PBQP cannot guarantee to deliver a solution that satisfies all constraints modeled in terms of ∞ costs. This would be a NP-complete problem. One way to work around this limitation is to include a small set of rules that cover each node individually and that can be used as a fallback rule in situations where no feasible solution has been obtained, which is similar to macro substitution techniques and ensures a correct but possibly non-optimal matching. These limitations do not apply to exact PBQP solvers such as the branch-&-bound algorithm. It is also straight-forward to extend the heuristic algorithm with a backtracking scheme on RN reductions, which would of course also be exponential in the worst case.

20.4 Summary and Further Reading

Aggressive optimizations for the instruction code selection problem are enabled by the use of SSA graph. The whole flow of a function is taken into account rather a local scope. The move from basic tree-pattern matching [5] to SSA-based DAG matching is a relative small step as long as a PBQP library and some basic infrastructure (graph grammar translator, etc.) is provided. The complexity of the approach is hidden in the discrete optimization problem called PBQP. Free PBQP libraries are available from the web-pages of the authors and a library is implemented as part of the LLVM [188] framework.

Many aspects of the PBQP formulation presented in this chapter could not be covered in detail. The interested reader is referred to the relevant literature [108, 106] for an in-depth discussion.

As we move from acyclic linear code regions to whole-functions, it becomes less clear in which basic block, the selected machine instructions should be emitted. For chain rules, the obvious choices are often non-optimal. In [244],

a polynomial-time algorithm based on generic network flows is introduced that allows a more efficient placement of chain rules across basic block boundaries. This technique is orthogonal to the generalization to complex patterns.

if-conversion
Very Long Instruction Word, VLIW
Instruction Level Parallelism, ILP
scheduling
predicated execution
speculative execution
dependence, control-dependence, data-parallelism
scheduling
select, instruction

CHAPTER 21

If-Conversion

C. Bruel

Progress: 90%

Text and figures formating in progress

21.1 Introduction

*

Very Large Instruction Word (VLIW)^{*} or Explicitly Parallel Instruction Computing (EPIC) architectures make Instruction Level Parallelism (ILP)^{*} visible within the Instruction Set Architecture (ISA), relying on static schedulers^{*} to organize the compiler output such that multiple instructions can be issued in each cycle.

If-conversion is the process of transforming a control-flow region with conditional branches, into an equivalent predicated^{*} or speculated^{*} sequence of instructions (into a region of Basic Blocks (possibly single) referred to as an Hyper-block). If-converted code replaces control dependencies^{*} by data dependencies^{*}, and thus exposes Instruction Level Parallelism^{*} very naturally within the new region at the software level.

Removing control hazards improves performance in several ways: By removing the misprediction penalty, the instruction fetch throughput is increased and the instruction cache locality improved. Enlarging the size of basic blocks allows earlier execution of long latency operations and the merging of multiple control-flow paths into a single flow of execution, that can later be exploited by scheduling frameworks such as VLIW scheduling, hyperblock scheduling or modulo scheduling^{*}.

Consider the simple example given in figure 21.1, that represents the execution of an `if-then-else-end` statement on a 4-issue processor with non-biased branches. In this figure, $r = q ? r_1 : r_2$ stands for a `select`^{*} instruction where r

speculation

is assigned r_1 if q is true, and r_2 otherwise. With standard basic block ordering, assuming that all instructions have a one cycle latency, the schedule height goes from five cycles in the most optimistic case, to six cycles. After if-conversion the execution path is reduced to four cycles with no branches, regardless of the test outcome, and assuming a very optimistic one cycle branch penalty. But the main benefit here is that it can be executed without branch disruption.

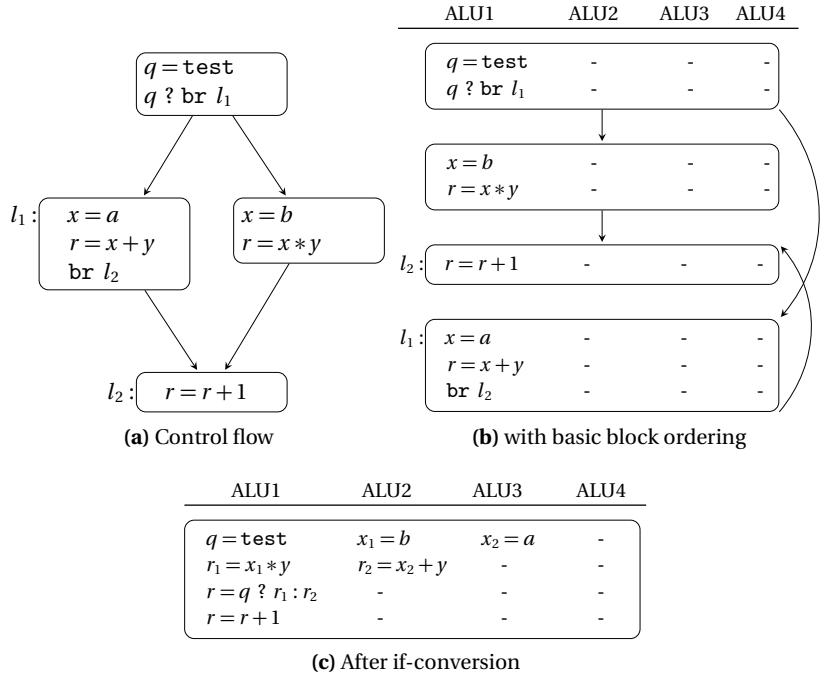


Fig. 21.1 Static schedule of a simple region on a 4-issue processor.

From this introductory example, we can observe that:

- the two possible execution paths have been merged into a single execution path, implying a better exploitation of the available resources;
- the schedule height has been reduced, because instructions can be control speculated before the branch;^{*}
- the variables have been renamed, and a *merge* pseudo-instruction have been introduced.

Thanks to SSA, the merging point is already materialized in the original control flow as a ϕ pseudo-instruction, and register renaming was performed by SSA construction. Given this, the transformation to generate if-converted code seems natural locally. Still exploiting those properties on larger scale control-flow regions requires a framework that we will develop further.

21.1.1 Architectural requirements

The *merge* pseudo operations needs to be mapped to a conditional form of execution in the target's architecture. As illustrated by Figure 21.2 we differentiate the three following models of conditional execution:

- *Fully predicated execution*: Any instruction can be executed conditionally on the value of a predicate operand.*
- *(Control) speculative execution*: The instruction is executed unconditionally, and then committed using conditional moves (`cmov` or `select`) instructions.*
- *Partially predicated execution*: Only a subset of the ISA is predicated, usually memory operations that are not easily speculated; other instructions are speculated.

*

In this figure, we use the notation $r = c ? r_1 : r_2$, to represent a `select`* like operation. Its semantic is identical to the gated ϕ_{if} -function* presented in Chapter 18: r takes the value of r_1 if c is true, r_2 otherwise. Similarly, we also use the notation $c ? r = op$ to represent the predicated execution of op if the predicate c is true; $\bar{c} ? r = op$ if the predicate c is false.

$p ? x = a + b$	$t_1 = a + b$	$x = a + b$
$\bar{p} ? x = a * b$	$t_2 = a * b$	$t = a * b$
	$x = p ? t_1 : t_2$	$x = \text{cmov } p, t$

(a) fully predicated

(b) speculative using `select`(c) speculative using `cmov`

Fig. 21.2 Conditional execution using different models

To be speculated, an instruction must not have any side effects, or hazards. For instance a memory load must not trap because of an invalid address. Memory operations are a major impediment to if-conversion. This is regrettable, because as any other long latency instructions, speculative loads can be very effective to fetching data earlier in the instruction stream, reducing stalls. Modern architectures provide architectural support to dismiss invalid address exceptions. Examples are the `1dw.d` dismissible load operation in the Multiflow Trace series of computers, or in the STMicroelectronics ST231 processor, but also the speculative load of the Intel IA64. The main difference is that with a dismissible model, invalid memory access exceptions are not delivered, which can be problematic in embedded or kernel environment that relies on memory exception for correct behavior. A speculative model allows to catch the exception thanks to the token bit check instruction. Some architectures, such as the IA64, offer both speculative and predicated memory operations.* Stores can also be executed conditionally by speculating part of their address value, with additional constraints on the ordering on the memory operations due to possible alias between the

predication, full speculation
predication, partial select, instruction gated instruction
 ϕ_{if} -function
dismissible model

*gated instruction
conditional move*

two paths. Figure 21.3 shows examples of various forms of speculative memory operations.

Note that the `select` instruction is an architecture instruction that does not need to be replaced during the SSA destruction phase. If the target architecture does not provide such gated instruction,^{*} it can be emulated using two conditional moves.^{*} This translation can be done afterward, and the `select` instruction still be used as an intermediate form. It allows the program to stay in full SSA form where all the data dependencies are made explicit, and can thus be fed to all SSA optimizers.

```
t = ld.s(addr)
chk.s (t)
*p ? x = t
```

(a) IA64 speculative load

```
t = ldw.d(addr)
x = select p ? t : x
```

(b) Multiflow/ST231 dismissible load

```
x = select p ? addr: dummy
stw(x, value)
```

(c) base store hoisting

```
index = select p ? i : j
stw(x[index], value)
```

(d) index store hoisting

Fig. 21.3 Examples of speculated memory operations

This chapter is organized as follow: we start to describe the SSA techniques to transform a CFG region in SSA form to produce an if-converted SSA representation using speculation. We then describe how this framework is extended to use predicated instructions, using the ψ -SSA form presented in Chapter 17. Finally we propose a global framework to pull together those techniques, incrementally enlarging the scope of the if-converted region to its maximum beneficial size.

21.2 Basic Transformations

Unlike global approaches, that identify a control-flow region and if-convert it in one shot, the technique described in this chapter is based on incremental reductions. To this end, we consider basic SSA transformations whose goal is to isolate a simple diamond-DAG structure (informally an `if-then-else-end`) that can be easily if-converted. The complete framework, that identifies and incrementally performs the transformation, is described in Section 21.3.

21.2.1 SSA operations on Basic Blocks

The basic transformation that actually if-converts the code is the ϕ removal^{*} that takes a simple diamond-DAG as an input, i.e., a single-entry-node/single-exit-node (SESE)^{*} DAG with only two distinct forward paths from its entry-node to its exit-node. The ϕ removal consists in (1) speculating the code of both branches in the entry basic block (denoted *head*); (2) then replacing the ϕ -function by a *select*; (3) finally simplifying the control flow to a single basic block. This transformation is illustrated by the example of Figure 21.4.

ϕ removal, reduction
single-entry-node/single-exit-node,
SESE
 ϕ reduction, reduction
 ϕ_{if} -function
Gated-SSA form

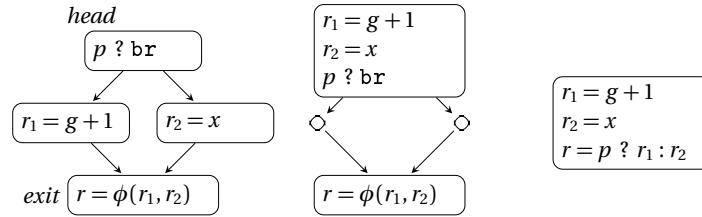


Fig. 21.4 ϕ removal

The goal of the ϕ reduction^{*} transformation is to isolate a diamond-DAG from a structure that resemble a diamond-DAG but has side entries to its exit block. This diamond-DAG can then be reduced using the ϕ removal transformation. Nested if-then-else-end in the original code can create such control flow. One can notice the similarity with the nested arity-two ϕ_{if} -functions^{*} used for gated-SSA^{*} (see Chapter 18). In the most general case, the join node of the considered region has n predecessors with ϕ -functions of the form $B_0 : r = \phi(B_1 : r_1, B_2 : r_2, \dots, B_n : r_n)$, and is such that removing edges from B_3, \dots, B_n would give a diamond-DAG. After the transformation, B_1 and B_2 point to a freshly created basic block, say B_{12} , that itself points to B_0 ; a new variable $r_{12} : r_{12} = \phi(B_1 : r_1, B_2 : r_2)$ is created in this new basic block; the ϕ -function in B_0 is replaced by $B_0 : r = \phi(B_{12} : r_{12}, \dots, B_n : r_n)$. This is illustrated through the example of Figure 21.5.

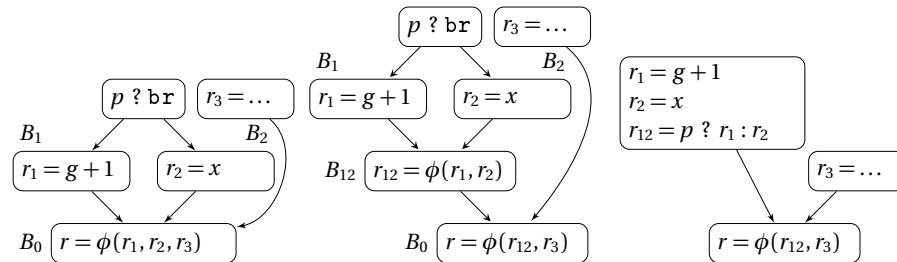


Fig. 21.5 ϕ reduction

*path duplication
renaming, of variables
copy folding
conjunctive predicate
merge
minimal SSA form
pruned SSA form
speculation
predicated execution*

The objective of *path duplication*^{*} is to get rid of all side entry edges that avoid a single-exit-node region to be a diamond-DAG. Through path duplication, all edges that point to a node different than the exit node or to the willing entry node, are “redirected” to the exit node. ϕ reduction can then be applied to the obtained region. More formally, consider two distinguished nodes, named *head* and the single exit node of the region *exit*, such that there are exactly two different control-flow paths from *head* to *exit*; consider (if exists), the first node *side_i* on one of the forward path *head* → *side₀* → ... *side_p* → *exit* that has at least two predecessors. The transformation duplicates the path $P = \text{side}_i \rightarrow \dots \rightarrow \text{side}_p \rightarrow \text{exit}$ into $P' = \text{side}'_i \rightarrow \dots \rightarrow \text{side}'_p \rightarrow \text{exit}$ and redirects *side_{i-1}* (or *head* if $i = 0$) to *side'_i*. All the ϕ -functions that are along P and P' for which the number of predecessors have changed have to be updated accordingly. Hence, a $r = \phi(\text{side}_p : r_1, B_2 : r_2, \dots, B_n : r_n)$ in *exit* will be updated into $r = \phi(\text{side}'_p : r_1, B_2 : r_2, \dots, B_n : r_n, \text{side}_p : r_1)$; a $r = \phi(\text{side}_{i-1} : r_0, r_1, \dots, r_m)$ originally in *side_i* will be updated into $r = \phi(r_1, \dots, r_m)$ in *side'_i* and into $r = \phi(r_0)$ i.e., $r = r_0$ in *side'_i*. Variables renaming^{*} (see Chapter 5) along with copy-folding can then be performed on P and P' . All steps are illustrated through the example of Figure 21.6.

The last transformation, namely the *Conjunctive predicate merge*^{*}, concerns the if-conversion of a control-flow pattern that sometimes appears on codes to represent logical and or or conditional operations. As illustrated by Figure 21.7 the goal is to get rid of side exit edges that avoid a single-entry-node region to be a diamond-DAG. As opposed to path duplication, the transformation is actually restricted to a very simple pattern highlighted in Figure 21.7 made up of three distinct basic block, *head*, that branches with predicate *p* to *side*, or *exit*. *side*, which is empty, branches itself with predicate *q* to another basic block outside of the region or to *exit*. Conceptually the transformation can be understood as first isolating the outgoing path *p* → *q* and then if-converting the obtained diamond-DAG.

Implementing the same framework on a non-SSA form program, would require more efforts: The ϕ reduction would require variable renaming, involving either a global data-flow analysis or the insertion of copies at the *exit* node of the diamond-DAG; inferring the minimum amount of select operations would require having and updating liveness information. SSA form solves the renaming issue for free, and as illustrated by Figure 21.8 the minimality and the pruned flavor of the SSA form allows to avoid inserting useless select operations.

21.2.2 Handling of predicated execution model

The ϕ removal transformation described above considered a speculative execution^{*} model. As we will illustrate hereafter, in the context of a predicated execution model, the choice of speculation versus predication^{*} is an optimization decision that should not be imposed by the intermediate representation. Also, transforming specu-

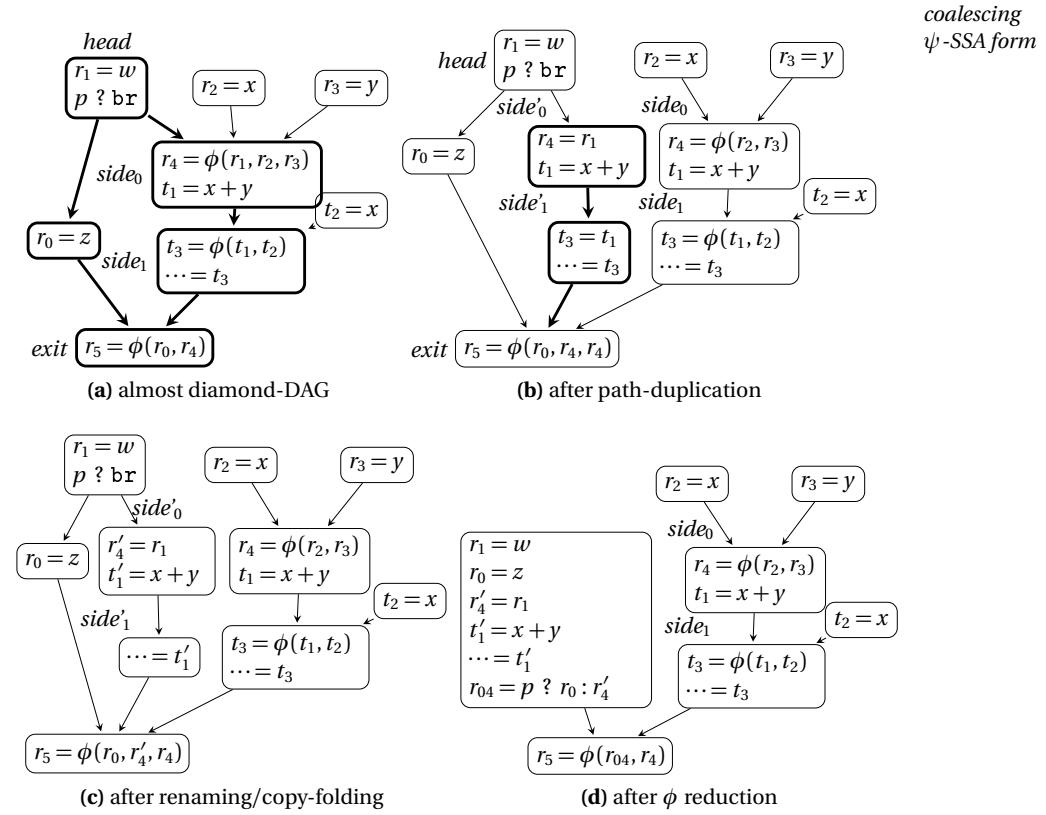


Fig. 21.6 Path duplication

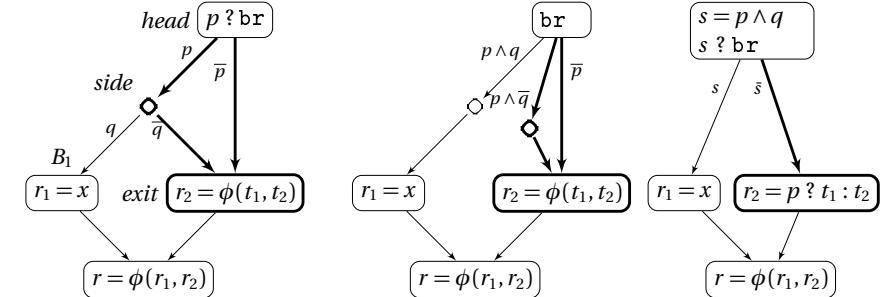


Fig. 21.7 convergent conjunctive merge

lated code into predicated code can be viewed as a coalescing^{*} problem. The use of ψ -SSA^{*} (see Chapter 17), as the intermediate form of if-conversion, allows to postpone the decision of speculating some code, while the coalescing problem is naturally handled by the ψ -SSA destruction phase.

*dependence, control-dependence, data-dependence, anti-parallel execution
 ψ -function*

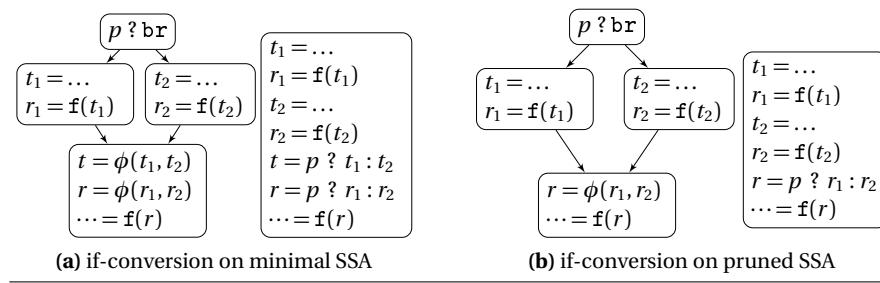


Fig. 21.8 SSA predicate minimality

Just as (control) speculating an operation on a control-flow graph corresponds to ignore the control dependence* with the conditional branch, speculating an operation on an if-converted code corresponds to remove the data dependence* with the corresponding predicate. On the other-hand on register allocated code, speculation adds anti-dependencies*. This trade-off can be illustrated through the example of Figure 21.9: For the fully predicated version of the code, the computation of p has to be done before the computations of x_1 and x_2 ; speculating the computation of x_1 removes the dependence with p and allows to execute it in parallel* with the test ($a \neq b$); if both the computation of x_1 and x_2 are speculated, they cannot be coalesced and when destruction ψ -SSA, the ψ -function* will give rise to some select instruction; if only the computation of x_1 is speculated, then x_1 and x_2 can be coalesced to x , but then an anti-dependence from $x = a + b$ and $p ? x = c$ appears that forbid its execution in parallel.

$p = (a \neq b)$			
$p ? x_1 = a + b$	$x_1 = a + b$	$x_1 = a + b$	$x = a + b$
$\bar{p} ? x_2 = c$	$x_2 = c$	$\bar{p} ? x_2 = c$	$\bar{p} ? x = c$
$x = \psi(p ? x_1, \bar{p} ? x_2)$	$x = \psi(p ? x_1, \bar{p} ? x_2)$	$x = \psi(p ? x_1, \bar{p} ? x_2)$	$x = \psi(p ? x_1, \bar{p} ? x_2)$

(a) predicated code

(b) fully speculated

(c) partially speculated

(d) after coalescing

Fig. 21.9 Speculation removes the dependency with the predicate but adds anti-dependencies between concurrent computations.

In practice, speculation is performed during the ϕ removal transformation, whenever it is possible (operations with side effect cannot be speculated) and considered as beneficial. As illustrated by Figure 21.10b, only the operations part of *one* of the diamond-DAG branch are actually speculated. This partial speculation leads to manipulating code made up of a mixed of predication and speculation.

Speculating code is the easiest part as it could be done prior to the actual if-conversion by simply hoisting the code above the conditional branch. Still we

			<i>def-use chains projection ψ-projection guard</i>
if (q)		$p = (a \otimes b)$	
then			
$p = (a \otimes b)$		$p = (a \otimes b)$	
$x_1 = a + b$		$x_1 = a + b$	
$\bar{p} ? x_2 = c$		$\bar{p} ? x_2 = c$	
$x = \psi(x_1, \bar{p} ? x_2)$		$x = \psi(x_1, \bar{p} ? x_2)$	
$d_1 = f(x)$		$d_1 = f(x)$	
else			
$d_2 = 3$		$\bar{q} ? d_2 = 3$	
end			
$d = \phi(d_1, d_2)$		$d = \psi(q ? d_1, \bar{q} ? d_2)$	
(a) nested if	(b) speculating the then branch	(c) predicating both branches	

Fig. 21.10 Inner region ψ

would like to outline that since ψ -functions are part of the intermediate representation, they can be considered for inclusion in a candidate region for if-conversion, and in particular for speculation. However, the strength of ψ -SSA allows to treat ψ -functions just as any other operation. Consider the code of Figure 21.10a containing a sub region already processed. To speculate the operation $d_1 = f(x)$, the operation defining x , i.e., the ψ -function, also has to be speculated. Similarly, all the operations defining the operands x_1 and x_2 should also be speculated. If one of them can produce hazardous execution, then the ψ -function cannot be speculated, which forbids in turn the operation $d_1 = f(x)$ to be speculated. Marking operations that cannot be speculated can be done easily using a forward propagation along def-use chains.^{*}

All operations that cannot be speculated, including possibly some ψ -functions, must be predicated. Suppose we are considering a none-speculated operation we aim at if-converting, that is part of the then branch on predicate q . Just as for $x_2 = c$ in Figure 21.10a, this operation might be already predicated (on \bar{p} here) prior to the if-conversion. In that case, a *projection*^{*} on q is performed, meaning that instead of predicing $x_2 = c$ by \bar{p} it gets prediced by $q \wedge \bar{p}$. A ψ -function can also be projected on a predicate q as described in Chapter 17: All gates of each operand are individually projected on q . As an example, originally non-gated operand x_1 gets gated by q , while the \bar{p} gated operand x_2 gets gated by $s = q \wedge \bar{p}$. Note that as opposed to speculating it, predicing a ψ -function does not impose to predicate the operation that define its operands. The only subtlety related to projection is related to generating the new predicate as the logical conjunction of the original guard (e.g., \bar{p})^{*} and the current branch predicate (e.g., q). Here s needs to be computed at some point. Our heuristic consists in first listing the set of all necessary predicates and then emitting the corresponding code at the earlier place. Here, used predicates are q , \bar{q} , and $q \wedge \bar{p}$. q and \bar{q} are already available. The earlier place where $q \wedge \bar{p}$ can be computed is just after calculating p .

*incremental update, of
SSA*

Once operations have been speculated or projected (on q for the `then` branch, on \bar{q} for the `else`), each ϕ -functions at the merge point is replaced by a ψ -function: operands of speculated operations are placed first and guarded by true; operands of projected operations follow, guarded by the predicate of the corresponding branch.

21.3 Global Analysis and Transformations

Hot regions are rarely just composed of simple `if-then-else-end` control-flow regions but processors have limited resources: the number of registers will determine the acceptable level of data dependencies to minimize register pressure; the number of predicate registers will determine the depth of the if-conversion so that the number of conditions does not exceed the number of available predicates; the number of processing units will determine the number of instructions that can be executed simultaneously. The inner-outer incremental process advocated in this chapter allows to evaluate finely the profitability of if-conversion.

21.3.1 SSA Incremental if-conversion algorithm

The algorithm takes as input a CFG in SSA form and applies incremental reductions on it using a list of candidate conditional basic blocks sorted in post-order. Each basic block in the list designates the head of a sub-graph that can be if-converted using the transformations described in Section 21.2. Post-order traversal allows to process the regions from inner to outer. When the if-converted region cannot grow anymore because of resources, or because a basic block cannot be if-converted, then another candidate is considered in the post-order list until all the CFG is explored. Note that as the reduction proceeds, maintaining SSA^{*} can be done using the general technique described in Chapter 5. However, basic ad-hoc updates can be implemented instead.

Consider for example the CFG from the gnu `wc` (word count) program reported in Figure 21.11a. The exit node $BB7$, and basic block $BB3$ that contains a function call cannot be if-converted (represented in gray). The post-order list of conditional blocks (represented in bold) is [$BB11$, $BB17$, $BB16$, $BB14$, $BB10$, $BB9$, $BB6$, $BB2$]. The first candidate region is composed of $\{BB11, BB2, BB12\}$. ϕ reduction can be applied, promoting the instructions of $BB12$ in $BB11$. $BB2$ becomes the single successor of $BB11$. The region headed by $BB17$ is then considered. $BB19$ cannot yet be promoted because of the side entries coming both from $BB15$ and $BB16$. $BB19$ is duplicated into a $BB19'$ with $BB2$ as successor. $BB19'$ can then be promoted into $BB17$. The region headed by $BB16$ that have

now $BB17$ and $BB19$ as successors is considered. $BB19$ is duplicated into $BB19''$, so as to promote $BB17$ and $BB19'$ into $BB16$ through ϕ reduction. $BB19'$ already contains predicated operations from the previous transformation. So, a new merging predicate is computed and inserted. After the ϕ removal is fully completed $BB16$ has a unique successor, $BB2$. $BB14$ is the head of the new candidate region where $BB15$, $BB16$ can be promoted. Again since $BB16$ contains predicated and predicate setting operations, a newly predicate must be created to hold the merged conditions. $BB10$ is then considered; $BB14$ needs to be duplicated to $BB14'$. The process finished with the region head by $BB9$.

*tail duplication
path duplication
hyperblock formation*

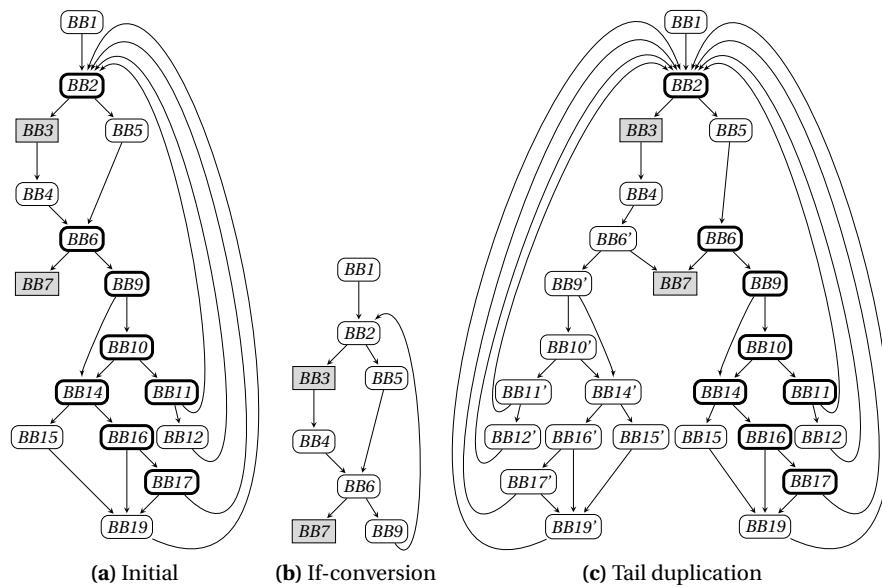


Fig. 21.11 If-conversion of wc (word count program). basic blocks in gray cannot be if-converted. Tail duplication can be used to exclude $BB3$ from the to-be-if-converted region.

21.3.2 Tail Duplication

*

Just as for the example of Figure 21.11, some basic blocks (such as $BB3$) may have to be excluded from the region to if-convert. *Tail duplication* can be used for this purpose. Similar to path duplication* described in Section 21.2, the goal of tail duplication is to get rid of incoming edges of a region to if-convert. This is usually done in the context of hyperblock formation*, which technique consists in, as opposed to the inner-outer incremental technique described in this

branch coalescing

chapter, to if-convert a region in “one shot”. Consider again the example of Figure 21.11a, and suppose the set of selected basic blocks defining the region to if-convert consists of all basic blocks from BB2 to BB19 excluding BB3, BB4, and BB7. Getting rid of the incoming edge from BB4 to BB6 is possible by duplicating all basic blocks of the region reachable from BB6 as shown in Figure 21.11b.

Formally, consider a region \mathcal{R} made up of a set of basic blocks, a distinguished one *entry* and the others denoted $(B_i)_{2 \leq i \leq n}$, such that any B_i is reachable from *entry* in \mathcal{R} . Suppose a basic block B_s has some predecessors out_1, \dots, out_m that are not in \mathcal{R} . Tail duplication consists in: (1) for all B_j (including B_s) reachable from B_s in \mathcal{R} , create a basic block B'_j as a copy of B_j ; (2) any branch from B'_j that points to a basic block B_k of the region is rerouted to its duplicate B'_k ; (3) any branch from a basic block out_k to B_s is rerouted to B'_s . In our example, we would have *entry* = BB2, B_s = BB6, and *out* = BB4.

A global approach would just do as in Figure 21.11c: first select the region; second, get-rid of side incoming edges using tail duplication; finally perform if-conversion of the whole region in one shot. We would like to point out that there is no phasing issue with tail duplication. To illustrate this point, consider the example of Figure 21.12a where BB2 cannot be if-converted. The selected region is made up of all other basic blocks. Using a global approach as in standard hyperblock formation, tail duplication would be performed prior to any if-conversion. This would lead to the CFG of Figure 21.12b. Note that a new node, BB7, has been added here after the tail duplication by a process called branch coalescing.^{*} Applying if-conversion on the two disjoint regions respectively head by BB4 and BB4' would lead to the final code shown in Figure 21.12c. Our incremental scheme would first perform if-conversion of the region head by BB4, leading to the code depicted in Figure 21.12e. Applying tail duplication to get rid of side entry from BB2 would lead to exactly the same final code as shown in Figure 21.12f.

21.3.3 Profitability

Fusing execution paths can over commit the architectural ability to execute in parallel the multiple instructions: Data dependencies and register renaming introduce new register constraints. Moving operations earlier in the instruction stream increases live-ranges. Aggressive if-conversion can easily exceed the processor resources, leading to excessive register pressure or moving infrequently used long latencies instructions into the critical path. The prevalent idea is that a region can be if-converted if the cost of the resulting if-converted basic block is smaller than the cost of each basic block of the region taken separately weighted by its execution probability. To evaluate those costs, we consider all possible paths impacted by the if-conversion.

For all transformations but the conjunctive predicate merge, there are two such paths starting at basic block *head*. For the code of Figure 21.13, we would

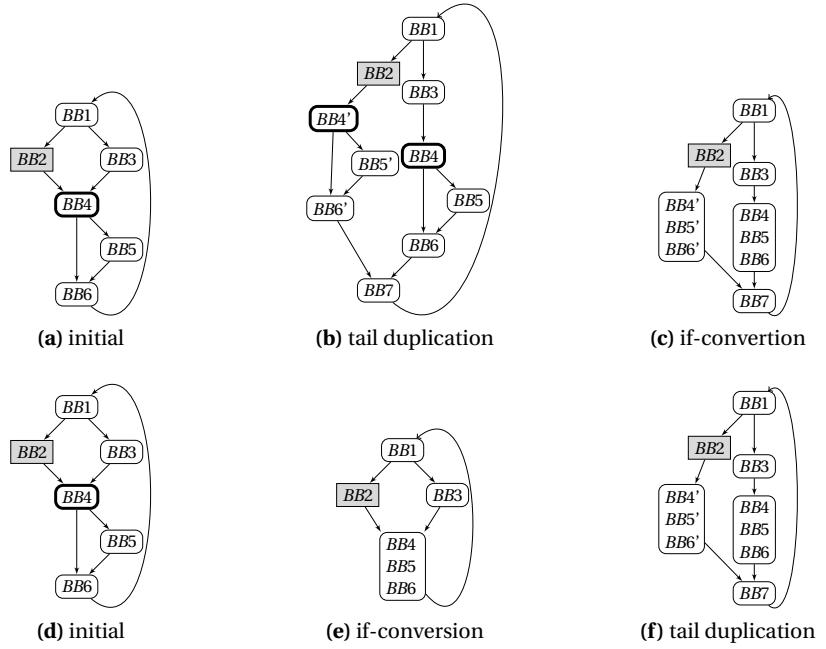


Fig. 21.12 Absence of phasing issue for tail duplication

have $\text{path}_p = [\text{head}, B_1, \text{exit}]$ and $\text{path}_{\bar{p}} = [\text{head}, B'_1, B'_2, \text{exit}]$ of respective probability $\text{prob}(p)$ and $\text{prob}(\bar{p})$. For a path $P_q = [B_0, B_1, \dots, B_n]$ of probability $\text{prob}(q)$, its cost is given by $\widehat{P}_q = \text{prob}(q) \times \sum_{i=0}^{n-1} [\widehat{B_i}, \widehat{B_{i+1}}]$ where $[\widehat{B_i}, \widehat{B_{i+1}}]$ represents the cost of basic block $\widehat{[B_i]}$ estimated using its schedule height plus the branch latency br_lat , if the edge (B_i, B_{i+1}) corresponds to a conditional branch, 0 otherwise. Note that if B_i branches to S_q on predicate q , and falls through to $S_{\bar{q}}$, $\widehat{B_i} = \text{prob}(q) \times ([\widehat{B_i}, \widehat{S_q}]) + \text{prob}(\bar{q}) \times ([\widehat{B_i}, \widehat{S_{\bar{q}}}] = [\widehat{B_i}] + \text{prob}(q) \times \text{br_lat}$.

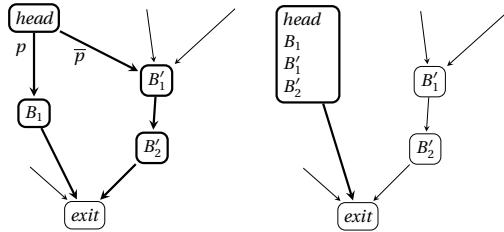
Let $\text{path}_p = [\text{head}, B_1, \dots, B_n, \text{exit}]$ and $\text{path}_{\bar{p}} = [\text{head}, B'_1, \dots, B'_m, \text{exit}]$ be the two paths with the condition taken branch on p . Then, the overall cost before if-conversion simplifies to

$$\text{cost}_{\text{control}} = \widehat{\text{path}_p} + \widehat{\text{path}_{\bar{p}}} = [\widehat{\text{head}}] + \text{prob}(p) \times \left(\text{br_lat} + \sum_{i=0}^n [\widehat{B_i}] \right) + \text{prob}(\bar{p}) \times \sum_{i=0}^m [\widehat{B'_i}]$$

This is to be compared to the cost after if-conversion

$$\text{cost}_{\text{predicated}} = [\text{head} \circ (\bigcirc_{i=1}^n B_i) \circ (\bigcirc_{i=1}^m B'_i)]$$

where \circ is the composition function that merges basic blocks together, removes associated branches and creates the predicate operations.

Fig. 21.13 ϕ reduction

The profitability for the logical conjunctive merge of Figure 21.14 can be evaluated similarly. There are three paths impacted by the transformation: $path_{p \wedge q} = [head, side, B_1[], path_{p \wedge \bar{q}} = [head, side, exit[]$, and $path_{\bar{p}} = [head, exit[]$ of respective probability $\text{prob}(p \wedge q)$, $\text{prob}(p \wedge \bar{q})$, and $\text{prob}(\bar{p})$. The overall cost before the transformation (if branches are on p and q) $\widehat{path}_{p \wedge q} + \widehat{path}_{p \wedge \bar{q}} + \widehat{path}_{\bar{p}}$ simplifies to

$$\text{cost}_{\text{control}} = \widehat{\text{head}} + \widehat{\text{side}} = [\widehat{\text{head}}] + \text{prob}(p) \times (1 + \text{prob}(q)) \times \text{br_lat}$$

which should be compared to (if the branch on the new *head* block is on $p \wedge q$)

$$\text{cost}_{\text{predicated}} = \widehat{\text{head} \circ \text{side}} = [\widehat{\text{head} \circ \text{side}}] + \text{prob}(p) \times \text{prob}(q) \times \text{br_lat}$$

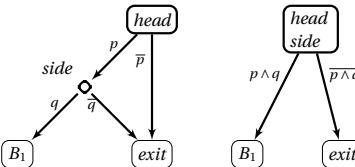


Fig. 21.14 conjunctive predicate merge

Note that if $\text{prob}(p) \ll 1$, emitting a conjunctive merge might not be beneficial. In that case, another strategy such as path duplication from the *exit* block will be evaluated. Profitability for any conjunctive predicate merge (disjunctive or conjunctive; convergent or not) is evaluated similarly.

A speed oriented objective function needs the target machine description to derive the instruction latencies, resource usage and scheduling constraints. The local dependencies computed between instructions are used to compute the dependence height. The branch probability is obtained either from static branch prediction heuristics, profile information or user inserted directives. Naturally, this heuristic can be either pessimistic, because it does not take into account new optimization opportunities introduced by the branch removals or explicit new dependencies, or optimistic because of bad register pressure es-

timation leading to register spilling on the critical path, or uncertainty in the branch prediction. But since the SSA incremental if-conversion framework reduces the scope for the decision function to a localized part of the CFG, the size and complexity of the inner region under consideration makes the profitability a comprehensive process. This cost function is fast enough to be reapplied to each region during the incremental processing, with the advantage that all the instructions introduced by the if-conversion process in the inner regions, such as new predicate merging instructions or new temporary pseudo registers, can be taken into account.

21.4 Conclusion

We presented in this chapter how an if-conversion algorithm can take advantage of the SSA properties to efficiently assign predicates and lay out the new control flow in an incremental, inner-outer process. As opposed to the alternative top-down approach, the region selection can be reevaluated at each nested transformation, using local analysis. Basic block selection and if-conversion are performed as a single process, hyperblocks being created lazily, using well known techniques such as tail duplication or branch coalescing only when the benefit is established. Predication and speculation are often presented as two different alternatives for if-conversion. While it is true that they both require different hardware support, they should coexist in an efficient if-conversion process such that every model of conditional execution is accepted. Thanks to conditional moves and ψ transformations, they are now generated together in the same framework.

21.5 Additional reading

[233], exposes ILP in VLIW architectures using trace scheduling and local if-converted if-then-else regions using the `select` and dismissible load operations. The idea behind was to enable the compiler to statically reorganize the instruction. In this respect, predictability [122] becomes a major criteria for profitability.

To overcome the hard to predict profitability in conventional if-conversion algorithms, Reverse if-conversion was proposed in [20], reconstructing the control flow at schedule time, after application of more aggressive region selection criteria.

Hyperblocks [195] was proposed as the primary if-converted scheduling framework, excluding basic blocks which do not justify their inclusion into the if-converted flow of control.

The duality between SSA like ϕ s and predicate dependencies have been used in other works. In SSA-PS [148], Predicated Switching operations are used to realize the conditional assignments using aggressive speculation techniques with conditional moves. Phi-Predication [78] uses a modified version of the RK algorithm, to map phi-predication with phi-lists, holding guard and topological information. In both works, the use of SSA aims at solving the multiple definition problem exploiting variable renaming and join points, but they are based on speculation using conditional moves.

In [264], ψ instructions are inserted while in SSA using a modified version of the classical Fang algorithm [112], enabling support for a fully predicated ISA. Those works established that standard if-conversion techniques can be applied to a SSA form using the ψ -SSA representation, or light weight ϕ -SSA generation, but do not yet exploit the native SSA properties to build up the if-converted region. A global SSA framework was presented [57] to support *select* moves using aggressive speculation techniques, further extended to ψ -SSA [58] allowing a mix of speculative and predicated techniques.

[193] evaluates how predicated operations can be performed using an equivalent sequences of speculative and conditional moves, starting from an if-converted region fully predicated.

CHAPTER 22

SSA destruction for machine code

F. Rastello

Progress: 90%



Text and figures formating in progress

*

*



22.1 Introduction

Chapter 3 provides a basic algorithm for destructing SSA that suffers from several limitations and drawbacks: first, it works under implicit assumptions that are not necessarily fulfilled at machine level; second, it must rely on subsequent phases to remove the numerous copy operations it inserts; finally, it increases subsequently the size of the intermediate representation, thus making it not suitable for just-in-time compilation.

Correctness SSA at machine level complicates the process of destruction that can potentially lead to bugs if not performed carefully. The algorithm described in Section 3.2 involves the splitting of every critical edges*. Unfortunately, because of specific architectural constraints, region boundaries, or exception handling code, the compiler might not allow the splitting of a given edge. Fortunately, as we will see further, this obstacle can easily be overcame. But then it becomes essential to be able to append a copy operation at the very end of a basic block which might neither be possible. Also, care must be taken with duplicated edges, i.e., when the same basic block appears twice in the list of predecessors. This can occur after control-flow graph structural optimizations like dead code elimination or empty block elimination. In such case, the edges should be considered as critical and then split.

SSA imposes a strict discipline on variable naming: every “name” must be associated to only one definition which is obviously most of the time not compatible with the instruction set of the targeted architecture. As an example, a two-address mode instruction, such as auto-increment ($x = x + 1$) would enforce its definition to use the same resource than one of its arguments (defined elsewhere), thus imposing two different definitions for the same temporary variable. This is why some prefer using, for SSA construction, the notion of *versioning*^{*} in place of renaming. Implicitly, two versions of the same original variable should not interfere, while two names can. The former simplifies the SSA destruction phase, while the latter simplifies and allows more transformations to be performed under SSA. Apart from dedicated registers for which optimizations are usually very careful in managing their live-range, register constraints related to calling conventions or instruction set architecture might be handled by the register allocation phase. However, as we will see, enforcement of register constraints impacts the register pressure as well as the number of copy operations. For those reasons we may want those constraints to be expressed earlier (such as for the pre-pass scheduler), in which case the SSA destruction phase might have to cope with them.

Code quality The natural way of lowering ϕ -functions and expressing register constraints is through the insertion of copies (when edge-splitting is not mandatory as discussed above). If done carelessly, the resulting code will contain many temporary-to-temporary copy operations. In theory, reducing the amount of these copies is the role of the coalescing during the register allocation phase. A few memory and time-consuming existing coalescing heuristics mentioned in Chapter 23 are quite effective in practice. The difficulty comes both from the size of the interference graph (the information of colorability is spread out) and the presence of many overlapping live-ranges that carry the same value (so non-interfering). Coalescing can also, with less effort, be performed prior to the register allocation phase. As opposed to a (so-called conservative) coalescing during register allocation, this aggressive coalescing would not cope with the interference graph colorability. As we will see, strict SSA form is really helpful for both computing and representing equivalent variables. This makes the SSA destruction phase the good candidate for eliminating (or not inserting) those copies.

Speed and Memory Footprint The cleanest and simplest way to perform SSA destruction with good code quality is to first insert copy instructions to make the SSA form conventional, then take advantage of the SSA form to run efficiently aggressive coalescing (without breaking the conventional property), before eventually renaming ϕ -webs^{*} and getting rid of ϕ -functions. Unfortunately this approach will lead, in a transitional stage, to an intermediate representation with a substantial number of variables: the size of the liveness sets and interference graph classically used to perform coalescing become prohibitively large for dynamic compilation. To overcome this difficulty one can compute liveness and interference on demand which, as we already mentioned, is made simpler by the use of SSA form. Remains the process of copy insertion itself that might

still take a substantial amount of time. To fulfill memory and time constraints imposed by just-in-time compilation, one idea is to *virtually* insert those copies, and only *effectively* insert the non-coalesced ones.

This chapter addresses those three issues: handling of machine level constraints, code quality (elimination of copies), and algorithm efficiency (speed and memory footprint). The layout falls into three corresponding sections.

22.2 Correctness

Isolating ϕ -node using copies In most cases, edge splitting can be avoided by treating symmetrically ϕ -uses and definition-operand: instead of just inserting copies on the incoming control-flow edges of the ϕ -node (one for each use operand), a copy is also inserted on the outgoing edge (one for its defining operand). This has the effect of isolating the value associated to the ϕ -node thus avoiding (as discussed further) SSA destruction issues such as the well known lost-copy problem. The process of ϕ -node isolation is illustrated by Figure 22.1. The corresponding pseudo-code is given in Algorithm 22.1. If, because of different ϕ -functions, several copies are introduced at the same place, they should be viewed as parallel copies. For that reason, an empty parallel copy is inserted both at the beginning (i.e., right after ϕ -functions, if any) and at the end of each basic-block (i.e., just before the branching operation, if any). Note that, as far as correctness is concerned, copies can be sequentialized in any order, as they concern different variables.



Fig. 22.1 Isolation of a ϕ -node

When incoming edges are not split, inserting a copy not only for each argument of the ϕ -function, but also for its result is important: without the copy $a'_0 \leftarrow a_0$, the ϕ -function defines directly a_0 whose live range can be long enough to intersect the live range of some a'_i , $i > 0$, if a_0 is live out of the block B_i where a'_i is defined. Two cases are possible: either a_0 is used in a successor of $B_i \neq B_0$, in which case the edge from B_i to B_0 is *critical* (as in the “lost copy problem”), or a_0 is used in B_0 as a ϕ -function argument (as in the “swap problem”). In this latter case, if parallel copies are used, a_0 is dead before a'_i is defined but, if copies are

sequentialized blindly, the live range of a_0 can go beyond the definition point of a'_i and lead to incorrect code after renaming a_0 and a'_i with the same name. ϕ -node isolation allows to solve most of the issues that can be faced at machine level. However, there remains subtleties listed below.

Limitations There is a tricky case, when the basic block contains variables *defined after* the point of copy insertion. This is for example the case for the PowerPC `bclr` branch instructions with a behavior similar to hardware loop. In addition to the condition, a counter u is decremented by the instruction itself. If u is used in a ϕ -function in a direct successor block, no copy insertion can split its live range. It must then be given the same name as the variable defined by the ϕ -function. If both variables interfere, this is just impossible! To solve the problem, the SSA optimization could be designed with more care, or the counter variable must not be promoted to SSA, or some instruction must be changed, or the control-flow edge must be split somehow. SSA destruction by copy insertion alone is not always possible, depending on the branch instructions and the particular case of interferences.

For example, suppose that for the code of Figure 22.2a, the instruction selection chooses a branch with decrement (denoted `br_dec`) for Block B_1 (Figure 22.2b). Then, the ϕ -function of Block B_2 , which uses u , cannot be translated out of SSA by standard copy insertion because u interferes with t_1 and its live range cannot be split. To go out of SSA, one could add $t_1 \leftarrow u - 1$ in Block B_1 to anticipate the branch. Or one could split the critical edge between B_1 and B_2 as in Figure 22.2c. In other words, simple copy insertions is not enough in this case.

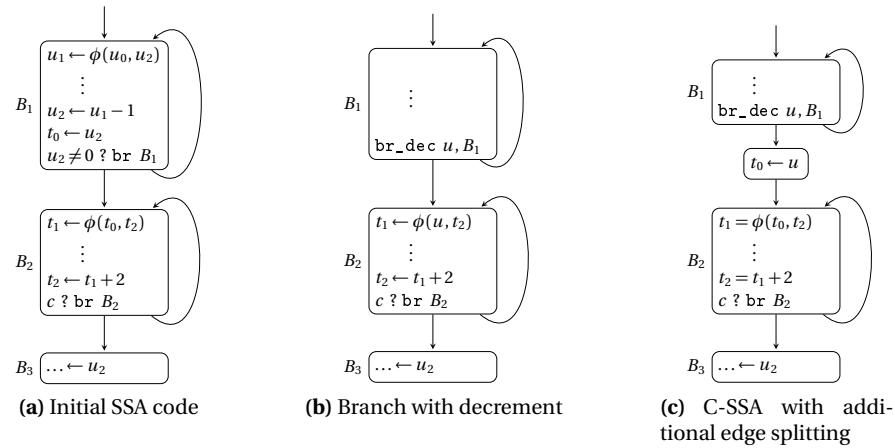


Fig. 22.2 Copy insertion may not be sufficient. `br_dec u, B1` decrements u , then branches to B_1 if $u \neq 0$.

There is another tricky case when a basic-block have twice the same predecessor block. This can result from consecutively applying copy-folding and

control-flow graph structural optimizations like dead code elimination or empty block elimination. This is the case for the example of Figure 22.3 where copy-folding^{*} would remove the copy $a_2 \leftarrow b$ in Block B_2 . If B_2 is eliminated, there is no way to implement the control dependence of the value to be assigned to a_3 other than through predicated code (see chapters 17 and 18) or through the re-insertion of a basic-block between B_1 and B_0 by the split of one of the edges.

*copy-folding
instruction set
architecture
application binary
interface
pinning*

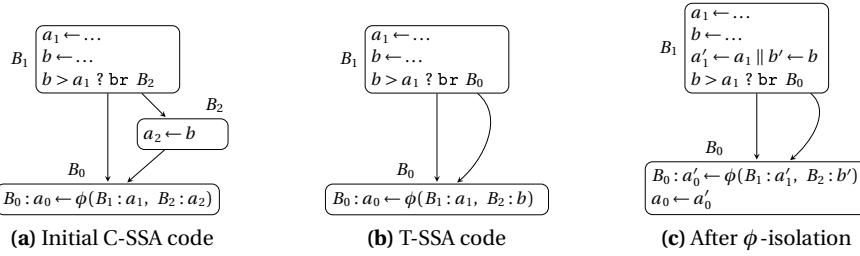


Fig. 22.3 Copy-folding followed by empty block elimination can lead to SSA code for which destruction is not possible through simple copy insertion

The last difficulty SSA destruction faces when performed at machine level is related to register constraints such as instruction set architecture (ISA)^{*} or application binary interface (ABI)^{*} constraints. For the sake of the discussion we differentiate two kinds of resource constraints that we will refer as *operand pinning* and *live-range pinning*. The live-range pinning of a variable v to resource R will be represented R_v , just as if v were a version of temporary R . An operand pinning to a resource R will be represented using the exponent ${}^{\dagger R}$ on the corresponding operand. Live-range pinning expresses the fact that the *entire* live-range of a variable must reside in a given resource (usually a dedicated register). An example of live-range pinning are versions of the stack-pointer temporary that must be assigned back to register SP . On the other hand the pinning of an operation's operand to a given resource does not impose anything on the live-range of the corresponding variable. The scope of the constraint is restricted to the operation. Examples of operand pinning are operand constraints such as *2-address-mode* where two operands of one instruction must use the same resource; or where an operand must use a given register. This last case encapsulates ABI constraints.

Note that looser constraints where the live-range or the operand can reside in more than one resource are not handled here. We assume this to always be the responsibility of register allocation. We first simplify the problem by transforming any operand pinning to a live-range pinning as sketched in Figure 22.4: parallel copies with new variables pinned to the corresponding resource are inserted just before (for use-operand pinning) and just after (for definition-operand pinning) the operation.

$p_2^{\uparrow T} \leftarrow p_1^{\uparrow T} + 1$ (a) Operand pinning of an auto-increment	$T_{p_1} \leftarrow p_1$ $T_{p_2} \leftarrow T_{p_1} + 1$ $p_2 \leftarrow T_{p_2}$ (b) Corresponding live-range pinning
$a^{\uparrow R0} \leftarrow f(b^{\uparrow R0}, c^{\uparrow R1})$ (c) Operand pinning of a function call	$R0_{b'} = b \parallel R1_{c'} = c$ $R0_{a'} = f(R0_{b'}, R1_{c'})$ $a = R0_{a'}$ (d) Corresponding live-range pinning

Fig. 22.4 Operand pinning and corresponding live-range pinning

Detection of strong interferences The scheme we propose in this section to perform SSA destruction that deals with machine level constraints does not address compilation cost (in terms of speed and memory footprint). It is designed to be simple. It first inserts parallel copies to isolate ϕ -functions and operand pinning. Then it checks for interferences that would persist. We will denote such interferences as strong, as they cannot be tackled through the simple insertion of temporary-to-temporary copies in the code. We consider that fixing strong interferences should be done on a case by case basis and restrict the discussion here on their detection.

As far as correctness is concerned, Algorithm 22.1 splits the data flow between variables and ϕ -nodes through the insertion of copies. For a given ϕ -function $a_0 \leftarrow \phi(a_1, \dots, a_n)$, this transformation is correct as long as the copies can be inserted close enough to the ϕ -function. It might not be the case if the insertion point (for a use-operand) of copy $a'_i \leftarrow a_i$ is not dominated by the definition point of a_i (such as for argument u of the ϕ -function $t_1 \leftarrow \phi(u, t_2)$ for the code of Figure 22.2b); symmetrically, it will not be correct if the insertion point (for the definition-operand) of copy $a_0 \leftarrow a'_0$ does not dominate all the uses of a_0 . Precisely this leads to inserting in Algorithm 22.1 the following tests:

- line 11: “**if** the definition of a_i does not dominate PC_i **then** continue;”
- line 19: “**if** one use of a_0 is not dominated by PC_0 **then** continue;”

For the discussion, we will denote as *split operands* the newly created local variables to differentiate them to the ones concerned by the two previous cases (designed as *non-split operands*). We suppose a similar process have been performed for operand pinning to express them in terms of live-range pinning with very short (when possible) live-ranges around the concerned operations.

At this point, the code is still under SSA and the goal of the next step is to check that it is conventional: this will obviously be the case only if all the variables of a ϕ -web can be coalesced together. But not only: the set of all variables pinned to a common resource must also be interference free. We say that x and y are *pin- ϕ -related* to one another if they are ϕ -related or if they are pinned to a common resource. The transitive closure of this relation defines an equivalence relation

that partitions the variables defined locally in the procedure into equivalence classes, the pin- ϕ -webs. Intuitively, the pin- ϕ -equivalence class of a resource represents a set of resources “connected” via ϕ -functions and resource pining. The computation of ϕ -webs given by Algorithm 3.4 can be generalized easily to compute pin- ϕ -webs. The resulting pseudo-code is given by Algorithm 22.2.

Algorithm 22.1: Algorithm making non-conventional SSA form conventional by isolating ϕ -nodes

```

1 begin
2   foreach  $B$ : basic block of the CFG do
3     | insert an empty parallel copy at the beginning of  $B$ ;
4     | insert an empty parallel copy at the end of  $B$ ;
5   end
6   foreach  $B_0$ : basic block of the CFG do
7     | foreach  $\phi$ -function at the entry of  $B_0$  of the form  $a_0 = \phi(B_1 : a_1, \dots, B_n : a_n)$  do
8       |   | foreach  $a_i$  (argument of the  $\phi$ -function corresponding to  $B_i$ ) do
9         |   |   | let  $PC_i$  be the parallel-copy at the end of  $B_i$ ;
11
12         |   |   | let  $a'_i$  be a freshly created variable;
13         |   |   | add copy  $a'_i \leftarrow a_i$  to  $PC_i$ ;
14         |   |   | replace  $a_i$  by  $a'_i$  in the  $\phi$ -function;
15       |   | end
16       |   begin
17         |   | let  $PC_0$  be the parallel-copy at the beginning of  $B_0$ ;
19
20         |   | let  $a'_0$  be a freshly created variable;
21         |   | add copy  $a'_0 \leftarrow a'_i$  to  $PC_0$ ;
22         |   | replace  $a_0$  by  $a'_0$  in the  $\phi$ -function;
23       |   | end
24       |   | /* all  $a'_i$  can be coalesced and the  $\phi$ -function removed */
25     |   end
26   end
27 end

```

Now, one need to check that each web is interference free. A web contains variables and resources. The notion of interferences between two variables is the one discussed in Section 2.7 for which we will propose an efficient implementation later in this chapter. A variable and a physical resource do not interfere while two distinct physical resources will interfere with one another.

If any interference have been discovered, it has to be fixed on a case by case basis. Note that some interferences such as the one depicted in Figure 22.3 can be detected and handled initially (through edge splitting if possible) during the copy insertion phase.

Algorithm 22.2: The pin- ϕ -webs discovery algorithm, based on the union-find pattern

```

1 begin
2   for each resource  $R$   $\text{web}(R) \leftarrow \{R\}$ ;
3   for each variable  $v$   $\text{web}(v) \leftarrow \{v\}$ ;
4   if  $v$  pinned to a resource  $R$   $\text{union}(\text{web}(R), \text{web}(v))$  for each instruction of the form
       $a_{\text{dest}} = \phi(a_1, \dots, a_n)$  for each source operand  $a_i$  in instruction
       $\text{union}(\text{web}(a_{\text{dest}}), \text{web}(a_i))$ 
5 end

```

22.3 Code Quality

Once the code is in conventional SSA, the correctness problem is solved: de-structuring it is by definition straightforward, as it relies in renaming all variables in each ϕ -web into a unique representative name and then remove all ϕ -functions. To improve the code however, it is important to remove as many copies as possible.

Aggressive coalescing

This can be treated with standard non-SSA coalescing technique as conventional SSA allows to get rid of ϕ -functions: the set of variables of a SSA web can be coalesced leading to a single non-SSA variable; liveness and interferences can then be defined as for regular code (with parallel copies). An interference graph (as depicted in Figure 22.5e) can be used. A solid edge between two nodes (e.g., between x_2 and x_3) materialize the presence of an interference between the two corresponding variables, i.e., expressing the fact that they cannot be coalesced and share the same resource. A dashed edge between two nodes materializes an affinity between the two corresponding variables, i.e., the presence of a copy (e.g., between x_2 and x'_2) that could be removed by their coalescing.

This process is illustrated by Figure 22.5: the isolation of the ϕ -function leads to inserting the three copies that respectively define x'_1 , x'_3 , and uses x'_2 ; the corresponding ϕ -web $\{x'_1, x'_2, x'_3\}$ is coalesced into a representative variable x ; according to the interference graph of Figure 22.5e, x_1 , x_3 can then be coalesced with x leading to the code of Figure 22.5d.

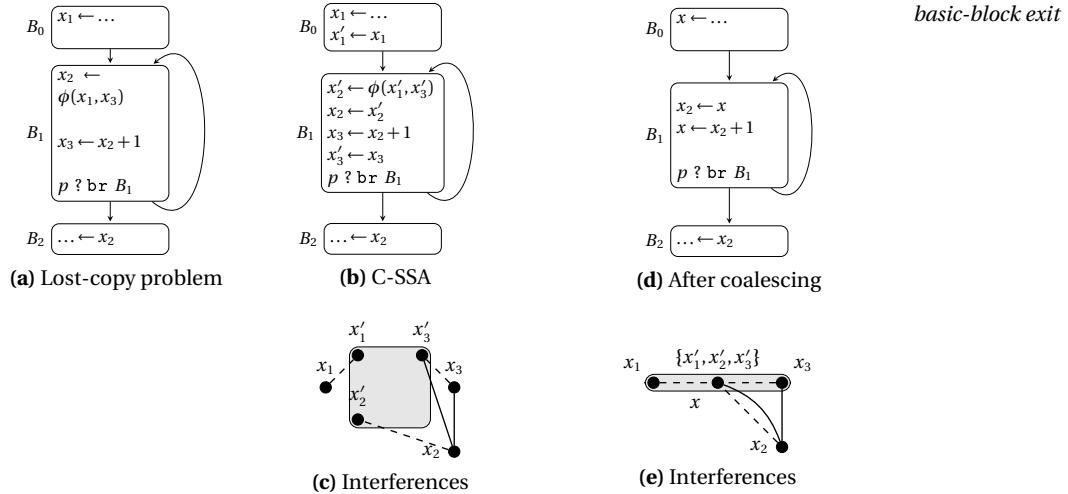


Fig. 22.5 SSA destruction for the lost-copy problem.

Liveness under SSA

If the goal is not to destruct SSA completely but remove as many copies as possible while maintaining the conventional property, liveness of ϕ -function operands should reproduce the behavior of the corresponding non-SSA code as if the variables of the ϕ -web were coalesced all together. The semantic of the ϕ -operator in the so called *multiplexing* mode fits the requirements. The corresponding interference graph on our example is depicted in Figure 22.5c.

Definition 4 (multiplexing mode). Let a ϕ -function $B_0 : a_0 = \phi(B_1 : a_1, \dots, B_n : a_n)$ be in *multiplexing* mode, then its liveness follows the following semantic: its definition-operand is considered to be at the entry of B_0 , in other words variable a_0 is live-in of B_0 ; its use-operands are at the exit^{*} of the corresponding predecessor basic-blocks, in other words variable a_i for $i > 0$ is live-out of basic-block B_i .

Value-based interference

As said earlier, after the ϕ -isolation phase and the treatment of operand pinning constraints, the code contains many overlapping live-ranges that carry the same value. Because of this, to be efficient coalescing must use an accurate notion of interference. As already mentioned in Chapter 2 the ultimate notion of interference contains two dynamic (i.e., related to the execution) notions: the notion of liveness and the notion of value. Analyzing statically if a variable is live at a given execution point or if two variables carry identical values is a difficult problem. The scope of variable coalescing is usually not so large, and graph coloring based

register allocation commonly take the following conservative test: *two variables interfere if one is live at a definition point of the other and this definition is not a copy between the two variables.*

One can notice that, with this conservative interference definition, when a and b are coalesced, the set of interferences of the new variable may be strictly smaller than the union of interferences of a and b . Thus, simply merging the two corresponding nodes in the interference graph is an over-approximation with respect to the interference definition. For example, in a block with two successive copies $b = a$ and $c = a$ where a is defined before, and b and c (and possibly a) are used after, it is considered that b and c interfere but that none of them interfere with a . However, after coalescing a and b , c should not interfere anymore with the coalesced variable. Hence the interference graph would have to be updated or rebuilt.

However, in SSA, each variable has, statically, a *unique* value, given by its unique definition. Furthermore, the “has-the-same-value” binary relation defined on variables is, if the SSA form fulfills the dominance property, an equivalence relation. The *value* of an equivalence class¹ is the variable whose definition dominates the definitions of all other variables in the class. Hence, using the same scheme as in SSA copy folding, finding the value of a variable can be done by a simple topological traversal of the dominance tree: when reaching an assignment of a variable b , if the instruction is a copy $b = a$, $V(b)$ is set to $V(a)$, otherwise $V(b)$ is set to b . The interference test is now both simple and accurate (no need to rebuild/update after a coalescing): if $\text{live}(x)$ denotes the set of program points where x is live,

$$a \text{ interfere with } b \text{ if } \text{live}(a) \text{ intersects } \text{live}(b) \text{ and } V(a) \neq V(b)$$

The first part reduces to $\text{def}(a) \in \text{live}(b)$ or $\text{def}(b) \in \text{live}(a)$ thanks to the dominance property. In the previous example, a , b , and c have the same value $V(c) = V(b) = V(a) = a$, thus they do not interfere.

Note that our notion of values is limited to the live ranges of SSA variables, as we consider that each ϕ -function defines a new variable. We could propagate information through a ϕ -function when its arguments are equivalent (same value). But, we would face the complexity of general value numbering. By comparison, our equality test in SSA comes for free.

Shared copies

It turns out that after the ϕ -isolation phase, and the treatment of operand pinning constraints the code also contains what we design as shared copies. A *shared copy* corresponds precisely to the previous example of two successive copies $b = a$ and $c = a$ i.e., the presence of two copies from the same source. We have seen that, thanks to our definition of value, the fact that b is live at the definition of c does not imply that b and c interfere. Suppose however that a (after

¹ Dominance property is required here, e.g., consider the following loop body if($i \neq 0$) { $b \leftarrow a;$ } $c \leftarrow \dots; \dots \leftarrow b; a \leftarrow c$; the interference between b and c is actual.

some other coalescing) interferes with b and c . Then, no coalescing can occur although coalescing b and c would save one copy, by “sharing” the copy of a . This sharing problem is difficult to model and optimize (the problem of placing copies is even worse), but we can optimize it a bit. We coalesce two variables b and c if they are both copies of the same variable a and if their live ranges intersect. This can be done in a second pass after all standard affinities have been treated. Note that if their live ranges are disjoint, such a coalescing may be incorrect as it would increase the live range of the dominating variable, possibly creating some interference not taken into account.

22.4 Speed and Memory Footprint

Implementing the technique of the previous section may be considered too costly. First, it inserts many instructions before realizing most are useless, and copy insertion is already by itself time-consuming. It introduces many new variables, too. The size of the variable universe has an impact on the liveness analysis and the interference graph construction. Also, if a general coalescing algorithm is used, a graph representation with adjacency lists (in addition to the bit matrix) and a working graph to explicitly merge nodes when coalescing variables, would be required. All these constructions, updates, manipulations are time-consuming and memory-consuming. We may improve the whole process by: a) avoiding the use of any interference graph and liveness sets; b) avoid the quadratic complexity of interference check between two sets of variables by an optimistic approach that first coalesces even interfering variables, then traverses each set of coalesced variables and un-coalesce one by one all the interfering ones; c) emulating (“virtualizing”) the introduction of the ϕ -related copies.

Interference check

Liveness sets and interference graph are the major source of memory usage. This motivates, in the context of JIT compilation, not to build any interference graph at all, and rely on the liveness check described in Chapter 9 to test if two live-ranges intersect or not. Let us suppose for this purpose that a “has-the-same-value” equivalence relation, is available thanks to a mapping V of variables to symbolic values:

$$\text{variables } a \text{ and } b \text{ have the same value} \Leftrightarrow V(a) = V(b)$$

As explained in Paragraph 22.3 this can be done linearly (without requiring any hash map-table) on a single traversal of the program if under strict SSA form. We also suppose that liveness check is available, meaning that for a given variable

a and program point p , one can answer if a is live at this point through the boolean value of $a.islive(p)$. This can directly be used, under strict SSA form, to check if two variables live-ranges, say a and b intersect:

$$\begin{aligned} \text{intersect}(a, b) &\Leftrightarrow \text{liverange}(a) \cap \text{liverange}(b) \neq \emptyset \\ &\Leftrightarrow \begin{cases} a.\text{def.op} = b.\text{def.op} \\ a.\text{def.op} \text{ dominates } b.\text{def.op} \wedge a.\text{islive}(\text{out}(b.\text{def.op})) \\ b.\text{def.op} \text{ dominates } a.\text{def.op} \wedge b.\text{islive}(\text{out}(a.\text{def.op})) \end{cases} \end{aligned}$$

Which leads to our refined notion of interference:

$$\text{interfere}(a, b) \Leftrightarrow \text{intersect}(a, b) \wedge V(a) \neq V(b)$$

De-coalescing in linear time

The interference checking outlined in the previous paragraph allows to avoid building an interference graph of the SSA form program. However, coalescing has the effect of merging vertices and interference queries are actually to be done between sets of vertices. To overcome this complexity issue, the technique proposed here is based on a de-coalescing scheme. The idea is to first merge all copy and ϕ -function related variables together. A merged-set might contain interfering variables at this point. The principle is to identify some variables that interfere with some other variables within the merged-set, and remove them (along with the one they are pinned with) from the merged-set. As we will see, thanks to the dominance property, this can be done linearly using a single traversal of the set.

In reference with register allocation, and graph coloring, we will associate the notion of colors to merged-sets: every variables of the same set are assigned the same color, and different sets are assigned different colors. The process of *de-coalescing* a variable is to extract it from its set; it is not put in another set, just isolated. We will say *un-colored*. Actually, variables pinned together have to stay together. We denote the (interference free) set of variables pinned to a common resource that contains variable v , $\text{atomic-merged-set}(v)$. So the process of un-coloring a variable might have the effect of un-coloring some others. In other words, a colored variable is to be coalesced with variables of the same color, and any un-colored variable v is to be coalesced only with the variables it is pinned with, $\text{atomic-merged-set}(v)$.

We suppose that variables have already been colored and the goal is to un-color some of them (preferably not all of them) so that each merged-set become interference free. We suppose that if two variables are pinned together they have been assigned the same color, and that a merged-set cannot contain variables pinned to different physical resources. Here we focus on a single merged-set and the goal is to make it interference free within a single traversal. The idea exploits

the tree shape of variables live-ranges under strict SSA. To this end, variables are identified by their definition point and ordered using dominance accordingly.

Algorithm 22.3 performs a traversal of this set along the dominance order, enforcing at each step the sub-set of already considered variables to be interference free. From now, we will abusively design as the dominators of a variable v , the set of variables of *color identical to* v which definition dominates the definition of v . Variables defined at the same program point are arbitrarily ordered, so as to use the standard definition of immediate dominator (denoted $v.\text{idom}$, set to \perp if not exists, updated lines 10-11). To illustrate the role of $v.\text{eanc}$ in Algorithm 22.3, let us consider the example of Figure 22.6 where all variables are assumed to be originally in the same merged-set: $v.\text{eanc}$ (updated line 15) represents the immediate intersecting dominator with the same value than v ; so we have $b.\text{eanc} = \perp$ and $d.\text{eanc} = a$. When line 15 is reached, cur_anc (if not \perp) represents a dominating variable interfering with v and with the same value than $v.\text{idom}$: when v is set to c ($c.\text{idom} = b$), as b does not intersect c and as $b.\text{eanc} = \perp$, $\text{cur_anc} = \perp$ which allows to conclude that there is no dominating variable that interfere with c ; when v is set to e , d does not intersect e but as a intersects and has the same value than d (otherwise a or d would have been un-colored), we have $d.\text{eanc} = a$ and thus $\text{cur_anc} = a$. This allows to detect on line 17 the interference of e with a .

Algorithm 22.3: De-coalescing of a merged-set

```

1 begin
2   |   cur_idom =  $\perp$ ;
3   |   foreach variable  $v$  of the merged-set in DFS pre-order of the dominance tree do
4   |     |   DeCoalesce( $v$ , cur_idom);
5   |     |   cur_idom  $\leftarrow v$ ;
6   |   end
7 end

8 Procedure DeCoalesce( $v$ ,  $u$ )
9 begin
10  |   while ( $u \neq \perp$ )  $\wedge (\neg(u.\text{idom} \text{ dominates } v) \vee \text{uncolored}(u))$  do  $u \leftarrow u.\text{idom}$ ;
11  |   |    $v.\text{idom} \leftarrow u$ ;
12  |   |    $v.\text{eanc} \leftarrow \perp$ ;
13  |   |    $\text{cur\_anc} \leftarrow v.\text{idom}$ ;
14  |   |   while  $\text{cur\_anc} \neq \perp$   $\wedge \text{cur\_anc} \neq \perp \wedge \neg(\text{colored}(\text{cur\_anc}) \wedge \text{intersect}(\text{cur\_anc}, v))$ 
15  |   |   |    $\text{cur\_anc} \leftarrow \text{cur\_anc.eanc}$ ;
16  |   |   |   if  $\text{cur\_anc} \neq \perp$  if  $V(\text{cur\_anc}) = V(v)$   $v.\text{eanc} \leftarrow \text{cur\_anc}$ ;
17  |   |   |   break;
18  |   |   |   else /*  $\text{cur\_anc}$  and  $v$  interfere */           /* if preferable to uncolor  $v$  uncolor
19  |   |   |   |   atomic-merged-set( $v$ );
20  |   |   |   break;
21  |   |   |   else uncolor atomic-merged-set( $\text{cur\_anc}$ );
22  |   |   |    $\text{cur\_anc} \leftarrow \text{cur\_anc.eanc}$ ;
23 end

```

basic-block early
basic-block late

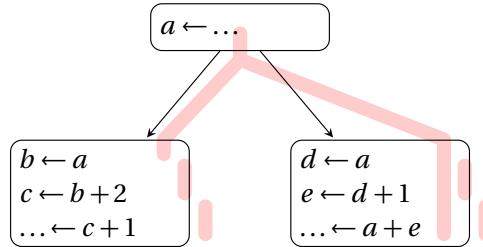


Fig. 22.6 Variables live-ranges are sub-trees of the dominator tree

Virtualizing ϕ -related copies

The last step toward a memory friendly and fast SSA-destruction algorithm consists in emulating the initial introduction of copies and only actually insert them on the fly when they appear to be required. We use *exactly the same algorithms as for the solution without virtualization*, and use a special location in the code, identified as a “virtual” parallel copy, where the real copies, if any, will be placed.

Because of this we consider a different semantic for ϕ -functions than the multiplexing mode previously defined. To this end we differentiate ϕ -operands for which a copy cannot be inserted (such as for the `br_dec` of Figure 22.2b) to the others. We use the term non-split and split operands introduced in Section 22.2. For a ϕ -function $B_0 : a_0 = \phi(B_1 : a_1, \dots, B_n : a_n)$ and a split operand $B_i : a_i$, we denote the program point where the corresponding copy would be inserted as the *early point* of B_0 ($\text{early}(B_0)$ – right after the ϕ -functions of B_0) for the definition-operand, and as the *late point* of B_i ($\text{late}(B_i)$ – just before the branching instruction) for a use-operand.

Definition 5 (copy mode). Let a ϕ -function $B_0 : a_0 = \phi(B_1 : a_1, \dots, B_n : a_n)$ be in *copy mode*, then the liveness for any split-operand follows the following semantic: its definition-operand is considered to be at the early point^{*} of B_0 , in other words variable a_0 is *not* live-in of B_0 ; its use-operands are at the late point^{*} of the corresponding predecessor basic-blocks, in other words variable a_i for $i > 0$ is (unless used further) *not* live-out of basic-block B_i . The liveness for non-split operands follows the multiplexing mode semantic.

When the algorithm decides that a virtual copy $a'_i \leftarrow a_i$ (resp. $a_0 \leftarrow a'_0$) cannot be coalesced, it is *materialized* in the parallel copy and a'_i (resp. a'_0) becomes explicit in its merged-set. The corresponding ϕ -operator is replaced and the use of a'_i (resp. def of a'_0) is now assumed, as in the multiplexing mode, to be on the corresponding control-flow edge. This way, only copies that the first approach would finally leave un-coalesced are introduced. We choose to postpone the materialization of all copies along a single traversal of the program at the really end of the de-coalescing process. Because of the virtualization of ϕ -related copies, the de-coalescing scheme given by Algorithm 22.3 have to be adapted

to emulate live-ranges of split-operands. The pseudo-code for processing a local virtual variable is given by Algorithm 22.5. The trick is to use the ϕ -function itself as a placeholder for its set of local virtual variables. As the live-range of a local virtual variable is very small, the cases to consider is quite limited: a local virtual variable can interfere with another virtual variable (lines 3-4) or with a “real” variable (lines 5-12).

Algorithm 22.4: De-coalescing with virtualization of ϕ -related copies

```

1 foreach  $c \in \text{COLORS}$  do  $c.\text{cur\_idom} = \perp$ ;
2 foreach basic-block  $B$  in CFG in DFS pre-order of the dominance tree do
3   foreach program point  $l$  of  $B$  in topological order do
4     if  $l = \text{late}(B)$  foreach  $c \in \text{COLORS}$  do  $c.curphi = \perp$ 
5     foreach basic-block  $B'$  successor of  $B$  do
6       foreach operation  $\Phi$ : “ $B': a_0 = \phi(\dots, B : v, \dots)$ ” in  $B'$  do
7         if  $\neg\text{colored}(\Phi)$  then continue else  $c \leftarrow \text{color}(\Phi)$ ;
8          $\text{DeCoalesce\_virtual}(\Phi, B : v, c.curphi, c.cur\_idom);$ 
9         if  $\text{colored}(\Phi)$  then  $c.curphi \leftarrow \Phi$ ;
10        end
11      end
12    else foreach operation  $OP$  at  $l$  (including  $\phi$ -functions) do
13      foreach variable  $v$  defined by  $OP$  do
14        if  $\neg\text{colored}(v)$  then continue else  $c \leftarrow \text{color}(v)$ ;
15         $\text{DeCoalesce}(v, c.cur\_idom);$ 
16        if  $\text{colored}(v)$  then  $c.cur\_idom \leftarrow v$ ;
17      end
18    end
19  end
20 end

```

The overall scheme works as follow: (1) every copy related (even virtual) variables are first coalesced (unless pinning to physical resources forbid it); (2) then merged-sets (identified by their associated color) are traversed and interfering variables are de-coalesced; (3) finally, materialization of remaining virtual copies is performed through a single traversal of all ϕ -functions of the program: whenever one of the two operands of a virtual copy is uncolored, or whenever the colors are different, a copy is inserted.

A key implementation aspect is related to the handling of pinning. In particular, for correctness purpose, coalescing is performed in two separated steps. First pin- ϕ -webs have to be coalesced. Detection and fix of strong interferences is handled at this point. The so obtained merged sets (that contain local virtual variables) have to be identified as atomic i.e., they cannot be separated. After the second step of coalescing, atomic merged sets will compose larger merged sets. A variable cannot be de-coalesced from a set without de-coalescing its atomic merged-set from the set also. Non singletons atomic merged-sets have to be represented somehow. For ϕ -functions, the trick is again to use the ϕ -function itself as a placeholder for its set of local virtual variables: the pinning of the virtual

Algorithm 22.5: Process (de-coalescing) a virtual variable

```

1 Procedure DeCoalesce_virtual( $\Phi_i$ ,  $B : v$ ,  $\Phi'_i$ ,  $u$ );
2 begin
3   if /* Interference */  $\Phi'_i \neq \perp \wedge V(\text{operand\_from\_B}(\Phi'_i)) \neq V(a')$  uncolor
4     atomic-merged-set( $\Phi'_i$ );
5   return;
6   while ( $u \neq \perp \wedge (\neg(u \text{ dominates } B) \vee \neg \text{colored}(u))$ )  $u \leftarrow u.\text{idom}$ ;
7    $v.\text{idom} \leftarrow u$ ;
8    $v.eanc \leftarrow \perp$ ;  $cur\_anc \leftarrow u$ ;
9   while  $cur\_anc \neq \perp$  while  $cur\_anc \neq \perp \wedge \neg(\text{colored}(cur\_anc) \wedge cur\_anc.\text{islive}(\text{out}(B)))$ 
10     $cur\_anc \leftarrow cur\_anc.eanc$ ;
11    if /* interference */  $cur\_anc \neq \perp$  if preferable to uncolor  $\Phi_i$  uncolor
12      atomic-merged-set( $\Phi'_i$ );
13    break;
14    else uncolor atomic-merged-set( $cur\_anc$ );
15     $cur\_anc \leftarrow cur\_anc.eanc$ ;
16 end

```

variables is represented through the pinning of the corresponding ϕ -function. As a consequence, any ϕ -function will be pinned to all its non-split operands.

Without any virtualization, the process of transforming operand pinning into live-range pinning also introduces copies and new local variables pinned together. This systematic copy insertion can also be avoided and managed lazily just as for ϕ -nodes isolation. We do not address this aspect of the virtualization here: to simplify we consider any operand pinning to be either ignored (handled by register allocation) or expressed as live-range pinning.

Sequentialization of parallel copies

During the whole algorithm, we treat the copies placed at a given program point as *parallel copies*, which are indeed the semantics of ϕ -functions. This gives several benefits: a simpler implementation, in particular for defining and updating liveness sets, a more symmetric implementation, and fewer constraints for the coalescer. However, at the end of the process, we need to go back to standard code, i.e., write the final copies in some sequential order.

As explained in Section 3.2 (Algorithm 3.6) the sequentialization of a parallel copy can be done using a simple traversal of a windmill farm shaped graph from the tip of the blades to the wheels. Algorithm 22.6 emulates a traversal of this graph (without building it), allowing to overwrite a variable as soon as it is saved in some other variable.

When a variable a is copied in a variable b , the algorithm remembers b as the last location where the initial value of a is available. This information is stored into $\text{loc}(a)$. The initial value that must be copied into b is stored in $\text{pred}(b)$. The initialization consists in identifying the variables whose values are not needed

(tree leaves), which are stored in the list `ready`. The list `to_do` contains the destination of all copies to be treated. Copies are first treated by considering leaves (while loop on the list `ready`). Then, the `to_do` list is considered, ignoring copies that have already been treated, possibly breaking a circuit with no duplication, thanks to an extra copy into the fresh variable n .

Algorithm 22.6: Parallel copy sequentialization algorithm

Data: Set P of parallel copies of the form $a \mapsto b$, $a \neq b$, one extra fresh variable n
Output: List of copies in sequential order

```

1 ready ← [] ; to_do ← [] ; pred( $n$ ) ← ⊥ ;
2 forall the  $(a \mapsto b) \in P$  do
3   | loc( $b$ ) ← ⊥ ; pred( $a$ ) ← ⊥ ;                                /* initialization */
4 end
5 forall the  $(a \mapsto b) \in P$  do
6   | loc( $a$ ) ←  $a$  ;                                         /* needed and not copied yet */
7   | pred( $b$ ) ←  $a$  ;                                         /* (unique) predecessor */
8   | to_do.push( $b$ ) ;                                         /* copy into  $b$  to be done */
9 end
10 forall the  $(a \mapsto b) \in P$  do
11   | if loc( $b$ ) = ⊥ then ready.push( $b$ ) ; /*  $b$  is not used and can be overwritten
12   | */,
12 end
13 while to_do ≠ [] while ready ≠ []  $b$  ← ready.pop() ;           /* pick a free location */
14   |  $a$  ← pred( $b$ ) ;  $c$  ← loc( $a$ ) ;                         /* available in  $c$  */
15   | emit_copy( $c \mapsto b$ ) ;                                     /* generate the copy */
16   | loc( $a$ ) ←  $b$  ;                                         /* now, available in  $b$  */
17   | if  $a = c$  and pred( $a$ ) ≠ ⊥ then ready.push( $a$ ) ; /* just copied, can be overwritten
18   | */,
18   |  $b$  ← to_do.pop() ;                                       /* look for remaining copy */
19   | if  $b = loc(pred(b))$                                      /* break circuit with copy */
20   | emit_copy( $b \mapsto n$ ) ;                                 /* now, available in  $n$  */
21   | loc( $b$ ) ←  $n$  ;                                         /*  $b$  can be overwritten */
22 ready.push( $b$ ) ;
```

.....

22.5 Further Readings

SSA destruction was first addressed by Cytron et al. [92] who propose to simply replace each ϕ -function by copies in the predecessor basic-block. Although this naive translation seems, at first sight, correct, Briggs et al. [53] pointed subtle errors due to parallel copies and/or critical edges in the control-flow graph. Two typical situations are identified, namely the “lost copy problem” and the “swap problem”. The first solution, both simple and correct, was proposed by Sreedhar et al. [256]. They address the associated problem of coalescing and describe three solutions. The first one, consists in three steps: a) translate SSA into

CSSA, by isolating ϕ -functions; b) eliminate redundant copies; c) eliminate ϕ -functions and leave CSSA. The third solution that turns out to be nothing else than the first solution that would virtualizes the isolation of ϕ -functions shows to introduce less copies. The reason for that, identified by Boissinot et al., is due to the fact that in the presence of many copies the code contains many intersecting variables that do not actually interfere. Boissinot et al. [40] revisited Sreedhar et al.'s approach in the light of this remark and proposed the value based interference described in this chapter.

The ultimate notion of interference was discussed by Chaitin et al. [67] in the context of register allocation. They proposed a simple conservative test: *two variables interfere if one is live at a definition point of the other and this definition is not a copy between the two variables*. This interference notion is the most commonly used, see for example how the interference graph is computed in [17]. Still they noticed that, with this conservative interference definition, after coalescing some variables the interference graph has to be updated or rebuilt. A counting mechanism to update the interference graph was proposed, but it was considered to be too space consuming. Recomputing it from time to time was preferred [67, 66].

The value based technique described here can also obviously be used in the context of register allocation even if the code is not under SSA form. The notion of value may be approximated using data-flow analysis on specific lattices [12] and under SSA form simple global value numbering [240] can be used.

Leung and George [182] addressed SSA destruction for machine code. Register renaming constraints, such as calling conventions or dedicated registers, are treated with pinned variables. Simple data-flow analysis scheme is used to place repairing copies. By revisiting this approach to address the coalescing of copies Rastello et al. [232] pointed out and fixed a few errors present in the original algorithm. While being very efficient in minimizing the introduced copies, this algorithm is quite complicated to implement and not suited for just in time compilation.

The first technique to address speed and memory footprint was proposed by Budimlić et al. [60]. It proposes the de-coalescing technique, revisited in this chapter, that exploits the underlying tree structure of dominance relation between variables of the same merged-set.

Last, this chapter describes a fast sequentialization algorithm that requires the minimum number of copies. A similar algorithm has already been proposed by C. May [198].

CHAPTER 23

Register Allocation

*F. Bouchez
S. Hack*

Progress: 40%



Material gathering in progress

Register allocation maps the variables of a program to physical memory locations. The compiler determines the location for each variable and each program point. Ideally, as many operations as possible should draw their operands from processor registers without loading them from memory beforehand. Due to the huge latency of the memory hierarchy (even loading from L1 cache takes three to ten times longer than accessing a register), register allocation is one of the most important optimizations in a compiler.

There is only a small number of registers available in a CPU, with usual values ranging from 8 to 128. Hence, the task of register allocation is not only assigning the variables to registers but also deciding which variables should be evicted from registers and when to store and load them from memory (spilling). Furthermore, register allocation has to remove spurious copy operations (copy coalescing) inserted by previous phases in the compilation process, and to deal with allocation restrictions that the instruction set architecture and the runtime system impose (register targeting).

23.1 Introduction

Register allocation is usually performed per procedure. In each procedure, a liveness analysis determines for each variable the program points where the variable is live.

A variable is live at a program point if there exists a path from this point on which the variable will be used later on and is not overwritten before that use.

Hence, storage needs to be allocated for that variable; ideally a register. The set of all program points where a variable is live is called the *live-range* of the variable. The resource conflict of two variables is called *interference* and is usually defined via liveness: Two variables interfere if (and only if) there exists a program point where they are simultaneously live.¹ The number of live variables at a program point is called the *register pressure* at that program point. The maximum register pressure over all program points in a procedure is called the register pressure of that procedure, or “Maxlive.”

It is helpful to think of interferences as an undirected *interference graph*: The nodes are the variables of the program, and two nodes are connected if they interfere. The set of variables live at some program point form a *clique* in this graph: Any two variables in this set interfere, hence their nodes are all mutually connected. The size of the largest clique in the interference graph is called its *clique number*, denoted ω . A *coloring* of this graph with k colors corresponds to a valid register allocation with k registers. Hence, we will use the terms “register” and “color” interchangeably in this chapter. A k -coloring is a mapping from the nodes of the graph to the first k natural numbers (called colors) such that two neighbouring nodes have different colors. The smallest k for which a coloring exists is called the *chromatic number*, denoted χ . Of course $\chi \geq \omega$ holds for every graph, because the nodes of a clique must all have different colors.

Chaitin et al. [65] showed that for every undirected graph, there is a program whose interference graph is this graph. From the NP-completeness of graph coloring then follows the NP-completeness of register allocation. Hence, determining the chromatic number of an interference graph is not possible in polynomial time unless P=NP.

This is the major nuisance of classical register allocation: The compiler cannot efficiently determine how many registers are needed. This means one might need more registers than Maxlive, Maxlive being only a lower bound on the chromatic number. The somewhat unintuitive result is that, even if at every program point there are no more than Maxlive variables live, we still might need more than Maxlive registers for a correct register allocation! The causes of this problem are control-flow merges, as can be seen in Figure 23.1.

In that example, the register pressure is at most two at every program point. However, the interference graph cannot be colored with two colors. Its chromatic number is three. The inequality between Maxlive and the chromatic number is caused by the cycle in the interference graph.²

This situation changes if we permit *live-range splitting*. That is, inserting a copy (move) instruction at a program point that creates a new variable. Thus, the value of the variables is allowed to reside in different registers at different

¹ This definition of interference by liveness is an overapproximation. There are refined definitions that create less interferences. However, in this chapter we will restrict ourselves to this definition.

² In theory, the gap between the chromatic number and Maxlive can even be arbitrarily large. This can be shown using by applying Chaitin’s proof of the NP-completeness of register allocation to Mycielski graphs.

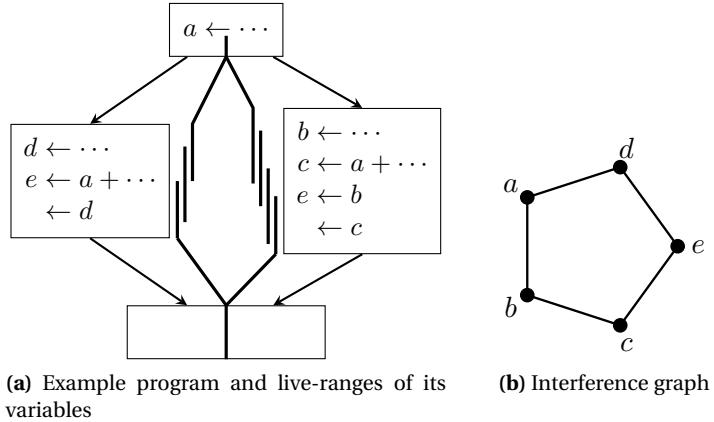


Fig. 23.1 Example program and its interference graph

times. In Figure 23.1a, assume we split the live-range of e in the left block: we rename it by e' and add a copy $e \leftarrow e'$ at the end of the block. Then, the node e in the graph is split into two nodes, e and e' , that do not interfere: this breaks the cycle in the graph, making its chromatic number equal to two, i.e., Maxlive.

In an extreme setting, one could split the live-ranges of all variables after every instruction. The interference graph degenerates into several disconnected components and its chromatic number drops to Maxlive. Of course, such an extreme live-range splitting introduces *a lot* of shuffle code, which degrades the runtime of the compiled program. This interplay between live-range splitting and colorability is the key issue in register allocation.

23.1.1 Non-SSA Register Allocators

As finding a valid k -coloring is NP-complete, assigning registers is performed by some heuristic algorithm. If that algorithm fails to assign a register to a variable, then it either spills this variable to memory, or frees a register by spilling the variable it contains. The problem here is that the spilling decision is made to revive the coloring heuristic, and not because the variable that gets spilled is in itself a good candidate for spilling. Even worse, we might spill a variable because the heuristic is “not good enough,” and not because we are actually out of registers.

At the same time, classical algorithms perform copy coalescing (i.e., undoing live-range splitting) during coloring. However, if done aggressively, coalescing may increase the chromatic number of the graph, or make the graph harder to color for the heuristic. Both cases generate additional spills. This is of course unacceptable because we never want to add memory accesses in favor of an

eliminated copy. Thus, existing techniques often apply *conservative* coalescing approaches which are guaranteed not to increase the chromatic number of the graph, at the expense of the quality of the coalescing.

23.1.2 SSA form to the rescue of register allocation

The live-ranges in an SSA form program all have a certain property: SSA requires that all uses of a variable are dominated by its definition. Hence, the whole live-range is dominated by the definition of the variable. Dominance, however, induces a tree on the control-flow graph. Thus, the live-ranges of SSA variables are all tree-shaped [44, 55, 137]. They can branch downwards on the dominance tree but have a single root: the program point where the variable is defined. Hence a situation like in Figure 23.1 can no longer occur: e had two “roots” because it was defined twice. Under SSA form, the live-range of e is split by a ϕ -function. The argument and result variables of that ϕ -function constitute new live-ranges, giving more freedom to the register allocator since they can be assigned to different registers.

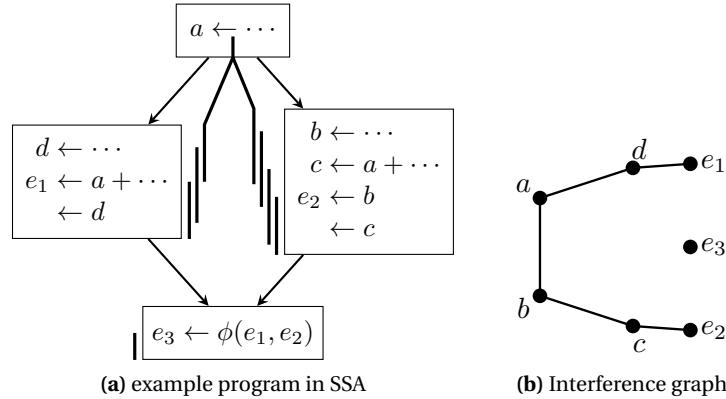


Fig. 23.2 SSA version of the example program

This special structure of the live-ranges leads to a special class of interference graphs: Gavril [126] showed that the intersection graphs of subtrees are the *chordal graphs*. Chordal graphs can be optimally colored in linear time with respect to the number of edges in the graph. Furthermore, they are *perfect*, which means that for every subgraph, the clique number equals the chromatic number. This is a very important property for register allocation because

it means that Maxlive equals the chromatic number of the graph;³ It has even been shown [44, 137] that, for every clique in the interference graph, there is one program point where all the variables of that clique are live. This allows to decouple spilling from coloring: First, lower the register pressure to k everywhere in the program; Then, color the interference graph with k colors in polynomial time.

23.1.3 Maxlive and colorability of graphs

TODO: Florent: This part does not seem to be useful, but double-check my words !

TODO: seb: do we need this? doing a standard liveness analysis does the job. we can also refer to Appel's liveness algo.

- *definition of Maxlive*
- *show that Maxlive is the minimum number of colors required*
- *scan of the dominance tree sufficient to know Maxlive: pseudo-code*
- *Fab: MAXLIVE minimum if interference graph==intersection graph. Clearly we do not want to go into such details but be careful to use precise and correct terms.*

The register pressure at a program point is the number of variables alive at this point.⁴ The greatest register pressure over all a program is called Maxlive, since it is the maximum number of simultaneously alive variables. Obviously, spilling is mandatory if Maxlive is strictly greater than R , the number of registers. Under SSA, the good news is that this is the only case where spilling is required, i.e., if $\text{Maxlive} \leq R$, there is no need for spilling. Indeed, Maxlive is the coloring number of the chordal interference graph of the program (**TODO: ref to chapter 2?**). So, we can devise a polynomial test to know whether spilling is necessary or not by computing Maxlive. This can be done by checking on every program point the number of variables that are alive. If liveness information is available on the whole program, the order does not matter. However, in the case of SSA, this is not required as Maxlive can easily be computed without liveness information by performing a scan of the dominance tree, starting from the root. We only need to know which uses are “last-uses.” The pseudo-code is given on Figure 23.3. **TODO: check Maxlive scan needs last uses.**

³ This also implies that chordal graphs do not have “holes” like the interference graph in Figure 23.1. Every cycle is spanned by *chords*.

⁴ Considering all simultaneously alive variables interfere.

```

function compute-branch (p, live)
    live <- live - { uses in p that are last uses }
    live <- live + { definitions of p }
    maxlive <- max( maxlive, #live )
    foreach child c of p
        compute-branch (c, live)

function compute-maxlive
    maxlive <- 0
    compute-branch (root, {})

```

Fig. 23.3 Pseudo-code for Maxlive

23.2 Spilling

The term spilling is heavily overloaded: In the graph-coloring setting it often means to spill a node of the interference graph and thus the *entire* live-range of a variable. This implies inserting load instructions in front of every use and store instructions after each definition of the (non-SSA) variables. Other approaches, like some linear scan allocators, are able to keep parts of the live-range of a variable in a register such that multiple uses can potentially reuse the same reload.

23.2.1 Spilling under the SSA Form

As stated above, lowering Maxlive to R ensures that a register allocation with R registers can be found in polynomial time for SSA programs. Thus, spilling should take place before registers are assigned *and* yield a program in SSA form. For the moment, let us ignore *how* to choose “what” to spill and “where,” and concentrate on the implications of spilling in SSA programs. Because many compilers use SSA heavily, it is a reasonable assumption that the program was in SSA before spilling. Hence, we consider spilling as an SSA program transformation that establishes $\text{Maxlive} \leq R$ by inserting loads and stores into the program.

23.2.2 Finding Spilling Candidates

We need to find program points to place loads and stores such that $\text{Maxlive} \leq R$. Let us start with a simplistic setting where the location of loads and stores is not optimized: loads are placed directly in front of uses and stores directly after the definition. Furthermore we consider a *spill everywhere* setting: a variable is either spilled completely or not at all; If spilled, the live-range of a variable then

degenerates into small intervals: one from the definition and the store, and one from each load to its subsequent use. **TODO: figure? Flo: if space allows it, yes.** However, even in this simplistic setting, it is NP-complete to find the minimum number of nodes to establish $\text{Maxlive} \leq R$. It is even NP-complete to find the minimum number of nodes to spill to decrease Maxlive just by one! Finding spill candidates is not facilitated by SSA.

Moreover, formulations like *spill everywhere* are often not appropriate for practical purposes, and putting the whole live-range to memory is too simplistic. A variable might be spilled because at some program point the pressure is too high; However, if that same variable is later used in a loop where the register pressure is low, a spill everywhere will place a superfluous (and costly!) load in that loop. Spill everywhere approaches try to minimize this behavior by adding costs to variables, to make the spilling of such a variable less likely. These bring in a flow-insensitive information that a variable reside an frequently executed area. However, such approximations are often too coarse to give good performance. Hence, it is imperative to cleverly split the live-range of the variable according to the *program structure* and spill only parts of it.

In the context of a basic block, a simple algorithm that works well is the “furthest first” algorithm. The idea is to scan the block from top to bottom, whenever where the register pressure is too high, we spill the variable whose next use is the furthest away; And we spill it only *up to this next use*. Spilling this variable frees a register for the longest time, hence diminishing the chances to have to spill other variables later. This algorithm is not optimal since the first time a variable is spilled, a load and a store are added; For subsequent spilling of the same variable (but on a different part of its live-range), only a load must be added. Hence this algorithm may produce more stores than necessary (e.g., by spilling two variables while spilling an other part of the first one would have been enough), but it gives good results on “straight-line code,” i.e., basic blocks.

We will present in the next section an algorithm that extends this “furthest first” algorithm to general control-flow graphs. The idea is to use a data-flow analysis to approximately compute the next use values on the control-flow graph. In this next use calculation, loops are virtually unrolled by a large factor such that code “behind” loops appears far away from code in front of the loop. This forces the spilling algorithm to prefer variables that are not used in loops and to leave in registers the parts of live-ranges that are used inside loops. We then have to find how and in which order to process the basic blocks of the control-flow graph in order to determine a reasonable initial register occupation for each block.

23.2.3 Spilling under SSA

Still, we give a possibility to perform spilling under SSA

greedy coloring scheme

- Hack's algorithm: simplified version?
- à la Belady following dominance tree?
- Fab: The last Algorithm from Sebastian is Belady like. So yes try to simplify it, provide the pseudo code and give the intuitions on how the full version improves it.
- Seb: I think it's ok to briefly outline the approaches that are out there and refer to the papers.

TODO: Ask Sebastian to do it.

[?] **TODO:** file braun09cc.pdf in biblio

23.3 Coloring and coalescing

Due to the decoupled register allocation and the use of SSA, the interference graph after spilling has two interesting properties. First, it still is chordal, hence easily colorable in linear time; Second, we know that R colors are sufficient to color it since $\text{Maxlive} \leq R$. In this section, we will first present the traditional graph coloring heuristic of Chaitin et al. and show how it successfully color programs under SSA form. We will then present a lighter coloring algorithm based on the scanning on the dominance tree. Finally, we will explain how to add *coalescing* to the coloring phase to reduce the number of copies in a program.

23.3.1 Greedy coloring scheme

In traditional graph coloring register allocation algorithms, the assignment of registers to variables is done by coloring the interference graph using the greedy heuristic of Chaitin et al. This scheme is based on the observation that given R colors—representing the registers—, if a node in the graph has at most $R - 1$ neighbours, there will always be one color available for this node whatever colors the remaining nodes have. Such a node can be *simplified*, i.e., removed from the graph and placed on a stack. This process can be iterated with the remaining nodes, whose degree may have decreased (if the simplified node was one of their neighbours). If the graph becomes empty, we know it is possible to color the graph with R colors; In particular, a valid R -coloring can be obtained by assigning colors to nodes in the reverse order of their simplification, i.e., popping nodes from the stack and assigning them one available color. This is always possible since they have at most $R - 1$ colored neighbours. We call this algorithm the *greedy coloring scheme*.^{*}

The greedy scheme is a coloring heuristic, and as such, it can get stuck; it happens whenever all remaining nodes have degree at least R . In that case, we do not know whether the graph is R -colorable or not. In traditional register allo-

cation, this is the trigger for spilling some variables so as to unstuck the simplification process. However, we will see that under the SSA form, if spilling has already been done so that the maximum register pressure is at most R , the greedy coloring scheme can never get stuck.

```

Function Simplify(G)
  Data: Undirected graph G = (V, E);
  For all v, degree[v] = #neighbours of v in G,
    k number of colors
  stack = {} ;
  worklist = {v in V | degree[v] < k} ;
  while worklist != {} do
    let v in worklist ;
    foreach w neighbour of v do
      degree[w] = degree[w]-1 ;
      if degree[w] = k - 1 then worklist = worklist U {w}
    push v on stack ;
    worklist = worklist \ {v} ; /* Remove v from G */
  if V != {} then Failure "The graph is not simplifiable"
  return stack ;

```

Fig. 23.4 Iskgreedy function

23.3.2 SSA graphs are colorable with a greedy scheme

One of the interesting properties of SSA form for register allocation is that the live-ranges of variables are subtrees of the dominance tree (see Chapter 2, Section 2.4). The interference graph is hence the intersection graph of a family of subtrees. In the graph theory, this means this graph is chordal and as such colorable in polynomial time, but it is in fact quite easy to understand why without resorting to this theory.

Consider the greedy coloring scheme described in the previous section, and picture a “tree” representation in your head. At the end of each dangling branch there is a “leaf” variable, that is defined last in this branch. We can visually see that this variable will not have many intersecting variables; it is in fact a candidate for simplification, and once removed another variable will become the new leaf. We see that simplification can always happen at the end of the branches of the dominance tree, moving upwards until the whole tree is simplified.

We will now prove more thoroughly that if the register pressure is at most R , the interference graph can be colored with R colors using the greedy scheme.

Proof. Let us consider such a subtree representation of the interference graph (see Figure 23.5 TODO) and try to apply the greedy coloring scheme to it. In this representation, a node candidate for simplification corresponds to a live-range

which intersect with at most $R - 1$ other live-ranges. Consider the “leaf” of a branch of the dominance tree, i.e., the live-range on this branch which starts last, say at point d (the instruction that defines the leaf). If this leaf does not intersect any other live-range, it can be simplified. If it does, then all its neighbour live-ranges start before d —by definition of the leaf—and end after d —because they intersect. So, all neighbours of the leaf are live at d ; Since spilling has already been done, we know that the register pressure at d is at most R , hence there is at most $R - 1$ intersecting live-ranges: the corresponding node in the interference graph has at most $R - 1$ neighbours and can be simplified. Removing a live-range does not change the shape of the graph (the remaining live-ranges are still subtrees of the dominance tree): as long as the graph is not empty, there will still be “leaf” live-ranges, i.e., nodes that can be simplified. Eventually, the simplification process will empty the whole graph, and colors can be assigned in the reverse order of the simplification.

Insert subtree figure HERE.

Fig. 23.5 Subtree

We just proved that if we are under SSA form and the spilling has already been done so that $\text{Maxlive} \leq R$, the classical greedy coloring scheme of Chaitin et al. is *guaranteed* to perform register allocation with R colors without any additional spilling. This is practical for instance when working on an existing compiler that uses a classical register allocation algorithm, for instance, the Iterated Register Coalescing.

It is also possible to assign colors directly after the greediness has been checked during the spilling phase (see Section ??). Indeed, the function `Is_k_Greedy` maintains a stack of simplified variables. This stack can be directly used for coloring by making a call to `Assign_colors` given in Figure 23.6.

Among the good things, this also means that existing register allocators can be easily modified to handle SSA code and that register allocation can benefit from powerful coalescing strategies such as aggressive or conservative ones. Moreover, the fact that register allocation can be decoupled into a phase of spilling first and then a phase of coloring/coalescing allows the writing of more involved coalescing strategies, as we have reduced the interplay (hence complexity) between those two phases.

One disadvantage remains compared to lighter algorithms, we still use the interference graph, a structure that takes time to build and uses memory up. By building on our current understanding of why the greedy scheme works for SSA graphs, we will show next how to design a lighter coloring algorithm.

```

Function Assign_colors(G)
    available = new array of size R and values True
    while stack != {} do
        v = pop stack ;
        for each neighbour w in G
            available[color(w)] = False
        col = 0;
        for each color c from 1 to R
            if available[c]
                col = c
        available[c] = True /* prepare for next round */
        color(v) = col
        add v to G

```

Fig. 23.6 Assign_color function

23.3.3 A light tree-scan coloring algorithm

- On the dominance tree, possible to scan from top and assign colors to variables as they come.
- Biased coloring is used to coalesce variables linked by ϕ -functions.
- see next section for more involved coalescing
- Fab: I think that you already talked about scan coloring. Ok to talk about biased.
- Fab: You can talk about biased coloring that uses the result of an aggressive coalescing (see with Quentin). It is not more complicated and it improves results.

One of the advantages of SSA is to make things simpler, faster, and use less memory. For the register assignment problem, we now know that the existing greedy scheme based on simplification still works; However, it requires constructing and maintaining the interference graph, a structure judged to big and cumbersome to be used in a JIT context, where compilation time and memory prints matter more. We will now show how to perform a fast register assignment for SSA after spilling that does not need more than the dominance tree and def-use chains.

The algorithm we propose scans the dominance tree, coloring the variables from the root to the leaves in a top-down order. This means the variables are simply colored in the order of their definitions, first come first served! We give the pseudo-code in procedure Tree_Scan, Figure 23.7; It works because it mimics the effect of the greedy scheme (“leaves” are simplified first hence colored *last*), without actually performing any simplification.

Intuitively, when the scanning arrives at the definition of a variable, the only colored variables are “above” it and since there is at most $R - 1$ other variables live at the definition, there is always a free color. We will now prove more formally that this method always work; The key observation is that the order in which variables are considered corresponds to a reverse order of simplifica-

tion, i.e., a valid order of the nodes on the stack after the greedy simplification scheme.

Proof. Let us consider an ordering v_1, v_2, \dots, v_n of the nodes of the nodes based on the dominance: if v_i dominates v_j , then v_j appear before v_i in the ordering ($j < i$). We will show that this order correspond to a greedy simplification. Suppose v_1, \dots, v_{i-1} have been simplified by the greedy scheme (for $i = 1$, this is the initial graph). If v_i dominates any variable, say v_j , then $j < i$ and the node has already been simplified; So, there is no variable defined after v_i on the dominance tree: the live-range of v_i is a “leaf” of the subtree representation (see proof in Section 23.3.2). So, v_i can be chosen for simplification, and by induction the ordering v_1, \dots, v_n is a valid order of simplification. As we have seen previously in Section 23.3.1, this means the reverse order v_n, v_{n-1}, \dots, v_1 is a valid order for coloring.

It is important to understand why this method does not work in the general non-SSA case. Under SSA, the variables are “split” at join points by the use of ϕ -functions. If this was not the case, we would have live-ranges that spans on multiple branches of the dominance tree, creating cycles in the representation. In that case, coloring a branch would constrain the coloring on leaves on a different branch; Under SSA form, such live-ranges are split and their colorings are independent. See Figure ?? for an example.

```

Tree_Scan(T)

function assign_color(p, available)
    for each v last use at p
        available[color(v)] = True      /* colors not used anymore */

    for each v defined at p
        c = choose_color(v, available)  /* choose available color */
        available[c] = False
        color(v) = c

    for each child p' of p
        assign_color(p', available)

assign_color(root(T), [True, True, ..., True])

```

Fig. 23.7 Tree scan coloring algorithm for SSA.

The tree-scan coloring algorithm is really fast as it only needs one traversal of the dominance tree. Since the number of colors is fixed and small, it can be implemented as a bit-set to speed-up updates of the `available` array. The pseudo-code for function `choose_color` is deliberately not given yet. A very basic implementation could just scan the `available` array until it finds one that is not

taken. We will see in the next section how to bias the choosing of colors so as to perform some coalescing.

23.3.4 Coalescing under SSA form

The goal of *coalescing* is to minimize the number of register-to-register move instructions in the final code. While there may not be so many such “copies” at the high-level (e.g., instruction “ $a = b$;” in C)—especially after a phase of copy propagation under SSA ([TODO: See chapter X](#))—, many such instructions are added by different compiler phases by the time compilation reaches the register allocation phase. For instance, adding copies is a common way to deal with register constraints (see Section 23.4). An even more obvious and unavoidable reason in our case is the presence of ϕ -functions due to the SSA form.

A ϕ -function represents in fact parallel copies on incoming edges of basic blocks. For instance, $a \leftarrow \phi(b, c)$ means that instruction $a \leftarrow b$ should be executed on the edge coming from the left, and $a \leftarrow c$ on the one from the right. Suppose now that register allocation decided to put a and b in register R_1 , but c in register R_2 ; Then, the copy $a \leftarrow b$ does not need to be executed anymore, but $a \leftarrow c$ still does: the value of variable c will be contained in R_2 and needs to transferred to R_1 , which means the final code will contain an instruction “`mov R1, R2`” or the like.

In the case of SSA, it is obviously better to assign variables linked by a ϕ -function, to the same register, so as to remove copies between subscripts of the same variable (a_1, a_2, \dots). This is possible if we are under conventional SSA (CSSA, see Chapter ??), but optimizations can break this property and interferences can appear between subscripted variables with the same origin. Such two variables cannot be assigned to the same register without breaking the program.

To bypass this problem and also allow for more involved coalescing between SSA variables of different origins, we define a notion of *affinity*, acting as the converse of the relation of interference and expressing how much two variables “want” to share the same register. By adding a metric to this notion, it measures the benefit one could get if the two variables were assigned to the same register: the weight represents how much instructions we would save at execution if the two variables share the same register. For instance, given two variables a and b , the weight of the affinity between them is the number of occurrence of copy instructions involving them ($a \leftarrow b$ or $b \leftarrow a$) or ϕ -functions ($a \leftarrow \phi(\dots, b, \dots)$ or $b \leftarrow \phi(\dots, a, \dots)$). Actually, some parts of the program are more often executed than others, and a more accurate cost for each copy can be calculated using actual execution frequencies based on profiling; It is also possible to use empirical parameters, for instance $\times 10$ for each level of nested loop, $\times 0.5$ if in a conditional branch, etc.

23.3.4.1 Biased coalescing

Since the number of affinities is sparse compared to the number of possible pairs of variables, it is best to keep for each variable v an “affinity list” where each element is a pair (u, w) meaning that v has an affinity with u of weight w . We propose a version of function `choose_color` that bias the choosing of colors based on these affinity lists, shown on Figure 23.8. The bias is very simple: each choice of color maximizes locally—under the current knowledge—the coalescing for the variable being colored. It is not of course optimal since the general problem of coalescing is NP-complete, and biased algorithms are known to be easily misguided.

```

global variables
    count = array of size #colors and containing only zeros

choose_color(v, available)
    for each (u,weight) in affinity_list(v)
        count[color(u)] += weight
    max_weight = -1
    col = 0
    for each color c
        if available[c] == true /* color is not used by a neighbour of v */
            if count[c] > max_weight
                col = c
                max_weight = count[c]
        count[c] = 0 /* prepare array for next call to choose_color */

```

Fig. 23.8 Choosing a color with a bias for coalescing.

23.3.4.2 Aggressive coalescing to improve biased coalescing.

There is another simple and interesting idea to improve biased coalescing. We still consider a JIT context where compilation speed is very important. **TODO:** what kind of aggressive? Ask fab...

23.3.4.3 More involved coalescing strategies

It is beyond the scope of this book to describe in details deeply involved strategies for coalescing. So we will just give a short overview of other coalescing schemes that might be used in a decoupled register allocator.

- Iterated Register Coalescing: conservative coalescing that works on the interference graph

- Brute force strategy, working on all “greedy- k -colorable” graphs

23.4 Practical and advanced discussions

23.4.1 Handling registers constraints

In theory, register allocation algorithms are always nicely working with nodes and colors. In practice, however, not all variables or registers are equivalent. Depending on the architecture, some registers might be dedicated to perform memory operations, some instructions expects their operands to reside in particular registers, and in general conventions are defined to simplify for example the writing of libraries, by specifying how parameters to functions are passed and how results are returned. This adds constraints to the register allocation, usually by restricting the coloring possibilities of variables. Some examples are given in Figure 23.9.

Context	Constraint	Instruction	Effect on reg. alloc.
x86	Division in registers R_x and R_y	a/b	a in R_x , b in R_y
st200	Memory load cannot use R_z	$a = \text{load}(b)$	b cannot be in R_z
ABI	Functions arguments in $R_1, R_2\dots$ result in R_0	$a = f(b, c)$	b in R_1 , c in R_2 a in R_0

Fig. 23.9 Examples of register constraints.

The problem with such constraints is that they cannot be expressed directly in the register allocation problem under SSA form. For instance, if variable a and b must absolutely reside in register R_1 , it is not possible to pre-color them with color R_1 as the interference graph could not be chordal anymore. Furthermore, a and b maybe interfere, for instance if they are both the first parameter of successive function calls. In that case putting them in the same register would break the program. We propose two solutions to deal with this problem. The first is classic in the literature, and the second is newer and promising.

Splitting variables to handle register constraints.

Traditionally, variables involved in constraining operations are split before and after the operation. For instance, suppose a is involved in an instruction that requires it to be in register R_x . Then the instructions $a' \leftarrow a$ and $a \leftarrow a'$ are issued before and after the division, a' is forced to be in R_x and is used instead of a in the instruction. If a happens to be in R_x after the register allocation, the copies can be removed, if not, the copies ensure the constraint is respected.

In SSA form, the problem is that variable a' is R_x , and if this constraint appears somewhere else in the program, it breaks the SSA property since there would be multiple definition of the same variable. Additionally, a also is redefined which also breaks the SSA form.

A workaround solution is to split *all* variables alive before and after the operation using a *parallel copy*, as shown on Figure 23.10. Moreover, the second copy must define new variables to keep the SSA property (and subsequent uses must be changed accordingly).

The use of parallel copies assures that the locally created variables, with very short live-ranges (span of only one instruction), are completely disconnected from the rest of the interference graph. Their only relation with the other, “normal” variable are affinities that each variable share with the two other parts coming from the same original variable.

The only problem with this method is that many copy instructions can remain in the program if the coalescing is not good enough to assign most of the created variables to the same register.

Before	After
Live in = {a, b, c, d, e}	Live in = {a, b, c, d, e}
$f = a/b$	$(a', b', c', d', e') = (a, b, c, d, e)$
Live out = {a, c, d, e, f}	$f' = a'/b'$
	Live out = {a'', c'', d'', e'', f} = {a', c', d', e', f'}
	Live out = {a'', c'', d'', e'', f}

Fig. 23.10 Examples of full parallel splitting.

Repairing problems afterwards.

Another possibility is to let register allocation do its job, and then intervene to repair the coloring whenever it does not fit the constraints, by adding copies afterwards around mismatches. To minimize the number of conflicts on constraints, it is however important to drive the coloring so that it still gives the right color whenever possible. To do so, we add a clique (complete graph) of R nodes, one for each register, to the interference graph. This clique is completely disconnected from the graph, and serves only to materialize affinities.

For instance, if variable a must reside in register R_1 we add R_1 to the affinity list of a . On the contrary, if a cannot be assigned to a register, say R_2 , an affinity of *negative weight* is added to the list, representing the added cost it will required if a is still put in register R_2 .

Existing coalescing algorithms in the literature usually do not support negative weight. However, it is possible to emulate a negative weight between two variables, say a and b with weight $w < 0$, by using the following trick: we create a new “dummy” variable, v_{ab} , and add an interference between a and v_{ab} and an affinity of weight $|w|$ between v_{ab} and b . Hence, coalescing v_{ab} and b ensures a and b will not be in the same register.

The problem with this trick is that it adds many useless nodes and interferences in the graph. In practice, it is best to modify the coalescing rules to act “as if” there were such a dummy node in case of negative weight affinities, hence not requiring an actual modification of the graph.

23.4.2 Out-of-SSA and edge splitting

ϕ -functions are not actual machine instructions, they must be replaced by actual instructions when going “out-of-SSA.” Traditionally, this translation out-of-SSA happens before register allocation. The classical “Sreedhar” method involves inserting copies before *and after* the ϕ -functions, i.e., at the end the previous basic blocks, and at the beginning of the current block. This creates a new short-lived variable for each ϕ -function, for instance, consider $a \leftarrow \phi(b, c)$, then copies $a' \leftarrow b$ and $a' \leftarrow c$ are created on the predecessors, and $a \leftarrow a'$ replaces the ϕ -function. This elegant method ensures the programs stays semantically correct, and fixes problems of earlier methods, the so-called “swap problem” and “lost-copy problem.”

When going out-of-SSA after register allocation, however, variables are already allocated to memory or registers, hence arguments and results of ϕ -functions are not variables anymore but registers. Shreedar’s method in this context would generate local, uncolored variables. Trying to color these variables could lead to a dead-end, as we come back to regular, non-SSA coloring.

We propose then a different approach. First, if the definition and arguments of a ϕ -function have been allocated to the same register, then it can be safely removed. We then replace the remaining ϕ -functions by adding copies directly *on the control-flow edges*. Of course, it is not actually possible to place code on edges, so we have to create new basic blocks on these edges to hold the added code. Since the semantics of ϕ -functions is parallel, we have to replace them with *parallel copies* on the edges. These copies can then be sequentialized using actual simple move instructions.

Sequentializing is not a difficult task, but the order of the final instructions is important. A very simple algorithm to sequentialize copies consist of first dealing with “cycles” in the parallel copy. For instance, in Figure 23.11, registers R_1 and R_2 form a cycle since their values must be swapped. We use the currently available register R_3 to act as a temporary register during the swap, which requires three copies. Then, R_3 gets the value that was contained in R_2 and which is in R_1 after the swap.

ϕ -functions	Parallel copy	Sequential copies
$R_1 \leftarrow \phi(R_2, \dots)$		$R_3 \leftarrow R_1$
$R_2 \leftarrow \phi(R_1, \dots)$	$(R_1, R_2, R_3) \leftarrow (R_2, R_1, R_2)$	$R_1 \leftarrow R_2$
$R_3 \leftarrow \phi(R_2, \dots)$		$R_2 \leftarrow R_3$ $R_3 \leftarrow R_1$

Fig. 23.11 Replacing ϕ -functions by actual copies.

In our case, we could save a copy by choosing to save R_2 into R_3 instead of R_1 for the swap. Better algorithms can be found in the literature but we prefer not to describe them here.

An important observation is that cycles can appear in parallel copies. These cycles need a “free” register to be sequentialized. What happens if there is no such free register, this happens if there is a cycle and every other register contains a live variable. In that case, there is no available register to “break” the cycle so we need to store one value in memory (spill),⁵ for instance R_1 , then perform all copies starting with $R_1 \leftarrow R_2$ and ending by loading the stored value to R_n .

23.4.3 Parallel copy motion

We have seen in the previous section that we can split control-flow edges to add copy code to get rid of ϕ -functions. Splitting edges has multiple disadvantages, it adds a new indirection in the code (such as a goto instruction), code on this block cannot be scheduled with code elsewhere, or it might prevent the use of hardware loop accelerators to name a few.

We can see however that splitting is not always necessary, it depends on the notion of *criticality* of the edge. We say that an edge is *critical* if it goes from a block with multiple successors to a block with multiple predecessors, see Figure 23.12. For instance, edges in a regular “if” construct are not critical, while the back-edge of a “do...until” loop is critical.

Code added on a non-critical edge can in fact be moved either at the bottom the predecessor block (*move up*) or at the top of the successor block (*move down*), depending on which one has only one exit or entry edge. For instance, if there is a ϕ -function after the join point of the *if*-construct of Figure 23.12, the move instructions can be placed at the end of the corresponding *true* and *false* branches.

However, in the case of a critical edge, it is dangerous to move up or down the copies, as such code would then be always executed, even if the control flow is not taken. For instance, Figure 23.12, it is not possible to move the copy $R_1 \leftarrow R_2$ up or down as this would change the behaviour of the program **TODO: figure**.

⁵ Some architecture, like VLIW, can swap the values in registers; It is also possible to swap integer values using three XOR instructions. In those cases, spilling is not necessary.

It is still possible in some cases to move copies out of critical edges. For instance, a parallel copy swaps the values in R_1 and R_2 , then it is possible to move this swap up on the preceding block, as long as the values are swapped back in their correct registers on all other outgoing edges, see Figure ?? . If all those edges are not critical, i.e., their successors have only one arriving control-flow edge in this case, then those swaps back can be moved down and no edge has been split. We call a critical edge with this property a *weak critical edge*. This one is *weak at the top* since all other edges leaving the predecessor are non-critical, meaning a copy on it can be moved up. It is also possible to have critical edges *weak at the bottom*, if all other edges arriving on the successor are non-critical, meaning we can move a parallel copy down.

In practice, not all parallel copies can be moved up or down on weak critical edges, since their effects are not always “reversible”: for instance, $R_1 \leftarrow R_2$ erases the value on R_1 and if this value is not saved elsewhere, it is not possible to get it back in R_1 on the other edges. We will not elaborate on this subject as our goal here is to show that optimizations are possible but not to cover them extensively. We advise the interested readers to read the existing literature on the subject, with some pointers given in the last section of this chapter. They will find there extended discussions on the “reversibility” of parallel copies, on the existence of “strong” critical edges (as opposed to “weak” ones), and what can be done for “abnormal” edges, i.e., control-flow edges that cannot be split. Here we will just conclude that in case we are stuck with our copies, there is always a means to solve our problem by spilling.

critical edge def todo	if-construct example todo	do...until loop example todo
---------------------------	------------------------------	---------------------------------

Fig. 23.12 Examples of critical and non-critical edges.

23.5 Further Reading & References

This section need to be completed/written:

Stuff to mention for interested readers:

- Examples of coalescing simpler / more involved because of split compilation.
- How to visualize dummy nodes with negative affinities: reference [?].
- swap and lost-copy problems: cite Briggs' thesis
- better sequentialization algorithm
- parallel copy motion

References to give:

- Chaitin et al. [65]

- IRC of George and Appel [129]
- Coalescing is NP-complete in most cases [45]
- Repairing ABI constraints: cite Fab's paper on tree-scan [?]
- Spilling techniques are based on page eviction, Belady's algorithm [28].
- Farrach and Liberatore [114] showed that Belady's algorithm gives good results on straight-line code.
- Braun and Hack [48] extended Belady's algorithm to general control-flow graphs.
- Gavril and Yannakakis [294] showed that, even in this simplistic setting (spilleverywhere “with holes”), it is NP-complete to find the minimum number of nodes to establish $\text{Maxlive} \leq R$.
- Bouchez et al. [46] showed that it is even NP-complete to find the minimum number of nodes to spill to decrease Maxlive just by one.

CHAPTER 24

Hardware Compilation using SSA

*Pedro C. Diniz
Philip Brisk*

Progress: 85%



Minor polishing in progress

Abstract This chapter describes the use of SSA-based high-level program representation for the realization of the corresponding computation using hardware circuits. We begin by highlighting the benefits of using a compiler SSA-based intermediate representation in this hardware mapping process using an illustrative example. The following sections describe hardware translation schemes for the core hardware logic or data-paths of hardware circuits. In this context we also outline several compiler transformations that benefit from the SSA-based representation of a computation. We conclude with a brief survey of various hardware compilation efforts both from academia as well as industry that have adopted SSA-based internal representations.

24.1 Brief history and overview

Hardware compilation is the process by which a high-level language, or behavioral, description of a computation is translated into a hardware-based implementation i.e., a circuit expressed in a hardware design language such as VHDL or Verilog which can be directly realized as an electrical (often digital) circuit.

While, initially, developers were forced to design hardware circuits using schematic-capture tools, high-level behavioral synthesis allowed them over the years to leverage the wealth of hardware-mapping and design exploration techniques to realize substantial productivity gains. As an example Figure 24.1 illustrates these concepts of hardware mapping for the computation expressed

as $x = (a * b) - (c * d) + f$. In figure 24.1(b) a graphical representation of a circuit that directly implements this computation is presented. Here there is a direct mapping of hardware operator such as adders and multipliers to the operations in the computation. Values are stored in register and the entire computation lasts a single (albeit long) clock cycle. The input data values are stored in the input registers on the left and a clock cycle later the corresponding results of the computation are ready and saved in the output registers on the right-hand-side of the figure. Overall this direct implementation uses two multipliers, two adders/subtractors and six registers. The execution using this design implementation can use a very simple control scheme, as it simply uses the data in the input registers, waits for a single clock cycle and stores the outputs of the operations in the output registers. In figure 24.1(c) we depict a different implementation variant of the same computation, this time using nine registers and the same amount of adders and subtractors. The increased number of registers allows for the circuit to be clocked at higher frequency as well as to be executed in a pipelined fashion. Finally, in figure 24.1(d) we depict yet another possible implementation of the same computation but using a single multiplier operator. This last version allows for the reuse in time of the multiplier operator and required thirteen register as well as multiplexers to route the inputs to the multiplier in two distinct control steps. As it is apparent, the reduction of number of operators, in this particular case the multipliers carries a penalty on increased number of registers, multiplexers¹ and increased complexity of the control scheme.

This example illustrates the many degrees of freedom in high-level behavioral hardware synthesis. Synthesis techniques perform the classical tasks of allocation, binding and scheduling of the various operations in a computation given specific target hardware resources. For instance, a designer can use behavioral synthesis tools (e.g., Mentor Graphics's Monet^(R)) to automatically derive an implementation for a computation as expressed in the example in Figure 24.1(a) by declaring that it pretends to use a single adder and a single multiplier automatically deriving an implementation that resembles the one depicts in figure 24.1(d). The tool then derives the control scheme required to route the data from registers to the selected units so as to meet the designers' goals.

Despite the introduction of high-level behavioral synthesis techniques in commercially available tools, hardware synthesis and thus hardware compilation has never enjoyed the same level of success as traditional, software compilation. Sequential programming paradigms popularized by programming languages such as C/C++ and more recently by Java, allow programmers to easily reason about program behavior as a sequence of program memory state tran-

¹ A 2×1 multiplexor is a combinatorial circuit with two data inputs, a single output and a control input, where the control input selects which of the two data inputs is transmitted to the output. It can be viewed as a hardware implementation of the C programming language selection operator: $out = (sel ? in1 : in2)$.

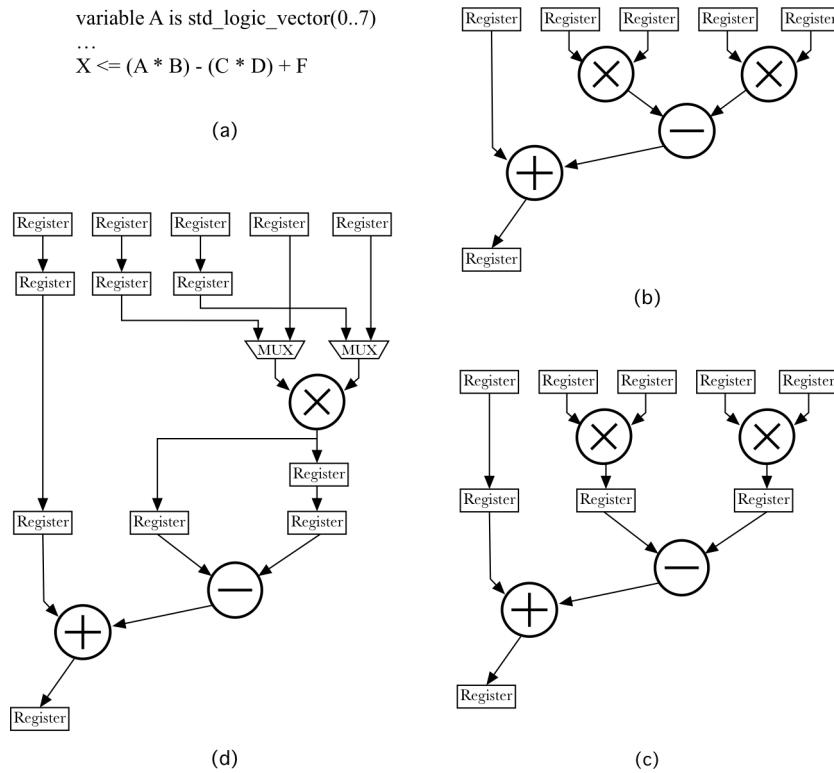


Fig. 24.1 Variants of mapping a computation to hardware: High-level source code (a); a simple hardware design (b); pipelined hardware design (c) design sharing resources (d).

sitions. The underlying processors and the corresponding system-level implementations present a number of simple unified abstractions – such as a unified memory model, a stack, and a heap – that do not exist (and often do not make sense) in customized hardware designs.

Hardware compilation, in contrast, has faced numerous obstacles that have impeded its progress and generality. When developing hardware solutions, designers must understand the concept of spatial concurrency that circuits do offer. Precise timing and synchronization between distinct hardware components are key abstractions in hardware. Solid and robust hardware design implies a detailed understanding of the precise timing of specific operations, including I/O, that simply cannot be expressed in language such as C, C++, or Java; for this reason, alternatives, such as SystemC have emerged in recent years, which give the programmer considerably more control over these issues. The inherent complexity of hardware designs has hampered the development of robust synthesis tools that can offer high-level programming abstractions enjoyed by tools that target traditional architecture and software systems, thus substantially raising

the barrier of entry for hardware designers in terms of productivity and robustness of the generated hardware solutions. At best, today hardware compilers can only handle certain subsets of mainstream high-level languages, and at worst, are limited to purely arithmetic sequences of operations with strong restrictions on control flow.

Nevertheless, the emergence of multi-core processing has led to the introduction of new parallel programming languages and parallel programming constructs that may be more amenable to hardware compilation than traditional languages and, similarly, abstractions such as a heap (e.g., pointer-based data structures) and a unified memory address space are proving to be bottlenecks with respect to effective parallel programming. For example, MapReduce, originally introduced by Google to spread parallel jobs across clusters of servers, has been an effective programming model for FPGAs; similarly, high-level languages based on parallel models of computation such as synchronous data flow, or functional single-assignment languages have also been shown to be good choices for hardware compilation.

Although the remainder of this chapter will be limited primarily to hardware compilation for imperative high-level programming languages with an obvious focus on SSA Form, many of the emerging parallel languages will retain sequential constructs, such as control-flow graphs, within a larger parallel framework. The extension of SSA Form, and SSA-like constructs, to these emerging languages, is an open area for future research; however, the fundamental uses of SSA Form for hardware compilation, as discussed in this chapter, are likely to remain generally useful.

24.2 Why use SSA for hardware compilation?

Hardware compilation, unlike its software counterpart, offers a spatially oriented computational infrastructure that presents opportunities that can leverage information exposed by the SSA representation. We illustrate the direct connection between SSA representation form and hardware compilation using the mapping of a computation example in Figure 24.2(a). Here the value of a variable v depends on the control flow of the computation as the temporary variable t can be assigned different values depending on the value of the p predicate. The representation of this computation is depicted in Figure 24.2(b) where a ϕ -function is introduced to capture the two possible assignments to the temporary variable t in both control branches of the `if-then-else` construct. Lastly, we illustrate in Figure 24.2(c) the corresponding mapping to hardware.

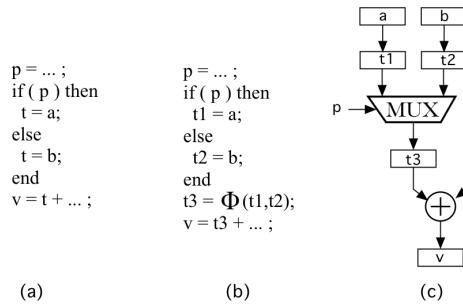


Fig. 24.2 Basic hardware mapping using SSA representation.

The basic observation is that the confluence of values for a given program variable leads to the use of a ϕ -function. This ϕ -function abstraction thus corresponds in terms of hardware implementation of the insertion of a multiplexer logic circuit. This logic circuit uses the Boolean value of a control input to select which of its input's value is to be propagated to its output. The selection or control input of a multiplexer thus acts as a gated transfer of value that parallels the actions of an *if-then-else* construct in software.

Equally important in this mapping to hardware is the notion that the computation in hardware can now take a spatial dimension. In the suggested hardware circuit in Figure 24.2(c) the computation derived from the statement in both branches of the *if-then-else* construct can be evaluated concurrently by distinct logic circuits. After the evaluation of both circuits the multiplexer will define which set of values are used based on the value of its control input, in this case of the value of the computation associated with p .

In a sequential software execution environment, the predicate p would be evaluated first, and then either branches of the *if-then-else* construct would be evaluated, based on the value of p ; as long as the register allocator is able to assign t_1 , t_2 , and t_3 to the same register, then the ϕ -function is executed implicitly; if not, it is executed as a register-to-register copy.

There have been some efforts that could automatically convert the sequential program above into a semi-spatial representation that could obtain some speedup if executed on a VLIW type of processor. For example, if-conversion (See Chapter 21) would convert the control dependency into a data dependency: statements from the *if-* and *else* blocks could be interleaved, as long as they do not overwrite one another's values, and the proper result (the ϕ -function) could be selected using a conditional-move instruction; however, in the worst case, this would effectively require the computation of both sides of the branch, rather than one, so it could actually lengthen the amount of time required to resolve the computation; in the spatial representation, in contrast, the correct result can be output as soon as two of the three inputs to the multi-

plexer are known (p , and one of t_1 or t_2 , depending on the value of p).

When targeting a hardware platform, one advantage of the SSA representation is that assigning a value to each scalar variable exactly once makes the sequences of definitions and uses of each variable both formal and explicit. A spatially-oriented infrastructure can leverage this information to perform the computation in ways that would make no sense in a traditional processor. For example, in an FPGA, one can use multiple registers in space and in time to hold the same variable, and even simultaneously assign distinct values to it; then, based on the outcome of specific predicates in the program, the hardware implementation can select the appropriate register to hold the outcome of the computation. In fact, this is precisely what was done using a multiplexer in the example above.

This example highlights the potential benefits of a SSA-based intermediate representation, when targeting hardware architectures that can easily exploit spatial computation, namely:

- Exposes the data dependences in each variable computation by explicitly incorporating into the representation each variable *def-use* chains. This allows a compiler to isolate the specific values and thus possible registers that contribute to each of the assumed values of the variable. Using separate registers for disjoint live ranges, allows hardware generation to reduce the amount of resources in multiplexer circuits, for example.
- Exposes the potential for the sharing of hardware register not only in time and but also in space providing insights for the high-level synthesis steps of allocation, binding and scheduling.
- Provides insight into control dependent regions where control predicates and thus the corresponding circuitry is identical and can thus be shared. This aspect has been so far neglected but might play an important role in the context of energy minimization.

While the Electronic Design Automation (EDA) community had for several decades now exploited similar information regarding data and control dependences for the generation of hardware circuits from increasingly higher-level representations (e.g., Behavioral HDL), SSA makes these dependences explicit in the intermediate representation itself. Similarly, the more classical compiler representations, using three-address instructions augmented with the *def-use* chains already exposes the data-flow information as for the SSA-based representation. The later however, and as we will explore in the next section, facilitates the mapping and selection of hardware resources.

24.3 Mapping a control-flow graph to hardware

In this section we are interested in hardware implementations or circuit that are spatial in nature and thus we do not address the mapping to architectures such as VLIW (Very Long Instruction Word) or Systolic Arrays. While these architectures pose very interesting and challenging issues, namely scheduling and resource use, we are more interested in exploring and highlighting the benefits of SSA representation which, we believe, are more naturally (although not exclusively) exposed in the context of spatial hardware computations.

24.3.1 Basic block mapping

As a basic block is a straight-line sequence of three-address instructions, a simple hardware mapping approach consists in composing or evaluating the operations in each instruction as a data-flow graph. The inputs and outputs of the instructions are transformed into registers² connected by nodes in the graph that represent the operators.

As a result of the "evaluation" of the instructions in the basic block this algorithm constructs a hardware circuit that has as input registers that will hold the values of the input variables to the various instructions and will have as outputs registers that hold only variables that are live outside the basic block.

24.3.2 Basic control-flow graph mapping

One can combine the various hardware circuits corresponding to a control-flow graph in two basic approaches, respectively, *spatial* and *temporal*. The spatial form of combining the hardware circuits consists in laying out the various circuits spatially by connecting variables that are live at the output of a basic block, and therefore the output registers of the corresponding hardware circuit, to the registers that will hold the values of those same variables in subsequent hardware circuits of the basic blocks that execute in sequence.

In the temporal approach the hardware circuits corresponding to the various CFG basic blocks are not directly interconnected. Instead their input and output registers are connected via dedicated buses to a local storage module. An

² As a first approach these registers are virtual and then after synthesis some of them are materialized to physical registers in a process similar to register allocation in software-oriented compilation.

execution controller "activates" a basic block or a set of basic blocks by transferring data between the storage and the input registers of the hardware circuits to be activated. Upon execution completion the controller transfers the data from the output registers of each hardware circuit to the storage module. These data transfers do not need necessarily to be carried out sequentially but instead can leverage the aggregation of the outputs of each hardware circuits to reduce transfer time to and from the storage module via dedicated wide buses.

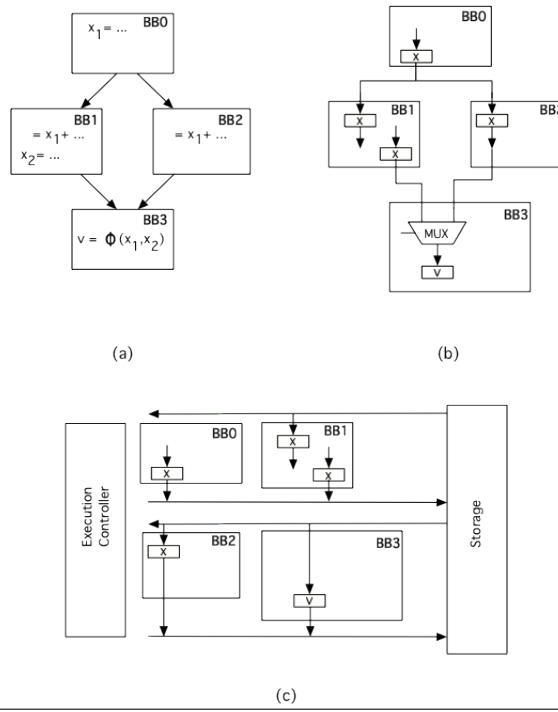


Fig. 24.3 Combination of hardware circuits for multiple basic blocks. CFG representation (a); Spatial mapping (b); Temporal mapping (c).

The temporal approach described above is well suited for the scenario where the target architecture does not have sufficient resources to simultaneously implement the hardware circuits corresponding to all basic blocks of interest as it trades off execution time for hardware resources.

In such a scenario, where hardware resources are very limited or the hardware circuit corresponding to a set of basic blocks is exceedingly large, one could opt for partitioning a basic block or set of basic blocks into smaller blocks until the space constraints for the realization of each hardware circuit are met. In reality this is the common approach in every processor today. It limits the hardware

resources to the resources required for each of the ISA instructions and schedules them in time at each step saving the state (registers) that were the output of the previous instruction. The computation thus proceed as described above by saving the values of the output registers of the hardware circuit corresponding to each smaller block.

These two approaches, illustrated in Figure 24.3, can obviously be merged in a hybrid implementation. As they lead to distinct control schemes for the orchestration of the execution of computation in hardware, their choice depends heavily on the nature and granularity of the target hardware architecture. For fine-grain hardware architectures such as FPGAs³ a spatial mapping can be favored, for coarse-grain architectures a temporal mapping is common.

While the overall execution control for the temporal mapping approach is simpler, as the transfers to and from the storage module are done upon the transfer of control between hardware circuits, a spatial mapping approach makes it more amenable to take advantage of pipelining execution techniques and speculation⁴. The temporal mapping approach can be, however, area-inefficient, as often only one basic block will execute at any point in time. This issue can, nevertheless, be mitigated by exposing additional amounts of instruction-level parallelism by merging multiple basic blocks into a single *hyper-block* and combining this aggregation with loop unrolling. Still, as these transformations and their combination, can lead to a substantial increase of the required hardware resources, a compiler can exploit resource sharing between the hardware units corresponding to distinct basic blocks to reduce the pressure on resource requirements and thus lead to feasible hardware implementation designs. As these optimizations are not specific to the SSA representation, we will not discuss them further here.

24.3.3 Control-flow graph mapping using SSA

In the case of the spatial mapping approach, the SSA form plays an important role in the minimization of multiplexers and thus in the simplification of the corresponding data-path logic and execution control.

Consider the illustrative example in Figure 24.4(a). Here basic block B0 defines a value for the variables x and y . One of the two subsequent basic blocks

³ Field-Programmable Gate Arrays are layers of two dimensional topology integrated circuits designed to be configurable after manufacturing

⁴ Speculation is also possible in the temporal mode by activating the inputs and execution of multiple hardware blocks and is only limited by the available storage bandwidth to restore the input context in each block which in the spatial approach is trivial.

BB1 redefines the value of x whereas the other basic block BB2 only reads them.

A naive implementation based exclusively on *livevariable* analysis would use for both variables x and y multiplexers to merge their values as inputs to the hardware circuit implementing basic block BB3 as depicted in Figure 24.4(b). As can be observed, however, the SSA-form representation captures the fact that such a multiplexer is only required for variable x . The value for the variable y can be propagated either from the output value in the hardware circuit for basic block BB0 (as shown in Figure 24.4(c)) or from any other register that has a valid copy of the y variable. The direct flow of the single definition point to all its uses, across the hardware circuits corresponding to the various basic blocks in the SSA form thus allows a compiler to use the minimal number of multiplexer strictly required⁵.

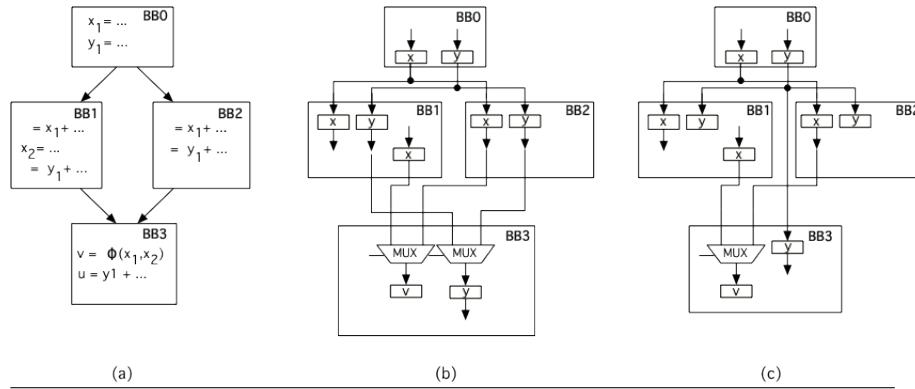


Fig. 24.4 Mapping of variable values across hardware circuit using spatial mapping; CFG in SSA form (a); naive hardware circuit using live variable information for multiplexer placement (b); hardware circuit using ϕ -function for multiplexer placement (c).

An important aspect regarding the implementation of a multiplexer associated with a ϕ -function is the definition and evaluation of the predicate associated with each multiplexer's control (or selection) input signal. In the basic SSA representation the selection predicates are not explicitly defined, as the execution of each ϕ -function is implicit when the control flow reaches it. When mapping a computation the ϕ -function to hardware, however, one has to clearly define which predicate to use to be included part of the hardware logic circuit that defines the value of the multiplexer circuit's selection input signal. The Gated-SSA form previously described in Chapter 18 of the SSA form explicitly captures the necessary information in the representation. As the variable name associated with each gating predicate fulfills the referential transparency as any SSA

⁵ Under the scenarios of a spatial mapping and with the common disclaimers about static control-flow analysis.

variable, this variant of the SSA form is particularly useful for hardware synthesis. The generation of the hardware circuit simply uses the register that holds the corresponding variable's version value of the predicate, as opposed to the base name of the variable as this might have been assigned a value that does not correspond to the actual value used in the gate. Figure 24.5 illustrates an example of a mapping using the information provided by the Gated-SSA form.

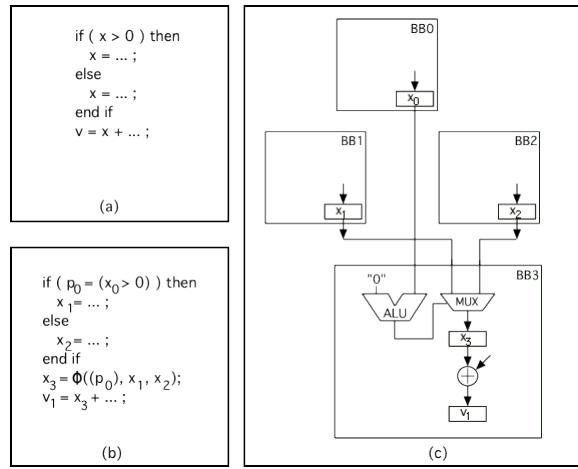


Fig. 24.5 Hardware generation example using Gated-SSA form;
(a) original code; (b) Gated-SSA representation; (c) hardware circuit implementation using spatial mapping.

When combining multiple predicates in the Gated-SSA form it is often desirable to leverage the control-flow representation in the form of the Program Dependence Graph (PDG) described in Chapter 18. In this representation, basic blocks that share common execution predicates (i.e., both execute under the same predicate conditions) are linked to the same *region* nodes. Nested execution conditions are easily recognized as the corresponding nodes are hierarchically organized in the PDG representation. As such, when generating code for a given basic block an algorithm will examine the various region nodes associated with a given basic block and compose (using AND operators) the outputs of the logic circuits that implement the predicates associated with these nodes. If an hardware circuit already exists for the same predicate, the implementation simply reuses its output signal. Otherwise, creates a new hardware circuit. This lazy code generation and predicate composition achieves the goal of hardware circuit sharing as illustrated by the example in Figure 24.6 where some of the details were omitted for simplicity.

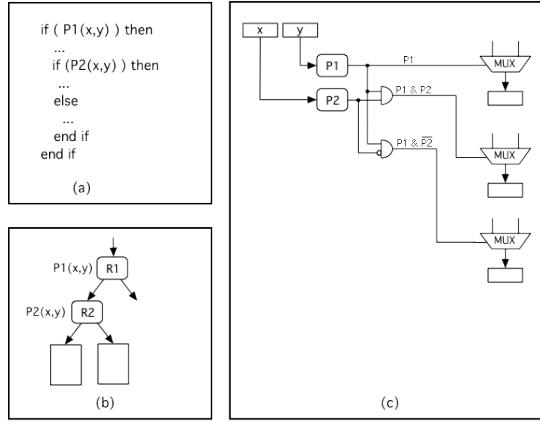


Fig. 24.6 Use of the predicates in region nodes of the PDG for mapping into the multiplexers associated with each ϕ -function; a) source code; b) the PDG representation skeleton; c) outline of hardware design implementation.

24.3.4 ϕ -function and multiplexer optimizations

We now describe a set of hardware-oriented transformations that can be applied to improve the amount of hardware resources devoted to multiplexer implementation. Although these transformations are not specific to the mapping of computations to hardware, the explicit representation of the selection constructs in SSA makes it very natural to map and therefore manipulate/transform the resulting hardware circuit using multiplexer. Other operations in the intermediate representation (e.g., predicated instructions) can also yield multiplexers in hardware without the explicit use of SSA Form.

A first transformation is motivated by a well-known result in computer arithmetic: integer addition scales with the number of operands. Building a large unified k -input integer addition circuit is more efficient than adding k integers two at a time. Moreover, hardware multipliers naturally contain multi-operand adders as building blocks: a partial product generator (a layer of AND gates) is followed by a multi-operand adder called a *partial product reduction tree*. For these reasons, there have been several efforts in recent years to apply high-level

algebraic transformations to source code with the goal of merging multiple addition operations with partial product reduction trees of multiplication operations. The basic flavor of these transformations is to push the addition operators toward the outputs of a data-flow graph, so that they can be merged at the bottom. Example of these transformations that use multiplexers are depicted in Figure 24.7(a,b). In the case of Figure 24.7(a) the transformation leads to the fact that an addition is always executed unlike in the original hardware design. Figure 24.7(c) depicts a similar transformation that merges two multiplexers sharing a common input, while exploiting the commutative property of the addition operator. The SSA-based representation facilitates these transformations as it explicitly indicates which values, and thus by tracing backwards in the representation, are involved in the computation of the corresponding values. For the example in Figure 24.7(b) the compiler would quickly detect a as a common variables in the two expressions associated with the ϕ -function.

A second transformation that can be applied to multiplexers is specific to FPGA whose basic building block consists of a k -input lookup table (LUT) logic element which can be programmed to implement any k -input logic function. For example, a 3-input LUT (3-LUT) can be programmed to implement a multiplexer with two data inputs and one selection bit; similarly, a 6-LUT can be programmed to implement a multiplexer with four data inputs and two selection bits. In particular, many FPGAs devices are organized using 4-LUTs, which are too small to implement a multiplexer with four data inputs, but leave one input unused when implementing multiplexers with two data inputs. These features can be explored to reduce the number of 4-LUTs required to implement a tree of multiplexers.

24.3.5 Implications of using SSA-form in floor-planing

For spatial oriented hardware circuits, moving a ϕ -function from one basic block to another can alter the length of the wires that are required to transmit data from the hardware circuits corresponding to the various basic blocks. As the boundaries of basic blocks are natural synchronization points, where values are captured in hardware registers, the length of wires dictate the maximum allowed hardware clock rate for synchronous designs. We illustrate this effect via an example as depicted in Figure 24.8. In this figure each basic block is mapped to a distinct hardware unit, whose spatial implementation is approximated by a rectangle. A floor-planning algorithm must place each of the units in a two-dimensional plane while ensuring that no two units overlap. As can be seen in Figure 24.8(a) placing the block 5 on the right-hand-side of the plane will result in several mid-range and one long-range wire connections. However, placing block 5 at the center of the design will virtually eliminate all mid-range con-

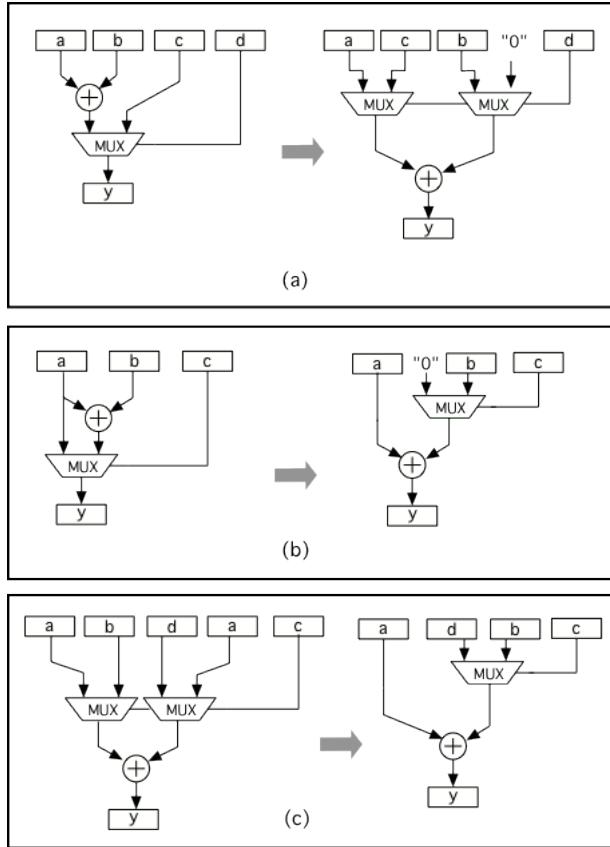


Fig. 24.7 Multiplexer-Operator transformations: juxtapose the positions of a multiplexer and an adder (a,b); reducing the number of multiplexers placed on the input of an adder (c).

nections as all connections corresponding to the transmission of the values for variable x are now next-neighboring connections.

As illustrated by this example, moving a multiplexer from one hardware unit to another can significantly change the dimensions of the resulting unit, which is not under the control of the compiler. Changing the dimensions of the hardware units fundamentally changes the placement of modules, so it is very difficult to predict whether moving a ϕ -function will actually be beneficial. For this reason, compiler optimizations that attempt to improve the physical layout must be performed using a feedback loop so that the results of the lower-level CAD tools that produce the layout can be reported back to the compiler.

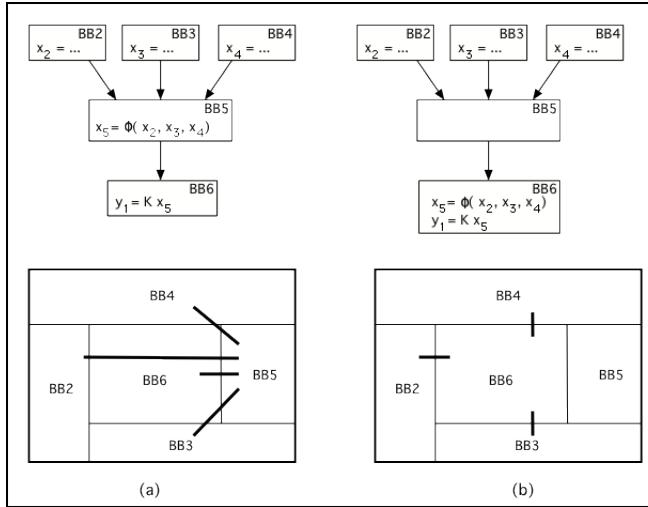


Fig. 24.8 Example of the impact of ϕ -function movement in reducing hardware wire length.

24.4 Existing SSA-based hardware compilation efforts

Several research projects have relied on SSA-based intermediate representation that leverage control- and data-flow information to exploit fine grain parallelism. Often, but not always, these efforts have been geared towards mapping computations to fine-grain hardware structure such as the ones offered by FPGAs.

The standard approach followed in these compilers has been to translate a high-level programming language such as Java in the case of the Sea Cucumber [272] compiler or C in the case of the ROCCC [135] to sequences of intermediate instructions. These sequences are then organized in basic blocks that compose the control-flow graph (CFG). For each basic block a data-flow graph (DFG) is typically extracted followed by conversion in to SSA representation, possibly using predicates associated with the control flow in the CFG thus explicitly using Predicated SSA representation (PSSA [64, 99, 261]).

As an approach to increase the potential amount of exploitable instruction-level parallelism (ILP) at the instruction level, many of these efforts (as well as others such as the earlier Garp compiler [63]) restructure the CFG into hyper-blocks [189]. An hyper-block consists on a single-entry multi-exit regions derived from the aggregation of multiple basic blocks thus serializing longer sequences of instruction. As not all instructions are executed in a hyper-block (due to early exit of a block), hardware circuit implementation must rely on predica-

tion to exploit the potential for additional ILP.

The CASH compiler [61] uses an augmented predicated SSA representation with tokens to explicitly express synchronization and handle *may-dependences* thus supporting speculative execution. This fine-grain synchronization mechanism is also used to serialize the execution of consecutive hyper-blocks, thus greatly simplifying the code generation.

I moved all the references found in the main body here. Please reorganize it. Put the reference for VHDL [18], Verilog [269], SystemC [218].

I guess a sentence summarizing what VHDL is would be helpful.

Add a reference for Mentor Graphics's if you think this is appropriate.

The following section taken from the main body contains many references that I removed. Might be put it here: *For example, MapReduce, originally introduced by Google to spread parallel jobs across clusters of servers [100], has been an effective programming model for FPGAs [295]; similarly, high-level languages based on parallel models of computation such as synchronous data flow [179], or functional single-assignment languages have also been shown to be good choices for hardware compilation [145, 138, 62].*

I removed this citation [8] concerning if-conversion and cross referred to the if-conversion chapter instead. Can be put here if you feel this is necessary.

Removed the reference for VLIW architecture [121] and for Systolic Arrays [173]. Please insert it here if appropriate.

Removed in the Section Basic CFG Mapping the references [145, 160].

I did not see any reference for FPGA that you talk about in the Basic CFG mapping section. I think you should give one and summaries in a sentence what FPGA is???

Removed this citation given for resource sharing in section Basic CFG mapping [199]).

Removed the reference [273] about gated-SSA in section SSA based CFG mapping.

Removed the reference to the PDG [116].

Removed the reference associated to the following sentence: Building a large unified k-input integer addition circuit is more efficient than adding k integers two at a time [258, 283].

Removed the reference associated with the following sentence: Figure 24.7(c) depicts a similar transformation that merges two multiplexers sharing a common input, while exploiting the commutative property of the addition operator [281].

Removed the reference associated with the following sentence: These features can be explored to reduce the number of 4-LUTs required to implement a tree of multiplexers [200].

Removed the reference [160] on the section of Floor planing.

CHAPTER 25

Building SSA form in a compiler for PHP

Paul Biggar and David Gregg, Lero@TCD, Trinity College Dublin

Progress: 55%

Review in progress

25.1 Introduction

Dynamic scripting languages such as PHP, Python and Javascript are among the most widely-used and fastest growing programming languages.

Scripting languages provide flexible high-level features, a fast modify-compile-test environment for rapid prototyping, strong integration with traditional programming languages, and an extensive standard library. Scripting languages are widely used in the implementation of many of the best known web applications of the last decade such as Facebook, Wikipedia, Gmail and Twitter. Many prominent web sites make use significant of scripting, not least because of strong support from browsers and simple integration with back-end databases.

One of the most widely used scripting languages is PHP, a general-purpose language that was originally designed for creating dynamic web pages. PHP has many of the features that are typical of dynamic scripting languages. These include simple syntax, dynamic typing, interpreted execution and run-time code generation. These simple, flexible features facilitate rapid prototyping, exploratory programming, and in the case of many non-professional websites, a copy-paste-and-modify approach to scripting by non-expert programmers.

Constructing SSA form for languages such as C/C++ and Java is a well-studied problem. Techniques exist to handle the most common features of static languages, and these solutions have been tried and tested in production level

compilers over many years. In these traditional languages it is not difficult to identify a set of scalar variables that can be safely renamed. Better analysis may lead to more such variables being identified, but significant numbers of such variables can be found with very simple analysis.

In our study of optimizing dynamic scripting languages, specifically PHP, we find this is not the case. The information required to build SSA form — that is, some conservatively complete set of unaliased scalars, and the locations of their uses and definitions—is not available directly from the program source, and can not be derived from a simple analysis. Instead, we find a litany of features whose presence must be ruled out, or heavily analysed, in order to obtain a non-pessimistic estimate.

Scripting languages commonly feature run-time code generation, built-in functions, and variable-variables, all of which may alter arbitrary unnamed variables. Less common — but still possible — features include the existence of object handlers which have the ability to alter any part of the program state, most dangerously a function’s local symbol-table.

Ruling out the presence of these features requires precise, inter-procedural, whole-program analysis. We discuss the futility of the pessimistic solution, the analyses required to provide a precise SSA form, and how the presence of variables of unknown types affect the precision of SSA.

The rest of this chapter describes our experiences with building SSA in PHC, our open source compiler for PHP. We identify the features of PHP that make building SSA difficult, outline the solutions we found to these some of these challenges, and draw some lessons about the use of SSA in analysis frameworks for PHP.

.....

25.2 SSA form in “traditional” languages

In traditional, non-scripting languages, such as C, C++ and Java, it is straightforward to identify which variables may be converted into SSA form. In Java, all scalars may be renamed. In C and C++, any scalar variables which do not have their address taken may be renamed. Figure 25.1 shows a simple C function, whose local variables may all be converted into SSA form. In Figure 25.1, it is trivial to convert each variable into SSA form. For each statement, the list of variables which are used and defined are immediately obvious from a simple syntactic check.

By contrast, Figure 25.2 contains variables which cannot be trivially converted into SSA form. On line 3, the variable x has its address taken. As a result, to convert x into SSA form, we must know that line 4 modifies x , and introduce a new version of x accordingly. Chapter hssa describes HSSA form, a powerful way to represent *indirect* modifications to scalars in SSA form.

```

1: int factorial(int num) {
2:     int result = 1, i;
3:
4:     while ( num > 1 ) {
5:         result = result * num;
6:         num--;
7:     }
8:     return result;
9: }
```

Fig. 25.1 Simple C code

```

1: void func() {
2:     int x=5, *y;
3:     y = &x;
4:     *y = 7;
5:     return x;
6: }
```

Fig. 25.2 C code with pointers

To discover the modification of *x* on line 4 requires an *alias analysis*. Alias analyses detect when multiple program names (*aliases*) represent the same memory location. In Figure 2, *x* and **y* alias each other.

There are many variants of alias analysis, of varying complexity. The most complex analyse the entire program text, taking into account control flow, function calls, and multiple levels of pointers. However, it is not difficult to perform a very simple alias analysis. *Address-taken alias analysis* identifies all variables whose addresses are taken by a referencing assignment. All variables whose address has been taken anywhere in the program are considered to alias each other (that is, all address-taken variables are in the same *alias set*).

When building SSA form with address-taken alias analysis, variables in the alias set are not renamed into SSA form. All other variables are converted. Variables not in SSA form do not possess any SSA properties, and pessimistic assumptions must be made. As a result of address-taken alias analysis, it is straightforward to convert any C program into SSA form, without sophisticated analysis. In fact, this allows more complex alias analyses to be performed on the SSA form.

.....

25.3 PHP and aliasing

In PHP, it is not possible to perform an address-taken alias analysis. The syntactic clues that C provides — notably as a result of static typing — are not available in PHP. Indeed, statements which may appear innocuous may perform complex operations behind the scenes, affecting any variable in the function. Figure 25.3

shows a small piece of PHP code. Note that in PHP variable names start with the \$ sign.

```
1: $x = 5;
2: $y =& $x;
3: $y = 7;
```

Fig. 25.3 Creating and using a reference in PHP

The most common appearance of aliasing in PHP is due to variables that store references. Creating references in PHP does not look so very different from C. In Figure 25.3, the variable *y* becomes a reference to the variable *x*. Once *y* has become an alias for *x*, the assignment to *y* (in line 3) also changes the value of *x*.

On first glance the PHP code in Figure 25.3 is not very different from similar C code in Figure 25.2. From the syntax it is easy to see that a reference to the variable *x* is taken. Thus, it is clear that *x* cannot be easily renamed. However, the problem is actually with the variable *y* which contains the reference to the variable *x*.

There is no type declaration to say that *y* in Figure 25.3 is a reference. In fact, due to dynamic typing, PHP variables may be references at one point in the program, and stop being references a moment later. Or at a given point in a program a given variable may be a reference variable or a non-reference variable depending upon the control flow that preceded that point. PHP's dynamic typing makes it difficult to simply identify when this occurs, and a sophisticated analysis over a larger region of code is essential to building a more precise conservative SSA form.

<pre>1: function snd(\$x, \$y) { 2: \$x = \$x + \$y; 3: return \$y; 4: }</pre>	<pre>1: int snd(int x, int y) { 2: x = x + y; 3: return y; 4: }</pre>
(a)	(b)

Fig. 25.4 Similar (a) PHP and (b) C functions with parameters

The size of this larger region of code is heavily influenced by the semantics of function parameters in PHP. Consider the PHP function in Figure 25.4(a). Superficially, it resembles the C function in Figure 25.4(b). In the C version, we know that *x* and *y* are simple integer variables and that no pointer aliasing relationship between them is possible. They are separate variables that can be safely renamed. In fact a relatively simple analysis can show that the assignment to *x* in line 2 of the C code can be optimized away because it is an assignment to a dead variable.

In the PHP version in Figure 25.4(a), *x* may alias *y* upon function entry. This can happen if *x* is a reference to *y*, or vice versa, or if both *x* and *y* are references

to a third variable. It is important to note, however, that the possibility of such aliasing is not apparent from the function prototype or any type declarations. Instead, whether the formal parameters x and/or y are references depends on the types of actual parameters that are passed when the function is invoked. If a reference is passed as a parameter to a function in PHP, the corresponding formal parameter in the function also becomes a reference.

The addition operation in line 2 of Figure 25.4(a) may therefore change the value of y , if x is a reference to y or vice versa. In addition, recall that dynamic typing in PHP means that whether or not a variable contains a reference can depend on control flow leading to different assignments. Therefore, on some executions of a function the passed parameters may be references, whereas on other executions they may not.

In the PHP version, there are no syntactic clues that the variables may alias. As a result, a simple conservative aliasing estimate — similar to C’s address-taken alias analysis — would need to place all variables in the alias set. This would leave no variables available for conversion to SSA form. Instead an interprocedural analysis is needed to track references between functions.

25.4 Our whole-program analysis

PHP’s dynamic typing means that program analysis cannot be performed a function at a time. As function signatures do not indicate whether parameters are references, this information must be determined by inter-procedural analysis. Furthermore, each function must be analysed with full knowledge of its calling context. This requires a whole-program analysis. We present an overview of the analysis below. A full description is beyond the scope of this chapter, and can be found elsewhere [32].

25.4.1 The analysis algorithm

The analysis is structured as a symbolic execution. This means the program is analysed by processing each statement in turn, and modelling the effect of the statement on the program state. This is similar in principle to the sparse conditional constant propagation (SCCP) algorithm sccp.¹

The SCCP algorithm models a function at a time. Instead, our algorithm models the entire program. The execution state of the program begins empty, and the analysis begins at the first statement in the program, which is placed in a worklist. The worklist is then processed a statement at a time.

¹ The analysis is actually based on a variation, conditional constant propagation [224].

For each analysed statement, the results of the analysis are stored. If the analysis results change, the statement's successors (in the control-flow graph) are added to the worklist. This is similar to CFG-edges in the SCCP algorithm. There is no parallel to the SSA edges, since the analysis is not performed on the SSA form. Instead, loops must be fully analysed if their headers change.

This analysis is therefore less efficient than the SCCP algorithm, in terms of time. It is also less efficient in terms of space. SSA form allows results to be compactly stored in a single array, using the SSA index as an array index. This is very space efficient. In our analysis, we must instead store a table of variable results at all points in the program.

Upon reaching a function or method call, the analysis begins analysing the callee function, pausing the caller's analysis. A new worklist is created, and initialized with the first statement in the callee function. The worklist is then run until it is exhausted. If another function call is analysed, the process recurses.

Upon reaching a callee function, the analysis results are copied into the scope of the callee. Once a worklist has ended, the analysis results for the exit node of the function are copied back to the calling statement's results.

25.4.2 Analysis results

Our analysis computes and stores three different kinds of results. Each kind of result is stored at each point in the program.

The first models the alias relationships in the program in a *points-to graph* [110]. The graph contains variable names as nodes, and the edges between them indicate aliasing relationships. An aliasing relationship indicates that two variables either *must*-alias, or *may*-alias. Two unconnected nodes cannot alias. A points-to graph is stored for each point in the program. Graphs are merged at CFG join points.

Secondly, our analysis also computes a conservative estimate of the types of variables in the program. Since PHP is an object-oriented language, polymorphic method calls are possible, and they must be analysed. As such, the set of possible types of each variable is stored at each point in the program. This portion of the analysis closely resembles using SCCP for type propagation [180], as described in Chapter sccp.

Finally, like the SCCP algorithm, constants are identified and propagated through the analysis of the program. Where possible, the algorithm resolves branches statically using propagated constant values. This is particularly valuable because our PHCahead-of-time compiler for PHP creates many branches in the intermediate representation during early stages of compilation. Resolving these branches statically eliminates unreachable paths, leading to significantly more precise results from the analysis algorithm.

25.4.3 Building def-use sets

To build SSA form we need to be able to identify the set of points in a program where a given variable is defined or used. Since we cannot easily identify these sets due to potential aliasing, we build them as part of our program analysis. Using our alias analysis, any variables which may be written to or read from during a statement's execution, are added to a set of defs and uses for that statement. These are then used during construction of the SSA form.

For an assignment by copy, $\$x = \y :

1. $\$x$'s value is defined.
2. $\$x$'s reference is used (by the assignment to $\$x$).
3. for each alias $\$x'$ of $\$x$, $\$x'$'s value is defined. If the alias is possible, $\$x'$'s value is may-defined instead of defined. In addition, $\$x'$'s reference is used.
4. $\$y$'s value is used.
5. $\$y$'s reference is used.

For an assignment by reference, $\$x =\& \y :

1. $\$x$'s value is defined.
2. $\$x$'s reference is defined (it is not used — $\$x$ does not maintain its previous reference relationships).
3. $\$y$'s value is used.
4. $\$y$'s reference is used.

25.4.4 HSSA

Once the set of locations where each variable is defined or used has been identified, we have the information needed to construct SSA. However, it is important to note that due to potential aliasing and potential side effects of some difficult-to-analyse PHP features (see section 25.5) many of the definitions we compute are may-definitions. Whereas a normal definition of a variable means that that variable will definitely be defined, a may-definition means that the variable may be defined at that point². In order to accommodate these may-definitions in our SSA form for PHP we use a variation of hashed SSA, described in Chapter hssa. As in HSSA, we use χ nodes, to model may-definitions.

The final distinction from traditional SSA form is that we do not only model scalar variables. All names in the program, such as fields of objects or the contents of arrays, can be represented in SSA form.

² Or more precisely, a may-definition means that a conservative analysis is unable to show that there does not exist some execution of the program where the variable is defined at that point.

25.5 Other challenging PHP features

For simplicity, the description so far of our algorithm has only considered the problems arising from aliasing due to PHP reference variables. However, in the process of construction SSA in our PHP compiler several other PHP language features that make it difficult to identify all the points in a program where a variable may be defined. In this section we briefly describe these language features and how they may be dealt with in order to conservatively identify all may-definitions of all variables in the program.

25.5.1 Run-time symbol tables

Figure 25.5 shows a program which accesses a variable indirectly. On line 2, a string value is read from the user, and stored in the variable *var_name*. On line 3, some variable—whose name is the value in *var_name*—is set to 5. That is, any variable can be updated, and the updated variable is chosen by the user at run-time. It is not possible to know whether the user has provided the value “*x*”, and so know whether *x* has the value 5 or 6.

```
1: $x = 5;
2: $var_name = readline();
3: $$var_name = 6;
4: print $x;
```

Fig. 25.5 PHP code using a variable-variable

This feature is known as *variable-variables*. They are possible because a function's symbol-table in PHP is available at run-time. Variables in PHP are not the same as variables in C. A C local variable is a name which represents a memory location on the stack. A PHP local variable is the domain of a map from strings to values. The same run-time value may be the domain of multiple string keys (references, discussed above). Similarly, variable-variables allow the symbol-table to be accessed dynamically at run-time, allowing arbitrary values to be read and written.

Upon seeing a variable-variable, all variables may be written to. In HSSA form, this creates a χ node for every variable in the program. In order to reduce the set of variables that might be updated by assigning to a variable-variable, the contents of the string stored in *var_name* may be modelled using string analysis [284]. String analysis is a static program analysis that models the structure of strings. For example, string analysis may be able to tell us that the name stored in the variable-variable (that is the name of the variable that will be written to)

begins with the letter “x”. In this case, all variables which do not begin with “x” do not require a χ node, leading to a more precise SSA form.

25.5.2 Execution of arbitrary code

PHP provides a feature that allows an arbitrary string to be executed as code. This *eval* statement simply executes the contents of any string as if the contents of the string appeared inline in the location of the *eval* statement. The resulting code can modify local or global variables in arbitrary ways. The string may be computed by the program or read in from the user or from a web form.

The PHP language also allows the contents of a file to be imported into the program text using the *include* statement. The name of the file is an arbitrary string which may be computed or read in at run time. The result is that any file can be included in the text of the program, even one that has just been created by the program. Thus, the *include* statement is potentially just as flexible as the *eval* statement for arbitrarily modifying program state.

Both of these may be modelled using the same string analysis techniques [284] as discussed in Section 25.5.1. The *eval* statement may also be handled using profiling [124], which restricts the set of possible *eval* statements to those which actually are used in practice.

25.5.3 Object handlers

PHP’s reference implementation allows classes to be partially written in C. Objects which are instantiations of these classes can have special behaviour in certain cases. For example, the objects may be dereferenced using array syntax. Or a special handler function may be called when the variable holding the object is read or written to.

The handler functions for these special cases are generally unrestricted, meaning that they can contain arbitrary C code that does anything the programmer chooses. Being written in C, they are given access to the entire program state, including all local and global variables.

There two characteristics of handlers make them very powerful, but break any attempts at function-at-a-time analysis. If one of these objects is passed in to a function, and is read, it may overwrite all variables in the local symbol table. The overwritten variables might then have the same handlers. These can then be returned from the function, or passed to any other called functions (indeed it can also call any function). This means that a single unconstrained variable in a function can propagate to any other point in the program, and we need to treat this case conservatively.

25.6 Discussion, implications and conclusion

In this chapter we have described our experience of building SSA in PHC, an ahead-of-time compiler for PHP. PHP is quite different to traditional static languages such as C/C++ and Java. In addition to dynamic typing it has other dynamic features such as very flexible and powerful reference types, variable-variables, and features that allow almost arbitrary code to be executed.

The main result of our experience is to show that SSA cannot be used and an end-to-end intermediate representation (IR) for a PHP compiler. The main reason is that in order to build SSA, significant analysis of the PHP program is needed to deal with aliasing and to rule out potential arbitrary updates of variables. We have found that in real PHP programs these features are seldom used in way that make analysis really difficult [32]. But analysis is nonetheless necessary to show the absence of the bad used of these features.

In principle our analysis could perform only the alias analysis prior to building SSA, and perform type analysis and constant propagation in SSA. But our experience is that combining all three analyses greatly improves the precision of alias analysis. In particular, type analysis significantly reduces the number of possible method callees in object-oriented PHP programs.

The complexity of building SSA for PHP suggests that we are unlikely to see SSA used in just-in-time (JIT) compilers for PHP soon. There is a great deal of interest in building JIT compilers for scripting languages, especially Javascript, but most of the effort is aimed at eliminating dynamic type checks using profiling and trace methods [125], rather than static analysis.

Once built the SSA representation forms a platform for further analysis and optimization. For example, we have used it to implement an aggressive dead-code elimination pass which can eliminate both regular assignments and reference assignments.

In recent years there has been increasing interest in the static analysis of scripting languages, with a particular focus on detecting potential security weaknesses in web-based PHP programs. For example, Jovanovic et al. [158] describe an analysis of PHP to detect vulnerabilities that allow the injection of executable code into a PHP program by a user or web form. Jensen et al. describe an alias analysis algorithm for Javascript which works quite differently to ours [32], but works well for Javascript [149].

To our knowledge we are the first to investigate building SSA in a PHP compiler. Our initial results show that building a reasonably precise SSA for PHP requires a great deal of analysis. Nonetheless, as languages such as PHP are used to build a growing number of web applications, we expect that there will be an increasing interest in tools for the static analysis of scripting languages.

Acknowledgements

This work was supported in part by Science Foundation Ireland grant 10/CE/I1855 to Lero – the Irish Software Engineering Research Centre (www.lero.ie).

*
*
*
*
*

*version|see variable
name|see variable
construction|see SSA
destruction|see SSA
out-of-SSA|see
SSA!destruction*

..... Index

Page numbers are underlined in the index when they represent the definition or the main source of information about whatever is being indexed. A page number is given in *italics* when that page contains an instructive example, use, or discussion of the concept in question.

- γ -function, 8
- ϕ reduction, reduction, 281
- ϕ removal, reduction, 281
- ϕ_{if} -function, 8, 279, 281
- ψ -SSA form, 283
- ψ -T-SSA, 220, 221
- ψ -permutation, 218
- ψ -projection, 217, 285
- ψ -promotion, 218
- ψ -reduction, 217
- χ -function, 184
- μ -function, 184
- ϕ -function, 7
- ϕ -function, as multiplexer, 7, 8
- ϕ -function, parallel semantic of, 8
- ϕ -web, 21, 36, 294
- ϕ -web, discovery, 37
- ψ -function, 214, 284
- ψ -web, 221–223
- ψ -Inlining, 217
- aggressive coalescing, 38
- aliasing, 183
- α -renaming, 74
- application binary interface, 297
- back-end, compiler, 12, 214
- basic-block early, 306
- basic-block exit, 301
- basic-block late, 306
- block
 - sinking, 82
- branch coalescing, 288
- C-SSA, 219
- coalescing, 221, 283
- coalescing, of variables, 36
- code selection, 263
- compositional equational reasoning, 74
- conditional move, 280
- conjunctive predicate merge, 282
- conservative coalescing, 38
- constant propagation, 217
- construction, *see* SA359
- construction, of ψ -SSA, 216
- continuation, 74
- control-flow graph, 7
- conventional SSA, 36
- conventional SSA form, 36
- conventional SSA form, CSSA, 36
- conventional SSA, C-SSA, 37
- conventional, making SSA, 39
- copy folding, 282
- copy-folding, 297
- copy-propagation, 41
- CPS, *see* continuation
- critical edge, 37, 293
- data flow, 186
- data-flow analysis, 10
- data-flow analysis, dense, 10
- data-flow analysis, non-zero value, 10
- data-flow analysis, sparse, 11
- data-flow tree, DFT, 263
- dead ϕ -functions, 40
- dead code elimination, 40, 187, 217
- def-use chains, 11, 32, 40, 216, 265, 285

- dependence, anti-, 284
 dependence, control-, 215, 277, 284
 dependence, data-, 277, 284
 dependence, flow-, 214
 destruction, *see* SA359
 SSA, 89
 direct style, 77
 dismissible model, 279
 DJ-graph, 32, 45
 dominance edge, D-edge, 32
 dominance frontier, 46
 dominance frontier edge, DF-edge,
 32
 dominance frontier, construction, 33
 dominance frontier, DF, 29
 dominance property, 30
 dominance property, SSA form with,
 36
 dominates, strictly, 29
 dropping
 λ , 82
 parameter, 84
 edge splitting, 37
 fixed-point arithmetic, 265
 formal parameters, 72
 forward control-flow graph, 16, 187
 free, *see* variable
 functional language, 12
 gate, 214
 gated instruction, 279, 280
 Gated-SSA form, 214, 281
 global vale numbering, 217
 Global Value Numbering, 190
 grammar, 263
 greedy coloring scheme, 318
 guard, 213, 214, 285
 Hashed-SSA form, H-SSA, 183
 hyperblock formation, 287
 if-conversion, 216, 277
 incremental update, of SSA, 286
 induction variable recognition, 217
 insertion of ϕ -function, pessimistic,
 41
 insertion, of ϕ -function, 28
 Instruction Level Parallelism, ILP, 277
 instruction set architecture, 297
 interference, 222, 223, 293
 interference graph, 19
 irreducible, 91
 iterated dominance frontier, DF⁺, 29
 join edge, J-edge, 32
 join set, 29
 just-in-time compiler, JIT, 12
 λ -dropping, 82
 λ -notation, 75
 let-binding, 73
 let-normal form, 79
 letrec, 91
 live-range, 18, 21, 24, 28, 224
 live-range splitting, 28, 33
 liveness, 19, 222
 liveness check, 224
 load, dismissible, 280
 load, speculative, 280
 local variable, 41
 loop nesting forest, 45, 91
 loop-closed, 86
 lost-copy problem, 89
 merge graph, 50
 minimal SSA form, 28, 36, 282
 minimal, making SSA, 41
 minimal, making SSA non-, 41
 monadic, 94
 myths, 13
 name, *see* variable
 normalized- ψ , 221
 parallel copy, 37
 parallel copy, sequentialization of,
 38
 parallel execution, 284
 parallel instruction, 185
 parallelism, 277

- parameter dropping, 84
- partial redundancy elimination, 217
- Partitioned Boolean Quadratic Programming, PBQP, 267
- path duplication, 282, 287
- pattern matching, 263
- phiweb, 219
- pining, 221, 297
- predicate, 213, 214
- predicated execution, 213, 277, 282
- predication, 217
- predication, full, 279
- predication, partial, 218, 279
- projection, 285
- pruned SSA form, 36, 282
- pruned SSA, building, 40
- reachable, by a definition, 28
- reaching definition., 34
- reducible, 90
- reducible CFG, 41
- referential transparency, 6, 73
- register allocation, 12, 38
- register promotion, 194
- renaming, of variable., 34
- renaming, of variables, 28, 33, 216, 282
- rule, base, 269
- rule, chain, 264
- rule, production, 269
- scheduling, 277
- scope
 - static, 73
- select, instruction, 277, 279
- semi-pruned SSA, 40
- single definition property, 28, 213
- single-entry-node/single-exit-node, SESE, 281
- speculation, 218, 219, 278, 279, 282, 284
- speculative execution, 277
- SSA
 - construction, 27
 - destruction, 27, 89, 219
 - myths, 13
- SSA destruction, 36, 293
- SSA graph, data-based, 265
- strict SSA form, 36
- strict, making SSA, 39
- strict, making SSA non-, 41
- swap problem, 89
- symbol, non-terminal, 264
- symbol, terminal, 264
- tail duplication, 287
- tail-recursive, 77
- top, T, 10
- transformed SSA, T-SSA, 37
- undefined variable, 36
- unrolling, 87
- use-def chains, 40, 186
- uses, of ψ -functions, 222, 223
- value numbering, 146
- variable
 - free, 73
 - name, 33
 - version, 33
- variable name, 7, 21, 28
- variable version, 186, 294
- version, *see* variable
- Very Long Instruction Word, WLIW, 277
- zero version, 186

References

1. ACM Press, 1993. 369
2. *Proceedings of the ACM SIGPLAN'96 Conference on Programming Language Design and Implementation (PLDI)*, volume 31. ACM Press, May 1996. 377
3. *POPL'98: Proceedings of the 25th ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages*. ACM Press, 1998. 374
4. *Proceedings of the third ACM SIGPLAN International Conference on Functional Programming (ICFP '98)*. ACM Press, 1998. SIGPLAN Notices 34(1), January 1999. 364
5. A. V. Aho and S. C. Johnson. Optimal code generation for expression trees. *J. ACM*, 23(3):488–501, 1976. 274
6. Alfred V. Aho, Ravi Sethi, and Jeffrey D. Ullman. *Compilers Principles Techniques and Tools*. Addison Wesley, 1986. 43
7. Frances E. Allen. Control flow analysis. *SIGPLAN Not.*, 5(7):1–19, 1970. 241
8. J. Allen, K. Kennedy, C. Porterfield, and J. Warren. Conversion of Control Dependence to Data Dependence. In *Proc. of the ACM Conf. on Principles of Programming Languages (POPL)*, pages 177–189. ACM Press, 1983. 346
9. J. R. Allen, Ken Kennedy, Carrie Porterfield, and Joe Warren. Conversion of control dependence to data dependence. In *Proceedings of the 10th ACM SIGACT-SIGPLAN symposium on Principles of programming languages*, POPL '83, pages 177–189, New York, NY, USA, 1983. ACM. 256
10. B. Alpern, M. Wegman, and F. Zadeck. Detecting equality of values in programs. In *Conference Record of the Fifteenth Annual ACM Symposium on Principles of Programming Languages*, pages 1–11, January 1988. 151
11. Bowen Alpern, Mark N. Wegman, and F. Kenneth Zadeck. Detecting equality of variables in programs. In *POPL '88: Proceedings of the 15th ACM SIGPLAN-SIGACT symposium on Principles of programming languages*, pages 1–11, New York, NY, USA, 1988. ACM. 11
12. Bowen Alpern, Mark N. Wegman, and F. Kenneth Zadeck. Detecting equality of variables in programs. In *Proceedings of the 15th ACM SIGPLAN-SIGACT symposium on Principles of programming languages*, POPL '88, pages 1–11. ACM, 1988. 242, 310
13. Scott Ananian. The static single information form. Master's thesis, MIT, September 1999. 180, 181
14. Andrew W. Appel. *Compiling with continuations*. Cambridge University Press, 1992. 93, 94
15. Andrew W. Appel. *Modern Compiler Implementation in ML*. Cambridge University Press, 1998. 93
16. Andrew W. Appel. SSA is functional programming. *SIGPLAN Notices*, 33(4):17–20, 1998. 93
17. Andrew W. Appel and Jens Palsberg. *Modern Compiler Implementation in Java*. Cambridge University Press, second edition, 2002. 310
18. P. Ashenden. *The Designer's Guide to VHDL*. Morgan Kaufmann Publishers Inc., San Francisco, CA, USA, 2001. 346

19. David I. August, Daniel A. Connors, Scott A. Mahlke, John W. Sias, Kevin M. Crozier, Ben-Chung Cheng, Patrick R. Eaton, Qudus B. Olaniran, and Wen-mei W. Hwu. Integrated predicated and speculative execution in the impact epic architecture. In *Proceedings of the 25th annual international symposium on Computer architecture*, ISCA '98, pages 227–237, Washington, DC, USA, 1998. IEEE Computer Society. 255
20. David I. August, Wen-Mei W. Hwu, and Scott A. Mahlke. The partial reverse if-conversion framework for balancing control flow and predication. *Int. J. Parallel Program.*, 27:381–423, October 1999. 291
21. John Aycock and Nigel Horspool. Simple generation of static single assignment form. In *Proceedings of the 9th International Conference in Compiler Construction*, volume 1781 of *Lecture Notes in Computer Science*, pages 110–125. Springer, 2000. 43
22. John Aycock and Nigel Horspool. Simple generation of static single-assignment form. *Compiler Construction*, pages 110–125, 2000. 70
23. Olaf Bachmann, Paul S. Wang, and Eugene V. Zima. Chains of recurrences a method to expedite the evaluation of closed-form functions. In *Proceedings of the international symposium on Symbolic and algebraic computation*, pages 242–249. ACM Press, 1994. 129
24. Rajkishore Barik and Vivek Sarkar. Interprocedural Load Elimination for Dynamic Optimization of Parallel Programs. In *PACT'09: Proceedings of the 18th International Conference on Parallel Architectures and Compilation Techniques*, pages 41–52, Washington, DC, USA, Sep 2009. IEEE Computer Society. 211
25. Denis Barthou, Jean-François Collard, and Paul Feautrier. Fuzzy array dataflow analysis. *J. of Parallel and Distributed Computing*, 40(2):210–226, 1997. 211
26. Samuel Bates and Susan Horwitz. Incremental program testing using program dependence graphs. In *POPL '93: Proceedings of the 20th ACM SIGPLAN-SIGACT symposium on Principles of programming languages*, pages 384–396, New York, NY, USA, 1993. ACM. 242
27. W. Baxter and H. R. Bauer, III. The program dependence graph and vectorization. In *POPL '89: Proceedings of the 16th ACM SIGPLAN-SIGACT symposium on Principles of programming languages*, pages 1–11, New York, NY, USA, 1989. ACM. 242
28. L. Belady. A Study of Replacement of Algorithms for a Virtual Storage Computer. *IBM Systems Journal*, 5:78–101, 1966. 330
29. Nick Benton, Andrew Kennedy, and George Russell. Compiling Standard ML to Java Bytecodes. In *Proceedings of the third ACM SIGPLAN International Conference on Functional Programming (ICFP '98)* [4], pages 129–140. SIGPLAN Notices 34(1), January 1999. 94
30. Lennart Beringer. Functional elimination of phi-instructions. *Electronic Notes in Theoretical Computer Science*, 176(3):3–20, 2007. 94
31. Lennart Beringer, Kenneth MacKenzie, and Ian Stark. Grail: a functional form for imperative mobile code. *Electronic Notes in Theoretical Computer Science*, 85(1), 2003. 95
32. Paul Biggar. *Design and Implementation of an Ahead-of-Time Compiler for PHP*. PhD thesis, Trinity College Dublin, 2010. 351, 356
33. Gianfranco Bilardi and Keshav Pingali. A framework for generalized control dependence. *SIGPLAN Not.*, 31(5):291–300, May 1996. 242
34. Jan Olaf Blech, Sabine Glesner, Johannes Leitner, and Steffen Mülling. Optimizing code generation from ssa form: A comparison between two formal correctness proofs in isabelle/hol. *Electronic Notes in Theoretical Computer Science*, 141(2):33–51, 2005. 43
35. David S. Blickstein, Peter W. Craig, Caroline S. Davidson, R. Neil Faiman, Jr., Kent D. Gossop, Richard B. Grove, Steven O. Hobbs, and William B. Noyce. The GEM optimizing compiler system. *Digital Technical Journal*, 4(4):121–136, Fall 1992. 253, 256
36. Rastislav Bodík, Rajiv Gupta, and Vivek Sarkar. ABCD: Eliminating array bounds checks on demand. In *PLDI '00: Proceedings of the Conference on Programming Language Design and Implementation*, pages 321–333, New York, NY, USA, 2000. ACM. 115
37. Rastislav Bodík, Rajiv Gupta, and Vivek Sarkar. ABCD: eliminating array bounds checks on demand. In *PLDI*, pages 321–333. ACM, 2000. 180, 181

38. Benoit Boissinot, Florian Brandner, Alain Darte, Benoît Dupont de Dinechin, and Fabrice Rastello. A non-iterative data-flow algorithm for computing liveness sets in strict ssa programs. In *Proceedings of the 9th Asian conference on Programming Languages and Systems*, APLAS'11, pages 137–154, Berlin, Heidelberg, 2011. Springer-Verlag. 254
39. Benoit Boissinot, Philip Brisk, Alain Darte, and Fabrice Rastello. SSI properties revisited. *ACM Transactions on Embedded Computing Systems*, 2012. Special Issue on Software and Compilers for Embedded Systems. 180, 254
40. Benoit Boissinot, Alain Darte, Benoît Dupont de Dinechin, Christophe Guillon, and Fabrice Rastello. Revisiting out-of-SSA translation for correctness, code quality, and efficiency. In *International Symposium on Code Generation and Optimization (CGO'09)*. IEEE Computer Society Press, March 2009. 310
41. Benoit Boissinot, Alain Darte, Fabrice Rastello, Benoît Dupont de Dinechin, and Christophe Guillon. Revisiting out-of-ssa translation for correctness, code quality and efficiency. In *Proceedings of the 7th annual IEEE/ACM International Symposium on Code Generation and Optimization*, pages 114–125, 2009. 43
42. Benoit Boissinot, Alain Darte, Fabrice Rastello, Benoît Dupont de Dinechin, and Christophe Guillon. Revisiting out-of-ssa translation for correctness, code quality and efficiency. In *Proceedings of the 7th annual IEEE/ACM International Symposium on Code Generation and Optimization*, CGO '09, pages 114–125, Washington, DC, USA, 2009. IEEE Computer Society. 252
43. Annalisa Bossi and Michael J. Maher, editors. *Proceedings of the 8th International ACM SIGPLAN Conference on Principles and Practice of Declarative Programming (PPDP'06)*. ACM Press, 2006. 373
44. Florent Bouchez, Alain Darte, Christophe Guillon, and Fabrice Rastello. Register allocation and spill complexity under ssa. Technical Report RR2005-33, LIP, ENS Lyon, France, August 2005. 314, 315
45. Florent Bouchez, Alain Darte, and Fabrice Rastello. On the complexity of register coalescing. In *CGO '07: Proceedings of the International Symposium on Code Generation and Optimization*, pages 102–114, Washington, DC, USA, March 2007. IEEE Computer Society Press. 330
46. Florent Bouchez, Alain Darte, and Fabrice Rastello. On the complexity of spill everywhere under ssa form. In *LCTES '07: Proceedings of the 2007 ACM SIGPLAN/SIGBED conference on Languages, compilers, and tools for embedded systems*, pages 103–112, New York, NY, USA, June 2007. ACM. 330
47. Marc M. Brandis and Hanspeter Mössenböck. Single-pass generation of static single-assignment form for structured languages. *ACM Transactions on Programming Languages and Systems*, 16(6):1684–1698, Nov 1994. 43
48. Matthias Braun and Sebastian Hack. Register Spilling and Live-Range Splitting for SSA-Form Programs. In *Compiler Construction 2009*, volume 5501 of *Lecture Notes In Computer Science*, pages 174–189. Springer, 2009. 330
49. P. Briggs, K. Cooper, and L. Simpson. Value numbering. *Software Practice and Experience*, 27(6):701–724, 1997. 151
50. P. Briggs, K. D. Cooper, T. J. Harvey, and L. Taylor Simpson. Practical Improvements to the Construction and Destruction of Static Single Assignment Form. *Software—Practice and Experience*, 28(8):859–881, 1998. 260
51. P. Briggs, K.D. Cooper, T.J. Harvey, and L.T. Simpson. Practical improvements to the construction and destruction of static single assignment form. *Software—Practice and Experience*, 28(8):859–882, 1998. 43
52. Preston Briggs, Keith D. Cooper, Timothy J. Harvey, and L. Taylor Simpson. Practical improvements to the construction and destruction of static single assignment form. *Softw., Pract. Exper.*, 28(8):859–881, 1998. 94
53. Preston Briggs, Keith D. Cooper, Timothy J. Harvey, and L. Taylor Simpson. Practical improvements to the construction and destruction of static single assignment form. *Software – Practice and Experience*, 28(8):859–881, July 1998. 309

54. Preston Briggs, Keith D. Cooper, and Linda Torczon. Rematerialization. In *PLDI '92: Proceedings of the ACM SIGPLAN 1992 conference on Programming language design and implementation*, pages 311–321, New York, NY, USA, 1992. ACM. 242
55. Philip Brisk, Foad Dabiri, Roozbeh Jafari, and Majid Sarrafzadeh. Optimal Register Sharing for High-Level Synthesis of SSA Form Programs. *IEEE Trans. on CAD of Integrated Circuits and Systems*, 25(5):772–779, 2006. 314
56. C. Bruel. If-conversion ssa framework for partially predicated vliw architectures. In *ODES 4*, pages 5–13. SIGPLAN, ACM and IEEE, March 2006. 257
57. Christian Bruel. If-conversion ssa framework. In *International Symposium on Code Generation and Optimization (CGO) ODES-04: Workshop on Optimizations for DSP and Embedded Systems*, 2006. 292
58. Christian Bruel. If-conversion for embedded vliw architectures. *IJES*, 4(1):2–16, 2009. 292
59. Zoran Budimlic, Keith D. Cooper, Timothy J. Harvey, Ken Kennedy, Timothy S. Oberg, and Steven W. Reeves. Fast copy coalescing and live-range identification. In *PLDI*, pages 25–32, 2002. 25
60. Zoran Budimlić, Keith D. Cooper, Timothy J. Harvey, Ken Kennedy, Timothy S. Oberg, and Steven W. Reeves. Fast copy coalescing and live-range identification. In *International Conference on Programming Language Design and Implementation (PLDI'02)*, pages 25–32. ACM Press, June 2002. 310
61. M. Budiu and S. Goldstein. Compiling application-specific hardware. In *Proc. of the Intl. Conf. on Field Programmable Logic and Applications (FPL)*, pages 853–863, 2002. 346
62. W. BÄÜhm, J. Hammes, B. Draper, M. Chawathe, C. Ross, R. Rinker, and W. Najjar. Mapping a single assignment programming language to reconfigurable systems. *The Journal of Supercomputing*, 21(2):117–130, February 2002. 346
63. T. Callahan, J. Hauser, and J. Wawrzynek. The Garp architecture and C compiler. *Computer*, 33(4):62–69, 2000. 345
64. L. Carter, B. Simon, B. Calder, L. Carter, and J. Ferrante. Predicated static single assignment. In *Proc. of the 1999 Intl. Conf. on Parallel Architectures and Compilation Techniques (PACT'99)*, page 245, Washington, DC, USA, 1999. IEEE Computer Society. 225, 345
65. G. J. Chaitin, M. A. Auslander, A. K. Chandra, J. Cocke, M. E. Hopkins, and P. W. Markstein. Register allocation via graph coloring. *Journal of Computer Languages*, 6:45–57, 1981. 312, 329
66. Gregory J. Chaitin. Register allocation & spilling via graph coloring. In *SIGPLAN Symp. on Compiler Construction (CC'82)*, pages 98–105. ACM Press, 1982. 310
67. Gregory J. Chaitin, Marc A. Auslander, Ashok K. Chandra, John Cocke, Martin E. Hopkins, and Peter W. Markstein. Register allocation via coloring. *Computer Languages*, 6:47–57, January 1981. 25, 310
68. Manuel M. T. Chakravarty, Gabriele Keller, and Patryk Zadarnowski. A functional perspective on SSA optimisation algorithms. *Electronic Notes in Theoretical Computer Science*, 82(2), 2003. 95
69. Craig Chambers and David Ungar. Customization: optimizing compiler technology for self, a dynamically-typed object-oriented programming language. *SIGPLAN Not.*, 24(7):146–160, 1989. 181
70. Jong Choi, Vivek Sarkar, and Edith Schonberg. Incremental computation of static single assignment form. In *Compiler Construction*, pages 223–237. IBM T.J. Watson Research Center P. O. Box 704 10598 Yorktown Heights NY P. O. Box 704 10598 Yorktown Heights NY, Springer Berlin / Heidelberg, 1996. 69
71. Jong-Deok Choi, Ron Cytron, and Jeanne Ferrante. Automatic construction of sparse data flow evaluation graphs. In *POPL '91: Proceedings of the 18th ACM SIGPLAN-SIGACT symposium on Principles of programming languages*, pages 55–66, New York, NY, USA, 1991. ACM. 25

72. Jong-Deok Choi, Ron Cytron, and Jeanne Ferrante. Automatic construction of sparse data flow evaluation graphs. In *POPL '91: Proceedings of the Symposium on Principles of Programming Languages*, pages 55–66, New York, NY, USA, 1991. ACM. 115
73. Jong-Deok Choi, Ron Cytron, and Jeanne Ferrante. Automatic construction of sparse data flow evaluation graphs. In *POPL*, pages 55–66, 1991. 181
74. F. Chow, S. Chan, R. Kennedy, S. Liu, R. Lo, and P. Tu. A new algorithm for partial redundancy elimination based on ssa form. In *Proceedings of the ACM SIGPLAN '97 Conference on Programming Language Design and Implementation*, pages 273–286, 1997. 150
75. Fred Chow, Sun Chan, Robert Kennedy, Shin-Ming Liu, Raymond Lo, and P. Tu. A new algorithm for partial redundancy elimination based on ssa form. *ACM SIGPLAN Notices*, 32(5):273 – 286, 1997. 225, 254
76. Fred C. Chow. Minimizing register usage penalty at procedure calls. In *Proceedings of the ACM SIGPLAN '88 Conference on Programming Language Design and Implementation*, pages 85–94, July 1988. 151
77. Weihaw Chuang, Brad Calder, and Jeanne Ferrante. Phi-predication for light-weight if-conversion. In *Proceedings of the international symposium on Code generation and optimization: feedback-directed and runtime optimization*, CGO '03, pages 179–190, Washington, DC, USA, 2003. IEEE Computer Society. 257
78. Weihaw Chuang, Brad Calder, and Jeanne Ferrante. Phi-predication for light-weight if-conversion. In *In Proceedings of the International Symposium on Code Generation and Optimization*, pages 179–190, 2003. 292
79. C. Click. Global code motion global value numbering. In *SIGPLAN International Conference on Programming Languages Design and Implementation*, pages 246 – 257, 1995. 225
80. Clifford Click. *Combining Analyses, Combining Optimizations*. PhD thesis, Rice University, February 1995. 69
81. J. Cocke and J. Schwartz. Programming languages and their compilers. Technical report, Courant Institute of Mathematical Sciences, New York University, April 1970. 151
82. Robert P. Colwell, Robert P. Nix, John J. O'Donnell, David B. Papworth, and Paul K. Rodman. A vliw architecture for a trace scheduling compiler. In *Proceedings of the second international conference on Architectural support for programming languages and operating systems*, ASPLOS-II, pages 180–192, Los Alamitos, CA, USA, 1987. IEEE Computer Society Press. 253
83. Keith D. Cooper, Timothy J. Harvey, and Ken Kennedy. An empirical study of iterative data-flow analysis. In *CIC '06: Proceedings of the Conference on Computing*, pages 266–276, Washington, DC, USA, 2006. IEEE Computer Society. 115
84. Keith D. Cooper, L. Taylor Simpson, and Christopher A. Vick. Operator strength reduction. *ACM Trans. Program. Lang. Syst.*, 23(5):603–625, 2001. 241, 242
85. P. Cousot and N.. Halbwachs. Automatic discovery of linear restraints among variables of a program. In *POPL*, pages 84–96. ACM, 1978. 181
86. R. Cytron, J. Ferrante, B. Rosen, and M. Wegman. Efficiently computing static single assignment form and the control dependence graph. *ACM Transactions on Programming Languages and Systems*, 13(4), 1991. 260
87. R. Cytron, J. Ferrante, B. K. Rosen, M. N. Wegman, and F. K. Zadek. Efficiently computing static single assignment form and the control dependence graph. *ACM Transactions on Programming Languages and Systems*, 13(4):451–490, October 1991. 69
88. Ron Cytron, Jeanne Ferrante, Barry K. Rosen, Mark N. Wegman, and F. Kenneth Zadeck. An efficient method of computing static single assignment form. In *POPL '89: Proceedings of the 16th ACM SIGPLAN-SIGACT symposium on Principles of programming languages*, pages 25–35, New York, NY, USA, 1989. ACM. 11, 42
89. Ron Cytron, Jeanne Ferrante, Barry K. Rosen, Mark N. Wegman, and F. Kenneth Zadeck. Efficiently computing static single assignment form and the control dependence graph. *ACM Trans. Program. Lang. Syst.*, 13(4):451–490, 1991. 24, 25

90. Ron Cytron, Jeanne Ferrante, Barry K. Rosen, Mark N. Wegman, and F. Kenneth Zadeck. Efficiently computing static single assignment form and the control dependence graph. *ACM Trans. Program. Lang. Syst.*, 13(4):451–490, 1991. 42
91. Ron Cytron, Jeanne Ferrante, Barry K. Rosen, Mark N. Wegman, and F. Kenneth Zadeck. Efficiently computing static single assignment form and the control dependence graph. *ACM Trans. Program. Lang. Syst.*, 13(4):451–490, 1991. 241
92. Ron Cytron, Jeanne Ferrante, Barry K. Rosen, Mark N. Wegman, and F. Kenneth Zadeck. Efficiently computing static single assignment form and the control dependence graph. *ACM Trans. Program. Lang. Syst.*, 13:451–490, October 1991. 309
93. Olivier Danvy, Kevin Millikin, and Lasse R. Nielsen. On one-pass CPS transformations. *Journal of Functional Programming*, 17(6):793–812, 2007. 94
94. Olivier Danvy and Ulrik Pagh Schultz. Lambda-dropping: transforming recursive equations into programs with block structure. *Theoretical Computer Science*, 248(1-2):243–287, 2000. 94
95. Alain Darte, Yves Robert, and Frederic Vivien. *Scheduling and Automatic Parallelization*. Birkhauser Boston, 1st edition, 2000. 242
96. B. Dupont de Dinechin, F. de Ferri, C. Guillon, and A. Stoutchinin. Code generator optimizations for the st120 dsp-mcu core. In *Proceedings of the 2000 international conference on Compilers, architecture, and synthesis for embedded systems*, CASES '00, pages 93–102, New York, NY, USA, 2000. ACM. 251
97. Benoît Dupont de Dinechin. A unified software pipeline construction scheme for modulo scheduled loops. In *Parallel Computing Technologies, 4th International Conference, PaCT-97*, volume 1277 of *Lecture Notes in Computer Science*, pages 189–200, 1997. 259
98. François de Ferrière. Improvements to the psi-ssa representation. In *Proceedings of the 10th international workshop on Software & compilers for embedded systems*, SCOPES '07, pages 111–121, New York, NY, USA, 2007. ACM. 225, 257
99. François de Ferrière. Improvements to the psi-SSA representation. In *Proc. of the 10th Intl. Workshop on Software and Compilers for Embedded Systems (SCOPES)*, pages 111–121, 2007. 345
100. J. Dean and S. Ghemawat. MapReduce: simplified data processing on large clusters. *Commun. ACM*, 51(1):107–113, 2008. 346
101. J. B. Dennis. First version of a data flow procedure language. In *Programming Symposium, Proceedings Colloque sur la Programmation*, pages 362–376, London, UK, 1974. Springer-Verlag. 241
102. J. B. Dennis. Data flow supercomputers. *Computer*, 13(11):48–56, 1980. 241
103. D. M. Dhamdhere. E-path_pre: partial redundancy made easy. *SIGPLAN Notices*, 37(8):53–65, 2002. 150
104. K. Drechsler and M. Stadel. A variation of knoop, rüthing and steffen's lazy code motion. *SIGPLAN Notices*, 28(5):29–38, 1993. 150
105. Evelyn Duesterwald, Rajiv Gupta, and Mary Lou Soffa. Reducing the cost of data flow analysis by congruence partitioning. In *CC '94: Proceedings of the Conference on Compiler Construction*, pages 357–373, London, UK, 1994. Springer-Verlag. 115
106. Dietmar Ebner, Florian Brandner, Bernhard Scholz, Andreas Krall, Peter Wiedermann, and Albrecht Kadlec. Generalized instruction selection using SSA -graphs. In Krisztián Flautner and John Regehr, editors, *Proceedings of the 2008 ACM SIGPLAN/SIGBED Conference on Languages, Compilers, and Tools for Embedded Systems (LCTES'08), Tucson, AZ, USA, June 12-13, 2008*, pages 31–40. ACM, 2008. 274
107. Dietmar Ebner, Florian Brandner, Bernhard Scholz, Andreas Krall, Peter Wiedermann, and Albrecht Kadlec. Generalized instruction selection using SSA-graphs. In *LCTES '08: Proceedings of the 2008 ACM SIGPLAN-SIGBED conference on languages, compilers, and tools for embedded systems*, pages 31–40, New York, NY, USA, 2008. ACM. 242
108. Erik Eckstein, Oliver König, and Bernhard Scholz. Code Instruction Selection Based on SSA-Graphs. In Krall [172], pages 49–65. 274

109. Alexandre E. Eichenberger and Edward S. Davidson. Register allocation for predicated code. In *Proceedings of the 28th annual international symposium on Microarchitecture*, MICRO 28, pages 180–191, Los Alamitos, CA, USA, 1995. IEEE Computer Society Press. 262
110. Maryam Emami, Rakesh Ghiya, and Laurie J. Hendren. Context-sensitive interprocedural points-to analysis in the presence of function pointers. In *Proceedings of the ACM SIGPLAN 1994 conference on Programming language design and implementation*, PLDI '94, pages 242–256, New York, NY, USA, 1994. ACM. 352
111. Jesse Zhixi Fang. Compiler algorithms on if-conversion, speculative predicates assignment and predicated code optimizations. In *Proceedings of the 9th International Workshop on Languages and Compilers for Parallel Computing*, LCPC '96, pages 135–153, London, UK, UK, 1997. Springer-Verlag. 256, 257
112. Jesse Zhixi Fang. Compiler algorithms on if-conversion, speculative predicates assignment and predicated code optimizations. In *Proceedings of the 9th International Workshop on Languages and Compilers for Parallel Computing*, LCPC '96, pages 135–153, London, UK, UK, 1997. Springer-Verlag. 292
113. Paolo Faraboschi, Geoffrey Brown, Joseph A. Fisher, Giuseppe Desoli, and Fred Homewood. Lx: a technology platform for customizable vliw embedded processing. *SIGARCH Comput. Archit. News*, 28(2):203–213, May 2000. 255
114. Martin Farach and Vincenzo Liberatore. On local register allocation. In *SODA '98: Proceedings of the ninth annual ACM-SIAM Symposium on Discrete Algorithms*, pages 564–573, Philadelphia, PA, USA, 1998. Society for Industrial and Applied Mathematics. 330
115. J. Ferrante, M. Mace, and B. Simons. Generating sequential code from parallel code. In *ICS '88: Proceedings of the 2nd international conference on Supercomputing*, pages 582–592, New York, NY, USA, 1988. ACM. 242
116. J. Ferrante, K. Ottenstein, and J. Warren. The program dependence graph and its use in optimization. *ACM Trans. on Programming Languages and Systems*, 9(3):319–349, 1987. 346
117. Jeanne Ferrante and Mary Mace. On linearizing parallel code. In *POPL '85: Proceedings of the 12th ACM SIGACT-SIGPLAN symposium on Principles of programming languages*, pages 179–190, New York, NY, USA, 1985. ACM. 242
118. Jeanne Ferrante, Karl J. Ottenstein, and Joe D. Warren. The program dependence graph and its use in optimization. *ACM Trans. Program. Lang. Syst.*, 9(3):319–349, 1987. 11
119. Jeanne Ferrante, Karl J. Ottenstein, and Joe D. Warren. The program dependence graph and its use in optimization. *ACM Trans. Program. Lang. Syst.*, 9(3):319–349, 1987. 242
120. Stephen J. Fink, Kathleen Knobe, and Vivek Sarkar. Unified analysis of array and object references in strongly typed languages. In *SAS '00: Proceedings of the 7th International Symposium on Static Analysis*, pages 155–174, London, UK, 2000. Springer-Verlag. 211
121. J. Fisher, P. Faraboschi, and C. Young. *Embedded computing: a VLIW approach to architecture, compilers and tools*. Elsevier Science Publishers B. V., 2005. 346
122. Joseph A. Fisher and Stefan M. Freudenberger. Predicting conditional branch directions from previous runs of a program. *SIGPLAN Not.*, 27:85–95, September 1992. 291
123. Cormac Flanagan, Amr Sabry, Bruce F. Duba, and Matthias Felleisen. The essence of compiling with continuations. In *Proceedings of the ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI'93)* [1], pages 237–247. 94
124. Michael Furr, Jong-hoon (David) An, Jeffrey S. Foster, and Michael Hicks. Static type inference for ruby. In *Proceedings of the 2009 ACM symposium on Applied Computing*, SAC '09, pages 1859–1866, New York, NY, USA, 2009. ACM. 355
125. Andreas Gal, Brendan Eich, Mike Shaver, David Anderson, David Mandelin, Mohammad R. Haghhighat, Blake Kaplan, Graydon Hoare, Boris Zbarsky, Jason Orendorff, Jesse Ruderman, Edwin W. Smith, Rick Reitmaier, Michael Bebenita, Mason Chang, and Michael Franz. Trace-based just-in-time type specialization for dynamic languages. In *Proceedings of the 2009 ACM SIGPLAN conference on Programming language design and implementation*, PLDI '09, pages 465–478, New York, NY, USA, 2009. ACM. 356

126. Fănică Gavril. The Intersection Graphs of Subtrees in Trees are exactly the Chordal Graphs. *J. Combin. Theory Ser. B*, (16):47–56, 1974. 314
127. Thomas Gawlitza, Jerome Leroux, Jan Reineke, Helmut Seidl, Gregoire Sutre, and Reinhard Wilhelm. Polynomial precise interval analysis revisited. *Efficient Algorithms*, 1:422 – 437, 2009. 181
128. Lal George and Andrew W Appel. Iterated register coalescing. *Transactions on Programming Languages and Systems (TOPLAS)*, 18(3):300 – 324, 1996. 260
129. Lal George and Andrew W. Appel. Iterated register coalescing. *ACM Transactions on Programming Languages and Systems*, 18(3):300–324, May 1996. 330
130. Lal George and Blu Matthias. Taming the ipx network processor. In *PLDI*, pages 26–37. ACM, 2003. 180
131. Michael P. Gerlek, Eric Stoltz, and Michael Wolfe. Beyond induction variables: detecting and classifying sequences using a demand-driven SSA form. *ACM Trans. Program. Lang. Syst.*, 17(1):85–122, 1995. 242
132. David M. Gillies, Dz-ching Roy Ju, Richard Johnson, and Michael Schlansker. Global predicate analysis and its application to register allocation. In *Proceedings of the 29th annual ACM/IEEE international symposium on Microarchitecture*, MICRO 29, pages 114–125, Washington, DC, USA, 1996. IEEE Computer Society. 251, 253
133. Sabine Glesner. An ASM semantics for SSA intermediate representations. In Zimmermann and Thalheim [301], pages 144–160. 95
134. J. R. Goodman and W.-C. Hsu. Code scheduling and register allocation in large basic blocks. In *Proceedings of the 2nd international conference on Supercomputing*, ICS ’88, pages 442–452, New York, NY, USA, 1988. ACM. 258
135. Z. Guo, B. Buyukkurt, J. Cortes, A. Mitra, and W. Najjar. A compiler intermediate representation for reconfigurable fabrics. *Int. J. Parallel Program.*, 36(5):493–520, 2008. 345
136. Yuri Gurevich. Sequential abstract-state machines capture sequential algorithms. *ACM Transactions on Computational Logic (TOCL)*, 1(1):77–111, 2000. 95
137. Sebastian Hack, Daniel Grund, and Gerhard Goos. Register Allocation for Programs in SSA form. In Andreas Zeller and Alan Mycroft, editors, *Compiler Construction 2006*, volume 3923. Springer, March 2006. 314, 315
138. A. Hagiescu, W.-F. Wong, D. Bacon, and R. Rabbah. A computing origami: folding streams in FPGAs. In *Proc. of the 2009 ACM/IEEE Design Automation Conf. (DAC)*, pages 282–287, 2009. 346
139. Rebecca Hasti and Susan Horwitz. Using static single assignment form to improve flow-insensitive pointer analysis. In *Proceedings of the ACM SIGPLAN 1998 Conference on Programming Language Design and Implementation*, pages 97–105, 1998. 13
140. W. Havanki, S. Banerjia, and T. Conte. Treegion scheduling for wide issue processors. *High-Performance Computer Architecture, International Symposium on*, 0:266, 1998. 258
141. Paul Havlak. Construction of thinned gated single-assignment form. In *In Proc. 6rd Workshop on Programming Languages and Compilers for Parallel Computing*, pages 477–499. Springer Verlag, 1993. 242
142. Paul Havlak. Nesting of Reducible and Irreducible Loops. *ACM Transactions on Programming Languages and Systems*, 19(4):557–567, 1997. 94, 254
143. Matthew S. Hecht and Jeffrey D. Ullman. Analysis of a simple algorithm for global data flow problems. In *POPL ’73: Proceedings of the Symposium on Principles of Programming Languages*, pages 207–217, New York, NY, USA, 1973. ACM. 114
144. Ralf Hinze and Norman Ramsey, editors. *Proceedings of the 12th ACM SIGPLAN International Conference on Functional Programming, ICFP 2007, Freiburg, Germany, October 1-3, 2007*. ACM Press, 2007. 371
145. A. Hormati, M. Kudlur, S. Mahlke, D. Bacon, and R. Rabbah. Optimus: Efficient Realization of Streaming Applications on FPGAs. In *Proc. of the 2008 Intl. Conf. on Compilers, Architecture, and Synthesis for Embedded Systems (CASES)*, pages 41–50, 2008. 346
146. Wen-Mei W. Hwu, Scott A. Mahlke, William Y. Chen, Pohua P. Chang, Nancy J. Warter, Roger A. Bringmann, Roland G. Ouellette, Richard E. Hank, Tokuzo Kiyohara, Grant E.

- Haab, John G. Holm, and Daniel M. Lavery. The superblock: an effective technique for vliw and superscalar compilation. *J. Supercomput.*, 7(1-2):229–248, May 1993. 258
147. Margarida F. Jacome, Gustavo de Veciana, and Satish Pillai. Clustered vliw architectures with predicated switching. In *Proceedings of the 38th Design Automation Conference, DAC*, pages 696–701, 2001. 256
148. Margarida F. Jacome, Gustavo De Veciana, and Satish Pillai. Clustered vliw architectures with predicated switching. In *In Proceedings of the Design Automation Conference, 2001*, pages 696–701. Online]. Available: citeseer.nj.nec.com/428688.html, 2001. 292
149. Simon Holm Jensen, Anders Møller, and Peter Thiemann. Type analysis for javascript. In *Proceedings of the 16th International Symposium on Static Analysis, SAS '09*, pages 238–255, Berlin, Heidelberg, 2009. Springer-Verlag. 356
150. Neil Johnson and Alan Mycroft. Combined code motion and register allocation using the value state dependence graph. In *In Proc. 12th International Conference on Compiler Construction (CC'03) (April 2003)*, pages 1–16, 2003. 243
151. Neil E. Johnson. Code size optimization for embedded processors. Technical Report UCAM-CL-TR-607, University of Cambridge, Computer Laboratory, November 2004. 243
152. R. Johnson, D. Pearson, and K. Pingali. The program tree structure. In *PLDI*, pages 171–185. ACM, 1994. 181
153. Richard Johnson and Keshav Pingali. Dependence-based program analysis. In *PLDI '93: Proceedings of the Conference on Programming Language Design and Implementation*, pages 78–89, New York, NY, USA, 1993. ACM. 115
154. Richard Johnson and Keshav Pingali. Dependence-based program analysis. In *PLDI*, pages 78–89. ACM, 1993. 181
155. Richard Johnson and Michael Schlansker. Analysis techniques for predicated code. In *Proceedings of the 29th annual ACM/IEEE international symposium on Microarchitecture, MICRO 29*, pages 100–113, Washington, DC, USA, 1996. IEEE Computer Society. 262
156. Thomas Johnsson. Lambda lifting: Transforming programs to recursive equations. In Jouannaud [157], pages 190–203. 94
157. Jean-Pierre Jouannaud, editor. *Functional Programming Languages and Computer Architecture, Proceedings*, volume 201 of *LNCS*. Springer, 1985. 371
158. Nenad Jovanovic, Christopher Kruegel, and Engin Kirda. Pixy: A static analysis tool for detecting web application vulnerabilities (short paper). In *2006 IEEE Symposium on Security and Privacy*, pages 258–263, 2006. 356
159. John B. Kam and Jeffrey D. Ullman. Global data flow analysis and iterative algorithms. *Journal of the ACM (JACM)*, 23(1):158–171, 1976. 114
160. A. Kaplan, P. Brisk, and R. Kastner. Data communication estimation and reduction for reconfigurable systems. In *Proc. of the 40th Design Automation Conf. (DAC'03)*, pages 616–621, 2003. 346
161. Richard Kelsey. A correspondence between continuation passing style and static single assignment form. In *Intermediate Representations Workshop*, pages 13–23, 1995. 93
162. Andrew Kennedy. Compiling with continuations, continued. In Hinze and Ramsey [144], pages 177–190. 94
163. R. Kennedy, S. Chan, S. Liu, R. Lo, P. Tu, and F. Chow. Partial redundancy elimination in ssa form. *ACM Trans. Program. Lang. Syst.*, 21(3):627–676, 1999. 150
164. R. Kennedy, F. Chow, P. Dahl, S. Liu, R. Lo, P. Tu, and M. Streich. Strength reduction via ssapre. In *Proceedings of the Seventh International Conference on Compiler Construction, 1988*. 151
165. Uday P. Khedker and Dhananjay M. Dhamdhere. Bidirectional data flow analysis: myths and reality. *SIGPLAN Not.*, 34(6):47–57, 1999. 180
166. V. Kislenkov, V. Mitrofanov, and E. Zima. Multidimensional chains of recurrences. In *Proceedings of the 1998 international symposium on symbolic and algebraic computation*, pages 199–206. ACM Press, 1998. 129

167. Kathleen Knobe and Vivek Sarkar. Array SSA form and its use in Parallelization. *Conf. Rec. Twenty-fifth ACM Symposium on Principles of Programming Languages, San Diego, California, January 1998.* 211
168. Kathleen Knobe and Vivek Sarkar. Conditional constant propagation of scalar and array references using array SSA form. In Giorgio Levi, editor, *Lecture Notes in Computer Science, 1503*, pages 33–56. Springer-Verlag, 1998. Proceedings from the 5th International Static Analysis Symposium. 211
169. J. Knoop, O. Rüthing, and B. Steffen. Lazy code motion. In *Proceedings of the ACM SIGPLAN'92 Conference on Programming Language Design and Implementation*, pages 224–234, 1992. 150
170. J. Knoop, O. Rüthing, and B. Steffen. Lazy strength reduction. *Journal of Programming Languages*, 1(1):71–91, 1993. 151
171. J. Knoop, O. Rüthing, and B. Steffen. Optimal code motion: theory and practice. *ACM Trans. Program. Lang. Syst.*, 16(4):1117–1155, 1994. 150
172. Andreas Krall, editor. *Software and Compilers for Embedded Systems, 7th International Workshop, SCOPES 2003, Vienna, Austria, September 24–26, 2003, Proceedings*, volume 2826 of *Lecture Notes in Computer Science*. Springer, 2003. 368
173. M. Lam. *A Systolic Array Optimizing Compiler*. Kluwer Academic Publishers, Norwell, MA, USA, 1989. 346
174. Peter Landin. A generalization of jumps and labels. Technical report, UNIVAC Systems Programming Research, August 1965. Reprinted in Higher Order and Symbolic Computation, 11(2):125–143, 1998, with a foreword by Hayo Thielecke. 93
175. Christopher Lapkowski and Laurie J. Hendren. Extended ssa numbering: introducing ssa properties to languages with multi-level pointers. In *Proceedings of the 1996 conference of the Centre for Advanced Studies on Collaborative research, CASCON '96*, pages 23–34. IBM Press, 1996. 254
176. Peeter Laud, Tarmo Uustalu, and Varmo Vene. Type systems equivalent to data-flow analyses for imperative languages. *Theoretical Computer Science*, 364(3):292–310, 2006. 95
177. Daniel M. Lavery and Wen-Mei W. Hwu. Unrolling-based optimizations for modulo scheduling. In *Proceedings of the 28th annual international symposium on Microarchitecture*, MICRO 28, pages 327–337, Los Alamitos, CA, USA, 1995. IEEE Computer Society Press. 258
178. Alan C. Lawrence. Optimizing compilation with the Value State Dependence Graph. Technical Report UCAM-CL-TR-705, University of Cambridge, Computer Laboratory, December 2007. 243
179. E. Lee and D. Messerschmitt. Synchronous data flow. *Proc. of the IEEE*, 75(9):1235–1245, 1987. 346
180. Alexandre Lenart, Christopher Sadler, and Sandeep K. S. Gupta. Ssa-based flow-sensitive type analysis: combining constant and type propagation. In *Proceedings of the 2000 ACM symposium on Applied computing - Volume 2, SAC '00*, pages 813–817, New York, NY, USA, 2000. ACM. 352
181. Allen Leung and Lal George. Static single assignment form for machine code. In *Proceedings of the ACM SIGPLAN 1999 conference on Programming language design and implementation, PLDI '99*, pages 204–214, New York, NY, USA, 1999. ACM. 251
182. Allen Leung and Lal George. Static single assignment form for machine code. In *International Conference on Programming Language Design and Implementation (PLDI'99)*, pages 204–214. ACM Press, 1999. 310
183. Rainer Leupers. Exploiting conditional instructions in code generation for embedded vliw processors. In *Proceedings of the conference on Design, automation and test in Europe, DATE '99*, New York, NY, USA, 1999. ACM. 256
184. R. Lo, F. Chow, R. Kennedy, S. Liu, and P. Tu. Register promotion by sparse partial redundancy elimination of loads and stores. In *Proceedings of the ACM SIGPLAN '98 Conference on Programming Language Design and Implementation*, pages 26–37, 1998. 150

185. Raymond Lo, Fred Chow, Robert Kennedy, Shin-Ming Liu, and Peng Tu. Register promotion by sparse partial redundancy elimination of loads and stores. In *PLDI*, pages 26–37. ACM, 1998. 180
186. Francesco Logozzo and Manuel Fahndrich. Pentagons: a weakly relational abstract domain for the efficient validation of array accesses. In *SAC*, pages 184–188. ACM, 2008. 181
187. P. Geoffrey Lowney, Stefan M. Freudenberger, Thomas J. Karzes, W. D. Lichtenstein, Robert P. Nix, John S. O'Donnell, and John Ruttenberg. The multiflow trace scheduling compiler. *J. Supercomput.*, 7(1-2):51–142, May 1993. 253, 255, 256, 258
188. LLVM Website. <http://llvm.cs.uiuc.edu>. 274
189. S. Mahlke, D. Lin, W. Chen, R. Hank, and R. Bringmann. Effective compiler support for predicated execution using the hyperblock. In *Proc. of the 25th Annual Intl. Symp. on Microarchitecture (MICRO)*, pages 45–54, Los Alamitos, CA, USA, 1992. IEEE Computer Society Press. 345
190. S. Mahlke, R. Ravindran, M. Schlansker, R. Schreiber, and T. Sherwood. Bitwidth cognizant architecture synthesis of custom hardware accelerators. *Computer-Aided Design of Integrated Circuits and Systems, IEEE Transactions on*, 20(11):1355–1371, 2001. 182
191. Scott A. Mahlke, William Y. Chen, Wen-mei W. Hwu, B. Ramakrishna Rau, and Michael S. Schlansker. Sentinel scheduling for vliw and superscalar processors. In *Proceedings of the fifth international conference on Architectural support for programming languages and operating systems*, ASPIOS-V, pages 238–247, New York, NY, USA, 1992. ACM. 255
192. Scott A. Mahlke, Richard E. Hank, James E. McCormick, David I. August, and Wen-Mei W. Hwu. A comparison of full and partial predicated execution support for ilp processors. In *Proceedings of the 22nd annual international symposium on Computer architecture*, ISCA '95, pages 138–150, New York, NY, USA, 1995. ACM. 253, 255
193. Scott A. Mahlke, Richard E. Hank, James E. McCormick, David I. August, and Wen mei W. Hwu. A comparison of full and partial predicated execution support for ilp processors. In *IN PROCEEDINGS OF THE 22TH INTERNATIONAL SYMPOSIUM ON COMPUTER ARCHITECTURE*, pages 138–150, 1995. 292
194. Scott A. Mahlke, David C. Lin, William Y. Chen, Richard E. Hank, and Roger A. Bringmann. Effective compiler support for predicated execution using the hyperblock. *SIGMICRO News*, 23(1-2):45–54, December 1992. 256, 258
195. Scott A. Mahlke, David C. Lin, William Y. Chen, Richard E. Hank, and Roger A. Bringmann. Effective compiler support for predicated execution using the hyperblock. *SIGMICRO News*, 23:45–54, December 1992. 291
196. Yutaka Matsuno and Atsushi Ohori. A type system equivalent to static single assignment. In Bossi and Maher [43], pages 249–260. 95
197. Cathy May. The parallel assignment problem redefined. *IEEE Trans. Softw. Eng.*, 15:821–824, June 1989. 94
198. Cathy May. The parallel assignment problem redefined. *IEEE Transactions on Software Engineering*, 15(6):821–824, June 1989. 310
199. S. Memik, G. Memik, R. Jafari, and E. Kursun. Global resource sharing for synthesis of control data flow graphs on FPGAs. In *Proc. of the 40th Design Automation Conf. (DAC'03)*, pages 604–609, 2003. 346
200. P. Metzgen and D. Nancekievill. Multiplexer restructuring for FPGA implementation cost reduction. In *Proc. of the ACM/IEEE Design Automation Conf. (DAC)*, pages 421–426, 2005. 346
201. Antoine Miné. The octagon abstract domain. *Higher Order Symbol. Comput.*, 19:31–100, 2006. 181
202. Eugenio Moggi. Notions of computation and monads. *Information and Computation*, 93(1):55–92, 1991. 94
203. E. Morel and C. Renvoise. Global optimization by suppression of partial redundancies. *Communications of the ACM*, 22(2):96–103, 1979. 150
204. Robert Morgan. Building an optimizing compiler, January 1998. 225

205. B. Murphy, V. Menon, F. Schneider, T. Shpeisman, and A. Adl-Tabatabai. Fault-safe code motion for type-safe languages. In *Proceedings of the 6th annual IEEE/ACM international symposium on Code generation and optimization*, pages 144–154, 2008. 150
206. Mangala Gowri Nanda and Saurabh Sinha. Accurate interprocedural null-dereference analysis for java. In *ICSE*, pages 133–143, 2009. 181
207. Flemming Nielson, Hanne R. Nielson, and Chris Hankin. *Principles of Program Analysis*. Springer-Verlag New York, Inc., Secaucus, NJ, USA, 1999. 114, 115
208. Flemming Nielson, Hanne Riis Nielson, and Chris Hankin. *Principles of program analysis*. Springer, 2005. 166, 180
209. Flemming Nielson, Hanne Riis Nielson, and Chris Hankin. *Principles of program analysis (2. corr. print)*. Springer, 2005. 95
210. Diego Novillo. A propagation engine for GCC. In *Proceedings of the GCC DevelopersâŽ Summit*, pages 175–184, 2005. 115
211. Ciaran O'Donnell. *High level compiling for low level machines*. PhD thesis, Ecole Nationale Supérieure des Télécommunications, 1994. Available from <ftp://enst.fr>. 93
212. Karl J. Ottenstein, Robert A. Ballance, and Arthur B. MacCabe. The program dependence web: a representation supporting control-, data-, and demand-driven interpretation of imperative languages. In *PLDI '90: Proceedings of the ACM SIGPLAN 1990 conference on Programming language design and implementation*, pages 257–271, New York, NY, USA, 1990. ACM. 11
213. Karl J. Ottenstein, Robert A. Ballance, and Arthur B. MacCabe. The program dependence web: a representation supporting control-, data-, and demand-driven interpretation of imperative languages. In *PLDI '90: Proceedings of the ACM SIGPLAN 1990 conference on Programming language design and implementation*, pages 257–271, New York, NY, USA, 1990. ACM. 242
214. Karl J. Ottenstein and Linda M. Ottenstein. The program dependence graph in a software development environment. In *SDE 1: Proceedings of the first ACM SIGSOFT/SIGPLAN software engineering symposium on Practical software development environments*, pages 177–184, New York, NY, USA, 1984. ACM. 242
215. David A. Padua, editor. *Encyclopedia of Parallel Computing*. Springer, 2011. 162
216. V.K. Paleri, Y.N. Srikant, and P. Shankar. Partial redundancy elimination: a simple, pragmatic and provably correct algorithm. *Science of Programming Programming*, 48(1):1–20, 2003. 150
217. Jens Palsberg. Type-based analysis and applications. In *Proceedings of the ACM SIGPLAN-SIGSOFT Workshop on Program analysis for software tools and engineering (PASTE'01)*, pages 20–27. ACM Press, 2001. See also <http://www.cs.ucla.edu/~palsberg/tba>. 95
218. P. Panda. SystemC: a modeling platform supporting multiple design abstractions. In *Proc. of the 14th Intl. Symp. on Systems Synthesis (ISSS'01)*, pages 75–80, New York, NY, USA, 2001. ACM. 346
219. J. C. Park and M. S. Schlansker. On predicated execution. Technical Report HPL-91-58, Hewlett Packard Laboratories, Palo Alto, California, 1991. 256
220. Fernando Magno Quintão Pereira and Jens Palsberg. Register allocation by puzzle solving. In *PLDI*, pages 216–226. ACM, 2008. 260
221. Simon Peyton Jones, Mark Shields, John Launchbury, and Andrew Tolmach. Bridging the gulf: a common intermediate language for ML and Haskell. In *POPL'98: Proceedings of the 25th ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages [3]*, pages 49–61. 94
222. Keshav Pingali and Gianfranco Bilardi. APT: A data structure for optimal control dependence computation. In *PLDI*, pages 211–222. ACM, 1995. 181
223. Keshav Pingali and Gianfranco Bilardi. Optimal control dependence computation and the roman chariots problem. In *TOPLAS*, pages 462–491. ACM, 1997. 181
224. Anthony Pioli and Michael Hind. Combining interprocedural pointer analysis and conditional constant propagation. Technical report, IBM T. J. Watson Research Center, 1999. 351

225. John Bradley Plevyak. *Optimization of Object-Oriented and Concurrent Programs*. PhD thesis, University of Illinois at Urbana-Champaign, 1996. 180
226. Gordon D. Plotkin. Call-by-name, call-by-value and the lambda-calculus. *Theoretical Computer Science*, 1(2):125–159, 1975. 93, 94
227. Sebastian Pop, Pierre Jouvelot, and Georges-Andre Silber. In and Out of SSA: a Denotational Specification. Technical report, July 2007. Available from <http://www.cri.ensmp.fr/classement/doc/E-285.pdf>. 95
228. G. Ramalingam. On Loops, Dominators, and Dominance Frontier. In *PLDI '00: Proceedings of the ACM SIGPLAN 2000 Conference on Programming Language Design and Implementation*, pages 233–241, New York, NY, USA, 2000. ACM Press. 254
229. G. Ramalingam. On loops, dominators, and dominance frontiers. *ACM Trans. Program. Lang. Syst.*, 24(5):455–490, 2002. 94
230. G Ramalingam and Thomas Reps. An Incremental Algorithm for Maintaining the Dominator Tree of a Reducible Flowgraph. In *Proceedings of the 21st ACM SIGPLAN-SIGACT symposium on Principles of programming languages - POPL '94*, pages 287–296, New York, New York, USA, 1994. ACM Press. 69
231. Ganesan Ramalingam. On sparse evaluation representations. *Theoretical Computer Science*, 277(1-2):119–147, 2002. 115, 181
232. Fabrice Rastello, Fran ois de Ferri re, and Christophe Guillon. Optimizing translation out of SSA using renaming constraints. In *International Symposium on Code Generation and Optimization (CGO'04)*, pages 265–278. IEEE Computer Society Press, 2004. 310
233. B. Ramakrishna Rau and Joseph A. Fisher. Instruction-level parallelism. In *Encyclopedia of Computer Science*, pages 883–887. John Wiley and Sons Ltd., Chichester, UK. 291
234. John H. Reppy. Optimizing nested loops using local CPS conversion. *Higher-Order and Symbolic Computation*, 15(2-3):161–180, 2002. 94
235. John C. Reynolds. Definitional interpreters for higher-order programming languages. In *Proceedings of the 25th ACM National Conference*, pages 717 – 714, 1972. Reprinted in Higher-Order and Symbolic Computation 11(4):363-397, 1998. 93
236. John C. Reynolds. On the relation between direct and continuation semantics. In *Proceedings of the 2nd Colloquium on Automata, Languages and Programming*, pages 141–156, London, UK, 1974. Springer. 93, 94
237. John C. Reynolds. The discoveries of continuations. *Lisp and Symbolic Computation*, 6(3-4):233–248, 1993. 93
238. Laurence Rideau, Bernard P. Serpette, and Xavier Leroy. Tilting at windmills with Coq: Formal verification of a compilation algorithm for parallel moves. *Journal of Automated Reasoning*, 40(4):307–326, 2008. 94
239. Andrei Alves Rimsa, Marcelo D'Amorim, and Fernando M. Q. Pereira. Tainted flow analysis on e-SSA-form programs. In *CC*, pages 124–143. Springer, 2011. 181
240. B. K. Rosen, M. N. Wegman, and F. K. Zadeck. Global value numbers and redundant computations. In *5th ACM SIGPLAN-SIGACT symposium on Principles of programming languages*, pages 12 – 27. ACM Press, 1988. 151, 310
241. Barry K. Rosen, Mark N. Wegman, and F. Kenneth Zadeck. Global value numbers and redundant computations. In *POPL '88: Proceedings of the 15th ACM SIGPLAN-SIGACT symposium on Principles of programming languages*, pages 12–27, New York, NY, USA, 1988. ACM. 11
242. Erik Ruf. Optimizing sparse representations for dataflow analysis. In *IR '95: Proceedings of the Workshop on Intermediate Representations*, pages 50–61, New York, NY, USA, 1995. ACM. 115
243. Vivek Sarkar and Stephen Fink. Efficient dependence analysis for java arrays. In Rizos Sakellariou, John Gurd, Len Freeman, and John Keane, editors, *Euro-Par 2001 Parallel Processing*, volume 2150 of *Lecture Notes in Computer Science*, pages 273–277. Springer Berlin / Heidelberg, 2001. 211
244. Stefan Sch fer and Bernhard Scholz. Optimal chain rule placement for instruction selection based on SSA graphs. In *SCOPES '07: Proceedings of the 10th international workshop*

- on Software & compilers for embedded systems*, pages 91–100, New York, NY, USA, 2007. ACM. 242, 271, 274
245. Michael Schlansker, Scott Mahlke, and Richard Johnson. Control cpr: a branch height reduction optimization for epic architectures. In *Proceedings of the ACM SIGPLAN 1999 conference on Programming language design and implementation*, PLDI '99, pages 155–168, New York, NY, USA, 1999. ACM. 259
 246. Robert M. Shapiro and Harry Saint. The representation of algorithms. Technical Report RADC-TR-69-313, Rome Air Development Center, September 1969. 43
 247. B. Simons, D. Alpern, and J. Ferrante. A foundation for sequentializing parallel code. In *SPAA '90: Proceedings of the second annual ACM symposium on Parallel algorithms and architectures*, pages 350–359, New York, NY, USA, 1990. ACM. 242
 248. Jeremy Singer. Sparse bidirectional data flow analysis as a basis for type inference. In *APPSEM '04: Web Proceedings of the Applied Semantics Workshop*, 2004. 115
 249. Jeremy Singer. *Static Program Analysis based on Virtual Register Renaming*. PhD thesis, Computer Laboratory, University of Cambridge, Cambridge, UK, 2005. 115
 250. Jeremy Singer. *Static Program Analysis Based on Virtual Register Renaming*. PhD thesis, University of Cambridge, 2006. 180, 181
 251. Vugranam C. Sreedhar and Guang R. Gao. A linear time algorithm for placing ϕ -nodes. In *Proceedings of the 22nd ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages*, pages 62–73, 1995. 43
 252. Vugranam C. Sreedhar, Guang R. Gao, and Yong-fong Lee. Incremental computation of dominator trees. In *Papers from the 1995 ACM SIGPLAN workshop on Intermediate representations*, pages 1–12, New York, NY, USA, 1995. ACM. 69
 253. Vugranam C. Sreedhar, Guang R. Gao, and Yong-fong Lee. A new framework for exhaustive and incremental data flow analysis using DJ graphs. In *Proceedings of the ACM SIGPLAN 1996 conference on Programming language design and implementation*, pages 278–290, New York, NY, USA, 1996. ACM. 69
 254. Vugranam C. Sreedhar, Guang R. Gao, and Yong-Fong Lee. Identifying loops using dj graphs. *ACM Transactions on Programming Languages and Systems*, 18(6):649–658, 1996. 94
 255. Vugranam C. Sreedhar, Roy Dz-Ching Ju, David M. Gillies, and Vatsa Santhanam. Translating out of static single assignment form. In *SAS*, pages 194–210, 1999. 25, 226, 254, 260
 256. Vugranam C. Sreedhar, Roy Dz-Ching Ju, David M. Gillies, and Vatsa Santhanam. Translating out of static single assignment form. In *Static Analysis Symposium (SAS'99)*, pages 194 – 204, 1999. 309
 257. James Stanier. Ph.D. thesis. Technical report, University of Sussex, School of Informatics, October 2011. 243
 258. P. Stelling, C. Martel, V. Oklobdzija, and R. Ravi. Optimal circuits for parallel multipliers. *IEEE Trans. on Computers*, 47(3):273–285, March 1998. 346
 259. Mark Stephenson, Jonathan Babb, and Saman Amarasinghe. Bidwidth analysis with application to silicon compilation. In *PLDI*, pages 108–120. ACM, 2000. 181
 260. Eric Stoltz, Michael P. Gerlek, and Michael Wolfe. Extended SSA with factored use-def chains to support optimization and parallelism. In *In Proceedings of the 27th Annual Hawaii International Conference on System Sciences*, pages 43–52, 1993. 242
 261. A. Stoutchinin and F. de Ferrière. Efficient static single assignment form for predication. In *Proc. of the 34th Annual Intl. Symp. on Microarchitecture (MICRO)*, pages 172–181, Dec. 2001. 345
 262. Arthur Stoutchinin and Francois de Ferriere. Efficient static single assignment form for predication. In *Proceedings of the 34th annual ACM/IEEE international symposium on Microarchitecture*, MICRO 34, pages 172–181, Washington, DC, USA, 2001. IEEE Computer Society. 225, 253, 257
 263. Arthur Stoutchinin and Guang R. Gao. If-conversion in ssa form. In *10th International Euro-Par Conference*, volume 3149 of *Lecture Notes in Computer Science*, pages 336–345. Springer, 2004. 257

264. Arthur Stoughton and Guang R Gao. If-conversion in ssa form. *EuroPar 2004 Parallel Processing*, pages 336–345, 2004. 292
265. Zhendong Su and David Wagner. A class of polynomially solvable range constraints for interval analysis without widenings. *Theoretical Computer Science*, 345(1):122–138, 2005. 181
266. Gerald Jay Sussman and Guy Lewis Steele Jr. Scheme: An interpreter for extended lambda calculus. Technical Report AI Lab Memo AIM-349, MIT AI Lab, December 1975. Reprinted in Higher-Order and Symbolic Computation 11(4):405-439, 1998. 94
267. David Tarditi, J. Gregory Morrisett, Perry Cheng, Christopher A. Stone, Robert Harper, and Peter Lee. Til: A type-directed optimizing compiler for ML. In *Proceedings of the ACM SIGPLAN'96 Conference on Programming Language Design and Implementation (PLDI) [2]*, pages 181–192. 94
268. Andre L. C. Tavares, Benoit Boissinot, Mariza A. S. Bigonha, Roberto Bigonha, Fernando M. Q. Pereira, and Fabrice Rastello. A program representation for sparse dataflow analyses. *Science of Computer Programming*, X(X):2–25, 201X. Invited paper with publication expected for 2012. 180
269. D. Thomas and P. Moorby. *The Verilog Hardware Description Language*. Kluwer Academic Publishers, Norwell, MA, USA, 1998. 346
270. Sam Tobin-Hochstadt and Matthias Felleisen. The design and implementation of typed scheme. *POPL*, pages 395–406, 2008. 182
271. Andrew P. Tolmach and Dino Oliva. From ML to Ada: Strongly-typed language interoperability via source translation. *Journal of Functional Programming*, 8(4):367–412, 1998. 94
272. J. Tripp, P. Jackson, and B. Hutchings. Sea Cucumber: A Synthesizing Compiler for FPGAs. In *Proc. of the 12th Intl. Conf. on Field-Programmable Logic and Applications (FPL'02)*, pages 875–885, London, UK, 2002. Springer-Verlag. 345
273. P. Tu and D. Padua. Gated SSA-based demand driven symbolic analysis for parallelizing compilers. In *Proc. of the ACM Intl. Conf. on Supercomputing (SC'95)*, pages 414–423, New York, NY, USA, 1995. ACM. 346
274. Peng Tu and David Padua. Efficient building and placing of gating functions. In *PLDI '95: Proceedings of the ACM SIGPLAN 1995 conference on Programming language design and implementation*, pages 47–55, New York, NY, USA, 1995. ACM. 242
275. Peng Tu and David Padua. Gated SSA-based demand-driven symbolic analysis for parallelizing compilers. In *Proceedings of the 9th international conference on Supercomputing, ICS '95*, pages 414–423, New York, NY, USA, 1995. ACM. 242
276. Eben Upton. Optimal sequentialization of gated data dependence graphs is NP-complete. In *PDPTA*, pages 1767–1770, 2003. 243
277. Eben Upton. Compiling with data dependence graphs, 2006. 243
278. Adriaan van Wijngaarden. Recursive definition of syntax and semantics. In T. B. Steel Jr., editor, *Formal Language Description Languages for Computer Programming*, pages 13 –24. North-Holland, 1966. 93
279. Thomas Vandrunen and Antony L. Hosking. Anticipation-based partial redundancy elimination for static single assignment form. *Software — Practice and Experience*, 34(15):1413–1439, December 2004. 151
280. Thomas Vandrunen and Antony L. Hosking. Value-based partial redundancy elimination. In *Proceedings of the 13th International Conference on Compiler Construction*, pages 167–184, 2004. 151
281. A. Verma, P. Brisk, and P. Ienne. Dataflow transformations to maximize the use of carry-save representation in arithmetic circuits. *IEEE Trans. on Computer-Aided Design of Integrated Circuits and Systems*, 27(10):1761–1774, October 2008. 346
282. Christopher P. Wadsworth. Continuations revisited. *Higher-Order and Symbolic Computation*, 13(1/2):131–133, 2000. 93
283. C.S. Wallace. A suggestion for a fast multiplier. *IEEE Trans. on Electronic Computers*, 13(14):14–17, February 1964. 346

284. Gary Wassermann and Zhendong Su. Sound and precise analysis of web applications for injection vulnerabilities. In *Proceedings of the 2007 ACM SIGPLAN conference on Programming language design and implementation*, PLDI '07, pages 32–41, New York, NY, USA, 2007. ACM. 354, 355
285. M. Wegman and K. Zadeck. Constant propagation with conditional branches. *ACM Transactions on Programming Languages and Systems*, 13(2):181–210, April 1991. 115
286. M. Wegman and K. Zadeck. Constant propagation with conditional branches. *ACM Transactions on Programming Languages and Systems*, 13(2):181 – 210, 1991. 225
287. Mark N. Wegman and F. Kenneth Zadeck. Constant propagation with conditional branches. *ACM Transactions on Programming Languages and Systems*, 13(2):181–210, 1991. 95
288. Daniel Weise, Roger F. Crew, Michael Ernst, and Bjarne Steensgaard. Value dependence graphs: representation without taxation. In *POPL '94: Proceedings of the 21st ACM SIGPLAN-SIGACT symposium on Principles of programming languages*, pages 297–310, New York, NY, USA, 1994. ACM. 243
289. M. Wolfe. Beyond induction variables. In *SIGPLAN International Conference on Programming Languages Design and Implementation*, pages 162 – 174, 1992. 225
290. Michael Wolfe. Beyond induction variables. *SIGPLAN Not.*, 27(7):162–174, 1992. 242
291. Michael Wolfe. $J+=j$. *SIGPLAN Not.*, 29(7):51–53, 1994. 25
292. J. Xue and Q. Cai. A lifetime optimal algorithm for speculative pre. *ACM Transactions on Architecture and Code Optimization*, 3(2):115–155, 2006. 150
293. J. Xue and J. Knoop. A fresh look at pre as a maximum flow problem, 2006. 150
294. Mihalis Yannakakis and Fanica Gavril. The maximum k-colorable subgraph problem for chordal graphs. *Information Processing Letters*, 24(2):133–137, 1987. 330
295. J. Yeung, C. Tsang, K. Tsoi, B. Kwan, C. Cheung, A. Chan, and P. Leong. Map-reduce as a programming model for custom computing machines. In *Proc. of the 2008 16th Intl. Symp. on Field-Programmable Custom Computing Machines (FCCM'08)*, pages 149–159, Washington, DC, USA, 2008. IEEE Computer Society. 346
296. Cliff Young and Michael D. Smith. Better global scheduling using path profiles. In *Proceedings of the 31st annual ACM/IEEE international symposium on Microarchitecture*, MICRO 31, pages 115–123, Los Alamitos, CA, USA, 1998. IEEE Computer Society Press. 259
297. Frank Kenneth Zadeck. *Incremental Data Flow Analysis in a Structured Program Editor*. PhD thesis, Rice University, 1984. 181
298. H. Zhou, W.G. Chen, and F. Chow. An ssa-based algorithm for optimal speculative code motion under an execution profile. In *Proceedings of the ACM SIGPLAN '11 Conference on Programming Language Design and Implementation*, 2011. 150
299. Hucheng Zhou, Wenguang Chen, and Fred C. Chow. An ssa-based algorithm for optimal speculative code motion under an execution profile. In *Proceedings of the 32nd ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI 2011)*, pages 98–108. ACM, 2011. 195
300. Eugene V. Zima. On computational properties of chains of recurrences. In *Proceedings of the 2001 international symposium on symbolic and algebraic computation*, pages 345–352. ACM Press, 2001. 129
301. Wolf Zimmermann and Bernhard Thalheim, editors. *Abstract State Machines 2004. Advances in Theory and Practice, 11th International Workshop (ASM 2004), Proceedings*, volume 3052 of *LNCS*. Springer, 2004. 370